

Cardano.BM - benchmarking and logging

Alexander Diemand

Andreas Triantafyllos

November 2018

Abstract

This is a framework that combines logging, benchmarking and monitoring. Complex evaluations of STM or monadic actions can be observed from outside while reading operating system counters before and after, and calculating their differences, thus relating resource usage to such actions. Through interactive configuration, the runtime behaviour of logging or the measurement of resource usage can be altered. Further reduction in logging can be achieved by redirecting log messages to an aggregation function which will output the running statistics with less frequency than the original message.

Contents

1	Cardano BM	3
1.1	Overview	3
1.2	Introduction	3
1.2.1	Logging with <i>Trace</i>	3
1.2.2	Measuring <i>Observables</i>	3
1.2.3	Monitoring	3
1.2.4	Information reduction in <i>Aggregation</i>	3
1.2.5	Output selection	3
1.2.6	Setup procedure	3
1.3	Examples	3
1.3.1	Observing evaluation of a STM action	3
1.3.2	Observing evaluation of a monad action	3
1.4	Code listings	3
1.4.1	Cardano.BM.Observer.STM	3
1.4.2	Cardano.BM.Observer.Monadlic	6
1.4.3	BaseTrace	7
1.4.4	Cardano.BM.Trace	8
1.4.5	Cardano.BM.Setup	12
1.4.6	Cardano.BM.Counters	13
1.4.7	Cardano.BM.Counters.Common	14
1.4.8	Cardano.BM.Counters.Dummy	14
1.4.9	Cardano.BM.Counters.Linux	15
1.4.10	Cardano.BM.Data.Aggregated	21
1.4.11	Cardano.BM.Data.Backend	22
1.4.12	Cardano.BM.Data.Configuration	23
1.4.13	Cardano.BM.Data.Counter	24
1.4.14	Cardano.BM.Data.LogItem	26
1.4.15	Cardano.BM.Data.Observable	27
1.4.16	Cardano.BM.Data.Output	27
1.4.17	Cardano.BM.Data.Severity	28
1.4.18	Cardano.BM.Data.SubTrace	28
1.4.19	Cardano.BM.Data.Trace	28
1.4.20	Cardano.BM.Configuration	29
1.4.21	Cardano.BM.Configuration.Model	29
1.4.22	Cardano.BM.Output.Switchboard	34
1.4.23	Cardano.BM.Output.Log	36
1.4.24	Cardano.BM.Output.EKGView	42

1.4.25 Cardano.BM.Output.Aggregation	43
--	----

Chapter 1

Cardano BM

1.1 Overview

In figure 1.1 we display the relationships among modules in *Cardano.BM*. The arrows indicate import of a module. The arrows with a triangle at one end would signify "inheritance", but we use it to show that one module replaces the other in the namespace, thus refines its interface.

1.2 Introduction

1.2.1 Logging with *Trace*

1.2.2 Measuring *Observables*

1.2.3 Monitoring

1.2.4 Information reduction in *Aggregation*

1.2.5 Output selection

1.2.6 Setup procedure

1.3 Examples

1.3.1 Observing evaluation of a STM action

1.3.2 Observing evaluation of a monad action

1.4 Code listings

1.4.1 Cardano.BM.Observer.STM

$$\begin{aligned} stmWithLog &:: STM.STM(t, [LogObject]) \rightarrow STM.STM(t, [LogObject]) \\ stmWithLog \ action &= action \end{aligned}$$

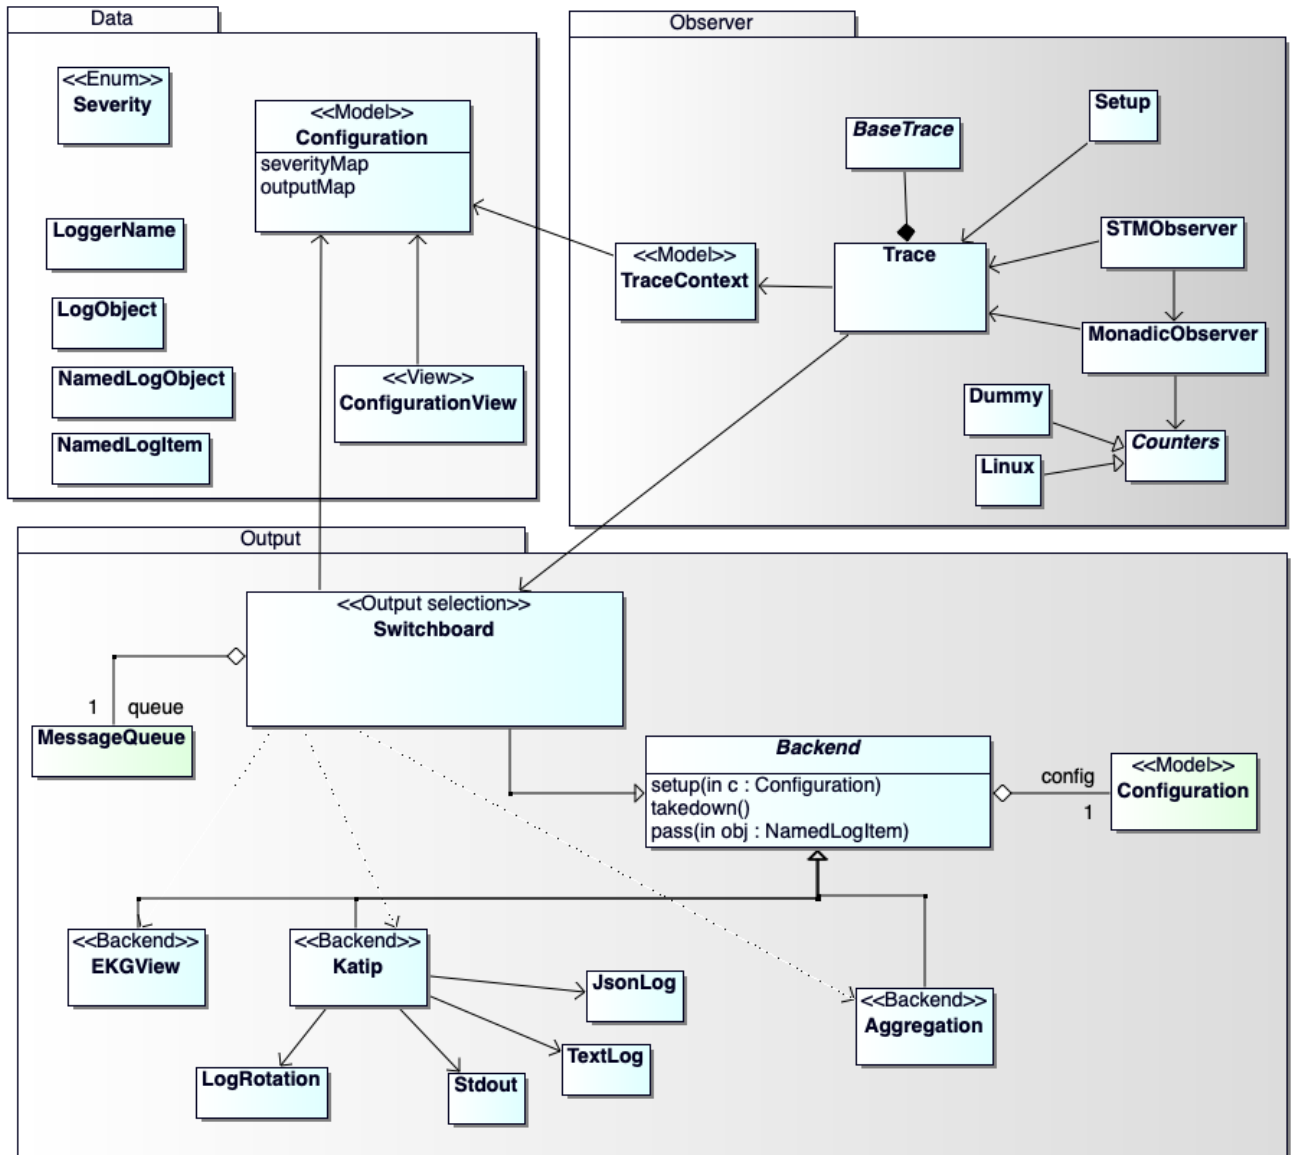


Figure 1.1: Overview of module relationships

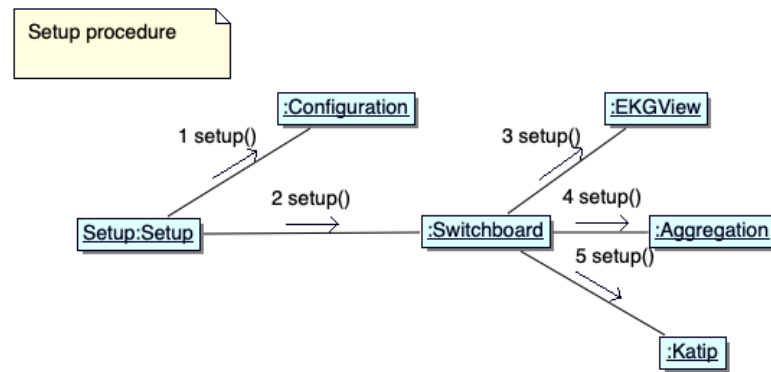


Figure 1.2: Setup procedure

Observe STM action in a named context

With given name, create a *SubTrace* according to *Configuration* and run the passed STM action on it.

```

bracketObserveIO :: Trace IO → Text → STM.STM t → IO t
bracketObserveIO logTrace0 name action = do
  logTrace ← subTrace name logTrace0
  let subtrace = typeofTrace logTrace
  bracketObserveIO' subtrace logTrace action
where
  bracketObserveIO' :: SubTrace → Trace IO → STM.STM t → IO t
  bracketObserveIO' NoTrace _ act =
    STM.atomically act
  bracketObserveIO' subtrace logTrace act = do
    countersid ← observeOpen subtrace logTrace
    -- run action, returns result only
    t ← STM.atomically act
    observeClose subtrace logTrace countersid []
    pure t
  
```

Observe STM action in a named context and output captured log items

The STM action might output messages, which after "success" will be forwarded to the logging trace. Otherwise, this function behaves the same as Observe STM action in a named context.

```

bracketObserveLogIO :: Trace IO → Text → STM.STM (t, [LogObject]) → IO t
bracketObserveLogIO logTrace0 name action = do
  logTrace ← subTrace name logTrace0
  let subtrace = typeofTrace logTrace
  bracketObserveLogIO' subtrace logTrace action
where
  
```

```

bracketObserveLogIO' :: SubTrace → Trace IO → STM.STM (t, [LogObject]) → IO t
bracketObserveLogIO' NoTrace _ act = do
  (t, _) ← STM.atomically $ stmWithLog act
  pure t
bracketObserveLogIO' subtrace logTrace act = do
  countersid ← observeOpen subtrace logTrace
  -- run action, return result and log items
  (t, as) ← STM.atomically $ stmWithLog act
  observeClose subtrace logTrace countersid as
  pure t

```

1.4.2 Cardano.BM.Observer.Monad

Monadic.bracketObserverIO

Observes an IO action and adds a name to the logger name of the passed in *Trace*. An empty *Text* leaves the logger name untouched.

```

bracketObserveIO :: Trace IO → Text → IO t → IO t
bracketObserveIO logTrace0 name action = do
  logTrace ← subTrace name logTrace0
  bracketObserveIO' (typeofTrace logTrace) logTrace action
where
  bracketObserveIO' :: SubTrace → Trace IO → IO t → IO t
  bracketObserveIO' NoTrace _ act = act
  bracketObserveIO' subtrace logTrace act = do
    countersid ← observeOpen subtrace logTrace
    -- run action
    t ← act
    observeClose subtrace logTrace countersid []
    pure t

```

observerOpen

```

observeOpen :: SubTrace → Trace IO → IO CounterState
observeOpen subtrace logTrace = do
  identifier ← newUnique
  logInfo logTrace $ "Opening: " <> pack (show $ hashUnique identifier)
  -- take measurement
  counters ← readCounters subtrace
  let state = CounterState identifier counters
  -- send opening message to Trace
  traceNamedObject logTrace $ ObserveOpen state
  return state

```


observeClose

```

observeClose :: SubTrace → Trace IO → CounterState → [LogObject] → IO ()
observeClose subtrace logTrace initState logObjects = do
  let identifier = csIdentifier initState
      initialCounters = csCounters initState
  logInfo logTrace $ "Closing: " <> pack (show $ hashUnique identifier)
  -- take measurement
  counters ← readCounters subtrace
  -- send closing message to Trace
  traceNamedObject logTrace $ ObserveClose (CounterState identifier counters)
  -- send diff message to Trace
  traceNamedObject logTrace $
    ObserveDiff (CounterState identifier (diffCounters initialCounters counters))
  -- trace the messages gathered from inside the action
  forM_logObjects $ traceNamedObject logTrace

```

1.4.3 BaseTrace**Contravariant**

A covariant is a functor: $F A \rightarrow F B$

A contravariant is a functor: $F B \rightarrow F A$

$Op\ a\ b$ implements the inverse to 'arrow' " $getOp :: b \rightarrow a$ ", which when applied to a *BaseTrace* of type " $Op\ (m\ ())\ s$ ", yields " $s \rightarrow m\ ()$ ". In our case, *Op* accepts an action in a monad *m* with input type *LogNamed LogObject* (see 'Trace').

```

newtype BaseTrace m s = BaseTrace { runTrace :: Op (m ()) s }

```

contramap

A covariant functor defines the function " $fmap :: (a \rightarrow b) \rightarrow f\ a \rightarrow f\ b$ ". In case of a contravariant functor, it is the dual function " $contramap :: (a \rightarrow b) \rightarrow f\ b \rightarrow f\ a$ " which is defined.

In the following instance, *runTrace* extracts type " $Op\ (m\ ())\ s$ " to which *contramap* applies *f*, thus " $f\ s \rightarrow m\ ()$ ". The constructor *BaseTrace* restores " $Op\ (m\ ())\ (f\ s)$ ".

```

instance Contravariant (BaseTrace m) where
  contramap f = BaseTrace ∘ contramap f ∘ runTrace

```

traceWith

Accepts a *Trace* and some payload *s*. First it gets the contravariant from the *Trace* as type " $Op\ (m\ ())\ s$ " and, after " $getOp :: b \rightarrow a$ " which translates to " $s \rightarrow m\ ()$ ", calls the action on the *LogNamed LogObject*.

```

traceWith :: BaseTrace m s → s → m ()
traceWith = getOp ∘ runTrace

```

natTrace

Natural transformation from monad m to monad n .

$$\begin{aligned} \text{natTrace} &:: (\text{forall } x \circ m \ x \rightarrow n \ x) \rightarrow \text{BaseTrace } m \ s \rightarrow \text{BaseTrace } n \ s \\ \text{natTrace } \text{nat } (\text{BaseTrace } (\text{Op } tr)) &= \text{BaseTrace } \$ \text{Op } \$ \text{nat } \circ tr \end{aligned}$$
noTrace

A *Trace* that discards all inputs.

$$\begin{aligned} \text{noTrace} &:: \text{Applicative } m \Rightarrow \text{BaseTrace } m \ a \\ \text{noTrace} &= \text{BaseTrace } \$ \text{Op } \$ \text{const } (\text{pure } ()) \end{aligned}$$
1.4.4 Cardano.BM.Trace**Utilities**

Natural transformation from monad m to monad n .

$$\begin{aligned} \text{natTrace} &:: (\text{forall } x \circ m \ x \rightarrow n \ x) \rightarrow \text{Trace } m \rightarrow \text{Trace } n \\ \text{natTrace } \text{nat } (\text{ctx}, \text{trace}) &= (\text{ctx}, \text{BaseTrace.natTrace } \text{nat } \text{trace}) \end{aligned}$$

Access type of *Trace*.

$$\begin{aligned} \text{typeofTrace} &:: \text{Trace } m \rightarrow \text{SubTrace} \\ \text{typeofTrace } (\text{ctx}, _) &= \text{tracetype } \text{ctx} \end{aligned}$$

Update type of *Trace*.

$$\begin{aligned} \text{updateTracetype} &:: \text{SubTrace} \rightarrow \text{Trace } m \rightarrow \text{Trace } m \\ \text{updateTracetype } \text{subtr } (\text{ctx}, tr) &= (\text{ctx } \{ \text{tracetype} = \text{subtr} \}, tr) \end{aligned}$$
Enter new named context

The context name is created and checked that its size is below a limit (currently 50 chars). The minimum severity that a log message must be labelled with is looked up in the configuration and recalculated.

$$\begin{aligned} \text{appendName} &:: \text{MonadIO } m \Rightarrow \text{LoggerName} \rightarrow \text{Trace } m \rightarrow m (\text{Trace } m) \\ \text{appendName } \text{name } (\text{ctx}, \text{trace}) &= \mathbf{do} \\ &\quad \mathbf{let} \ \text{prevLoggerName} = \text{loggerName } \text{ctx} \\ &\quad \text{prevMinSeverity} = \text{minSeverity } \text{ctx} \\ &\quad \text{newLoggerName} = \text{appendWithDot } \text{prevLoggerName } \text{name} \\ &\quad \text{globMinSeverity} \leftarrow \text{liftIO } \$ \text{Config.minSeverity } (\text{configuration } \text{ctx}) \\ &\quad \text{namedSeverity} \leftarrow \text{liftIO } \$ \text{Config.inspectSeverity } (\text{configuration } \text{ctx}) \ \text{newLoggerName} \\ &\quad \mathbf{case} \ \text{namedSeverity} \ \mathbf{of} \\ &\quad \quad \text{Nothing} \rightarrow \text{return } (\text{ctx } \{ \text{loggerName} = \text{newLoggerName} \}, \text{trace}) \\ &\quad \quad \text{Just } \text{sev} \rightarrow \text{return } (\text{ctx } \{ \text{loggerName} = \text{newLoggerName} \}, \text{trace}) \end{aligned}$$

```

    ,minSeverity = max (max sev prevMinSeverity) globMinSeverity}
    ,trace)
appendWithDot :: LoggerName → LoggerName → LoggerName
appendWithDot "" newName = T.take 50 newName
appendWithDot xs "" = xs
appendWithDot xs newName = T.take 50 $ xs <> " ." <> newName

-- return a BaseTrace from a TraceNamed
named :: BaseTrace.BaseTrace m (LogNamed i) → LoggerName → BaseTrace.BaseTrace m i
named trace name = contramap (LogNamed name) trace

```

TODO remove *locallock*

```

locallock :: MVar ()
locallock = unsafePerformIO $ newMVar ()

```

Trace that forwards to the Switchboard

Every *Trace* ends in the Switchboard which then takes care of dispatching the messages to outputs

```

mainTrace :: Switchboard.Switchboard → TraceNamed IO
mainTrace sb = BaseTrace.BaseTrace $ Op $ \lognamed → do
    Switchboard.effectuate sb lognamed

```

Concrete Trace on stdout

This function returns a trace with an action of type `"(LogNamed LogObject) → IO ()"` which will output a text message as text and all others as JSON encoded representation to the console.

```

stdoutTrace :: TraceNamed IO
stdoutTrace = BaseTrace.BaseTrace $ Op $ \lognamed →
    case lnItem lognamed of
        LP (LogMessage logItem) →
            withMVar locallock $ \_ →
                output (lnName lognamed) $ liPayload logItem
        obj →
            withMVar locallock $ \_ →
                output (lnName lognamed) $ toStrict (encodeToLazyText obj)
    where
        output nm msg = TIO.putStrLn $ nm <> " :: " <> msg

```

Concrete Trace into a TVar

```

traceInTVar :: STM.TVar [a] → BaseTrace.BaseTrace STM.STM a
traceInTVar tvar = BaseTrace.BaseTrace $ Op $ \a → STM.modifyTVar tvar ((:) a)

```

```

traceInTVarIO :: STM.TVar [LogObject] → TraceNamed IO
traceInTVarIO tvar = BaseTrace.BaseTrace $ Op $ \ln →
  STM.atomically $ STM.modifyTVar tvar ((:) (lnItem ln))
traceNamedInTVarIO :: STM.TVar [LogNamed LogObject] → TraceNamed IO
traceNamedInTVarIO tvar = BaseTrace.BaseTrace $ Op $ \ln →
  STM.atomically $ STM.modifyTVar tvar ((:) ln)

```

Check a log item's severity against the *Trace's* minimum severity

do we need three different *minSeverity* defined?

We do a lookup of the global *minSeverity* in the configuration. And, a lookup of the *minSeverity* for the current named context. These values might have changed in the meanwhile. A third filter is the *minSeverity* defined in the current context.

```

traceConditionally
  :: (MonadIO m)
  ⇒ TraceContext → BaseTrace.BaseTrace m LogObject → LogObject
  → m ()
traceConditionally ctx logTrace msg@(LP (LogMessage item)) = do
  globminsev ← liftIO $ Config.minSeverity (configuration ctx)
  globnamesev ← liftIO $ Config.inspectSeverity (configuration ctx) (loggerName ctx)
  let minsev = max (minSeverity ctx) $ max globminsev (fromMaybe Debug globnamesev)
  flag = (liSeverity item) ≥ minsev
  when flag $ BaseTrace.traceWith logTrace msg
  traceConditionally _ logTrace logObject = BaseTrace.traceWith logTrace logObject

```

Enter message into a trace

The function *traceNamedItem* creates a *LogObject* and threads this through the action defined in the *Trace*.

```

traceNamedItem
  :: (MonadIO m)
  ⇒ Trace m
  → LogSelection
  → Severity
  → T.Text
  → m ()
traceNamedItem (ctx, logTrace) p s m =
  let logmsg = LP $ LogMessage $ LogItem {liSelection = p
    ,liSeverity = s
    ,liPayload = m
    }
  in
    traceConditionally ctx (named logTrace (loggerName ctx)) $ logmsg
logDebug, logInfo, logNotice, logWarning, logError

```

```

    :: (MonadIO m) => Trace m -> T.Text -> m ()
logDebug logTrace = traceNamedItem logTrace Both Debug
logInfo logTrace  = traceNamedItem logTrace Both Info
logNotice logTrace = traceNamedItem logTrace Both Notice
logWarning logTrace = traceNamedItem logTrace Both Warning
logError logTrace  = traceNamedItem logTrace Both Error
logDebugS, logInfoS, logNoticeS, logWarningS, logErrorS
    :: (MonadIO m) => Trace m -> T.Text -> m ()
logDebugS logTrace = traceNamedItem logTrace Private Debug
logInfoS logTrace  = traceNamedItem logTrace Private Info
logNoticeS logTrace = traceNamedItem logTrace Private Notice
logWarningS logTrace = traceNamedItem logTrace Private Warning
logErrorS logTrace  = traceNamedItem logTrace Private Error
logDebugP, logInfoP, logNoticeP, logWarningP, logErrorP
    :: (MonadIO m) => Trace m -> T.Text -> m ()
logDebugP logTrace = traceNamedItem logTrace Public Debug
logInfoP logTrace  = traceNamedItem logTrace Public Info
logNoticeP logTrace = traceNamedItem logTrace Public Notice
logWarningP logTrace = traceNamedItem logTrace Public Warning
logErrorP logTrace  = traceNamedItem logTrace Public Error
logDebugUnsafeP, logInfoUnsafeP, logNoticeUnsafeP, logWarningUnsafeP, logErrorUnsafeP
    :: (MonadIO m) => Trace m -> T.Text -> m ()
logDebugUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Debug
logInfoUnsafeP logTrace  = traceNamedItem logTrace PublicUnsafe Info
logNoticeUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Notice
logWarningUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Warning
logErrorUnsafeP logTrace  = traceNamedItem logTrace PublicUnsafe Error

traceNamedObject
    :: Trace m
    -> LogObject
    -> m ()
traceNamedObject (ctx, logTrace) = BaseTrace.traceWith (named logTrace (loggerName ctx))

```

subTrace

Transforms the input *Trace* according to the *Configuration* using the logger name of the current *Trace* appended with the new name. If the empty *Text* is passed, then the logger name remains untouched.

```

subTrace :: MonadIO m => T.Text -> Trace m -> m (Trace m)
subTrace name tr@(ctx, _) = do
    let newName = appendWithDot (loggerName ctx) name
    subtrace0 <- liftIO $ Config.findSubTrace (configuration ctx) newName
    let subtrace = case subtrace0 of Nothing -> Neutral; Just str -> str
    case subtrace of

```

```

Neutral      → do
                tr' ← appendName name tr
                return $ updateTracetype subtrace tr'
UntimedTrace → do
                tr' ← appendName name tr
                return $ updateTracetype subtrace tr'
NoTrace      → return $ updateTracetype subtrace (ctx, BaseTrace.BaseTrace $ Op $ \_ → pure ())
DropOpening  → return $ updateTracetype subtrace (ctx, BaseTrace.BaseTrace $ Op $ \lognamed → do
    case lnItem lognamed of
        ObserveOpen _ → return ()
        obj            → traceNamedObject tr obj)
ObservableTrace _ → do
                tr' ← appendName name tr
                return $ updateTracetype subtrace tr'

```

1.4.5 Cardano.BM.Setup

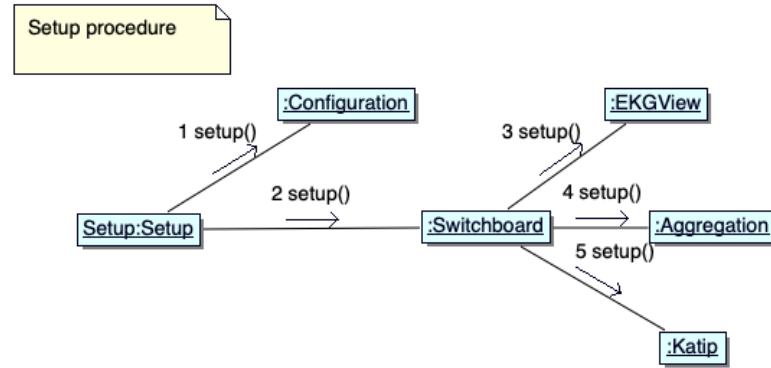


Figure 1.3: Setup procedure

setupTrace

Setup a new *Trace* (Trace) with either a given *Configuration* (Configuration.Model) or a *FilePath* to a configuration file.

```

setupTrace :: MonadIO m ⇒ Either FilePath Config.Configuration → Text → m (Trace m)
setupTrace (Left cfgFile) name = do
    c ← liftIO $ Config.setup cfgFile
    setupTrace_ c name
setupTrace (Right c) name = setupTrace_ c name
setupTrace_ :: MonadIO m ⇒ Config.Configuration → Text → m (Trace m)
setupTrace_ c name = do
    sb ← liftIO $ Switchboard.realize c

```

```

sev ← liftIO $ Config.minSeverity c
ctx ← liftIO $ newContext name c sev
let logTrace = natTrace liftIO (ctx, mainTrace sb)
logTrace' ← subTrace " " logTrace
return logTrace'

```

withTrace

```

withTrace :: MonadIO m ⇒ Config.Configuration → Text → (Trace m → m t) → m t
withTrace cfg name action = do
  logTrace ← setupTrace (Right cfg) name
  action logTrace

```

newContext

```

newContext :: LoggerName → Config.Configuration → Severity → IO TraceContext
newContext name cfg sev = do
  return $ TraceContext {
    loggerName = name
    ,configuration = cfg
    ,minSeverity = sev
    ,tracetype = Neutral
  }

```

1.4.6 Cardano.BM.Counters

Here the platform is chosen on which we compile this program.

Currently, we only support *Linux* with its 'proc' filesystem.

```

{-# LANGUAGE CPP #-}
# if defined (linux_HOST_OS)
# define LINUX
# endif
module Cardano.BM.Counters
(
  Platform.readCounters
  ,diffTimeObserved
  ,getMonoClock
) where
# ifdef LINUX
import qualified Cardano.BM.Counters.Linux as Platform
# else
import qualified Cardano.BM.Counters.Dummy as Platform
# endif

```

```

import Cardano.BM.Counters.Common (getMonoClock)
import Cardano.BM.Data.Counter
import Data.Time.Units (Microsecond)

```

Calculate difference between clocks

```

diffTimeObserved :: CounterState → CounterState → Microsecond
diffTimeObserved (CounterState id0 startCounters) (CounterState id1 endCounters) =
  let
    startTime = getMonotonicTime startCounters
    endTime = getMonotonicTime endCounters
  in
    if (id0 ≡ id1)
    then endTime - startTime
    else error "these clocks are not from the same experiment"
where
  getMonotonicTime counters = case (filter isMonotonicClockCounter counters) of
    [(Counter MonotonicClockTime _ micros)] → fromInteger micros
    _ → error "A time measurement is missing!"
  isMonotonicClockCounter :: Counter → Bool
  isMonotonicClockCounter = (MonotonicClockTime ≡) ∘ cType

```

1.4.7 Cardano.BM.Counters.Common

Common functions that serve *readCounters* on all platforms.

```

nominalTimeToMicroseconds :: Word64 → Microsecond
nominalTimeToMicroseconds = fromMicroseconds ∘ toInteger ∘ ('div' 1000)

```

```

getMonoClock :: IO [Counter]
getMonoClock = do
  t ← getMonotonicTimeNSec
  return [Counter MonotonicClockTime "monoclock" $ toInteger $ nominalTimeToMicroseconds t]

```

1.4.8 Cardano.BM.Counters.Dummy

This is a dummy definition of *readCounters* on platforms that do not support the 'proc' filesystem from which we would read the counters.

The only supported measurement is monotonic clock time for now.

we could well imagine that some day we support all platforms

```

readCounters :: SubTrace → IO [Counter]
readCounters NoTrace = return []
readCounters Neutral = return []

```



```

readCounters UntimedTrace = return []
readCounters DropOpening = return []
readCounters (ObservableTrace tts) = foldrM (\(sel,fun) a →
  if sel `member` tts
  then (fun >>= \xs → return $ a ++ xs)
  else return a) [] selectors
where
  selectors = [(MonotonicClock, getMonoClock)
    -- , (MemoryStats, readProcStatM)
    -- , (ProcessStats, readProcStats)
    -- , (IOStats, readProcIO)
  ]

```

1.4.9 Cardano.BM.Counters.Linux

we have to expand the `readMemStats` function
to read full data from `proc`

```

readCounters :: SubTrace → IO [Counter]
readCounters NoTrace      = return []
readCounters Neutral      = return []
readCounters UntimedTrace = return []
readCounters DropOpening = return []
readCounters (ObservableTrace tts) = foldrM (\(sel,fun) a →
  if sel `member` tts
  then (fun >>= \xs → return $ a ++ xs)
  else return a) [] selectors
where
  selectors = [(MonotonicClock, getMonoClock)
    , (MemoryStats, readProcStatM)
    , (ProcessStats, readProcStats)
    , (IOStats, readProcIO)
  ]

pathProc :: FilePath
pathProc = "/proc/"
pathProcStat :: ProcessID → FilePath
pathProcStat pid = pathProc </> (show pid) </> "stat"
pathProcStatM :: ProcessID → FilePath
pathProcStatM pid = pathProc </> (show pid) </> "statm"
pathProcIO :: ProcessID → FilePath
pathProcIO pid = pathProc </> (show pid) </> "io"

```

Reading from a file in `/proc/<pid>`

```

readProcList :: FilePath → IO [Integer]
readProcList fp = do

```

```
cs ← readFile fp
return $ map (\s → maybe 0 id $ (readMaybe s :: Maybe Integer)) (words cs)
```

readProcStatM - /proc/<pid>/statm

```
/proc/[pid]/statm
Provides information about memory usage, measured in pages. The columns are:
size      (1) total program size
           (same as VmSize in /proc/[pid]/status)
resident  (2) resident set size
           (same as VmRSS in /proc/[pid]/status)
shared    (3) number of resident shared pages (i.e., backed by a file)
           (same as RssFile+RssShmem in /proc/[pid]/status)
text      (4) text (code)
lib       (5) library (unused since Linux 2.6; always 0)
data      (6) data + stack
dt        (7) dirty pages (unused since Linux 2.6; always 0)
```

```
readProcStatM :: IO [Counter]
readProcStatM = do
  pid ← getProcessID
  ps0 ← readProcList (pathProcStatM pid)
  let ps = zip colnames ps0
      psUseful = filter (("unused" ≠) ∘ fst) ps
  return $ map (\(n,i) → Counter MemoryCounter n i) psUseful
where
  colnames :: [Text]
  colnames = ["size", "resident", "shared", "text", "unused", "data", "unused"]
```

readProcStats - //proc//<pid>//stat

```
/proc/[pid]/stat
Status information about the process. This is used by ps(1). It is defined in the kernel source file
fs/proc/array.c.

The fields, in order, with their proper scanf(3) format specifiers, are listed below. Whether or not
certain of these fields display valid information is governed by a ptrace access mode
PTTRACE_MODE_READ_FSCREDS | PTTRACE_MODE_NOAUDIT check (refer to ptrace(2)). If the check denies access,
then the field value is displayed as 0. The affected fields are indicated with the marking [PT].

(1) pid %d
    The process ID.

(2) comm %s
    The filename of the executable, in parentheses. This is visible whether or not the exe-
    cutable is swapped out.

(3) state %c
    One of the following characters, indicating process state:

    R Running
    S Sleeping in an interruptible wait
    D Waiting in uninterruptible disk sleep
    Z Zombie
```

T Stopped (on a signal) or (before Linux 2.6.33) trace stopped
 t Tracing stop (Linux 2.6.33 onward)
 W Paging (only before Linux 2.6.0)
 X Dead (from Linux 2.6.0 onward)
 x Dead (Linux 2.6.33 to 3.13 only)
 K Wakekill (Linux 2.6.33 to 3.13 only)
 W Waking (Linux 2.6.33 to 3.13 only)
 P Parked (Linux 3.9 to 3.13 only)

(4) ppid %d
 The PID of the parent of this process.

(5) pgrp %d
 The process group ID of the process.

(6) session %d
 The session ID of the process.

(7) tty_nr %d
 The controlling terminal of the process. (The minor device number is contained in the combination of bits 31 to 20 and 7 to 0; the major device number is in bits 15 to 8.)

(8) tpgid %d
 The ID of the foreground process group of the controlling terminal of the process.

(9) flags %u
 The kernel flags word of the process. For bit meanings, see the PF_* defines in the Linux kernel source file include/linux/sched.h. Details depend on the kernel version.
 The format for this field was %lu before Linux 2.6.

(10) minflt %lu
 The number of minor faults the process has made which have not required loading a memory page from disk.

(11) cminflt %lu
 The number of minor faults that the process's waited-for children have made.

(12) majflt %lu
 The number of major faults the process has made which have required loading a memory page from disk.

(13) cmajflt %lu
 The number of major faults that the process's waited-for children have made.

(14) utime %lu
 Amount of time that this process has been scheduled in user mode, measured in clock ticks (divide by sysconf(_SC_CLK_TCK)). This includes guest time, guest_time (time spent running a virtual CPU, see below), so that applications that are not aware of the guest time field do not lose that time from their calculations.

(15) stime %lu
 Amount of time that this process has been scheduled in kernel mode, measured in clock ticks (divide by sysconf(_SC_CLK_TCK)).

(16) cutime %ld
 Amount of time that this process's waited-for children have been scheduled in user mode, measured in clock ticks (divide by sysconf(_SC_CLK_TCK)). (See also times(2).) This includes guest time, cguest_time (time spent running a virtual CPU, see below).

(17) cstime %ld

Amount of time that this process's waited-for children have been scheduled in kernel mode, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).

- (18) `priority` %ld
(Explanation for Linux 2.6) For processes running a real-time scheduling policy (policy below; see `sched_setscheduler(2)`), this is the negated scheduling priority, minus one; that is, a number in the range -2 to -100, corresponding to real-time priorities 1 to 99. For processes running under a non-real-time scheduling policy, this is the raw nice value (`setpriority(2)`) as represented in the kernel. The kernel stores nice values as numbers in the range 0 (high) to 39 (low), corresponding to the user-visible nice range of -20 to 19.
- (19) `nice` %ld
The nice value (see `setpriority(2)`), a value in the range 19 (low priority) to -20 (high priority).
- (20) `num_threads` %ld
Number of threads in this process (since Linux 2.6). Before kernel 2.6, this field was hard coded to 0 as a placeholder for an earlier removed field.
- (21) `itrealvalue` %ld
The time in jiffies before the next SIGALRM is sent to the process due to an interval timer. Since kernel 2.6.17, this field is no longer maintained, and is hard coded as 0.
- (22) `starttime` %llu
The time the process started after system boot. In kernels before Linux 2.6, this value was expressed in jiffies. Since Linux 2.6, the value is expressed in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).

The format for this field was %lu before Linux 2.6.
- (23) `vsize` %lu
Virtual memory size in bytes.
- (24) `rss` %ld
Resident Set Size: number of pages the process has in real memory. This is just the pages which count toward text, data, or stack space. This does not include pages which have not been demand-loaded in, or which are swapped out.
- (25) `rsslim` %lu
Current soft limit in bytes on the rss of the process; see the description of `RLIMIT_RSS` in `getrlimit(2)`.
- (26) `startcode` %lu [PT]
The address above which program text can run.
- (27) `endcode` %lu [PT]
The address below which program text can run.
- (28) `startstack` %lu [PT]
The address of the start (i.e., bottom) of the stack.
- (29) `kstkesp` %lu [PT]
The current value of ESP (stack pointer), as found in the kernel stack page for the process.
- (30) `kstkeip` %lu [PT]
The current EIP (instruction pointer).
- (31) `signal` %lu
The bitmap of pending signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (32) `blocked` %lu
The bitmap of blocked signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (33) `sigignore` %lu
The bitmap of ignored signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.

- (34) `sigcatch %lu`
The bitmap of caught signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (35) `wchan %lu [PT]`
This is the "channel" in which the process is waiting. It is the address of a location in the kernel where the process is sleeping. The corresponding symbolic name can be found in `/proc/[pid]/wchan`.
- (36) `nswap %lu`
Number of pages swapped (not maintained).
- (37) `cnswap %lu`
Cumulative `nswap` for child processes (not maintained).
- (38) `exit_signal %d` (since Linux 2.1.22)
Signal to be sent to parent when we die.
- (39) `processor %d` (since Linux 2.2.8)
CPU number last executed on.
- (40) `rt_priority %u` (since Linux 2.5.19)
Real-time scheduling priority, a number in the range 1 to 99 for processes scheduled under a real-time policy, or 0, for non-real-time processes (see `sched_setscheduler(2)`).
- (41) `policy %u` (since Linux 2.5.19)
Scheduling policy (see `sched_setscheduler(2)`). Decode using the `SCHED_*` constants in `linux/sched.h`.

The format for this field was `%lu` before Linux 2.6.22.
- (42) `delayacct_blkio_ticks %llu` (since Linux 2.6.18)
Aggregated block I/O delays, measured in clock ticks (centiseconds).
- (43) `guest_time %lu` (since Linux 2.6.24)
Guest time of the process (time spent running a virtual CPU for a guest operating system), measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (44) `cguest_time %ld` (since Linux 2.6.24)
Guest time of the process's children, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (45) `start_data %lu` (since Linux 3.3) [PT]
Address above which program initialized and uninitialized (BSS) data are placed.
- (46) `end_data %lu` (since Linux 3.3) [PT]
Address below which program initialized and uninitialized (BSS) data are placed.
- (47) `start_brk %lu` (since Linux 3.3) [PT]
Address above which program heap can be expanded with `brk(2)`.
- (48) `arg_start %lu` (since Linux 3.5) [PT]
Address above which program command-line arguments (`argv`) are placed.
- (49) `arg_end %lu` (since Linux 3.5) [PT]
Address below which program command-line arguments (`argv`) are placed.
- (50) `env_start %lu` (since Linux 3.5) [PT]
Address above which program environment is placed.
- (51) `env_end %lu` (since Linux 3.5) [PT]
Address below which program environment is placed.
- (52) `exit_code %d` (since Linux 3.5) [PT]
The thread's exit status in the form reported by `waitpid(2)`.

```

readProcStats :: IO [Counter]
readProcStats = do
    pid ← getProcessID
    ps0 ← readProcList (pathProcStat pid)
    let ps = zip colnames ps0
    psUseful = filter (("unused"  $\neq$ )  $\circ$  fst) ps
    return $ map ( $\lambda(n,i) \rightarrow$  Counter StatInfo n i) psUseful
where
    colnames :: [Text]
    colnames = [ "pid", "unused", "unused", "ppid", "pggrp", "session", "ttynr", "tpgid", "flags", "minflt",
        "cminflt", "majflt", "cmajflt", "utime", "stime", "cutime", "cstime", "priority", "nice", "num",
        "itrealvalue", "starttime", "vsize", "rss", "rsslim", "startcode", "endcode", "startstack",
        "signal", "blocked", "sigignore", "sigcatch", "wchan", "nswap", "cnsnap", "exitsignal", "proc",
        "policy", "blkio", "guesttime", "cguesttime", "startdata", "enddata", "startbrk", "argstart",
        "envend", "exitcode"
    ]

```

readProcIO - //proc//<pid >//io

/proc/[pid]/io (since kernel 2.6.20)

This file contains I/O statistics for the process, for example:

```

# cat /proc/3828/io
rchar: 323934931
wchar: 323929600
syscr: 632687
syscw: 632675
read_bytes: 0
write_bytes: 323932160
cancelled_write_bytes: 0

```

The fields are as follows:

rchar: characters read

The number of bytes which this task has caused to be read from storage. This is simply the sum of bytes which this process passed to read(2) and similar system calls. It includes things such as terminal I/O and is unaffected by whether or not actual physical disk I/O was required (the read might have been satisfied from pagecache).

wchar: characters written

The number of bytes which this task has caused, or shall cause to be written to disk. Similar caveats apply here as with rchar.

syscr: read syscalls

Attempt to count the number of read I/O operations—that is, system calls such as read(2) and pread(2).

syscw: write syscalls

Attempt to count the number of write I/O operations—that is, system calls such as write(2) and pwrite(2).

read_bytes: bytes read

Attempt to count the number of bytes which this process really did cause to be fetched from the storage layer. This is accurate for block-backed filesystems.

write_bytes: bytes written

Attempt to count the number of bytes which this process caused to be sent to the storage layer.

cancelled_write_bytes:

The big inaccuracy here is truncate. If a process writes 1MB to a file and then deletes the

file, it will in fact perform no writeout. But it will have been accounted as having caused 1MB of write. In other words: this field represents the number of bytes which this process caused to not happen, by truncating pagecache. A task can cause "negative" I/O too. If this task truncates some dirty pagecache, some I/O which another task has been accounted for (in its write_bytes) will not be happening.

Note: In the current implementation, things are a bit racy on 32-bit systems: if process A reads process B's /proc/[pid]/io while process B is updating one of these 64-bit counters, process A could see an intermediate result.

Permission to access this file is governed by a ptrace access mode `PTRACE_MODE_READ_FSCREDS` check; see `ptrace(2)`.

```
readProcIO :: IO [Counter]
readProcIO = do
  pid ← getProcessID
  ps0 ← readProcList (pathProcIO pid)
  let ps = zip colnames ps0
  return $ map (\(n,i) → Counter IOCounter n i) ps
where
  colnames :: [Text]
  colnames = [ "rchar", "wchar", "syscr", "syscw", "rbytes", "wbytes", "cxwbytes" ]
```

1.4.10 Cardano.BM.Data.Aggregated

Stats

```
data Stats = Stats {
  fmin :: Integer,
  fmax :: Integer,
  fcount :: Integer,
  fsum_A :: Integer,
  fsum_B :: Integer
} deriving (Show, Eq, Generic, ToJSON)
```

Aggregated

```
data Aggregated = Aggregated {
  fstats :: Stats,
  flast :: Integer,
  fdelta :: Stats
} deriving (Show, Eq, Generic, ToJSON)
```

Update aggregation

We distinguish an uninitialized from an already initialized aggregation:

```
updateAggregation :: Integer → Maybe Aggregated → Maybe Aggregated
updateAggregation v Nothing =
```

```

Just $
  Aggregated {fstats = Stats {
    fmin = v, fmax = v, fcount = 1
    , fsum_A = v, fsum_B = v * v}
    , flast = v
    , fdelta = Stats {
    fmin = 0, fmax = 0, fcount = 0
    , fsum_A = 0, fsum_B = 0}
  }
updateAggregation v (Just (Aggregated (Stats _min _max _count _sumA _sumB)
  _last
  (Stats _dmin _dmax _dcount _dsumA _dsumB)
  )) =
  let delta = v - _last
  in
  Just $
    Aggregated {fstats = Stats {
      fmin = (min _min v)
      , fmax = (max _max v)
      , fcount = (_count + 1)
      , fsum_A = (_sumA + v)
      , fsum_B = (_sumB + v * v)
    }
      , flast = v
      , fdelta = Stats {
        fmin = (min _dmin delta)
        , fmax = (max _dmax delta)
        , fcount = (_dcount + 1)
        , fsum_A = (_dsumA + delta)
        , fsum_B = (_dsumB + delta * delta)
      }
    }

```

1.4.11 Cardano.BM.Data.Backend

Accepts a `NamedLogItem`

```

class IsEffectuator t where
  effectuate :: t → NamedLogItem → IO ()

```

Declaration of a *Backend*

```

class (IsEffectuator t) ⇒ IsBackend t where
  typeof :: t → BackendKind
  realize  :: Configuration → IO t
  realizefrom :: forall s. (IsEffectuator s) ⇒ Configuration → s → IO t

```



```

default realizefrom :: forall s o (IsEffectuator s) => Configuration -> s -> IO t
realizefrom c _ = realize c
unrealize :: t -> IO ()
effectuatefrom :: forall s o (IsEffectuator s) => t -> NamedLogItem -> s -> IO ()
default effectuatefrom :: forall s o (IsEffectuator s) => t -> NamedLogItem -> s -> IO ()
effectuatefrom t nli _ = effectuate t nli

```

Backend

A backend is referenced through the function *bPass* which accepts a *NamedLogItem* and a terminating function *bTerminate* which is responsible for closing the specific backend.

```

data Backend = MkBackend
  { bPass :: NamedLogItem -> IO ()
  , bTerminate :: IO ()
  }

```

1.4.12 Cardano.BM.Data.Configuration

Data structure to help parsing configuration files.

Representation

```

type Port = Int
data Representation = Representation
  { minSeverity    :: Severity
  , rotation      :: RotationParameters
  , setupScribes  :: [ ScribeDefinition ]
  , defaultScribes :: [ (ScribeKind, Text) ]
  , setupBackends :: [ BackendKind ]
  , defaultBackends :: [ BackendKind ]
  , hasEKG        :: Maybe Port
  , hasGUI        :: Maybe Port
  , options       :: HM.HashMap Text Object
  }
deriving (Generic, Show, ToJSON, FromJSON)

```

parseRepresentation

```

parseRepresentation :: FilePath -> IO Representation
parseRepresentation fp = do
  repr :: Representation <- decodeFileThrow fp
  return $ implicit_fill_representation repr

```

after parsing the configuration representation we implicitly correct it.

```

implicit_fill_representation :: Representation → Representation
implicit_fill_representation =
  remove_ekgview_if_not_defined ∘
  filter_duplicates_from_backends ∘
  filter_duplicates_from_scribes ∘
  union_setup_and_usage_backends ∘
  add_ekgview_if_port_defined ∘
  add_katip_if_any_scribes
where
  filter_duplicates_from_backends r =
    r {setupBackends = mkUniq $ setupBackends r}
  filter_duplicates_from_scribes r =
    r {setupScribes = mkUniq $ setupScribes r}
  union_setup_and_usage_backends r =
    r {setupBackends = setupBackends r <> defaultBackends r}
  remove_ekgview_if_not_defined r =
    case hasEKG r of
      Nothing → r {defaultBackends = filter (λbk → bk ≠ EKGViewBK) (defaultBackends r)
                  , setupBackends = filter (λbk → bk ≠ EKGViewBK) (setupBackends r)
                  }
      Just _ → r
  add_ekgview_if_port_defined r =
    case hasEKG r of
      Nothing → r
      Just _ → r {setupBackends = setupBackends r <> [EKGViewBK]}
  add_katip_if_any_scribes r =
    if (any ¬ [null $ setupScribes r, null $ defaultScribes r])
    then r {setupBackends = setupBackends r <> [KatipBK]}
    else r
  mkUniq :: Ord a ⇒ [a] → [a]
  mkUniq = Set.toList ∘ Set.fromList

```

1.4.13 Cardano.BM.Data.Counter

Counter

```

data Counter = Counter
  { cType :: CounterType
  , cName :: Text
  , cValue :: Integer
  }
  deriving (Eq, Show, Generic, ToJSON)

data CounterType = MonotonicClockTime
  | MemoryCounter
  | StatInfo
  | IOCounter
  | CpuCounter

```

```

deriving (Eq, Show, Generic, ToJSON)
nameCounter :: Counter → Text
nameCounter (Counter MonotonicClockTime _ _) = "Time"
nameCounter (Counter MemoryCounter _ _) = "Mem"
nameCounter (Counter StatInfo _ _) = "Stat"
nameCounter (Counter IOCounter _ _) = "IO"
nameCounter (Counter CpuCounter _ _) = "Cpu"
instance ToJSON Microsecond where
  toJSON = toJSON ∘ toMicroseconds
  toEncoding = toEncoding ∘ toMicroseconds

```

CounterState

```

data CounterState = CounterState {
  csIdentifier :: Unique
  , csCounters :: [Counter]
}
deriving (Generic, ToJSON)
instance ToJSON Unique where
  toJSON = toJSON ∘ hashUnique
  toEncoding = toEncoding ∘ hashUnique
instance Show CounterState where
  show cs = (show ∘ hashUnique) (csIdentifier cs)
    <> " => " <> (show $ csCounters cs)

```

Difference between counters

```

diffCounters :: [Counter] → [Counter] → [Counter]
diffCounters openings closings =
  getCountersDiff openings closings
where
  getCountersDiff :: [Counter]
    → [Counter]
    → [Counter]
  getCountersDiff as bs =
    let
      getName counter = nameCounter counter <> cName counter
      asNames = map getName as
      aPairs = zip asNames as
      bsNames = map getName bs
      bs' = zip bsNames bs
      bPairs = HM.fromList bs'
    in
      catMaybes $ (flip map) aPairs $ λ(name, Counter _ _ startValue) →

```

```

case HM.lookup name bPairs of
  Nothing    → Nothing
  Just counter → let endValue = cValue counter
                in Just counter {cValue = endValue – startValue}

```

1.4.14 Cardano.BM.Data.LogItem

LoggerName

```
type LoggerName = Text
```

NamedLogItem

```
type NamedLogItem = LogNamed LogObject
```

LogItem

TODO liPayload :: ToObject

```

data LogItem = LogItem
  { liSelection :: LogSelection
  , liSeverity  :: Severity
  , liPayload  :: Text -- TODO should become ToObject
  } deriving (Show, Generic, ToJSON)

```

```

data LogSelection =
  Public  -- only to public logs.
  | PublicUnsafe -- only to public logs, not console.
  | Private -- only to private logs.
  | Both   -- to public and private logs.
  deriving (Show, Generic, ToJSON, FromJSON)

```

LogObject

```

data LogPrims = LogMessage LogItem
  | LogValue Text Integer
  deriving (Generic, Show, ToJSON)
data LogObject = LP LogPrims
  | ObserveOpen CounterState
  | ObserveDiff CounterState
  | ObserveClose CounterState
  | AggregatedMessage Text Aggregated
  | KillPill
  deriving (Generic, Show, ToJSON)

```

LogNamed

A *LogNamed* contains of a context name and some log item.

```
data LogNamed item = LogNamed
  { lnName :: LoggerName
  , lnItem :: item
  } deriving (Show)
deriving instance Generic item  $\Rightarrow$  Generic (LogNamed item)
deriving instance (ToJSON item, Generic item)  $\Rightarrow$  ToJSON (LogNamed item)
```

1.4.15 Cardano.BM.Data.Observable**ObservableInstance**

```
data ObservableInstance = MonotonicClock
  | MemoryStats
  | ProcessStats
  | IOStats
  deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)
```

1.4.16 Cardano.BM.Data.Output**OutputKind**

```
data OutputKind = StdOut
  | TVarList (STM.TVar [LogObject])
  | TVarListNamed (STM.TVar [LogNamed LogObject])
  | Null
  deriving (Eq)
```

ScribeKind

This identifies katip's scribes by type.

```
data ScribeKind = FileTextSK
  | FileJsonSK
  | StdoutSK
  | StderrSK
  deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)
```

ScribeId

A scribe is identified by *ScribeKind* x *Filename*

```
type ScribeId = Text-- (ScribeKind, Filename)
```

ScribeDefinition

This identifies katip's scribes by type.

```
data ScribeDefinition = ScribeDefinition
  {
    scKind :: ScribeKind
  , scName :: Text
  , scRotation :: Maybe RotationParameters
  }
deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)
```

1.4.17 Cardano.BM.Data.Severity

Severity

```
data Severity = Debug | Info | Warning | Notice | Error
deriving (Show, Eq, Ord, Generic, ToJSON)
instance FromJSON Severity where
  parseJSON = withText "severity" $ \case
    "Debug"    → pure Debug
    "Info"     → pure Info
    "Notice"   → pure Notice
    "Warning"  → pure Warning
    "Error"    → pure Error
    _          → pure Info-- catch all
```

1.4.18 Cardano.BM.Data.SubTrace

SubTrace

```
data SubTrace = Neutral
  | UntimedTrace
  | NoTrace
  | DropOpening
  | ObservableTrace (Set ObservableInstance)
deriving (Generic, Show, FromJSON, ToJSON)
```

1.4.19 Cardano.BM.Data.Trace

Trace

A *Trace* consists of a *TraceContext* and a *TraceNamed* in *m*.

```
type Trace m = (TraceContext, TraceNamed m)
```

TraceNamed

A *TraceNamed* is a specialized Contravariant of type *LogNamed* with payload *LogObject*.

```
type TraceNamed m = BaseTrace m (LogNamed LogObject)
```

TraceContext

We keep the context's name and a reference to the *Configuration* in the *TraceContext*.

```
data TraceContext = TraceContext {
  loggerName  :: LoggerName
  ,configuration :: Configuration
  ,tracetype   :: SubTrace
  ,minSeverity :: Severity
}
```

1.4.20 Cardano.BM.Configuration

see *Cardano.BM.Configuration.Model* for the implementation.

```
getOptionOrDefault :: CM.Configuration → Text → Text → IO (Text)
getOptionOrDefault cg name def = do
  opt ← CM.getOption cg name
  case opt of
    Nothing → return def
    Just o  → return o
```

1.4.21 Cardano.BM.Configuration.Model**Configuration.Model**

```
type ConfigurationMVar = MVar ConfigurationInternal
newtype Configuration = Configuration
  {getCG :: ConfigurationMVar}
-- Our internal state; see -"Configuration model"-
data ConfigurationInternal = ConfigurationInternal
  {cgMinSeverity   :: Severity
  ,cgMapSeverity   :: HM.HashMap LoggerName Severity
  ,cgMapSubtrace   :: HM.HashMap LoggerName SubTrace
  ,cgOptions       :: HM.HashMap Text Object
  ,cgMapBackend    :: HM.HashMap LoggerName [BackendKind]
  ,cgDefBackendKs  :: [BackendKind]
  ,cgSetupBackends :: [BackendKind]
  ,cgMapScribe     :: HM.HashMap LoggerName [ScribeId]
  ,cgDefScribes    :: [ScribeId]
  ,cgSetupScribes  :: [ScribeDefinition]}
```

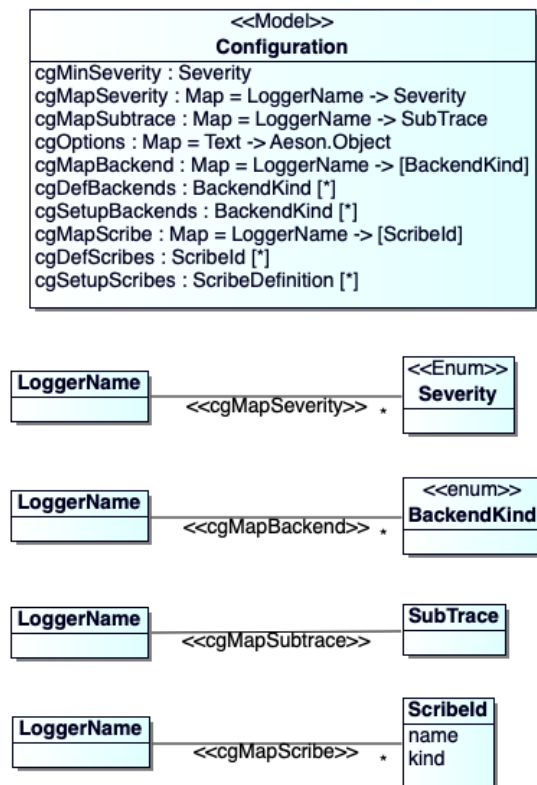


Figure 1.4: Configuration model


```

, cgPortEKG      :: Int
, cgPortGUI      :: Int
}

```

Backends configured in the *Switchboard*

For a given context name return the list of backends configured, or, in case no such configuration exists, return the default backends.

```

getBackends :: Configuration → LoggerName → IO [BackendKind]
getBackends configuration name =
  withMVar (getCG configuration) $ \cg → do
    let outs = HM.lookup name (cgMapBackend cg)
    case outs of
      Nothing → do
        return (cgDefBackendKs cg)
      Just os → return $ os -- TODO in (cgDefBackendKs cg)

getDefaultBackends :: Configuration → IO [BackendKind]
getDefaultBackends configuration =
  withMVar (getCG configuration) $ \cg → do
    return (cgDefBackendKs cg)

setDefaultBackends :: Configuration → [BackendKind] → IO ()
setDefaultBackends configuration bes = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgDefBackendKs = bes}

setBackend :: Configuration → LoggerName → Maybe [BackendKind] → IO ()
setBackend configuration name be = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMapBackend = HM.alter (\_ → be) name (cgMapBackend cg)}

```

Backends to be setup by the *Switchboard*

Defines the list of *Backends* that need to be setup by the *Switchboard*.

```

setSetupBackends :: Configuration → [BackendKind] → IO ()
setSetupBackends configuration bes = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgSetupBackends = bes}

```

Scribes configured in the *Log backend*

For a given context name return the list of scribes to output to, or, in case no such configuration exists, return the default scribes to use.

```

getScribes :: Configuration → LoggerName → IO [ScribeId]
getScribes configuration name =

```

```

withMVar (getCG configuration) $ \cg → do
  let outs = HM.lookup name (cgMapScribe cg)
  case outs of
    Nothing → do
      return (cgDefScribes cg)
    Just os → return $ os
setDefaultScribes :: Configuration → [ ScribeId ] → IO ()
setDefaultScribes configuration scs = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgDefScribes = scs}

```

Scribes to be setup in the Log backend

Defines the list of *Scribes* that need to be setup in the *Log* backend.

```

setSetupScribes :: Configuration → [ ScribeDefinition ] → IO ()
setSetupScribes configuration sds = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgSetupScribes = sds}
getSetupScribes :: Configuration → IO [ ScribeDefinition ]
getSetupScribes configuration =
  withMVar (getCG configuration) $ \cg → do
    return $ cgSetupScribes cg

```

Access port numbers of EKG, GUI

```

getEKGport :: Configuration → IO Int
getEKGport configuration =
  withMVar (getCG configuration) $ \cg → do
    return $ cgPortEKG cg
getGUIport :: Configuration → IO Int
getGUIport configuration =
  withMVar (getCG configuration) $ \cg → do
    return $ cgPortGUI cg

```

Options

```

getOption :: Configuration → Text → IO (Maybe Text)
getOption configuration name = do
  withMVar (getCG configuration) $ \cg →
    case HM.lookup name (cgOptions cg) of
      Nothing → return Nothing
      Just o → return $ Just $ pack $ show o

```

Global setting of minimum severity

```

minSeverity :: Configuration → IO Severity
minSeverity configuration = withMVar (getCG configuration) $ λcg →
    return $ cgMinSeverity cg
setMinSeverity :: Configuration → Severity → IO ()
setMinSeverity configuration sev = do
    cg ← takeMVar (getCG configuration)
    putMVar (getCG configuration) $ cg {cgMinSeverity = sev}

```

Relation of context name to minimum severity

```

inspectSeverity :: Configuration → Text → IO (Maybe Severity)
inspectSeverity configuration name = do
    withMVar (getCG configuration) $ λcg →
        return $ HM.lookup name (cgMapSeverity cg)
setSeverity :: Configuration → Text → Maybe Severity → IO ()
setSeverity configuration name sev = do
    cg ← takeMVar (getCG configuration)
    putMVar (getCG configuration) $ cg {cgMapSeverity = HM.alter (\_ → sev) name (cgMapSeverity cg)}

```

Relation of context name to SubTrace

A new context may contain a different type of *Trace*. The function *appendName* (Enter new named context) will look up the *SubTrace* for the context's name.

```

findSubTrace :: Configuration → Text → IO (Maybe SubTrace)
findSubTrace configuration name = do
    withMVar (getCG configuration) $ λcg →
        return $ HM.lookup name (cgMapSubtrace cg)
setSubTrace :: Configuration → Text → Maybe SubTrace → IO ()
setSubTrace configuration name trafo = do
    cg ← takeMVar (getCG configuration)
    putMVar (getCG configuration) $ cg {cgMapSubtrace = HM.alter (\_ → trafo) name (cgMapSubtrace cg)}

```

Parse configuration from file

Parse the configuration into an internal representation first. Then, fill in *Configuration* from it in a second step after refinement.

```

setup :: FilePath → IO Configuration
setup fp = do
    r ← R.parseRepresentation fp
    cgreg ← newEmptyMVar
    putMVar cgreg $ ConfigurationInternal

```

```

    {cgMinSeverity = R.minSeverity r
    ,cgMapSeverity = HM.empty
    ,cgMapSubtrace = HM.empty
    ,cgOptions = R.options r
    ,cgMapBackend = HM.empty
    ,cgDefBackendKs = R.defaultBackends r
    ,cgSetupBackends = R.setupBackends r
    ,cgMapScribe = HM.empty
    ,cgDefScribes = r_defaultScribes r
    ,cgSetupScribes = R.setupScribes r
    ,cgPortEKG = r_hasEKG r
    ,cgPortGUI = r_hasGUI r
    }
  return $ Configuration cgreg
where
  r_hasEKG r = case (R.hasEKG r) of
    Nothing → 0
    Just p → p
  r_hasGUI r = case (R.hasGUI r) of
    Nothing → 0
    Just p → p
  r_defaultScribes r = map (λ(k,n) → pack (show k) <> " :: " <> n) (R.defaultScribes r)

```

Setup empty configuration

```

empty :: IO Configuration
empty = do
  cgreg ← newEmptyMVar
  putMVar cgreg $ ConfigurationInternal Debug HM.empty HM.empty HM.empty HM.empty [ ] [ ] HM.empty [ ]
  return $ Configuration cgreg

```

1.4.22 Cardano.BM.Output.Switchboard

Switchboard

```

type SwitchboardMVar = MVar SwitchboardInternal
newtype Switchboard = Switchboard
  {getSB :: SwitchboardMVar}
data SwitchboardInternal = SwitchboardInternal
  {sbQueue :: TBQ.TBQueue NamedLogItem
  }

```

Process incoming messages

Incoming messages are put into the queue, and then processed by the dispatcher. The queue is initialized and the message dispatcher launched.

```

instance IsEffectuator Switchboard where
  effectuate switchboard item = do
    let writequeue :: TBQ.TBQueue NamedLogItem → NamedLogItem → IO ()
    writequeue q i = do
      nocapacity ← atomically $ TBQ.isFullTBQueue q
      if nocapacity
      then return ()
      else atomically $ TBQ.writeTBQueue q i
    withMVar (getSB switchboard) $ \sb →
      writequeue (sbQueue sb) item

```

Switchboard **implements** Backend functions

Switchboard is an Declaration of a Backend

```

instance IsBackend Switchboard where
  typeof _ = SwitchboardBK
  realize cfg = do
    let spawnDispatcher :: IsBackend e ⇒ Configuration → [(BackendKind, e)] → TBQ.TBQueue NamedLogItem
    spawnDispatcher config backends queue =
      let qProc = do
        nli ← atomically $ TBQ.readTBQueue queue
        case lnItem nli of
          AggregatedMessage _ aggregated → do
            selectedBackends ← getBackends config (lnName nli)
            let dropAggrBackends = filter (≠ AggregationBK) selectedBackends
            forM_ backends (λ(bek, be) →
              when (bek ∈ dropAggrBackends) (dispatch nli be))
            qProc
          KillPill →
            forM_ backends $ λ(_, backend) →
              unrealize backend
        _ → do
          selectedBackends ← getBackends config (lnName nli)
          forM_ backends $ λ(bek, be) →
            when (bek ∈ selectedBackends) (dispatch nli be)
            qProc
      -- if Nothing is in the queue then every backend needs to be terminated
      -- and the dispatcher exits, the Switchboard stops.
    dispatch :: IsEffectuator e ⇒ NamedLogItem → e → IO ()
    dispatch nli backend = effectuate backend nli
  in
    Async.async qProc
  q ← atomically $ TBQ.newTBQueue 2048
  sbref ← newEmptyMVar
  putMVar sbref $ SwitchboardInternal q
  let sb = Switchboard sbref

```

```

backends ← getDefaultBackends cfg
bs ← setupBackends backends cfg sb []
_ ← spawnDispatcher cfg bs q
return sb
unrealize _ = pure ()

```

Realizing the backends according to configuration

```

setupBackends :: IsBackend e
              => [BackendKind]
              -> Configuration
              -> e
              -> [(BackendKind, e)]
              -> IO [(BackendKind, e)]
setupBackends [] _ _ acc = return acc
setupBackends (bk : bes) c sb acc = do
  be' ← setupBackend' bk c sb
  setupBackends bes c sb ((bk, be') : acc)
setupBackend' :: IsBackend e => BackendKind -> Configuration -> e -> IO e
setupBackend' EKGViewBK c _ = Cardano.BM.Output ◦ EKGView.realize c
setupBackend' AggregationBK c sb = Cardano.BM.Output ◦ Aggregation.realizefrom c sb
setupBackend' KatipBK c _ = Cardano.BM.Output ◦ Log.realize c
setupBackend' SwitchboardBK _ _ = error "cannot instantiate a further Switchboard"

```

1.4.23 Cardano.BM.Output.Log

Internal representation

```

type LogMVar = MVar LogInternal
newtype Log = Log
  {getK :: LogMVar}
data LogInternal = LogInternal
  {kLogEnv :: K.LogEnv
   ,configuration :: Config.Configuration}

```

Log implements effectuate

```

instance IsEffectuator Log where
  effectuate katip item = do
    c ← withMVar (getK katip) $ \k -> return (configuration k)
    selscribes ← getScribes c (lnName item)
    forM_ selscribes $ \sc -> passN sc katip item

```

Log implements backend functions

```

instance IsBackend Log where
  typeof _ = KatipBK
  realize config = do
    let updateEnv :: K.LogEnv → IO UTCTime → K.LogEnv
    updateEnv le timer =
      le { K._logEnvTimer = timer, K._logEnvHost = "hostname" }
    register :: [ScribeDefinition] → K.LogEnv → IO K.LogEnv
    register [ ] le = return le
    register (defsc : dscs) le = do
      let kind = scKind defsc
      name = scName defsc
      name' = pack (show kind) <> ":" <> name
      scr ← createScribe kind name
      register dscs ≡ K.registerScribe name' scr scribeSettings le
    mockVersion :: Version
    mockVersion = Version [0,1,0,0] []
    scribeSettings :: KC.ScribeSettings
    scribeSettings =
      let bufferSize = 5000 -- size of the queue (in log items)
      in
        KC.ScribeSettings bufferSize
    createScribe FileTextSK name = mkTextFileScribe (FileDescription $ unpack name) False
    createScribe FileJsonSK name = mkJsonFileScribe (FileDescription $ unpack name) False
    createScribe StdoutSK _ = mkStdoutScribe
    createScribe StderrSK _ = mkStderrScribe
    cfoKey ← Config.getOptionOrDefault config (pack "cfokey") (pack "<unknown>")
    le0 ← K.initLogEnv
      (K.Namespace [ "iohk" ])
      (fromString $ (unpack cfoKey) <> ":" <> showVersion mockVersion)
    -- request a new time 'getCurrentTime' at most 100 times a second
    timer ← mkAutoUpdate defaultUpdateSettings { updateAction = getCurrentTime, updateFreq = 10000 }
    let le1 = updateEnv le0 timer
    scribes ← getSetupScribes config
    le ← register scribes le1
    kref ← newEmptyMVar
    putMVar kref $ LogInternal le config
    return $ Log kref
  unrealize katip = do
    le ← withMVar (getK katip) $ λk → return (kLogEnv k)
    void $ K.closeScribes le

example :: IO ()
example = do
  config ← Config.setup "from_some_path.yaml"

```

```

k ← setup config
passN (pack (show StdoutSK)) k $ LogNamed
  {lnName = "test"
  ,lnItem = LP $ LogMessage $ LogItem
    {liSelection = Both
    ,liSeverity = Info
    ,liPayload = "Hello!"
    }
  }
passN (pack (show StdoutSK)) k $ LogNamed
  {lnName = "test"
  ,lnItem = LP $ LogValue "cpu-no" 1
  }

-- useful instances for katip
deriving instance K.ToObject LogObject
deriving instance K.ToObject LogItem
deriving instance K.ToObject (Maybe LogObject)
instance KC.LogItem LogObject where
  payloadKeys _ = KC.AllKeys
instance KC.LogItem LogItem where
  payloadKeys _ = KC.AllKeys
instance KC.LogItem (Maybe LogObject) where
  payloadKeys _ = KC.AllKeys

```

Log.passN

The following function copies the *NamedLogItem* to the queues of all scribes that match on their name. Compare start of name of scribe to *(show backend <> " : ")*. This function is non-blocking.

```

passN :: Text → Log → NamedLogItem → IO ()
passN backend katip namedLogItem = do
  env ← withMVar (getK katip) $ λk → return (kLogEnv k)
  forM_ (Map.toList $ K._logEnvScribes env) $
    λ(scName, (KC.ScribeHandle _ shChan)) →
      -- check start of name to match ScribeKind
      if backend 'isPrefixOf' scName
      then do
        let item = lnItem namedLogItem
        let (sev, msg, payload) = case item of
          (LP (LogMessage logItem)) →
            (liSeverity logItem, liPayload logItem, Nothing)
          (AggregatedMessage name aggregated) →
            (Info, pack (show name ++ " : " ++ show aggregated), Nothing)
          _ → (Info, "", (Nothing :: Maybe LogObject))
        if (msg ≡ "") ∧ (isNothing payload)
        then return ()

```



```

else do
  threadIdText ← KC.mkThreadIdText < $ > myThreadId
  let ns = lnName namedLogItem
  itemTime ← env ^. KC.logEnvTimer
  let itemKatip = K.Item {
    _itemApp = env ^. KC.logEnvApp
    , _itemEnv = env ^. KC.logEnvEnv
    , _itemSeverity = sev2klog sev
    , _itemThread = threadIdText
    , _itemHost = env ^. KC.logEnvHost
    , _itemProcess = env ^. KC.logEnvPid
    , _itemPayload = payload
    , _itemMessage = K.logStr msg
    , _itemTime = itemTime
    , _itemNamespace = (env ^. KC.logEnvApp) <> (K.Namespace [ ns ])
    , _itemLoc = Nothing
  }
  void $ atomically $ KC.tryWriteTBQueue shChan (KC.NewItem itemKatip)
else return ()

```

Scribes

```

mkStdoutScribe :: IO K.Scribe
mkStdoutScribe = mkTextFileScribeH stdout True
mkStderrScribe :: IO K.Scribe
mkStderrScribe = mkTextFileScribeH stderr True
mkTextFileScribeH :: Handle → Bool → IO K.Scribe
mkTextFileScribeH handler color = do
  mkFileScribeH handler formatter color
  where
    formatter h colorize verbosity item =
      TIO.hPutStrLn h $! toLazyText $ formatItem colorize verbosity item
mkFileScribeH
  :: Handle
  → (forall a ◦ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ())
  → Bool
  → IO K.Scribe
mkFileScribeH h formatter colorize = do
  hSetBuffering h LineBuffering
  locklocal ← newMVar ()
  let logger :: forall a ◦ K.LogItem a ⇒ K.Item a → IO ()
      logger item = withMVar locklocal $ \_ →
        formatter h colorize K.V0 item
  pure $ K.Scribe logger (hClose h)
mkTextFileScribe :: FileDescription → Bool → IO K.Scribe
mkTextFileScribe fdesc colorize = do

```

```

    mkFileScribe fdesc formatter colorize
  where
    formatter :: Handle → Bool → K.Verbosity → K.Item a → IO ()
    formatter hdl colorize' v' item = do
      let tmsg = toLazyText $ formatItem colorize' v' item
      TIO.hPutStrLn hdl tmsg
mkJsonFileScribe :: FileDescription → Bool → IO K.Scribe
mkJsonFileScribe fdesc colorize = do
  mkFileScribe fdesc formatter colorize
  where
    formatter :: (K.LogItem a) ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ()
    formatter h _ verbosity item = do
      let tmsg = case KC._itemMessage item of
        K.LogStr "" → K.itemJson verbosity item
        K.LogStr msg → K.itemJson verbosity $
          item { KC._itemMessage = K.logStr (" " :: Text)
              , KC._itemPayload = LogItem Both Info $ toStrict $ toLazyText msg
              }
      TIO.hPutStrLn h (encodeToLazyText tmsg)
mkFileScribe
  :: FileDescription
  → (forall a ◦ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ())
  → Bool
  → IO K.Scribe
mkFileScribe fdesc formatter colorize = do
  let prefixDir = prefixPath fdesc
  (createDirectoryIfMissing True prefixDir)
  'catchIO' (prtoutException ("cannot log prefix directory: " ++ prefixDir))
  let fpath = filePath fdesc
  h ← catchIO (openFile fpath WriteMode) $
    λe → do
      prtoutException ("error while opening log: " ++ fpath) e
      -- fallback to standard output in case of exception
      return stdout
  hSetBuffering h LineBuffering
  scribestate ← newMVar h
  let finalizer :: IO ()
      finalizer = withMVar scribestate hClose
  let logger :: forall a ◦ K.LogItem a ⇒ K.Item a → IO ()
      logger item =
        withMVar scribestate $ λhandler →
          formatter handler colorize K.V0 item
  return $ K.Scribe logger finalizer

formatItem :: Bool → K.Verbosity → K.Item a → Builder
formatItem withColor _verb K.Item {..} =
  fromText header <>

```

```

fromText " " <>
brackets (fromText timestamp) <>
fromText " " <>
KC.unLogStr _itemMessage
where
  header = colorBySeverity _itemSeverity $
    "[ " <> mconcat namedcontext <> ":" <> severity <> ":" <> threadid <> "]"
  namedcontext = KC.intercalateNs _itemNamespace
  severity = KC.renderSeverity _itemSeverity
  threadid = KC.getThreadIdText _itemThread
  timestamp = pack $ formatTime defaultTimeLocale tsformat _itemTime
  tsformat :: String
  tsformat = "%F %T%2Q %Z"
  colorBySeverity s m = case s of
    K.EmergencyS → red m
    K.AlertS     → red m
    K.CriticalS  → red m
    K.ErrorS     → red m
    K.NoticeS    → magenta m
    K.WarningS   → yellow m
    K.InfoS      → blue m
    _            → m
  red = colorize "31"
  yellow = colorize "33"
  magenta = colorize "35"
  blue = colorize "34"
  colorize c m
    | withColor = "\ESC[ " <> c <> "m" <> m <> "\ESC[0m"
    | otherwise = m
-- translate Severity to Log.Severity
sev2klog :: Severity → K.Severity
sev2klog = λcase
  Debug → K.DebugS
  Info   → K.InfoS
  Notice → K.NoticeS
  Warning → K.WarningS
  Error  → K.ErrorS

data FileDescription = FileDescription {
  filePath :: !FilePath
  deriving (Show)
prefixPath :: FileDescription → FilePath
prefixPath = takeDirectory ∘ filePath
-- display message and stack trace of exception on stdout
prtoutException :: Exception e ⇒ String → e → IO ()
prtoutException msg e = do
  putStrLn msg

```

```
putStrLn ("exception: " ++ displayException e)
```

1.4.24 Cardano.BM.Output.EKGView

Structure of EKGView

```
type EKGViewMVar = MVar EKGViewInternal
newtype EKGView = EKGView
  {getEV :: EKGViewMVar}
data EKGViewInternal = EKGViewInternal
  {evGauges :: HM.HashMap Text Gauge.Gauge
  ,evLabels  :: HM.HashMap Text Label.Label
  ,evServer :: Server
  }
```

EKG view is an effectuator

```
instance IsEffectuator EKGView where
  effectuate ekgview item =
    let update :: LogObject → LoggerName → EKGViewInternal → IO (Maybe EKGViewInternal)
    update (LP (LogMessage logitem)) logname ekg@(EKGViewInternal _ labels server) =
      case HM.lookup logname labels of
        Nothing → do
          ekghdl ← getLabel logname server
          Label.set ekghdl (liPayload logitem)
          return $ Just $ ekg {evLabels = HM.insert logname ekghdl labels}
        Just ekghdl → do
          Label.set ekghdl (liPayload logitem)
          return Nothing
    update (LP (LogValue iname value)) logname ekg@(EKGViewInternal gauges _ server) =
      let name = logname <> " ." <> iname
      in
      case HM.lookup name gauges of
        Nothing → do
          ekghdl ← getGauge name server
          Gauge.set ekghdl (fromInteger value)
          return $ Just $ ekg {evGauges = HM.insert name ekghdl gauges}
        Just ekghdl → do
          Gauge.set ekghdl (fromInteger value)
          return Nothing
    update _ _ _ = return Nothing
  in do
    ekg ← takeMVar (getEV ekgview)
    upd ← update (lnItem item) (lnName item) ekg
    case upd of
```

```
Nothing → putMVar (getEV ekgview) ekg
Just ekg' → putMVar (getEV ekgview) ekg'
```

EKGView implements Backend functions

EKGView is an Declaration of a *Backend*

```
instance IsBackend EKGView where
  typeof _ = EKGViewBK
  realize config = do
    evref ← newEmptyMVar
    evport ← getEKGport config
    ehdl ← forkServer "127.0.0.1" evport
    ekghdl ← getLabel "iohk-monitoring version" ehdl
    Label.set ekghdl $ pack (showVersion version)
    putMVar evref $ EKGViewInternal
      { evGauges = HM.empty
      , evLabels = HM.empty
      , evServer = ehdl
      }
    return $ EKGView evref
  unrealize ekgview = do
    ekg ← takeMVar $ getEV ekgview
    killThread $ serverThreadId $ evServer ekg
```

Interactive testing EKGView

```
test :: IO ()
test = do
  c ← Cardano.BM.Configuration.setup "test/config.yaml"
  ev ← Cardano.BM.Output ◦ EKGView.realize c
  effectuate ev $ LogNamed "test.questions" (LP (LogValue "answer" 42))
  effectuate ev $ LogNamed "test.monitor023" (LP (LogMessage (LogItem Public Warning "!!!! ALARM !!!")))
```

1.4.25 Cardano.BM.Output.Aggregation

Internal representation

```
type AggregationMVar = MVar AggregationInternal
newtype Aggregation = Aggregation
  { getAg :: AggregationMVar }
data AggregationInternal = AggregationInternal
  { agQueue :: TBQ.TBQueue (Maybe NamedLogItem)
  , agDispatch :: Async.Async ()
  }
}
```

Relation from context name to aggregated statistics

We keep the aggregated values (Aggregated) for a named context in a *HashMap*.

```
type AggregationMap = HM.HashMap Text Aggregated
```

Aggregation **implements** *effectuate*

Aggregation is an *Accepts a NamedLogItem* Enter the log item into the *Aggregation* queue.

```
instance IsEffectuator Aggregation where
  effectuate agg item = do
    ag ← readMVar (getAg agg)
    atomically $ TBQ.writeTBQueue (agQueue ag) $ Just item
```

Aggregation **implements** *Backend functions*

Aggregation is an Declaration of a *Backend*

```
instance IsBackend Aggregation where
  typeof _ = AggregationBK
  realize _ = error "Aggregation cannot be instantiated by 'realize'"
  realizefrom _ switchboard = do
    aggref ← newEmptyMVar
    aggregationQueue ← atomically $ TBQ.newTBQueue 2048
    dispatcher ← spawnDispatcher HM.empty aggregationQueue switchboard
    putMVar aggref $ AggregationInternal aggregationQueue dispatcher
    return $ Aggregation aggref
  unrealize aggregation = do
    let clearMVar :: MVar a → IO ()
        clearMVar = void ∘ tryTakeMVar
    (dispatcher, queue) ← withMVar (getAg aggregation) (λag →
      return (agDispatch ag, agQueue ag))
    -- send terminating item to the queue
    atomically $ TBQ.writeTBQueue queue Nothing
    -- wait for the dispatcher to exit
    res ← Async.waitCatch dispatcher
    either throwM return res
    (clearMVar ∘ getAg) aggregation
```

Asynchronously reading log items from the queue and their processing

```
spawnDispatcher :: IsEffectuator e
  ⇒ AggregationMap
  → TBQ.TBQueue (Maybe NamedLogItem)
  → e
```

```

    → IO (Async.Async ())
spawnDispatcher aggMap aggregationQueue switchboard = Async.async $ qProc aggMap
where
  qProc aggregatedMap = do
    maybeItem ← atomically $ TBQ.readTBQueue aggregationQueue
    case maybeItem of
      Just item → do
        let (updatedMap, msgs) =
            update (lnItem item) (lnName item) aggregatedMap
        sendAggregated msgs switchboard (lnName item)
        qProc updatedMap
      Nothing → return ()
  update :: LogObject
    → LoggerName
    → HM.HashMap Text Aggregated
    → (HM.HashMap Text Aggregated, [LogObject])
  update (LP (LogValue iname value)) logname agmap =
    let name = logname <> "." <> iname
        maybeAggregated = updateAggregation value $ HM.lookup name agmap
        aggregatedMessage = case maybeAggregated of
            Nothing →
                []
            Just aggregated →
                [AggregatedMessage iname aggregated]
    in
      -- use of HM.alter so that in future we can clear the Aggregated
      -- by using as alter's arg a function which returns Nothing.
      (HM.alter (const $ maybeAggregated) name agmap, aggregatedMessage)
  update (ObserveDiff counterState) logname agmap =
    let
      counters = csCounters counterState
      (mapNew, msgs) = updateCounter counters logname agmap [ ]
    in
      (mapNew, reverse msgs)
  -- TODO for text messages aggregate on delta of timestamps
  update _ _ agmap = (agmap, [ ])
  updateCounter :: [Counter]
    → LoggerName
    → HM.HashMap Text Aggregated
    → [LogObject]
    → (HM.HashMap Text Aggregated, [LogObject])
  updateCounter [ ] _ aggrMap msgs = (aggrMap, msgs)
  updateCounter (counter : cs) logname aggrMap msgs =
    let
      name = cName counter
      fullname = logname <> "." <> name
      maybeAggregated = updateAggregation (cValue counter) $ HM.lookup fullname aggrMap

```

```

    aggregatedMessage = case maybeAggregated of
      Nothing →
        error "This should not have happened!"
      Just aggregated →
        AggregatedMessage ((nameCounter counter) <> "." <> name) aggregated
    updatedMap = HM.alter (const $ maybeAggregated) fullname aggrMap
in
    updateCounter cs logname updatedMap (aggregatedMessage : msgs)
sendAggregated :: IsEffectuator e ⇒ [LogObject] → e → Text → IO ()
sendAggregated [] _ _ = return ()
sendAggregated (aggregatedMsg@(AggregatedMessage _ _):ms) sb logname = do
  -- forward the aggregated message to Switchboard
  effectuate sb $
    LogNamed
      { lnName = logname <> ".aggregated"
      , lnItem = aggregatedMsg
      }
  sendAggregated ms sb logname
-- ignore all other messages that are not of type AggregatedMessage
sendAggregated (_:ms) sb logname =
  sendAggregated ms sb logname

```