Cardano.BM - benchmarking and logging

Alexander Diemand   Andreas Triantafyllos

November 2018

Abstract

This is a framework that combines logging, benchmarking and monitoring. Complex evaluations of STM or monadic actions can be observed from outside while reading operating system counters before and after, and calculating their differences, thus relating resource usage to such actions. Through interactive configuration, the runtime behaviour of logging or the measurement of resource usage can be altered. Further reduction in logging can be achieved by redirecting log messages to an aggregation function which will output the running statistics with less frequency than the original message.

Contents

Chapter 1

Cardano BM

1.1 Overview

In figure 1.1 we display the relationships among modules in *Cardano.BM*. The arrows indicate import of a module. The arrows with a triangle at one end would signify "inheritance" in object-oriented programming, but we use it to show that one module replaces the other in the namespace, thus refines its interface.

1.2 Introduction

1.2.1 Logging with Trace

1.2.2 Setup procedure

Hierarchy of Traces

1.2.3 Measuring *Observables*

1.2.4 Information reduction in Aggregation

1.2.5 Output selection

1.2.6 Monitoring

1.3 Examples

1.3.1 Observing evaluation of a STM action

1.3.2 Observing evaluation of a monad action

1.3.3 Simple example showing plain logging

```
{-# LANGUAGE OverloadedStrings #-}
module Main
  (main)
  where
import Control.Concurrent (threadDelay)
import Cardano.BM.Configuration.Static (defaultConfigStdout)
```
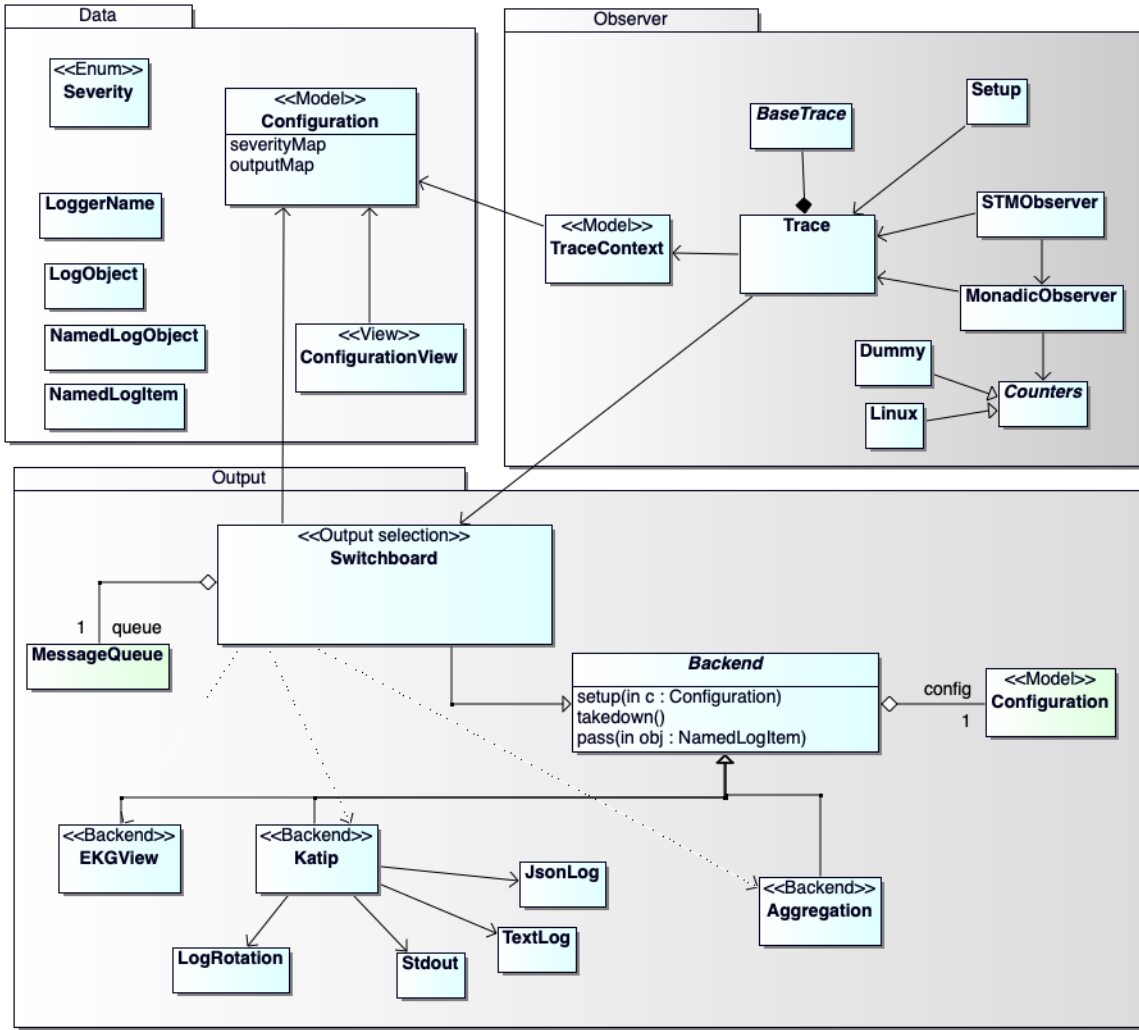
Figure 1.1: Overview of module relationships

import *Cardano.BM.Setup* (setupTrace)
import *Cardano.BM.Trace* (logDebug, logError, logInfo, logNotice,
          logWarning)

main :: *IO* ()
main = **do**
  c ← *defaultConfigStdout*
  tr ← setupTrace (*Right c*) "simple"

  logDebug *tr* "this is a debug message"
  logInfo   *tr* "this is an information."
  logNotice *tr* "this is a notice!"
  logWarning *tr* "this is a warning!"
  logError   *tr* "this is an error!"

  threadDelay 80000

Figure 1.2: Setup procedure

return ()

1.3.4 Complex example showing logging, aggregation of log items, and observing *IO* actions

Module header and import directives

{-# LANGUAGE OverloadedStrings #-}

module *Main*
  (*main*)
  where

import *Control.Concurrent* (*threadDelay*)
import *qualified Control.Concurrent.Async as Async*
import *Control.Monad* (*forM*, *forM_*)
import *GHC.Conc.Sync* (*STM*, *TVar*, *atomically*, *newTVar*, *readTVar*, *writeTVar*)
import *Data.Text* (*pack*)
import *System.Random*

import *qualified Cardano.BM.Configuration.Model as CM*
import *Cardano.BM.Data.Aggregated* (*Measurable* (..))
import *Cardano.BM.Data.AggregatedKind*
import *Cardano.BM.Data.BackendKind*
import *Cardano.BM.Data.LogItem*
import *Cardano.BM.Data.Observable*
import *Cardano.BM.Data.Output*
import *Cardano.BM.Data.Severity*
import *Cardano.BM.Data.SubTrace*
import *Cardano.BM.Observer.Monadic* (*bracketObserveIO*)
import *qualified Cardano.BM.Observer.STM as STM*
import *Cardano.BM.Setup*
import *Cardano.BM.Trace*

Define configuration

The output can be viewed in EKG on `http://localhost:12789`.

```
config :: IO CM.Configuration
config = do
  c ← CM.empty
  CM.setMinSeverity c Debug
  CM.setSetupBackends c [KatipBK, AggregationBK, EKGViewBK]
  CM.setDefaultBackends c [KatipBK]
  CM.setSetupScribes c [ScribeDefinition {
        scName = "stdout"
        ,scKind = StdoutSK
        ,scRotation = Nothing
        }
      ,ScribeDefinition {
        scName = "out.odd.json"
        ,scKind = FileJsonSK
        ,scRotation = Nothing
        }
      ,ScribeDefinition {
        scName = "out.even.json"
        ,scKind = FileJsonSK
        ,scRotation = Nothing
        }
      ,ScribeDefinition {
        scName = "out.txt"
        ,scKind = FileTextSK
        ,scRotation = Nothing
        }
      ]
  CM.setDefaultScribes c ["StdoutSK::stdout"]
  CM.setScribes c "complex.random" (Just ["StdoutSK::stdout", "FileTextSK::out.txt"])
  CM.setScribes c "#aggregated.complex.random" (Just ["StdoutSK::stdout"])
  forM_ [(1 :: Int)..10] $ λx →
    if odd x
    then
      CM.setScribes c ("#aggregation.complex.observeSTM." <> (pack $ show x)) $ Just ["FileJsonSK::ou
    else
      CM.setScribes c ("#aggregation.complex.observeSTM." <> (pack $ show x)) $ Just ["FileJsonSK::ou
  CM.setSubTrace c "complex.random" (Just $ TeeTrace "ewma")
  CM.setSubTrace c "#ekgview"
    (Just $ FilterTrace [(Drop (StartsWith "#ekgview.#aggregation.complex.random"),
        Unhide [(EndsWith ".count"),
          (EndsWith ".avg"),
          (EndsWith ".mean")]),
      (Drop (StartsWith "#ekgview.#aggregation.complex.observeIO"),
        Unhide [(Contains "diff.RTS.cpuNs.timed.")]),
```

```
        (Drop (StartsWith "#ekgview.#aggregation.complex.observeSTM"),
          Unhide [(Contains "diff.RTS.gcNum.timed.")]),
        (Drop (StartsWith "#ekgview.#aggregation.complex.message"),
          Unhide [(Contains ".timed.m")])
        ])
  CM.setSubTrace c "complex.observeIO" (Just $ ObservableTrace [GhcRtsStats, MemoryStats])
  forM_ [(1 :: Int)..10] $ λx →
    CM.setSubTrace
      c
      ("complex.observeSTM." <> (pack $ show x))
      (Just $ ObservableTrace [GhcRtsStats, MemoryStats])
  CM.setBackends c "complex.message" (Just [AggregationBK, KatipBK])
  CM.setBackends c "complex.random" (Just [AggregationBK, KatipBK])
  CM.setBackends c "complex.random.ewma" (Just [AggregationBK])
  CM.setBackends c "complex.observeIO" (Just [AggregationBK])
  forM_ [(1 :: Int)..10] $ λx → do
    CM.setBackends c
      ("complex.observeSTM." <> (pack $ show x))
      (Just [AggregationBK])
    CM.setBackends c
      ("#aggregation.complex.observeSTM." <> (pack $ show x))
      (Just [EKGViewBK])
  CM.setAggregatedKind c "complex.random.rr" (Just StatsAK)
  CM.setAggregatedKind c "complex.random.ewma.rr" (Just (EwmaAK 0.42))
  CM.setBackends c "#aggregation.complex.message" (Just [EKGViewBK])
  CM.setBackends c "#aggregation.complex.observeIO" (Just [EKGViewBK])
  CM.setBackends c "#aggregation.complex.random" (Just [EKGViewBK])
  CM.setBackends c "#aggregation.complex.random.ewma" (Just [EKGViewBK])
  CM.setEKGport c 12789
  return c
```

Thread that outputs a random number to a Trace

```
randomThr :: Trace IO → IO (Async.Async ())
randomThr trace = do
  logInfo trace "starting random generator"
  trace' ← subTrace "random" trace
  proc ← Async.async (loop trace')
  return proc
  where
    loop tr = do
      threadDelay 500000 -- 0.5 second
      num ← randomRIO (42 − 42, 42 + 42) :: IO Double
      lo ← LogObject < $ > mkLOMeta < ∗ > pure (LogValue "rr" (PureD num))
      traceNamedObject tr lo
      loop tr
```

Thread that observes an *IO* action

```
observeIO :: Trace IO → IO (Async.Async ())
observeIO trace = do
   logInfo trace "starting observer"
   proc ← Async.async (loop trace)
   return proc
   where
      loop tr = do
         threadDelay 5000000-- 5 seconds
         _ ← bracketObserveIO tr "observeIO" $ do
            num ← randomRIO (100000, 200000) :: IO Int
            ls ← return $ reverse $ init $ reverse $ 42 : [1..num]
            pure $ const ls ()
         loop tr
```

Thread that observes an *IO* action which downloads a txt in order to observe the I/O statistics

> disabled for now! on Mac OSX this function was blocking all IO.

```
observeDownload :: Trace IO → IO ()
observeDownload trace = loop trace
   where
      loop tr = do
         threadDelay 10000000-- 10 seconds
         tr' ← appendName "observeDownload" tr
         bracketObserveIO tr' "" $ do
            license ← openURI "http://www.gnu.org/licenses/gpl.txt"
            case license of
               Right bs → logNotice tr' $ pack $ take 100 $ BS8.unpack bs
               Left _ → return ()
            threadDelay 500000-- .5 second
            pure ()
         loop tr
```

Threads that observe *STM* actions on the same TVar

```
observeSTM :: Trace IO → IO [Async.Async ()]
observeSTM trace = do
   logInfo trace "starting STM observer"
   tvar ← atomically $ newTVar ([1..1000] :: [Int])
   -- spawn 10 threads
   proc ← forM [(1 :: Int)..10] $ λx → Async.async (loop trace tvar (pack $ show x))
   return proc
   where
      loop tr tvarlist name = do
```

```
        threadDelay 10000000-- 10 seconds
        STM.bracketObserveIO tr ("observeSTM." <> name) (stmAction tvarlist)
        loop tr tvarlist name
stmAction :: TVar [Int] → STM ()
stmAction tvarlist = do
  list ← readTVar tvarlist
  writeTVar tvarlist $ reverse $ init $ reverse $ list
  pure ()
```

Thread that periodically outputs a message

```
msgThr :: Trace IO → IO (Async.Async ())
msgThr trace = do
  logInfo trace "start messaging .."
  trace' ← subTrace "message" trace
  Async.async (loop trace')
  where
    loop tr = do
      threadDelay 3000000-- 3 seconds
      logNotice tr "N O T I F I C A T I O N ! ! !"
      logDebug tr "a detailed debug message."
      logError tr "Boooommm .."
      loop tr
```

Main entry point

```
main :: IO ()
main = do
  -- create configuration
  c ← config
  -- create initial top-level Trace
  tr ← setupTrace (Right c) "complex"
  logNotice tr "starting program; hit CTRL-C to terminate"
  logInfo tr "watch its progress on http://localhost:12789"
    {-start thread sending unbounded sequence of random numbers to a trace which aggregates them into a
  procRandom ← randomThr tr
  -- start thread endlessly reversing lists of random length
  procObsvIO ← observeIO tr
  -- start threads endlessly observing STM actions operating on the same TVar
  procObsvSTMs ← observeSTM tr
  -- start a thread to output a text messages every n seconds
  procMsg ← msgThr tr
  -- wait for message thread to finish, ignoring any exception
```

```
_ ← Async.waitCatch procMsg
-- wait for observer thread to finish, ignoring any exception
_ ← forM procObsvSTMs Async.waitCatch
-- wait for observer thread to finish, ignoring any exception
_ ← Async.waitCatch procObsvIO
-- wait for random thread to finish, ignoring any exception
_ ← Async.waitCatch procRandom
return ()
```

1.4   Code listings

1.4.1   Cardano.BM.Observer.STM

$$stmWithLog :: STM.STM\,(t, [\text{LogObject}]) \rightarrow STM.STM\,(t, [\text{LogObject}])$$
$$stmWithLog\ action = action$$

Observe *STM* action in a named context

With given name, create a SubTrace according to Configuration and run the passed *STM* action
on it.

```
bracketObserveIO :: Trace IO → Text → STM.STM t → IO t
bracketObserveIO logTrace0 name action = do
    logTrace ← subTrace name logTrace0
    let subtrace = typeofTrace logTrace
    bracketObserveIO' subtrace logTrace action
  where
    bracketObserveIO' :: SubTrace → Trace IO → STM.STM t → IO t
    bracketObserveIO' NoTrace _ act =
      STM.atomically act
    bracketObserveIO' subtrace logTrace act = do
      mCountersid ← observeOpen subtrace logTrace

      -- run action; if an exception is caught will be logged and rethrown.
      t ← (STM.atomically act) ‘catch‘ (λ(e :: SomeException) → (logError logTrace (pack (show e)) ≫ throwM e)
      case mCountersid of
        Left openException →
          -- since observeOpen faced an exception there is no reason to call observeClose
          -- however the result of the action is returned
          logNotice logTrace ("ObserveOpen: " <> pack (show openException))
        Right countersid → do
          res ← observeClose subtrace logTrace countersid []
          case res of
            Left ex → logNotice logTrace ("ObserveClose: " <> pack (show ex))
            _ → pure ()
      pure t
```

Observe *STM* action in a named context and output captured log items

The *STM* action might output messages, which after "success" will be forwarded to the logging trace. Otherwise, this function behaves the same as Observe *STM* action in a named context.

```
bracketObserveLogIO :: Trace IO → Text → STM.STM (t, [LogObject]) → IO t
bracketObserveLogIO logTrace0 name action = do
    logTrace ← subTrace name logTrace0
    let subtrace = typeofTrace logTrace
    bracketObserveLogIO' subtrace logTrace action
  where
    bracketObserveLogIO' :: SubTrace → Trace IO → STM.STM (t, [LogObject]) → IO t
    bracketObserveLogIO' NoTrace _ act = do
        (t, _) ← STM.atomically $ stmWithLog act
        pure t
    bracketObserveLogIO' subtrace logTrace act = do
        mCountersid ← observeOpen subtrace logTrace

        -- run action, return result and log items; if an exception is
        -- caught will be logged and rethrown.
        (t, as) ← (STM.atomically $ stmWithLog act) 'catch'
            (λ(e :: SomeException) → (logError logTrace (pack (show e)) ≫ throwM e))

        case mCountersid of
          Left openException →
            -- since observeOpen faced an exception there is no reason to call observeClose
            -- however the result of the action is returned
            logNotice logTrace ("ObserveOpen: " <> pack (show openException))
          Right countersid → do
            res ← observeClose subtrace logTrace countersid as
            case res of
              Left ex → logNotice logTrace ("ObserveClose: " <> pack (show ex))
              _ → pure ()
        pure t
```

1.4.2 Cardano.BM.Observer.Monadic

Monadic.bracketObserverIO

Observes an *IO* action and adds a name to the logger name of the passed in Trace. An empty *Text* leaves the logger name untouched.

Microbenchmarking steps:

1. Create a *trace* which will have been configured to observe things besides logging.

```
import qualified Cardano.BM.Configuration.Model as CM
∘ ∘ ∘
c ← config
trace@(ctx, _) ← setupTrace (Right c) "demo-playground"
```

```
where
  config :: IO CM.Configuration
  config = do
    c ← CM.empty
    CM.setMinSeverity c Debug
    CM.setSetupBackends c [KatipBK, AggregationBK]
    CM.setDefaultBackends c [KatipBK, AggregationBK]
    CM.setSetupScribes c [ScribeDefinition {
      scName = "stdout"
      , scKind = StdoutSK
      , scRotation = Nothing
      }
      ]
    CM.setDefaultScribes c ["StdoutSK::stdout"]
    return c
```

2. *c* is the Configuration of *trace*. In order to enable the collection and processing of measurements (min, max, mean, std-dev) *AggregationBK* is needed.

CM.setDefaultBackends c [*KatipBK, AggregationBK*]

in a configuration file (YAML) means

defaultBackends :
  − *KatipBK*
  − *AggregationBK*

3. Set the measurements that you want to take by changing the configuration of the *trace* using *setSubTrace*, in order to declare the namespace where we want to enable the particular measurements and the list with the kind of measurements.

```
CM.setSubTrace
    (configuration ctx)
    "demo-playground.submit-tx"
    (Just $ ObservableTrace observablesSet)
  where
    observablesSet = [MonotonicClock, MemoryStats]
```

4. Find an action to measure. e.g.:

runProtocolWithPipe x hdl proto 'catch' (*λProtocolStopped → return* ())

and use bracketObserveIO. e.g.:

```
bracketObserveIO trace "submit-tx" $
    runProtocolWithPipe x hdl proto 'catch' (λProtocolStopped → return ())
```

————————-

```
bracketObserveIO :: Trace IO → Text → IO t → IO t
bracketObserveIO logTrace0 name action = do
```

```
logTrace ← subTrace name logTrace0
bracketObserveIO' (typeofTrace logTrace) logTrace action
where
  bracketObserveIO' :: SubTrace → Trace IO → IO t → IO t
  bracketObserveIO' NoTrace _ act = act
  bracketObserveIO' subtrace logTrace act = do
    mCountersid ← observeOpen subtrace logTrace

    -- run action; if an exception is caught will be logged and rethrown.
    t ← act 'catch' (λ(e :: SomeException) → (logError logTrace (pack (show e)) ≫ throwM e))

    case mCountersid of
      Left openException →
        -- since observeOpen faced an exception there is no reason to call observeClose
        -- however the result of the action is returned
        logNotice logTrace ("ObserveOpen: " <> pack (show openException))
      Right countersid → do
        res ← observeClose subtrace logTrace countersid [ ]
        case res of
          Left ex → logNotice logTrace ("ObserveClose: " <> pack (show ex))
          _ → pure ()
    pure t
```

Monadic.bracketObserverM

Observes a *MonadIO m ⇒ m* action and adds a name to the logger name of the passed in Trace. An empty *Text* leaves the logger name untouched.

```
bracketObserveM :: (MonadCatch m, MonadIO m) ⇒ Trace IO → Text → m t → m t
bracketObserveM logTrace0 name action = do
    logTrace ← liftIO $ subTrace name logTrace0
    bracketObserveM' (typeofTrace logTrace) logTrace action
  where
    bracketObserveM' :: (MonadCatch m, MonadIO m) ⇒ SubTrace → Trace IO → m t → m t
    bracketObserveM' NoTrace _ act = act
    bracketObserveM' subtrace logTrace act = do
      mCountersid ← liftIO $ observeOpen subtrace logTrace

      -- run action; if an exception is caught will be logged and rethrown.
      t ← act 'catch'
          (λ(e :: SomeException) → (liftIO (logError logTrace (pack (show e)) ≫ throwM e)))

      case mCountersid of
        Left openException →
          -- since observeOpen faced an exception there is no reason to call observeClose
          -- however the result of the action is returned
          liftIO $ logNotice logTrace ("ObserveOpen: " <> pack (show openException))
        Right countersid → do
          res ← liftIO $ observeClose subtrace logTrace countersid [ ]
          case res of
            Left ex → liftIO (logNotice logTrace ("ObserveClose: " <> pack (show ex)))
```

$$\_ \rightarrow pure\,()$$
$$pure\ t$$

observerOpen

observeOpen :: SubTrace → Trace *IO* → *IO* (*Either SomeException* CounterState)
observeOpen subtrace logTrace = (**do**
  identifier ← *newUnique*

  -- take measurement
  counters ← *readCounters subtrace*
  let *state* = CounterState *identifier counters*
  if *counters* ≡ [ ]
  then *return* ()
  else do
    -- send opening message to Trace
    traceNamedObject *logTrace* ⋙
      LogObject < $ > *mkLOMeta* < ∗ > *pure* (*ObserveOpen state*)
  return (*Right state*)) '*catch*' (*return* ∘ *Left*)

observeClose

observeClose :: SubTrace → Trace *IO* → CounterState → [LogObject] → *IO* (*Either SomeException* ())
observeClose subtrace logTrace initState logObjects = (**do**
  let *identifier* = *csIdentifier initState*
    initialCounters = *csCounters initState*

  -- take measurement
  counters ← *readCounters subtrace*
  if *counters* ≡ [ ]
  then *return* ()
  else do
    mle ← *mkLOMeta*
    -- send closing message to Trace
    traceNamedObject *logTrace* $
      LogObject *mle* (*ObserveClose* (CounterState *identifier counters*))
    -- send diff message to Trace
    traceNamedObject *logTrace* $
      LogObject *mle* (*ObserveDiff* (CounterState *identifier* (diffCounters *initialCounters counters*)))
  -- trace the messages gathered from inside the action
  forM_ logObjects $ traceNamedObject *logTrace*
  return (*Right* ())) '*catch*' (*return* ∘ *Left*)

1.4.3 BaseTrace

Contravariant

A covariant is a functor: F A → F B
A contravariant is a functor: F B → F A

Op a b implements the inverse to 'arrow' *"getOp :: b → a"*, which when applied to a BaseTrace of type *"Op (m ()) s"*, yields *"s → m ()"*. In our case, *Op* accepts an action in a monad *m* with input type LogNamed LogObject (see 'Trace').

> **newtype** BaseTrace *m s* = BaseTrace {*runTrace :: Op (m ()) s*}

contramap

A covariant functor defines the function *"fmap :: (a → b) → f a → f b"*. In case of a contravariant functor, it is the dual function *"contramap :: (a → b) → f b → f a"* which is defined.

In the following instance, *runTrace* extracts type *"Op (m ()) s"* to which contramap applies *f*, thus *"f s → m ()"*. The constructor BaseTrace restores *"Op (m ()) (f s)"*.

> **instance** *Contravariant* (BaseTrace *m*) **where**
>   *contramap f* = BaseTrace ∘ *contramap f* ∘ *runTrace*

traceWith

Accepts a Trace and some payload *s*. First it gets the contravariant from the Trace as type *"Op (m ()) s"* and, after *"getOp::b → a"* which translates to *"s → m ()"*, calls the action on the LogNamed LogObject.

> traceWith :: BaseTrace *m s → s → m* ()
> traceWith = *getOp* ∘ *runTrace*

natTrace

Natural transformation from monad *m* to monad *n*.

> natTrace :: (*forall x ∘ m x → n x*) → BaseTrace *m s* → BaseTrace *n s*
> natTrace *nat* (BaseTrace (*Op tr*)) = BaseTrace $ *Op* $ *nat* ∘ *tr*

noTrace

A Trace that discards all inputs.

> noTrace :: *Applicative m* ⇒ BaseTrace *m a*
> noTrace = BaseTrace $ *Op* $ *const* (*pure* ())

1.4.4   Cardano.BM.Trace

Utilities

Natural transformation from monad *m* to monad *n*.

> natTrace :: (*forall x ∘ m x → n x*) → Trace *m* → Trace *n*
> natTrace *nat* (*ctx, trace*) = (*ctx*, BaseTrace.natTrace *nat trace*)

Access type of Trace.

```
typeofTrace :: Trace m → SubTrace
typeofTrace (ctx, _) = tracetype ctx
```

Update type of Trace.

```
updateTracetype :: SubTrace → Trace m → Trace m
updateTracetype subtr (ctx, tr) = (ctx {tracetype = subtr}, tr)
```

Enter new named context

The context name is created and checked that its size is below a limit (currently 80 chars). The minimum severity that a log message must be labelled with is looked up in the configuration and recalculated.

```
appendName :: MonadIO m ⇒ LoggerName → Trace m → m (Trace m)
appendName name (ctx, trace) = do
   let prevLoggerName = loggerName ctx
       prevMinSeverity = minSeverity ctx
       newLoggerName   = appendWithDot prevLoggerName name
   globMinSeverity ← liftIO $ Config.minSeverity (configuration ctx)
   namedSeverity ← liftIO $ Config.inspectSeverity (configuration ctx) newLoggerName
   case namedSeverity of
      Nothing → return (ctx {loggerName = newLoggerName}, trace)
      Just sev → return (ctx {loggerName = newLoggerName
        , minSeverity = max (max sev prevMinSeverity) globMinSeverity}
        , trace)
appendWithDot :: LoggerName → LoggerName → LoggerName
appendWithDot "" newName = T.take 80 newName
appendWithDot xs "" = xs
appendWithDot xs newName = T.take 80 $ xs <> "." <> newName
```

Contramap a trace and produce the naming context

```
named :: BaseTrace.BaseTrace m (LogNamed i) → LoggerName → BaseTrace.BaseTrace m i
named trace name = contramap (LogNamed name) trace
```

Trace a LogObject through

```
traceNamedObject
   :: MonadIO m
   ⇒ Trace m
   → LogObject
   → m ()
traceNamedObject trace@(ctx, logTrace) lo@(LogObject _ lc) = do
```

```
    let lname = loggerName ctx
    doOutput ← case (typeofTrace trace) of
      FilterTrace filters →
        case lc of
          LogValue loname _ →
            return $ evalFilters filters (lname <> "." <> loname)
          _ →
            return $ evalFilters filters lname
      TeeTrace secName → do
        -- create a newly named copy of the LogObject
        BaseTrace.traceWith (named logTrace (lname <> "." <> secName)) lo
        return True
      _ → return True
    if doOutput
    then BaseTrace.traceWith (named logTrace lname) lo
    else return ()
```

Evaluation of FilterTrace

A filter consists of a *DropName* and a list of *UnhideNames*. If the context name matches the *DropName* filter, then at least one of the *UnhideNames* must match the name to have the evaluation of the filters return *True*.

```
evalFilters :: [(DropName, UnhideNames)] → LoggerName → Bool
evalFilters fs nm =
    all (λ(no, yes) → if (dropFilter nm no) then (unhideFilter nm yes) else True) fs
  where
    dropFilter :: LoggerName → DropName → Bool
    dropFilter name (Drop sel) =  {-not -} (matchName name sel)
    unhideFilter :: LoggerName → UnhideNames → Bool
    unhideFilter _ (Unhide [ ]) = False
    unhideFilter name (Unhide us) = any (λsel → matchName name sel) us
    matchName :: LoggerName → NameSelector → Bool
    matchName name (Exact name') = name ≡ name'
    matchName name (StartsWith prefix) = T.isPrefixOf prefix name
    matchName name (EndsWith postfix) = T.isSuffixOf postfix name
    matchName name (Contains name') = T.isInfixOf name' name
```

Concrete Trace on stdout

This function returns a trace with an action of type "(LogNamed LogObject) → *IO* ()" which will output a text message as text and all others as JSON encoded representation to the console.

TODO remove *locallock*

```
locallock :: MVar ()
locallock = unsafePerformIO $ newMVar ()
```

```
stdoutTrace :: TraceNamed IO
stdoutTrace = BaseTrace.BaseTrace $ Op $ λ(LogNamed logname (LogObject _ lc)) →
    withMVar locallock $ \_ →
      case lc of
        (LogMessage logItem) →
          output logname $ liPayload logItem
        obj →
          output logname $ toStrict (encodeToLazyText obj)
  where
    output nm msg = TIO.putStrLn $ nm <> " :: " <> msg
```

Concrete Trace into a *TVar*

```
traceInTVar :: STM.TVar [a] → BaseTrace.BaseTrace STM.STM a
traceInTVar tvar = BaseTrace.BaseTrace $ Op $ λa → STM.modifyTVar tvar ((:) a)

traceInTVarIO :: STM.TVar [LogObject] → TraceNamed IO
traceInTVarIO tvar = BaseTrace.BaseTrace $ Op $ λln →
  STM.atomically $ STM.modifyTVar tvar ((:) (lnItem ln))

traceNamedInTVarIO :: STM.TVar [LogNamed LogObject] → TraceNamed IO
traceNamedInTVarIO tvar = BaseTrace.BaseTrace $ Op $ λln →
  STM.atomically $ STM.modifyTVar tvar ((:) ln)
```

Check a log item's severity against the Trace's minimum severity

do we need three different minSeverity defined?

We do a lookup of the global minSeverity in the configuration. And, a lookup of the minSeverity for the current named context. These values might have changed in the meanwhile. A third filter is the minSeverity defined in the current context.

```
traceConditionally
    :: MonadIO m
    ⇒ Trace m
    → LogObject
    → m ()
traceConditionally logTrace@(ctx, _) msg@(LogObject _ (LogMessage item)) = do
    globminsev ← liftIO $ Config.minSeverity (configuration ctx)
    globnamesev ← liftIO $ Config.inspectSeverity (configuration ctx) (loggerName ctx)
    let minsev = max (minSeverity ctx) $ max globminsev (fromMaybe Debug globnamesev)
        flag = (liSeverity item) ⩾ minsev
    when flag $ traceNamedObject logTrace msg
traceConditionally logTrace logObject =
    traceNamedObject logTrace logObject
```

Enter message into a trace

The function traceNamedItem creates a LogObject and threads this through the action defined in the Trace.

```
traceNamedItem
    :: MonadIO m
    ⇒ Trace m
    → LogSelection
    → Severity
    → T.Text
    → m ()
traceNamedItem trace p s m =
    traceConditionally trace ≪
        LogObject < $ > liftIO mkLOMeta
            < * > pure (LogMessage LogItem {liSelection = p
                , liSeverity = s
                , liPayload  = m
                })
```

Logging functions

```
logDebug, logInfo, logNotice, logWarning, logError, logCritical, logAlert, logEmergency
    :: MonadIO m ⇒ Trace m → T.Text → m ()
logDebug     logTrace = traceNamedItem logTrace Both Debug
logInfo      logTrace = traceNamedItem logTrace Both Info
logNotice    logTrace = traceNamedItem logTrace Both Notice
logWarning logTrace = traceNamedItem logTrace Both Warning
logError     logTrace = traceNamedItem logTrace Both Error
logCritical  logTrace = traceNamedItem logTrace Both Critical
logAlert     logTrace = traceNamedItem logTrace Both Alert
logEmergency logTrace = traceNamedItem logTrace Both Emergency

logDebugS, logInfoS, logNoticeS, logWarningS, logErrorS, logCriticalS, logAlertS, logEmergencyS
    :: MonadIO m ⇒ Trace m → T.Text → m ()
logDebugS    logTrace = traceNamedItem logTrace Private Debug
logInfoS     logTrace = traceNamedItem logTrace Private Info
logNoticeS   logTrace = traceNamedItem logTrace Private Notice
logWarningS logTrace = traceNamedItem logTrace Private Warning
logErrorS    logTrace = traceNamedItem logTrace Private Error
logCriticalS logTrace = traceNamedItem logTrace Private Critical
logAlertS    logTrace = traceNamedItem logTrace Private Alert
logEmergencyS logTrace = traceNamedItem logTrace Private Emergency

logDebugP, logInfoP, logNoticeP, logWarningP, logErrorP, logCriticalP, logAlertP, logEmergencyP
    :: MonadIO m ⇒ Trace m → T.Text → m ()
logDebugP    logTrace = traceNamedItem logTrace Public Debug
logInfoP     logTrace = traceNamedItem logTrace Public Info
logNoticeP   logTrace = traceNamedItem logTrace Public Notice
```

```
logWarningP logTrace = traceNamedItem logTrace Public Warning
logErrorP     logTrace = traceNamedItem logTrace Public Error
logCriticalP  logTrace = traceNamedItem logTrace Public Critical
logAlertP     logTrace = traceNamedItem logTrace Public Alert
logEmergencyP logTrace = traceNamedItem logTrace Public Emergency

logDebugUnsafeP, logInfoUnsafeP, logNoticeUnsafeP, logWarningUnsafeP, logErrorUnsafeP,
    logCriticalUnsafeP, logAlertUnsafeP, logEmergencyUnsafeP
    :: MonadIO m ⇒ Trace m → T.Text → m ()
logDebugUnsafeP     logTrace = traceNamedItem logTrace PublicUnsafe Debug
logInfoUnsafeP      logTrace = traceNamedItem logTrace PublicUnsafe Info
logNoticeUnsafeP    logTrace = traceNamedItem logTrace PublicUnsafe Notice
logWarningUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Warning
logErrorUnsafeP     logTrace = traceNamedItem logTrace PublicUnsafe Error
logCriticalUnsafeP  logTrace = traceNamedItem logTrace PublicUnsafe Critical
logAlertUnsafeP     logTrace = traceNamedItem logTrace PublicUnsafe Alert
logEmergencyUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Emergency
```

subTrace

Transforms the input Trace according to the Configuration using the logger name of the current
Trace appended with the new name. If the empty *Text* is passed, then the logger name remains
untouched.

```
subTrace :: MonadIO m ⇒ T.Text → Trace m → m (Trace m)
subTrace name tr@(ctx, _) = do
    let newName = appendWithDot (loggerName ctx) name
    subtrace0 ← liftIO $ Config.findSubTrace (configuration ctx) newName
    let subtrace = case subtrace0 of Nothing → Neutral; Just str → str
    case subtrace of
        Neutral          → do
                            tr' ← appendName name tr
                            return $ updateTracetype subtrace tr'
        UntimedTrace → do
                            tr' ← appendName name tr
                            return $ updateTracetype subtrace tr'
        TeeTrace _       → do
                            tr' ← appendName name tr
                            return $ updateTracetype subtrace tr'
        FilterTrace _    → do
                            tr' ← appendName name tr
                            return $ updateTracetype subtrace tr'
        NoTrace          → return $ updateTracetype subtrace (ctx, BaseTrace.BaseTrace $ Op $ \_ → pure ())
        DropOpening   → return $ updateTracetype subtrace (ctx, BaseTrace.BaseTrace $ Op $
                            λ(LogNamed _ lo@(LogObject _ lc)) → do
                                case lc of
                                    ObserveOpen _ → return ()
                                    _ → traceNamedObject tr lo)
```

```
ObservableTrace _ → do
             tr' ← appendName name tr
             return $ updateTracetype subtrace tr'
```

1.4.5 Cardano.BM.Setup
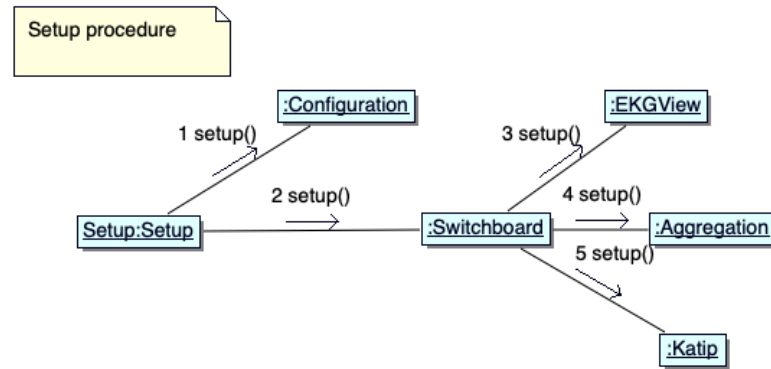
Figure 1.3: Setup procedure

setupTrace

Setup a new Trace (Trace) with either a given Configuration (Configuration.Model) or a *FilePath* to a configuration file.

```
setupTrace :: MonadIO m ⇒ Either FilePath Config.Configuration → Text → m (Trace m)
setupTrace (Left cfgFile) name = do
   c ← liftIO $ Config.setup cfgFile
   setupTrace_ c name
setupTrace (Right c) name = setupTrace_ c name

setupTrace_ :: MonadIO m ⇒ Config.Configuration → Text → m (Trace m)
setupTrace_ c name = do
   sb ← liftIO $ Switchboard.realize c
   sev ← liftIO $ Config.minSeverity c
   ctx ← liftIO $ newContext "" c sev sb
   subTrace name $ natTrace liftIO (ctx, Switchboard.mainTrace sb)
```

withTrace

```
withTrace :: MonadIO m ⇒ Config.Configuration → Text → (Trace m → m t) → m t
withTrace cfg name action = do
   logTrace ← setupTrace (Right cfg) name
   action logTrace
```

newContext

```
newContext :: LoggerName
    → Config.Configuration
    → Severity
    → Switchboard.Switchboard
    → IO TraceContext
newContext name cfg sev sb = do
   return $ TraceContext {
      loggerName = name
      , configuration = cfg
      , minSeverity = sev
      , tracetype = Neutral
      , shutdown = unrealize sb
      }
```

1.4.6   Cardano.BM.Counters

Here the platform is chosen on which we compile this program.
   Currently, we mainly support *Linux* with its 'proc' filesystem.

```
{-# LANGUAGE CPP #-}
# if defined (linux_HOST_OS)
# define LINUX
# endif
module Cardano.BM.Counters
   (
      Platform.readCounters
   , diffTimeObserved
   , getMonoClock
   ) where
# ifdef LINUX
import qualified Cardano.BM.Counters.Linux as Platform
# else
import qualified Cardano.BM.Counters.Dummy as Platform
# endif
import Cardano.BM.Counters.Common (getMonoClock)
import Cardano.BM.Data.Aggregated (Measurable (..))
import Cardano.BM.Data.Counter
```

Calculate difference between clocks

```
diffTimeObserved :: CounterState → CounterState → Measurable
diffTimeObserved (CounterState id0 startCounters) (CounterState id1 endCounters) =
    let
```

```
            startTime = getMonotonicTime startCounters
            endTime = getMonotonicTime endCounters
        in
        if (id0 ≡ id1)
            then endTime − startTime
            else error "these clocks are not from the same experiment"
    where
        getMonotonicTime counters = case (filter isMonotonicClockCounter counters) of
            [(Counter MonotonicClockTime _ mus)] → mus
            _ → error "A time measurement is missing!"

        isMonotonicClockCounter :: Counter → Bool
        isMonotonicClockCounter = (MonotonicClockTime ≡) ∘ cType
```

1.4.7 Cardano.BM.Counters.Common

Common functions that serve *readCounters* on all platforms.

```
nominalTimeToMicroseconds :: Word64 → Microsecond
nominalTimeToMicroseconds = fromMicroseconds ∘ toInteger ∘ (`div`1000)
```

Read monotonic clock

```
getMonoClock :: IO [Counter]
getMonoClock = do
    t ← getMonotonicTimeNSec
    return [Counter MonotonicClockTime "monoclock" $ Microseconds (t `div` 1000)]
```

Read GHC RTS statistics

Read counters from GHC's *RTS* (runtime system). The values returned are as per the last GC (garbage collection) run.

```
readRTSStats :: IO [Counter]
readRTSStats = do
    iscollected ← GhcStats.getRTSStatsEnabled
    if iscollected
        then ghcstats
        else return [ ]
    where
        ghcstats :: IO [Counter]
        ghcstats = do
            -- need to run GC?
            rts ← GhcStats.getRTSStats
            let getrts = ghcval rts
            return [getrts (Bytes ∘ fromIntegral ∘ GhcStats.allocated_bytes, "bytesAllocated")
                   ,getrts (Bytes ∘ fromIntegral ∘ GhcStats.max_live_bytes, "liveBytes")
```

```
              , getrts (Bytes ∘ fromIntegral ∘ GhcStats.max_large_objects_bytes, "largeBytes")
              , getrts (Bytes ∘ fromIntegral ∘ GhcStats.max_compact_bytes, "compactBytes")
              , getrts (Bytes ∘ fromIntegral ∘ GhcStats.max_slop_bytes, "slopBytes")
              , getrts (Bytes ∘ fromIntegral ∘ GhcStats.max_mem_in_use_bytes, "usedMemBytes")
              , getrts (Nanoseconds ∘ fromIntegral ∘ GhcStats.gc_cpu_ns, "gcCpuNs")
              , getrts (Nanoseconds ∘ fromIntegral ∘ GhcStats.gc_elapsed_ns, "gcElapsedNs")
              , getrts (Nanoseconds ∘ fromIntegral ∘ GhcStats.cpu_ns, "cpuNs")
              , getrts (Nanoseconds ∘ fromIntegral ∘ GhcStats.elapsed_ns, "elapsedNs")
              , getrts (PureI ∘ toInteger ∘ GhcStats.gcs, "gcNum")
              , getrts (PureI ∘ toInteger ∘ GhcStats.major_gcs, "gcMajorNum")
              ]
       ghcval :: GhcStats.RTSStats → ((GhcStats.RTSStats → Measurable), Text) → Counter
       ghcval s (f, n) = Counter RTSStats n $ (f s)
```

1.4.8   Cardano.BM.Counters.Dummy

This is a dummy definition of *readCounters* on platforms that do not support the 'proc' filesystem from which we would read the counters.

The only supported measurements are monotonic clock time and RTS statistics for now.

```
       readCounters :: SubTrace → IO [Counter]
       readCounters NoTrace            = return []
       readCounters Neutral            = return []
       readCounters (TeeTrace _)       = return []
       readCounters (FilterTrace _)    = return []
       readCounters UntimedTrace       = return []
       readCounters DropOpening        = return []
       readCounters (ObservableTrace tts) = foldrM (λ(sel, fun) a →
            if any (≡ sel) tts
            then (fun ≫= λxs → return $ a ++ xs)
            else return a)[] selectors
         where
           selectors = [(MonotonicClock, getMonoClock)
              -- , (MemoryStats, readProcStatM)
            -- , (ProcessStats, readProcStats)
                -- , (IOStats, readProcIO)
              , (GhcRtsStats, readRTSStats)
                            ]
```

1.4.9   Cardano.BM.Counters.Linux

> we have to expand the *readMemStats* function
> to read full data from *proc*

```
       readCounters :: SubTrace → IO [Counter]
       readCounters NoTrace            = return []
       readCounters Neutral            = return []
```

```
readCounters (TeeTrace _)     = return [ ]
readCounters (FilterTrace _) = return [ ]
readCounters UntimedTrace = return [ ]
readCounters DropOpening  = return [ ]
readCounters (ObservableTrace tts) = foldrM (λ(sel, fun) a →
     if any (≡ sel) tts
     then (fun ≫ λxs → return $ a ++ xs)
     else return a)[ ] selectors
   where
     selectors = [ (MonotonicClock, getMonoClock)
        , (MemoryStats, readProcStatM)
        , (ProcessStats, readProcStats)
        , (IOStats, readProcIO)
        ]
```

```
pathProc :: FilePath
pathProc = "/proc/"
pathProcStat :: ProcessID → FilePath
pathProcStat pid = pathProc < / > (show pid) < / > "stat"
pathProcStatM :: ProcessID → FilePath
pathProcStatM pid = pathProc < / > (show pid) < / > "statm"
pathProcIO :: ProcessID → FilePath
pathProcIO pid = pathProc < / > (show pid) < / > "io"
```

Reading from a file in /proc/<pid >

```
readProcList :: FilePath → IO [Integer ]
readProcList fp = do
   cs ← readFile fp
   return $ map (λs → maybe 0 id $ (readMaybe s :: Maybe Integer)) (words cs)
```

readProcStatM - /proc/<pid >/statm

```
/proc/[pid]/statm
      Provides information about memory usage, measured in pages.  The columns are:
            size         (1) total program size
                            (same as VmSize in /proc/[pid]/status)
            resident    (2) resident set size
                            (same as VmRSS in /proc/[pid]/status)
            shared      (3) number of resident shared pages (i.e., backed by a file)
                            (same as RssFile+RssShmem in /proc/[pid]/status)
            text         (4) text (code)
            lib          (5) library (unused since Linux 2.6; always 0)
            data         (6) data + stack
            dt           (7) dirty pages (unused since Linux 2.6; always 0)
```

```
readProcStatM :: IO [Counter ]
readProcStatM = do
   pid ← getProcessID
```

> $ps0 \leftarrow readProcList\ (pathProcStatM\ pid)$
> **let** $ps = zip\ colnames\ ps0$
> $psUseful = filter\ (("unused" \not\equiv) \circ fst)\ ps$
> $return\ \$\ map\ (\lambda(n,i) \rightarrow Counter\ MemoryCounter\ n\ (PureI\ i))\ psUseful$
> **where**
> $colnames :: [Text]$
> $colnames = ["size","resident","shared","text","unused","data","unused"]$

readProcStats - //proc//<pid >//stat

```
/proc/[pid]/stat
       Status  information about the process.  This is used by ps(1).  It is defined in the kernel source file
       fs/proc/array.c.

       The fields, in order, with their proper scanf(3) format specifiers, are listed below.  Whether  or  not
       certain  of  these  fields  display  valid information is governed by a ptrace access mode
       PTRACE_MODE_READ_FSCREDS | PTRACE_MODE_NOAUDIT check (refer to ptrace(2)).  If the check denies access,
       then the field value is displayed as 0.  The affected fields are indicated with the marking [PT].

       (1) pid  %d
                    The process ID.

       (2) comm  %s
                    The  filename  of  the  executable,  in parentheses.  This is visible whether or not the exe-
                    cutable is swapped out.

       (3) state  %c
                    One of the following characters, indicating process state:

                    R  Running

                    S  Sleeping in an interruptible wait

                    D  Waiting in uninterruptible disk sleep

                    Z  Zombie

                    T  Stopped (on a signal) or (before Linux 2.6.33) trace stopped

                    t  Tracing stop (Linux 2.6.33 onward)

                    W  Paging (only before Linux 2.6.0)

                    X  Dead (from Linux 2.6.0 onward)

                    x  Dead (Linux 2.6.33 to 3.13 only)

                    K  Wakekill (Linux 2.6.33 to 3.13 only)

                    W  Waking (Linux 2.6.33 to 3.13 only)

                    P  Parked (Linux 3.9 to 3.13 only)

       (4) ppid  %d
                    The PID of the parent of this process.

       (5) pgrp  %d
                    The process group ID of the process.

       (6) session  %d
                    The session ID of the process.

       (7) tty_nr  %d
                    The controlling terminal of the process.  (The minor device number is contained in the combi-
```

nation of bits 31 to 20 and 7 to 0; the major device number is in bits 15 to 8.)

(8) tpgid %d
          The ID of the foreground process group of the controlling terminal of the process.

(9) flags %u
          The kernel flags word of the process. For bit meanings, see the PF_* defines in the Linux
          kernel source file include/linux/sched.h. Details depend on the kernel version.

          The format for this field was %lu before Linux 2.6.

(10) minflt %lu
          The number of minor faults the process has made which have not required loading a memory page
          from disk.

(11) cminflt %lu
          The number of minor faults that the process's waited-for children have made.

(12) majflt %lu
          The number of major faults the process has made which have required loading a memory page
          from disk.

(13) cmajflt %lu
          The number of major faults that the process's waited-for children have made.

(14) utime %lu
          Amount of time that this process has been scheduled in user mode, measured in clock ticks
          (divide by sysconf(_SC_CLK_TCK)). This includes guest time, guest_time (time spent running a
          virtual CPU, see below), so that applications that are not aware of the guest time field do
          not lose that time from their calculations.

(15) stime %lu
          Amount of time that this process has been scheduled in kernel mode, measured in clock ticks
          (divide by sysconf(_SC_CLK_TCK)).

(16) cutime %ld
          Amount of time that this process's waited-for children have been scheduled in user mode, mea-
          sured in clock ticks (divide by sysconf(_SC_CLK_TCK)). (See also times(2).) This includes
          guest time, cguest_time (time spent running a virtual CPU, see below).

(17) cstime %ld
          Amount of time that this process's waited-for children have been scheduled in kernel mode,
          measured in clock ticks (divide by sysconf(_SC_CLK_TCK)).

(18) priority %ld
          (Explanation for Linux 2.6) For processes running a real-time scheduling policy (policy
          below; see sched_setscheduler(2)), this is the negated scheduling priority, minus one; that
          is, a number in the range -2 to -100, corresponding to real-time priorities 1 to 99. For
          processes running under a non-real-time scheduling policy, this is the raw nice value (set-
          priority(2)) as represented in the kernel. The kernel stores nice values as numbers in the
          range 0 (high) to 39 (low), corresponding to the user-visible nice range of -20 to 19.

(19) nice %ld
          The nice value (see setpriority(2)), a value in the range 19 (low priority) to -20 (high pri-
          ority).

(20) num_threads %ld
          Number of threads in this process (since Linux 2.6). Before kernel 2.6, this field was hard
          coded to 0 as a placeholder for an earlier removed field.

(21) itrealvalue %ld
          The time in jiffies before the next SIGALRM is sent to the process due to an interval timer.
          Since kernel 2.6.17, this field is no longer maintained, and is hard coded as 0.

(22) starttime %llu
          The time the process started after system boot. In kernels before Linux 2.6, this value was
          expressed in jiffies.   Since Linux 2.6, the value is expressed in clock ticks (divide by
          sysconf(_SC_CLK_TCK)).

```
               The format for this field was %lu before Linux 2.6.

(23) vsize  %lu
               Virtual memory size in bytes.

(24) rss  %ld
               Resident Set Size: number of pages the process has in real memory.  This is  just  the  pages
               which  count  toward  text, data, or stack space.  This does not include pages which have not
               been demand-loaded in, or which are swapped out.

(25) rsslim  %lu
               Current soft limit in bytes on the rss of the process; see the description of  RLIMIT_RSS  in
               getrlimit(2).

(26) startcode  %lu  [PT]
               The address above which program text can run.

(27) endcode  %lu  [PT]
               The address below which program text can run.

(28) startstack  %lu  [PT]
               The address of the start (i.e., bottom) of the stack.

(29) kstkesp  %lu  [PT]
               The current value of ESP (stack pointer), as found in the kernel stack page for the process.

(30) kstkeip  %lu  [PT]
               The current EIP (instruction pointer).

(31) signal  %lu
               The  bitmap of pending signals, displayed as a decimal number.  Obsolete, because it does not
               provide information on real-time signals; use /proc/[pid]/status instead.

(32) blocked  %lu
               The bitmap of blocked signals, displayed as a decimal number.  Obsolete, because it does  not
               provide information on real-time signals; use /proc/[pid]/status instead.

(33) sigignore  %lu
               The  bitmap of ignored signals, displayed as a decimal number.  Obsolete, because it does not
               provide information on real-time signals; use /proc/[pid]/status instead.

(34) sigcatch  %lu
               The bitmap of caught signals, displayed as a decimal number.  Obsolete, because it  does  not
               provide information on real-time signals; use /proc/[pid]/status instead.

(35) wchan  %lu  [PT]
               This  is  the  "channel" in which the process is waiting.  It is the address of a location in
               the kernel where the process is sleeping.  The corresponding symbolic name can  be  found  in
               /proc/[pid]/wchan.

(36) nswap  %lu
               Number of pages swapped (not maintained).

(37) cnswap  %lu
               Cumulative nswap for child processes (not maintained).

(38) exit_signal  %d  (since Linux 2.1.22)
               Signal to be sent to parent when we die.

(39) processor  %d  (since Linux 2.2.8)
               CPU number last executed on.

(40) rt_priority  %u  (since Linux 2.5.19)
               Real-time  scheduling priority, a number in the range 1 to 99 for processes scheduled under a
               real-time policy, or 0, for non-real-time processes (see sched_setscheduler(2)).

(41) policy  %u  (since Linux 2.5.19)
```

Scheduling policy (see sched_setscheduler(2)). Decode using the SCHED_* constants in linux/sched.h.

The format for this field was %lu before Linux 2.6.22.

(42) delayacct_blkio_ticks %llu (since Linux 2.6.18)
Aggregated block I/O delays, measured in clock ticks (centiseconds).

(43) guest_time %lu (since Linux 2.6.24)
Guest time of the process (time spent running a virtual CPU for a guest operating system), measured in clock ticks (divide by sysconf(_SC_CLK_TCK)).

(44) cguest_time %ld (since Linux 2.6.24)
Guest time of the process's children, measured in clock ticks (divide by sysconf(_SC_CLK_TCK)).

(45) start_data %lu (since Linux 3.3) [PT]
Address above which program initialized and uninitialized (BSS) data are placed.

(46) end_data %lu (since Linux 3.3) [PT]
Address below which program initialized and uninitialized (BSS) data are placed.

(47) start_brk %lu (since Linux 3.3) [PT]
Address above which program heap can be expanded with brk(2).

(48) arg_start %lu (since Linux 3.5) [PT]
Address above which program command-line arguments (argv) are placed.

(49) arg_end %lu (since Linux 3.5) [PT]
Address below program command-line arguments (argv) are placed.

(50) env_start %lu (since Linux 3.5) [PT]
Address above which program environment is placed.

(51) env_end %lu (since Linux 3.5) [PT]
Address below which program environment is placed.

(52) exit_code %d (since Linux 3.5) [PT]
The thread's exit status in the form reported by waitpid(2).

```
readProcStats :: IO [Counter]
readProcStats = do
    pid ← getProcessID
    ps0 ← readProcList (pathProcStat pid)
    let ps = zip colnames ps0
        psUseful = filter (("unused" ≢) ∘ fst) ps
    return $ map (λ(n,i) → Counter StatInfo n (PureI i)) psUseful
  where
    colnames :: [Text]
    colnames = ["pid","unused","unused","ppid","pgrp","session","ttynr","tpgid","flags","minfl
      ,"cminflt","majflt","cmajflt","utime","stime","cutime","cstime","priority","nice","num
      ,"itrealvalue","starttime","vsize","rss","rsslim","startcode","endcode","startstack","
      ,"signal","blocked","sigignore","sigcatch","wchan","nswap","cnswap","exitsignal","proc
      ,"policy","blkio","guesttime","cguesttime","startdata","enddata","startbrk","argstart"
      ,"envend","exitcode"
      ]
```

readProcIO - //proc//<pid >//io

```
/proc/[pid]/io (since kernel 2.6.20)
       This file contains I/O statistics for the process, for example:

                 # cat /proc/3828/io
                 rchar: 323934931
                 wchar: 323929600
                 syscr: 632687
                 syscw: 632675
                 read_bytes: 0
                 write_bytes: 323932160
                 cancelled_write_bytes: 0

       The fields are as follows:

       rchar: characters read
              The number of bytes which this task has caused to be read from storage.   This is simply the   sum
              of bytes which this process passed to read(2) and similar system calls.   It includes things such
              as terminal I/O and is unaffected by whether or not actual physical disk I/O was   required   (the
              read might have been satisfied from pagecache).

       wchar: characters written
              The   number  of bytes which this task has caused, or shall cause to be written to disk.   Similar
              caveats apply here as with rchar.

       syscr: read syscalls
              Attempt  to count the number of read I/O operations-that is, system calls   such   as   read(2)   and
              pread(2).

       syscw: write syscalls
              Attempt   to  count the number of write I/O operations-that is, system calls such as write(2) and
              pwrite(2).

       read_bytes: bytes read
              Attempt to count the number of bytes which this process really did cause to be fetched from  the
              storage layer.   This is accurate for block-backed filesystems.

       write_bytes: bytes written
              Attempt to count the number of bytes which this process caused to be sent to the storage layer.

       cancelled_write_bytes:
              The  big  inaccuracy  here  is  truncate.  If a process writes 1MB to a file and then deletes the
              file, it will in fact perform no writeout.   But it will have been accounted as having caused 1MB
              of  write.    In other words: this field represents the number of bytes which this process caused
              to not happen, by truncating pagecache.   A task can cause "negative"  I/O  too.    If  this  task
              truncates  some  dirty  pagecache,  some  I/O  which another task has been accounted for (in its
              write\_bytes) will not be happening.

       Note: In the current implementation, things are a bit racy  on  32-bit  systems:  if  process  A  reads
       process  B's  /proc/[pid]/io  while process B is updating one of these 64-bit counters, process A could
       see an intermediate result.

       Permission to access this file is governed by a ptrace access mode PTRACE\_MODE\_READ\_FSCREDS check;  see
       ptrace(2).
```

$readProcIO :: IO\,[Counter]$
$readProcIO = \textbf{do}$
$\quad pid \leftarrow getProcessID$
$\quad ps0 \leftarrow readProcList\,(pathProcIO\ pid)$
$\quad \textbf{let}\ ps = zip3\ colnames\ ps0\ units$
$\quad return\ \$\ map\ (\lambda(n,i,u) \rightarrow Counter\ IOCounter\ n\ (u\ i))\ ps$
$\quad\ \textbf{where}$
$\quad\quad colnames :: [\,Text\,]$

```
        colnames = ["rchar","wchar","syscr","syscw","rbytes","wbytes","cxwbytes"]
        units = [Bytes ∘ fromInteger, Bytes ∘ fromInteger, PureI, PureI, Bytes ∘ fromInteger, Bytes ∘ fromInteger, Bytes ∘
```

1.4.10   Cardano.BM.Data.Aggregated

Measurable

A Measurable may consist of different types of values. Time measurements are strict, so are *Bytes* which are externally measured. The real or integral numeric values are lazily linked, so we can decide later to drop them.

```
    data Measurable = Microseconds {-# UNPACK #-} ! Word64
      | Nanoseconds {-# UNPACK #-} ! Word64
      | Seconds      {-# UNPACK #-} ! Word64
      | Bytes        {-# UNPACK #-} ! Word64
      | PureD        Double
      | PureI        Integer
      deriving (Eq, Ord, Generic, ToJSON)
```

Measurable can be transformed to an integral value.

```
    getInteger :: Measurable → Integer
    getInteger (Microseconds a) = toInteger a
    getInteger (Nanoseconds a) = toInteger a
    getInteger (Seconds a)      = toInteger a
    getInteger (Bytes a)        = toInteger a
    getInteger (PureI a)        = a
    getInteger (PureD a)        = round a
```

Measurable can be transformed to a rational value.

```
    getDouble :: Measurable → Double
    getDouble (Microseconds a) = fromIntegral a
    getDouble (Nanoseconds a) = fromIntegral a
    getDouble (Seconds a)      = fromIntegral a
    getDouble (Bytes a)        = fromIntegral a
    getDouble (PureI a)        = fromInteger a
    getDouble (PureD a)        = a
```

It is a numerical value, thus supports functions to operate on numbers.

```
    instance Num Measurable where
      (+) (Microseconds a) (Microseconds b) = Microseconds (a + b)
      (+) (Nanoseconds a) (Nanoseconds b) = Nanoseconds (a + b)
      (+) (Seconds a)      (Seconds b)     = Seconds     (a + b)
      (+) (Bytes a)        (Bytes b)       = Bytes       (a + b)
      (+) (PureI a)        (PureI b)       = PureI       (a + b)
      (+) (PureD a)        (PureD b)       = PureD       (a + b)
      (+) _ _                              = error "Trying to add values with different units"
```

```
(∗) (Microseconds a) (Microseconds b) = Microseconds (a ∗ b)
(∗) (Nanoseconds a)  (Nanoseconds b) = Nanoseconds (a ∗ b)
(∗) (Seconds a)      (Seconds b)     = Seconds      (a ∗ b)
(∗) (Bytes a)        (Bytes b)       = Bytes        (a ∗ b)
(∗) (PureI a)        (PureI b)       = PureI        (a ∗ b)
(∗) (PureD a)        (PureD b)       = PureD        (a ∗ b)
(∗) _ _                              = error "Trying to multiply values with different units"

abs (Microseconds a) = Microseconds (abs a)
abs (Nanoseconds a)  = Nanoseconds (abs a)
abs (Seconds a)      = Seconds (abs a)
abs (Bytes a)        = Bytes    (abs a)
abs (PureI a)        = PureI    (abs a)
abs (PureD a)        = PureD    (abs a)

signum (Microseconds a) = Microseconds (signum a)
signum (Nanoseconds a)  = Nanoseconds (signum a)
signum (Seconds a)      = Seconds       (signum a)
signum (Bytes a)        = Bytes         (signum a)
signum (PureI a)        = PureI         (signum a)
signum (PureD a)        = PureD         (signum a)

negate (Microseconds a) = Microseconds (negate a)
negate (Nanoseconds a)  = Nanoseconds (negate a)
negate (Seconds a)      = Seconds       (negate a)
negate (Bytes a)        = Bytes         (negate a)
negate (PureI a)        = PureI         (negate a)
negate (PureD a)        = PureD         (negate a)

fromInteger = PureI
```

Pretty printing of Measurable.

```
instance Show Measurable where
  show (Microseconds a) = show a
  show (Nanoseconds a)   = show a
  show (Seconds a)       = show a
  show (Bytes a)         = show a
  show (PureI a)         = show a
  show (PureD a)         = show a

showUnits :: Measurable → String
showUnits (Microseconds _) = " s"
showUnits (Nanoseconds _)  = " ns"
showUnits (Seconds _)      = " s"
showUnits (Bytes _)        = " B"
showUnits (PureI _)        = ""
showUnits (PureD _)        = ""

-- show in S.I. units
showSI :: Measurable → String
showSI (Microseconds a) = show (fromFloatDigits ((fromIntegral a) / (1000000 :: Float))) ++
                          showUnits (Seconds a)
```

showSI (*Nanoseconds a*) = *show* (*fromFloatDigits* ((*fromIntegral a*) / (1000000000 :: *Float*))) ⧺
                       showUnits (*Seconds a*)
showSI v@(*Seconds a*)   = *show a* ⧺ *showUnits v*
showSI v@(*Bytes a*)   = *show a* ⧺ *showUnits v*
showSI v@(*PureI a*)   = *show a* ⧺ *showUnits v*
showSI v@(*PureD a*)   = *show a* ⧺ *showUnits v*

Stats

A Stats statistics is strictly computed.

data *BaseStats* = *BaseStats* {
   fmin :: !Measurable,
   fmax :: !Measurable,
   fcount :: !*Word64*,
   fsum_A :: !*Double*,
   fsum_B :: !*Double*
   } **deriving** (*Generic*, *ToJSON*, *Show*)
instance *Eq BaseStats* **where**
   (*BaseStats mina maxa counta sumAa sumBa*) ≡ (*BaseStats minb maxb countb sumAb sumBb*) =
     mina ≡ *minb* ∧ *maxa* ≡ *maxb* ∧ *counta* ≡ *countb* ∧
     abs (*sumAa* − *sumAb*) < 1.0*e*−4 ∧
     abs (*sumBa* − *sumBb*) < 1.0*e*−4
data Stats = Stats {
   flast :: !Measurable,
   fold :: !Measurable,
   fbasic :: !*BaseStats*,
   fdelta :: !*BaseStats*,
   ftimed :: !*BaseStats*
   } **deriving** (*Eq*, *Generic*, *ToJSON*, *Show*)

meanOfStats :: *BaseStats* → *Double*
meanOfStats = *fsum_A*

stdevOfStats :: *BaseStats* → *Double*
stdevOfStats s =
   if *fcount s* < 2
   then 0
   else *sqrt* $ (*fsum_B s*) / (*fromInteger* $ *fromIntegral* (*fcount s*) − 1)

> **instance** *Semigroup* Stats disabled for the moment, because not needed.

We use a parallel algorithm to update the estimation of mean and variance from two sample statistics. (see `https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance#Parallel_algorit`

instance *Semigroup* Stats **where**
   (<>) *a b* = **let** *counta* = *fcount a*

```
    countb = fcount b
    newcount = counta + countb
    delta = fsum_A b − fsum_A a
    in
    Stats {flast = flast b-- right associative
       ,fmin    = min (fmin a) (fmin b)
       ,fmax    = max (fmax a) (fmax b)
       ,fcount = newcount
       ,fsum_A = fsum_A a + (delta / fromInteger newcount)
       ,fsum_B = fsum_B a + fsum_B b + (delta ∗ delta) ∗ (fromInteger (counta ∗ countb) / fromInteger newcount)
       }
```

```
stats2Text :: Stats → Text
stats2Text (Stats slast _ sbasic sdelta stimed) =
    pack $
       "{ last=" ++ show slast ++
       ", basic-stats=" ++ showStats' (sbasic) ++
       ", delta-stats=" ++ showStats' (sdelta) ++
       ", timed-stats=" ++ showStats' (stimed) ++
       " }"
    where
      showStats' :: BaseStats → String
      showStats' s =
         ", { min=" ++ show (fmin s) ++
         ", max=" ++ show (fmax s) ++
         ", mean=" ++ show (meanOfStats s) ++ showUnits (fmin s) ++
         ", std-dev=" ++ show (stdevOfStats s) ++
         ", count=" ++ show (fcount s) ++
         " }"
```

Exponentially Weighted Moving Average (EWMA)

Following https://en.wikipedia.org/wiki/Moving_average#Exponential_moving_average we
calculate the exponential moving average for a series of values Y_t according to:

$$S_t = \begin{cases} Y_1, & t = 1 \\ \alpha \cdot Y_t + (1 - \alpha) \cdot S_{t-1}, & t > 1 \end{cases}$$

```
data EWMA = EmptyEWMA {alpha :: Double}
    | EWMA {alpha :: Double
       ,avg :: Measurable
       } deriving (Show, Eq, Generic, ToJSON)
```

Aggregated

```
data Aggregated = AggregatedStats Stats
  | AggregatedEWMA EWMA
  deriving (Eq, Generic, ToJSON)
```

instance *Semigroup* Aggregated disabled for the moment, because not needed.

```
instance Semigroup Aggregated where
  (<>) (AggregatedStats a) (AggregatedStats b) =
    AggregatedStats (a <> b)
  (<>) _ _ = error "Cannot combine different objects"


singletonStats :: Measurable → Aggregated
singletonStats a =
  let stats = Stats {flast = a
    ,fold                = 0
    ,fbasic = BaseStats
      {fmin = a
      ,fmax = a
      ,fcount = 1
      ,fsum_A = getDouble a
      ,fsum_B = 0}
    ,fdelta = BaseStats
      {fmin = 0
      ,fmax = 0
      ,fcount = 0
      ,fsum_A = 0
      ,fsum_B = 0}
    ,ftimed = BaseStats
      {fmin = Nanoseconds 999999999999
      ,fmax = Nanoseconds 0
      ,fcount = (−1)
      ,fsum_A = 0
      ,fsum_B = 0}
    }
  in
  AggregatedStats stats


instance Show Aggregated where
  show (AggregatedStats astats) =
    "{ stats = " ++ show astats ++ " }"
  show (AggregatedEWMA a) = show a
```

1.4.11   Cardano.BM.Data.Backend

Accepts a NamedLogItem

Instances of this type class accept a NamedLogItem and deal with it.

```
class IsEffectuator t where
    effectuate :: t → NamedLogItem → IO ()
    effectuatefrom :: forall s ∘ (IsEffectuator s) ⇒ t → NamedLogItem → s → IO ()
    default effectuatefrom :: forall s ∘ (IsEffectuator s) ⇒ t → NamedLogItem → s → IO ()
    effectuatefrom t nli _ = effectuate t nli
```

Declaration of a Backend

A backend is life-cycle managed, thus can be *realize*d and *unrealize*d.

```
class (IsEffectuator t) ⇒ IsBackend t where
    typeof     :: t → BackendKind
    realize    :: Configuration → IO t
    realizefrom :: forall s ∘ (IsEffectuator s) ⇒ Trace IO → s → IO t
    default realizefrom :: forall s ∘ (IsEffectuator s) ⇒ Trace IO → s → IO t
    realizefrom (ctx, _) _ = realize (configuration ctx)
    unrealize :: t → IO ()
```

Backend

This data structure for a backend defines its behaviour as an IsEffectuator when processing an incoming message, and as an IsBackend for unrealizing the backend.

```
data Backend = MkBackend
    {bEffectuate :: NamedLogItem → IO ()
    ,bUnrealize :: IO ()
    }
```

1.4.12   Cardano.BM.Data.Configuration

Data structure to help parsing configuration files.

Representation

```
type Port = Int
data Representation = Representation
    {minSeverity    :: Severity
    ,rotation       :: RotationParameters
    ,setupScribes   :: [ScribeDefinition]
    ,defaultScribes :: [(ScribeKind, Text)]
    ,setupBackends :: [BackendKind]
    ,defaultBackends :: [BackendKind]
```

```
    , hasEKG          :: Maybe Port
    , hasGUI          :: Maybe Port
    , options         :: HM.HashMap Text Object
    }
    deriving (Generic, Show, ToJSON, FromJSON)
```

parseRepresentation

```
parseRepresentation :: FilePath → IO Representation
parseRepresentation fp = do
    repr :: Representation ← decodeFileThrow fp
    return $ implicit_fill_representation repr
```

after parsing the configuration representation we implicitly correct it.

```
implicit_fill_representation :: Representation → Representation
implicit_fill_representation =
    remove_ekgview_if_not_defined ∘
    filter_duplicates_from_backends ∘
    filter_duplicates_from_scribes ∘
    union_setup_and_usage_backends ∘
    add_ekgview_if_port_defined ∘
    add_katip_if_any_scribes
  where
    filter_duplicates_from_backends r =
        r { setupBackends = mkUniq $ setupBackends r }
    filter_duplicates_from_scribes r =
        r { setupScribes = mkUniq $ setupScribes r }
    union_setup_and_usage_backends r =
        r { setupBackends = setupBackends r <> defaultBackends r }
    remove_ekgview_if_not_defined r =
        case hasEKG r of
        Nothing → r { defaultBackends = filter (λbk → bk ≢ EKGViewBK) (defaultBackends r)
            , setupBackends = filter (λbk → bk ≢ EKGViewBK) (setupBackends r)
            }
        Just _ → r
    add_ekgview_if_port_defined r =
        case hasEKG r of
        Nothing → r
        Just _ → r { setupBackends = setupBackends r <> [EKGViewBK] }
    add_katip_if_any_scribes r =
        if (any ¬ [null $ setupScribes r, null $ defaultScribes r])
        then r { setupBackends = setupBackends r <> [KatipBK] }
        else r
    mkUniq :: Ord a ⇒ [a] → [a]
    mkUniq = Set.toList ∘ Set.fromList
```

1.4.13   Cardano.BM.Data.Counter

Counter

```
data Counter = Counter
         {cType  :: CounterType
         ,cName :: Text
         ,cValue :: Measurable
         }
         deriving (Eq, Show, Generic, ToJSON)
data CounterType = MonotonicClockTime
  | MemoryCounter
  | StatInfo
  | IOCounter
  | CpuCounter
  | RTSStats
    deriving (Eq, Show, Generic, ToJSON)
instance ToJSON Microsecond where
  toJSON = toJSON ∘ toMicroseconds
  toEncoding = toEncoding ∘ toMicroseconds
```

Names of counters

```
nameCounter :: Counter → Text
nameCounter (Counter MonotonicClockTime _ _) = "Time-interval"
nameCounter (Counter MemoryCounter _ _) = "Mem"
nameCounter (Counter StatInfo       _ _) = "Stat"
nameCounter (Counter IOCounter      _ _) = "IO"
nameCounter (Counter CpuCounter     _ _) = "Cpu"
nameCounter (Counter RTSStats       _ _) = "RTS"
```

CounterState

```
data CounterState = CounterState {
  csIdentifier :: Unique
  ,csCounters :: [Counter]
  }
  deriving (Generic, ToJSON)
instance ToJSON Unique where
  toJSON = toJSON ∘ hashUnique
  toEncoding = toEncoding ∘ hashUnique
instance Show CounterState where
  show cs = (show ∘ hashUnique) (csIdentifier cs)
     <> " => " <> (show $ csCounters cs)
```

Difference between counters

```
diffCounters :: [Counter] → [Counter] → [Counter]
diffCounters openings closings =
    getCountersDiff openings closings
  where
    getCountersDiff :: [Counter]
              → [Counter]
              → [Counter]
    getCountersDiff as bs =
        let
          getName counter = nameCounter counter <> cName counter

          asNames = map getName as
          aPairs = zip asNames as

          bsNames = map getName bs
          bs' = zip bsNames bs
          bPairs = HM.fromList bs'
        in
          catMaybes $ (flip map) aPairs $ λ(name, Counter _ _ startValue) →
            case HM.lookup name bPairs of
              Nothing    → Nothing
              Just counter → let endValue = cValue counter
                          in Just counter {cValue = endValue − startValue}
```

1.4.14   Cardano.BM.Data.LogItem

LoggerName

A LoggerName has currently type *Text*.

```
type LoggerName = Text
```

NamedLogItem

```
type NamedLogItem = LogNamed LogObject
```

LogNamed

A LogNamed contains of a context name and some log item.

```
data LogNamed item = LogNamed
  {lnName :: LoggerName
  ,lnItem :: item
  } deriving (Show)

deriving instance Generic item ⇒ Generic (LogNamed item)
deriving instance (ToJSON item, Generic item) ⇒ ToJSON (LogNamed item)
```

Logging of outcomes with LogObject

```haskell
data LogObject = LogObject LOMeta LOContent
    deriving (Generic, Show, ToJSON)
```

Meta data for a LogObject:

```haskell
data LOMeta = LOMeta {
  tstamp :: {-# UNPACK #-} ! UTCTime
  , tid :: {-# UNPACK #-} ! ThreadId
  }
  deriving (Show)
instance ToJSON LOMeta where
  toJSON (LOMeta _tstamp _tid) =
    object ["tstamp" .= _tstamp, "tid" .= show _tid]
mkLOMeta :: IO LOMeta
mkLOMeta =
  LOMeta <$> getCurrentTime
    <*> myThreadId
```

Payload of a LogObject:

```haskell
data LOContent = LogMessage LogItem
  | LogValue Text Measurable
  | ObserveOpen CounterState
  | ObserveDiff CounterState
  | ObserveClose CounterState
  | AggregatedMessage [(Text, Aggregated)]
  | KillPill
    deriving (Generic, Show, ToJSON)
```

LogItem

> TODO liPayload :: ToObject

```haskell
data LogItem = LogItem
  {liSelection :: LogSelection
  , liSeverity :: Severity
  , liPayload :: Text-- TODO should become ToObject
  } deriving (Show, Generic, ToJSON)

data LogSelection =
  Public   -- only to public logs.
  | PublicUnsafe-- only to public logs, not console.
  | Private-- only to private logs.
  | Both   -- to public and private logs.
  deriving (Show, Generic, ToJSON, FromJSON)
```

1.4.15 Cardano.BM.Data.Observable

ObservableInstance

```
data ObservableInstance = MonotonicClock
  | MemoryStats
  | ProcessStats
  | IOStats
  | GhcRtsStats
    deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON, Read)
```

1.4.16 Cardano.BM.Data.Output

OutputKind

```
data OutputKind = TVarList (STM.TVar [LogObject])
  | TVarListNamed (STM.TVar [LogNamed LogObject])
  deriving (Eq)
```

ScribeKind

This identifies katip's scribes by type.

```
data ScribeKind = FileTextSK
  | FileJsonSK
  | StdoutSK
  | StderrSK
  deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)
```

ScribeId

A scribe is identified by ScribeKind *x Filename*

```
type ScribeId = Text-- (ScribeKind ::  Filename)
```

ScribeDefinition

This identifies katip's scribes by type.

```
data ScribeDefinition = ScribeDefinition
  {scKind  :: ScribeKind
  , scName :: Text
  , scRotation :: Maybe RotationParameters
  }
  deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)
```

1.4.17   Cardano.BM.Data.Severity

Severity

The intended meaning of severity codes:

Debug *detailed information about values and decision flow* Info *general information of events;* progressing properly Notice *needs attention;something ¬ progressing properly* Warning *may continue into an error condition if continued* Error *unexpected set* **of** *event or condition occured* Critical *error condition causing degrade of operation* Alert *a subsystem is no longer operating correctly,likely requires manual* at this point, the system can never progress without additional intervention

We were informed by the *Syslog* taxonomy: `https://en.wikipedia.org/wiki/Syslog#Severity_level`

```
data Severity = Debug
    | Info
    | Notice
    | Warning
    | Error
    | Critical
    | Alert
    | Emergency
      deriving (Show, Eq, Ord, Generic, ToJSON, Read)
instance FromJSON Severity where
  parseJSON = withText "severity" $ λcase
    "Debug"     → pure Debug
    "Info"      → pure Info
    "Notice"    → pure Notice
    "Warning"   → pure Warning
    "Error"     → pure Error
    "Critical"  → pure Critical
    "Alert"     → pure Alert
    "Emergency" → pure Emergency
    _           → pure Info-- catch all
```

1.4.18   Cardano.BM.Data.SubTrace

SubTrace

```
data NameSelector = Exact Text | StartsWith Text | EndsWith Text | Contains Text
                    deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)
data DropName     = Drop NameSelector
                    deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)
data UnhideNames = Unhide [NameSelector]
                    deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)
data SubTrace = Neutral
    | UntimedTrace
    | NoTrace
    | TeeTrace LoggerName
    | FilterTrace [(DropName, UnhideNames)]
```

```
    | DropOpening
    | ObservableTrace [ ObservableInstance ]
      deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)
```

1.4.19 Cardano.BM.Data.Trace

Trace

A Trace consists of a TraceContext and a TraceNamed in *m*.

```
    type Trace m = (TraceContext, TraceNamed m)
```

TraceNamed

A TraceNamed is a specialized Contravariant of type NamedLogItem, a LogNamed with payload LogObject.

```
    type TraceNamed m = BaseTrace m (NamedLogItem)
```

TraceContext

We keep the context's name and a reference to the Configuration in the TraceContext.

```
    data TraceContext = TraceContext
      { loggerName :: LoggerName
      , configuration :: Configuration
      , tracetype    :: SubTrace
      , minSeverity :: Severity
      , shutdown    :: IO ()
      }
```

1.4.20 Cardano.BM.Configuration

see Cardano.BM.Configuration.Model for the implementation.

```
    getOptionOrDefault :: CM.Configuration → Text → Text → IO (Text)
    getOptionOrDefault cg name def = do
      opt ← CM.getOption cg name
      case opt of
        Nothing → return def
        Just o → return o
```

Figure 1.4: Configuration model

1.4.21 Cardano.BM.Configuration.Model

Configuration.Model

```
type ConfigurationMVar = MVar ConfigurationInternal
newtype Configuration = Configuration
  {getCG :: ConfigurationMVar}
-- Our internal state; see -"Configuration model"-
data ConfigurationInternal = ConfigurationInternal
  {cgMinSeverity      :: Severity
  -- minimum severity level of every object that will be output
  ,cgMapSeverity      :: HM.HashMap LoggerName Severity
  -- severity filter per loggername
  ,cgMapSubtrace      :: HM.HashMap LoggerName SubTrace
  -- type of trace per loggername
  ,cgOptions          :: HM.HashMap Text Object
  -- options needed for tracing, logging and monitoring
  ,cgMapBackend       :: HM.HashMap LoggerName [BackendKind]
  -- backends that will be used for the specific loggername
  ,cgDefBackendKs     :: [BackendKind]
  -- backends that will be used if a set of backends for the
  -- specific loggername is not set
  ,cgSetupBackends    :: [BackendKind]
  -- backends to setup; every backend to be used must have
  -- been declared here
  ,cgMapScribe        :: HM.HashMap LoggerName [ScribeId]
  -- katip scribes that will be used for the specific loggername
  ,cgMapScribeCache   :: HM.HashMap LoggerName [ScribeId]
  -- map to cache info of the cgMapScribe
  ,cgDefScribes       :: [ScribeId]
  -- katip scribes that will be used if a set of scribes for the
  -- specific loggername is not set
  ,cgSetupScribes     :: [ScribeDefinition]
  -- katip scribes to setup; every scribe to be used must have
  -- been declared here
  ,cgMapAggregatedKind :: HM.HashMap LoggerName AggregatedKind
  -- kind of Aggregated that will be used for the specific loggername
  ,cgDefAggregatedKind :: AggregatedKind
  -- kind of Aggregated that will be used if a set of scribes for the
  -- specific loggername is not set
  ,cgPortEKG          :: Int
  -- port for EKG server
  ,cgPortGUI          :: Int
  -- port for changes at runtime (NOT IMPLEMENTED YET)
  } deriving (Show, Eq)
```

Backends configured in the Switchboard

For a given context name return the list of backends configured, or, in case no such configuration exists, return the default backends.

> *getBackends* :: Configuration → LoggerName → *IO* [*BackendKind*]
> *getBackends* configuration *name* =
>   *withMVar* (*getCG* configuration) $ *λcg* → **do**
>   **let** *outs* = *HM.lookup name* (*cgMapBackend cg*)
>   **case** *outs* **of**
>   *Nothing* → **do**
>   *return* (*cgDefBackendKs cg*)
>   *Just os* → *return os*
>
> *getDefaultBackends* :: Configuration → *IO* [*BackendKind*]
> *getDefaultBackends* configuration =
>   *withMVar* (*getCG* configuration) $ *λcg* → **do**
>   *return* (*cgDefBackendKs cg*)
>
> *setDefaultBackends* :: Configuration → [*BackendKind*] → *IO* ()
> *setDefaultBackends* configuration *bes* = **do**
>   *cg* ← *takeMVar* (*getCG* configuration)
>   *putMVar* (*getCG* configuration) $ *cg* {*cgDefBackendKs* = *bes*}
>
> *setBackends* :: Configuration → LoggerName → *Maybe* [*BackendKind*] → *IO* ()
> *setBackends* configuration *name be* = **do**
>   *cg* ← *takeMVar* (*getCG* configuration)
>   *putMVar* (*getCG* configuration) $ *cg* {*cgMapBackend* = *HM.alter* (_ → *be*) *name* (*cgMapBackend cg*)}

Backends to be setup by the Switchboard

Defines the list of Backends that need to be setup by the Switchboard.

> *setSetupBackends* :: Configuration → [*BackendKind*] → *IO* ()
> *setSetupBackends* configuration *bes* = **do**
>   *cg* ← *takeMVar* (*getCG* configuration)
>   *putMVar* (*getCG* configuration) $ *cg* {*cgSetupBackends* = *bes*}
>
> *getSetupBackends* :: Configuration → *IO* [*BackendKind*]
> *getSetupBackends* configuration =
>   *withMVar* (*getCG* configuration) $ *λcg* →
>   *return* $ *cgSetupBackends cg*

Scribes configured in the Log backend

For a given context name return the list of scribes to output to, or, in case no such configuration exists, return the default scribes to use.

> *getScribes* :: Configuration → LoggerName → *IO* [*ScribeId*]
> *getScribes* configuration *name* = **do**
>   (*updateCache, scribes*) ← *withMVar* (*getCG* configuration) $ *λcg* → **do**

```
        let defs = cgDefScribes cg
        let mapScribe = cgMapScribe cg
        let find_s lname = case HM.lookup lname mapScribe of
           Nothing →
              case dropToDot lname of
                 Nothing → defs
                 Just lname' → find_s lname'
           Just os → os
        let outs = HM.lookup name (cgMapScribeCache cg)
        -- look if scribes are already cached
        return $ case outs of
           -- if no cached scribes found; search the appropriate scribes that
           -- they must inherit and update the cached map
           Nothing → (True, find_s name)
           Just os → (False, os)
     when updateCache $ setCachedScribes configuration name $ Just scribes
     return scribes
   where
     dropToDot :: Text → Maybe Text
     dropToDot ts = dropToDot' (breakOnEnd "." ts)
     dropToDot' (_, "") = Nothing
     dropToDot' (name', _) = Just $ dropWhileEnd (≡ '.') name'
getCachedScribes :: Configuration → LoggerName → IO (Maybe [ScribeId])
getCachedScribes configuration name =
     withMVar (getCG configuration) $ λcg → do
        return $ HM.lookup name $ cgMapScribeCache cg
setScribes :: Configuration → LoggerName → Maybe [ScribeId] → IO ()
setScribes configuration name scribes = do
     cg ← takeMVar (getCG configuration)
     putMVar (getCG configuration) $
        cg {cgMapScribe = HM.alter (\_ → scribes) name (cgMapScribe cg)}
setCachedScribes :: Configuration → LoggerName → Maybe [ScribeId] → IO ()
setCachedScribes configuration name scribes = do
     cg ← takeMVar (getCG configuration)
     putMVar (getCG configuration) $
        cg {cgMapScribeCache = HM.alter (\_ → scribes) name (cgMapScribeCache cg)}
setDefaultScribes :: Configuration → [ScribeId] → IO ()
setDefaultScribes configuration scs = do
     cg ← takeMVar (getCG configuration)
     putMVar (getCG configuration) $ cg {cgDefScribes = scs}
```

Scribes to be setup in the Log backend

Defines the list of *Scribe*s that need to be setup in the Log backend.

```
setSetupScribes :: Configuration → [ScribeDefinition] → IO ()
setSetupScribes configuration sds = do
```

```
cg ← takeMVar (getCG configuration)
putMVar (getCG configuration) $ cg {cgSetupScribes = sds}
getSetupScribes :: Configuration → IO [ScribeDefinition]
getSetupScribes configuration =
   withMVar (getCG configuration) $ λcg → do
      return $ cgSetupScribes cg
```

AggregatedKind **to define the type of measurement**

For a given context name return its *AggregatedKind* or in case no such configuration exists, return the default *AggregatedKind* to use.

```
getAggregatedKind :: Configuration → LoggerName → IO AggregatedKind
getAggregatedKind configuration name =
   withMVar (getCG configuration) $ λcg → do
      let outs = HM.lookup name (cgMapAggregatedKind cg)
      case outs of
         Nothing → do
            return (cgDefAggregatedKind cg)
         Just os → return $ os
setDefaultAggregatedKind :: Configuration → AggregatedKind → IO ()
setDefaultAggregatedKind configuration defAK = do
   cg ← takeMVar (getCG configuration)
   putMVar (getCG configuration) $ cg {cgDefAggregatedKind = defAK}
setAggregatedKind :: Configuration → LoggerName → Maybe AggregatedKind → IO ()
setAggregatedKind configuration name ak = do
   cg ← takeMVar (getCG configuration)
   putMVar (getCG configuration) $ cg {cgMapAggregatedKind = HM.alter (\_ → ak) name (cgMapAggregatedKi
```

Access port numbers of EKG, GUI

```
getEKGport :: Configuration → IO Int
getEKGport configuration =
   withMVar (getCG configuration) $ λcg → do
      return $ cgPortEKG cg
setEKGport :: Configuration → Int → IO ()
setEKGport configuration port = do
   cg ← takeMVar (getCG configuration)
   putMVar (getCG configuration) $ cg {cgPortEKG = port}
getGUIport :: Configuration → IO Int
getGUIport configuration =
   withMVar (getCG configuration) $ λcg → do
      return $ cgPortGUI cg
setGUIport :: Configuration → Int → IO ()
setGUIport configuration port = do
```

```
cg ← takeMVar (getCG configuration)
putMVar (getCG configuration) $ cg {cgPortGUI = port}
```

Options

```
getOption :: Configuration → Text → IO (Maybe Text)
getOption configuration name = do
  withMVar (getCG configuration) $ λcg →
    case HM.lookup name (cgOptions cg) of
      Nothing → return Nothing
      Just o → return $ Just $ pack $ show o
```

Global setting of minimum severity

```
minSeverity :: Configuration → IO Severity
minSeverity configuration = withMVar (getCG configuration) $ λcg →
  return $ cgMinSeverity cg

setMinSeverity :: Configuration → Severity → IO ()
setMinSeverity configuration sev = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMinSeverity = sev}
```

Relation of context name to minimum severity

```
inspectSeverity :: Configuration → Text → IO (Maybe Severity)
inspectSeverity configuration name = do
  withMVar (getCG configuration) $ λcg →
    return $ HM.lookup name (cgMapSeverity cg)

setSeverity :: Configuration → Text → Maybe Severity → IO ()
setSeverity configuration name sev = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMapSeverity = HM.alter (\_ → sev) name (cgMapSeverity cg)}
```

Relation of context name to SubTrace

A new context may contain a different type of Trace. The function appendName (Enter new named context) will look up the SubTrace for the context's name.

```
findSubTrace :: Configuration → Text → IO (Maybe SubTrace)
findSubTrace configuration name = do
  withMVar (getCG configuration) $ λcg →
    return $ HM.lookup name (cgMapSubtrace cg)

setSubTrace :: Configuration → Text → Maybe SubTrace → IO ()
setSubTrace configuration name trafo = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMapSubtrace = HM.alter (\_ → trafo) name (cgMapSubtrace cg)}
```

Parse configuration from file

Parse the configuration into an internal representation first. Then, fill in Configuration after refinement.

```
setup :: FilePath → IO Configuration
setup fp = do
    r ← R.parseRepresentation fp
    setupFromRepresentation r

setupFromRepresentation :: R.Representation → IO Configuration
setupFromRepresentation r = do
    cgref ← newEmptyMVar
    let mapseverity = HM.lookup "mapSeverity" (R.options r)
        mapbackends = HM.lookup "mapBackends" (R.options r)
        mapsubtrace  = HM.lookup "mapSubtrace" (R.options r)
        mapscribes   = HM.lookup "mapScribes"  (R.options r)
        mapAggregatedKinds = HM.lookup "mapAggregatedkinds" (R.options r)
        mapScribe    = parseScribeMap mapscribes
    putMVar cgref $ ConfigurationInternal
        {cgMinSeverity = R.minSeverity r
        ,cgMapSeverity = parseSeverityMap mapseverity
        ,cgMapSubtrace = parseSubtraceMap mapsubtrace
        ,cgOptions = R.options r
        ,cgMapBackend = parseBackendMap mapbackends
        ,cgDefBackendKs = R.defaultBackends r
        ,cgSetupBackends = R.setupBackends r
        ,cgMapScribe = mapScribe
        ,cgMapScribeCache = mapScribe
        ,cgDefScribes = r_defaultScribes r
        ,cgSetupScribes = R.setupScribes r
        ,cgMapAggregatedKind = parseAggregatedKindMap mapAggregatedKinds
        ,cgDefAggregatedKind = StatsAK
        ,cgPortEKG = r_hasEKG r
        ,cgPortGUI = r_hasGUI r
        }
    return $ Configuration cgref
  where
    parseSeverityMap :: Maybe (HM.HashMap Text Value) → HM.HashMap Text Severity
    parseSeverityMap Nothing = HM.empty
    parseSeverityMap (Just hmv) = HM.mapMaybe mkSeverity hmv
    mkSeverity (String s) = Just (read (unpack s) :: Severity)
    mkSeverity _ = Nothing

    parseBackendMap Nothing = HM.empty
    parseBackendMap (Just hmv) = HM.map mkBackends hmv
    mkBackends (Array bes) = catMaybes $ map mkBackend $ Vector.toList bes
    mkBackends _ = [ ]
    mkBackend (String s) = Just (read (unpack s) :: BackendKind)
    mkBackend _ = Nothing
```

```
parseScribeMap Nothing = HM.empty
parseScribeMap (Just hmv) = HM.map mkScribes hmv
mkScribes (Array scs) = catMaybes $ map mkScribe $ Vector.toList scs
mkScribes (String s) = [(s :: ScribeId)]
mkScribes _ = [ ]
mkScribe (String s) = Just (s :: ScribeId)
mkScribe _ = Nothing

parseSubtraceMap :: Maybe (HM.HashMap Text Value) → HM.HashMap Text SubTrace
parseSubtraceMap Nothing = HM.empty
parseSubtraceMap (Just hmv) = HM.mapMaybe mkSubtrace hmv
mkSubtrace (String s) = Just (read (unpack s) :: SubTrace)
mkSubtrace (Object hm) = mkSubtrace' (HM.lookup "tag" hm) (HM.lookup "contents" hm)
mkSubtrace _ = Nothing
mkSubtrace' Nothing _ = Nothing
mkSubtrace' _ Nothing = Nothing
mkSubtrace' (Just (String tag)) (Just (Array cs)) =
    if tag ≡ "ObservableTrace"
    then Just $ ObservableTrace $ map (λ(String s) → (read (unpack s) :: ObservableInstance)) $ Vector.toLis
    else Nothing
mkSubtrace' _ _ = Nothing

r_hasEKG repr = case (R.hasEKG repr) of
    Nothing → 0
    Just p → p
r_hasGUI repr = case (R.hasGUI repr) of
    Nothing → 0
    Just p → p
r_defaultScribes repr = map (λ(k, n) → pack (show k) <> " :: " <> n) (R.defaultScribes repr)

parseAggregatedKindMap Nothing = HM.empty
parseAggregatedKindMap (Just hmv) =
    let
        listv = HM.toList hmv
        mapAggregatedKind = HM.fromList $ catMaybes $ map mkAggregatedKind listv
    in
        mapAggregatedKind
mkAggregatedKind (name, String s) = Just (name, read (unpack s) :: AggregatedKind)
mkAggregatedKind _ = Nothing
```

Setup empty configuration

```
empty :: IO Configuration
empty = do
    cgref ← newEmptyMVar
    putMVar cgref $ ConfigurationInternal Debug HM.empty HM.empty HM.empty HM.empty [ ] [ ] HM.empty H
    return $ Configuration cgref
```

1.4.22   Cardano.BM.Output.Switchboard

Switchboard

```
type SwitchboardMVar = MVar SwitchboardInternal
newtype Switchboard = Switchboard
  {getSB :: SwitchboardMVar}
data SwitchboardInternal = SwitchboardInternal
  {sbQueue :: TBQ.TBQueue NamedLogItem
  ,sbDispatch :: Async.Async ()
  }
```

Trace that forwards to the Switchboard

Every Trace ends in the Switchboard which then takes care of dispatching the messages to outputs

```
mainTrace :: Switchboard → TraceNamed IO
mainTrace sb = BaseTrace.BaseTrace $ Op $ λlognamed → do
  effectuate sb lognamed
```

Process incoming messages

Incoming messages are put into the queue, and then processed by the dispatcher.
The queue is initialized and the message dispatcher launched.

```
instance IsEffectuator Switchboard where
  effectuate switchboard item = do
    let writequeue :: TBQ.TBQueue NamedLogItem → NamedLogItem → IO ()
        writequeue q i = do
          nocapacity ← atomically $ TBQ.isFullTBQueue q
          if nocapacity
          then return ()
          else atomically $ TBQ.writeTBQueue q i
    withMVar (getSB switchboard) $ λsb →
      writequeue (sbQueue sb) item
```

Switchboard implements Backend functions

Switchboard is an Declaration of a Backend

```
instance IsBackend Switchboard where
  typeof _ = SwitchboardBK
  realize cfg =
    let spawnDispatcher
          :: Configuration
          → [(BackendKind, Backend)]
          → TBQ.TBQueue NamedLogItem
```

```
            → IO (Async.Async ())
      spawnDispatcher config backends queue =
         let sendMessage nli befilter = do
               selectedBackends ← getBackends config (lnName nli)
               let selBEs = befilter selectedBackends
               forM_ backends $ λ(bek, be) →
                  when (bek ∈ selBEs) (bEffectuate be $ nli)
            qProc = do
               nli ← atomically $ TBQ.readTBQueue queue
               case lnItem nli of
                  LogObject _ KillPill →
                     forM_ backends (λ(_, be) → bUnrealize be)
                  LogObject _ (AggregatedMessage _) → do
                     sendMessage nli (filter (≢ AggregationBK))
                     qProc
                  _ → sendMessage nli id ≫ qProc
         in
         Async.async qProc
      in do
      q ← atomically $ TBQ.newTBQueue 2048
      sbref ← newEmptyMVar
      putMVar sbref $ SwitchboardInternal q $ error "unitialized dispatcher"
      let sb :: Switchboard = Switchboard sbref

      backends ← getSetupBackends cfg
      bs ← setupBackends backends cfg sb [ ]
      dispatcher ← spawnDispatcher cfg bs q
      -- link the given Async to the current thread, such that if the Async
      -- raises an exception, that exception will be re-thrown in the current
      -- thread, wrapped in ExceptionInLinkedThread.
      Async.link dispatcher
      modifyMVar_ sbref $ λsbInternal → return $ sbInternal {sbDispatch = dispatcher}

      return sb

unrealize switchboard = do
   let clearMVar :: MVar a → IO ()
      clearMVar = void ∘ tryTakeMVar

   (dispatcher, queue) ← withMVar (getSB switchboard) (λsb → return (sbDispatch sb, sbQueue sb))
   -- send terminating item to the queue
   lo ← LogObject < $ > mkLOMeta < ∗ > pure KillPill
   atomically $ TBQ.writeTBQueue queue $ LogNamed "kill.switchboard" lo
   -- wait for the dispatcher to exit
   res ← Async.waitCatch dispatcher
   either throwM return res
   (clearMVar ∘ getSB) switchboard
```

Realizing the backends according to configuration

```
setupBackends :: [BackendKind]
          → Configuration
          → Switchboard
          → [(BackendKind, Backend)]
          → IO [(BackendKind, Backend)]
setupBackends [] _ _ acc = return acc
setupBackends (bk : bes) c sb acc = do
  be' ← setupBackend' bk c sb
  setupBackends bes c sb ((bk, be') : acc)
setupBackend' :: BackendKind → Configuration → Switchboard → IO Backend
setupBackend' SwitchboardBK _ _ = error "cannot instantiate a further Switchboard"
setupBackend' EKGViewBK c _ = do
  be :: Cardano.BM.Output ◦ EKGView.EKGView ← Cardano.BM.Output ◦ EKGView.realize c
  return MkBackend
    { bEffectuate = Cardano.BM.Output ◦ EKGView.effectuate be
    , bUnrealize = Cardano.BM.Output ◦ EKGView.unrealize be
    }
setupBackend' AggregationBK c sb = do
  let trace = mainTrace sb
    ctx = TraceContext { loggerName = ""
          , configuration = c
          , minSeverity = Debug
          , tracetype = Neutral
          , shutdown = pure ()
          }
  be :: Cardano.BM.Output ◦ Aggregation.Aggregation ← Cardano.BM.Output ◦ Aggregation.realizefrom (ctx,
  return MkBackend
    { bEffectuate = Cardano.BM.Output ◦ Aggregation.effectuate be
    , bUnrealize = Cardano.BM.Output ◦ Aggregation.unrealize be
    }
setupBackend' KatipBK c _ = do
  be :: Cardano.BM.Output ◦ Log.Log ← Cardano.BM.Output ◦ Log.realize c
  return MkBackend
    { bEffectuate = Cardano.BM.Output ◦ Log.effectuate be
    , bUnrealize = Cardano.BM.Output ◦ Log.unrealize be
    }
```

1.4.23 Cardano.BM.Output.Log

Internal representation

```
type LogMVar = MVar LogInternal
newtype Log = Log
  { getK :: LogMVar }
```

```
data LogInternal = LogInternal
  {kLogEnv :: K.LogEnv
  ,configuration :: Config.Configuration}
```

Log implements *effectuate*

```
instance IsEffectuator Log where
  effectuate katip item = do
    c ← withMVar (getK katip) $ λk → return (configuration k)
    selscribes ← getScribes c (lnName item)
    forM_ selscribes $ λsc → passN sc katip item
```

Log implements backend functions

```
instance IsBackend Log where
  typeof _ = KatipBK
  realize config = do
    let updateEnv :: K.LogEnv → IO UTCTime → K.LogEnv
        updateEnv le timer =
          le {K._logEnvTimer = timer, K._logEnvHost = "hostname"}
        register :: [ScribeDefinition] → K.LogEnv → IO K.LogEnv
        register [ ] le = return le
        register (defsc : dscs) le = do
          let kind = scKind defsc
              name = scName defsc
              name' = pack (show kind) <> "::" <> name
          scr ← createScribe kind name
          register dscs ≪ K.registerScribe name' scr scribeSettings le
        mockVersion :: Version
        mockVersion = Version [0, 1, 0, 0] [ ]
        scribeSettings :: KC.ScribeSettings
        scribeSettings =
          let bufferSize = 5000-- size of the queue (in log items)
          in
          KC.ScribeSettings bufferSize
        createScribe FileTextSK name = mkTextFileScribe (FileDescription $ unpack name) False
        createScribe FileJsonSK name = mkJsonFileScribe (FileDescription $ unpack name) False
        createScribe StdoutSK _ = mkStdoutScribe
        createScribe StderrSK _ = mkStderrScribe
      cfoKey ← Config.getOptionOrDefault config (pack "cfokey") (pack "<unknown>")
      le0 ← K.initLogEnv
            (K.Namespace ["iohk"])
            (fromString $ (unpack cfoKey) <> ":" <> showVersion mockVersion)
      -- request a new time 'getCurrentTime' at most 100 times a second
      timer ← mkAutoUpdate defaultUpdateSettings {updateAction = getCurrentTime, updateFreq = 10000}
```

```
        let le1 = updateEnv le0 timer
        scribes ← getSetupScribes config
        le ← register scribes le1

        kref ← newEmptyMVar
        putMVar kref $ LogInternal le config

        return $ Log kref
    unrealize katip = do
        le ← withMVar (getK katip) $ λk → return (kLogEnv k)
        void $ K.closeScribes le


    example :: IO ()
    example = do
        config ← Config.setup "from_some_path.yaml"
        k ← setup config
        passN (pack (show StdoutSK)) k $ LogNamed
          {lnName = "test"
          ,lnItem = LogMessage $ LogItem
              {liSelection = Both
              ,liSeverity = Info
              ,liPayload = "Hello!"
              }
          }
        passN (pack (show StdoutSK)) k $ LogNamed
          {lnName = "test"
          ,lnItem = LogValue "cpu-no" 1
          }
```

Needed instances for *katip*:

```
    deriving instance K.ToObject LogObject
    deriving instance K.ToObject LogItem
    deriving instance K.ToObject (Maybe LOContent)

    instance KC.LogItem LogObject where
        payloadKeys _ _ = KC.AllKeys
    instance KC.LogItem LogItem where
        payloadKeys _ _ = KC.AllKeys
    instance KC.LogItem (Maybe LOContent) where
        payloadKeys _ _ = KC.AllKeys
```

Log.passN

The following function copies the NamedLogItem to the queues of all scribes that match on their name. Compare start of name of scribe to (*show backend* <> "::"). This function is non-blocking.

```
    passN :: Text → Log → NamedLogItem → IO ()
    passN backend katip namedLogItem = do
        env ← withMVar (getK katip) $ λk → return (kLogEnv k)
```

```
forM_ (Map.toList $ K._logEnvScribes env) $
  λ(scName, (KC.ScribeHandle _ shChan)) →
    -- check start of name to match ScribeKind
      if backend 'isPrefixOf' scName
      then do
        let (LogObject lometa loitem) = lnItem namedLogItem
        let (sev, msg, payload) = case loitem of
            (LogMessage logItem) →
              (liSeverity logItem, liPayload logItem, Nothing)
            (ObserveDiff _) →
              let text = toStrict (encodeToLazyText loitem)
              in
              (Info, text, Just loitem)
            (ObserveOpen _) →
              let text = toStrict (encodeToLazyText loitem)
              in
              (Info, text, Just loitem)
            (ObserveClose _) →
              let text = toStrict (encodeToLazyText loitem)
              in
              (Info, text, Just loitem)
            (AggregatedMessage aggregated) →
              let text = T.concat $ (flip map) aggregated $ λ(name, agg) →
                "\n" <> name <> ": " <> pack (show agg)
              in
              (Info, text, Nothing)
            (LogValue name value) →
              (Debug, name <> " = " <> pack (showSI value), Nothing)
            KillPill →
              (Info, "Kill pill received!", Nothing)
        if (msg ≡ "") ∧ (isNothing payload)
        then return ()
        else do
          let threadIdText = KC.mkThreadIdText (tid lometa)
          let ns = lnName namedLogItem
          let itemTime = tstamp lometa
          let itemKatip = K.Item {
              _itemApp      = env ˆ. KC.logEnvApp
            , _itemEnv      = env ˆ. KC.logEnvEnv
            , _itemSeverity = sev2klog sev
            , _itemThread   = threadIdText
            , _itemHost     = env ˆ. KC.logEnvHost
            , _itemProcess  = env ˆ. KC.logEnvPid
            , _itemPayload  = payload
            , _itemMessage  = K.logStr msg
            , _itemTime     = itemTime
            , _itemNamespace = (env ˆ. KC.logEnvApp) <> (K.Namespace [ns])
            , _itemLoc      = Nothing
```

```
            }
        void $ atomically $ KC.tryWriteTBQueue shChan (KC.NewItem itemKatip)
    else return ()
```

Scribes

```
mkStdoutScribe :: IO K.Scribe
mkStdoutScribe = mkTextFileScribeH stdout True

mkStderrScribe :: IO K.Scribe
mkStderrScribe = mkTextFileScribeH stderr True

mkTextFileScribeH :: Handle → Bool → IO K.Scribe
mkTextFileScribeH handler color = do
    mkFileScribeH handler formatter color
  where
    formatter h colorize verbosity item =
        TIO.hPutStrLn h $! toLazyText $ formatItem colorize verbosity item

mkFileScribeH
    :: Handle
    → (forall a ∘ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ())
    → Bool
    → IO K.Scribe
mkFileScribeH h formatter colorize = do
    hSetBuffering h LineBuffering
    locklocal ← newMVar ()
    let logger :: forall a ∘ K.LogItem a ⇒ K.Item a → IO ()
        logger item = withMVar locklocal $ \_ →
            formatter h colorize K.V0 item
    pure $ K.Scribe logger (hClose h)

mkTextFileScribe :: FileDescription → Bool → IO K.Scribe
mkTextFileScribe fdesc colorize = do
    mkFileScribe fdesc formatter colorize
  where
    formatter :: Handle → Bool → K.Verbosity → K.Item a → IO ()
    formatter hdl colorize' v' item =
        case KC._itemMessage item of
          K.LogStr "" →
            -- if message is empty do not output it
            return ()
          _ → do
            let tmsg = toLazyText $ formatItem colorize' v' item
            TIO.hPutStrLn hdl tmsg

mkJsonFileScribe :: FileDescription → Bool → IO K.Scribe
mkJsonFileScribe fdesc colorize = do
    mkFileScribe fdesc formatter colorize
  where
    formatter :: (K.LogItem a) ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ()
```

```
    formatter h _ verbosity item = do
        let tmsg = case KC._itemMessage item of
            -- if a message is contained in item then only the
            -- message is printed and not the data
            K.LogStr "" → K.itemJson verbosity item
            K.LogStr msg → K.itemJson verbosity $
                item {KC._itemMessage = K.logStr ("" :: Text)
                    ,KC._itemPayload = LogItem Both Info $ toStrict $ toLazyText msg
                    }
        TIO.hPutStrLn h (encodeToLazyText tmsg)
mkFileScribe
    :: FileDescription
    → (forall a ∘ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ())
    → Bool
    → IO K.Scribe
mkFileScribe fdesc formatter colorize = do
    let prefixDir = prefixPath fdesc
    (createDirectoryIfMissing True prefixDir)
        `catchIO` (prtoutException ("cannot log prefix directory: " ++ prefixDir))
    let fpath = filePath fdesc
    h ← catchIO (openFile fpath WriteMode) $
        λe → do
            prtoutException ("error while opening log: " ++ fpath) e
            -- fallback to standard output in case of exception
            return stdout
    hSetBuffering h LineBuffering
    scribestate ← newMVar h
    let finalizer :: IO ()
        finalizer = withMVar scribestate hClose
    let logger :: forall a ∘ K.LogItem a ⇒ K.Item a → IO ()
        logger item =
            withMVar scribestate $ λhandler →
                formatter handler colorize K.V0 item
    return $ K.Scribe logger finalizer


formatItem :: Bool → K.Verbosity → K.Item a → Builder
formatItem withColor _verb K.Item {..} =
    fromText header <>
    fromText "  " <>
    brackets (fromText timestamp) <>
    fromText "  " <>
    KC.unLogStr _itemMessage
  where
    header = colorBySeverity _itemSeverity $
        "[" <> mconcat namedcontext <> ":" <> severity <> ":" <> threadid <> "]"
    namedcontext = KC.intercalateNs _itemNamespace
    severity = KC.renderSeverity _itemSeverity
```

```
      threadid = KC.getThreadIdText _itemThread
      timestamp = pack $ formatTime defaultTimeLocale tsformat _itemTime
      tsformat :: String
      tsformat = "%F %T%2Q %Z"
      colorBySeverity s m = case s of
        K.EmergencyS → red m
        K.AlertS     → red m
        K.CriticalS  → red m
        K.ErrorS     → red m
        K.NoticeS    → magenta m
        K.WarningS → yellow m
        K.InfoS      → blue m
        _ → m
      red = colorize "31"
      yellow = colorize "33"
      magenta = colorize "35"
      blue = colorize "34"
      colorize c m
        | withColor = "\ESC[" <> c <> "m" <> m <> "\ESC[0m"
        | otherwise = m
-- translate Severity to Log.Severity
sev2klog :: Severity → K.Severity
sev2klog = λcase
    Debug     → K.DebugS
    Info      → K.InfoS
    Notice    → K.NoticeS
    Warning → K.WarningS
    Error     → K.ErrorS
    Critical  → K.CriticalS
    Alert     → K.AlertS
    Emergency → K.EmergencyS



data FileDescription = FileDescription {
  filePath :: !FilePath}
  deriving (Show)
prefixPath :: FileDescription → FilePath
prefixPath = takeDirectory ∘ filePath
-- display message and stack trace of exception on stdout
prtoutException :: Exception e ⇒ String → e → IO ()
prtoutException msg e = do
  putStrLn msg
  putStrLn ("exception: " ++ displayException e)
```

1.4.24 Cardano.BM.Output.EKGView

Structure of EKGView

```
type EKGViewMVar = MVar EKGViewInternal
newtype EKGView = EKGView
    {getEV :: EKGViewMVar}

data EKGViewInternal = EKGViewInternal
    {evQueue :: TBQ.TBQueue (Maybe NamedLogItem)
    ,evLabels :: EKGViewMap
    ,evServer :: Server
    }
```

Relation from variable name to label handler

We keep the label handlers for later update in a *HashMap*.

```
type EKGViewMap = HM.HashMap Text Label.Label
```

Internal Trace

This is an internal Trace, named "#ekgview", which can be used to control the messages that are being displayed by EKG.

```
ekgTrace :: EKGView → Configuration → IO (Trace IO)
ekgTrace ekg c = do
    let trace = ekgTrace' ekg
        ctx = TraceContext {loggerName = ""
                ,configuration = c
                ,minSeverity = Debug
                ,tracetype = Neutral
                ,shutdown = pure ()
                }
    Trace.subTrace "#ekgview" (ctx,trace)
  where
    ekgTrace' :: EKGView → TraceNamed IO
    ekgTrace' ekgview = BaseTrace.BaseTrace $ Op $ λ(LogNamed lognamed lo) → do
        let setlabel :: Text → Text → EKGViewInternal → IO (Maybe EKGViewInternal)
            setlabel name label ekg_i@(EKGViewInternal _ labels server) =
                case HM.lookup name labels of
                    Nothing → do
                        ekghdl ← getLabel name server
                        Label.set ekghdl label
                        return $ Just $ ekg_i {evLabels = HM.insert name ekghdl labels}
                    Just ekghdl → do
                        Label.set ekghdl label
                        return Nothing
```

```haskell
        update :: LogObject → LoggerName → EKGViewInternal → IO (Maybe EKGViewInternal)
        update (LogObject _ (LogMessage logitem)) logname ekg_i =
            setlabel logname (liPayload logitem) ekg_i
        update (LogObject _ (LogValue iname value)) logname ekg_i =
          let logname′ = logname <> "." <> iname
          in
          setlabel logname′ (pack $ show value) ekg_i

        update _ _ _ = return Nothing
    ekgup ← takeMVar (getEV ekgview)
    let -- strip off some prefixes not necessary for display
      lognam1 = case stripPrefix "#ekgview.#aggregation." lognamed of
        Nothing → lognamed
        Just ln′ → ln′
      logname = case stripPrefix "#ekgview." lognam1 of
        Nothing → lognam1
        Just ln′ → ln′
    upd ← update lo logname ekgup
    case upd of
      Nothing → putMVar (getEV ekgview) ekgup
      Just ekgup′ → putMVar (getEV ekgview) ekgup′
```

EKG view is an effectuator

Function *effectuate* is called to pass in a NamedLogItem for display in EKG. If the log item is an *AggregatedStats* message, then all its constituents are put into the queue.

```haskell
    instance IsEffectuator EKGView where
      effectuate ekgview item = do
        ekg ← readMVar (getEV ekgview)
        let queue a = do
                nocapacity ← atomically $ TBQ.isFullTBQueue (evQueue ekg)
                if nocapacity
                then return ()
                else atomically $ TBQ.writeTBQueue (evQueue ekg) (Just a)
        case (lnItem item) of
          (LogObject lometa (AggregatedMessage ags)) → liftIO $ do
            let logname = lnName item
              traceAgg :: [(Text, Aggregated)] → IO ()
              traceAgg [ ] = return ()
              traceAgg ((n, AggregatedEWMA ewma) : r) = do
                queue $ LogNamed (logname <> "." <> n) $ LogObject lometa (LogValue "avg" $ avg ewma)
                traceAgg r
              traceAgg ((n, AggregatedStats stats) : r) = do
                let statsname = logname <> "." <> n
                  qbasestats s′ nm = do
                    queue $ LogNamed nm $ LogObject lometa (LogValue "mean" (PureD $ meanOfStats s′))
                    queue $ LogNamed nm $ LogObject lometa (LogValue "min" $ fmin s′)
```

queue $ LogNamed *nm* $ LogObject *lometa* (*LogValue* `"max"` $ *fmax s'*)
queue $ LogNamed *nm* $ LogObject *lometa* (*LogValue* `"count"` $ *PureI* $ *fromIntegral* $ *fcoun*
queue $ LogNamed *nm* $ LogObject *lometa* (*LogValue* `"stdev"` (*PureD* $ *stdevOfStats s'*))
queue $ LogNamed *statsname* $ LogObject *lometa* (*LogValue* `"last"` $ *flast stats*)
qbasestats (*fbasic stats*) $ *statsname* <> `".basic"`
qbasestats (*fdelta stats*) $ *statsname* <> `".delta"`
qbasestats (*ftimed stats*) $ *statsname* <> `".timed"`
traceAgg r
traceAgg ags
(LogObject _ (*LogMessage* _)) → *queue item*
(LogObject _ (*LogValue* _ _)) → *queue item*
_ → *return* ()

EKGView implements Backend functions

EKGView is an IsBackend

```
instance IsBackend EKGView where
  typeof _ = EKGViewBK
  realize config = do
    evref ← newEmptyMVar
    let ekgview = EKGView evref
    evport ← getEKGport config
    ehdl ← forkServer "127.0.0.1" evport
    ekghdl ← getLabel "iohk-monitoring version" ehdl
    Label.set ekghdl $ pack (showVersion version)
    ekgtrace ← ekgTrace ekgview config
    queue ← atomically $ TBQ.newTBQueue 512
    dispatcher ← spawnDispatcher queue ekgtrace
    -- link the given Async to the current thread, such that if the Async
    -- raises an exception, that exception will be re-thrown in the current
    -- thread, wrapped in ExceptionInLinkedThread.
    Async.link dispatcher
    putMVar evref $ EKGViewInternal
      {evLabels = HM.empty
      ,evServer = ehdl
      ,evQueue = queue
      }
    return ekgview
  unrealize ekgview = do
    ekg ← takeMVar $ getEV ekgview
    killThread $ serverThreadId $ evServer ekg
```

Asynchrouniously reading log items from the queue and their processing

spawnDispatcher :: *TBQ.TBQueue* (*Maybe* NamedLogItem)
          → Trace.Trace *IO*

```
                → IO (Async.Async ())
    spawnDispatcher evqueue trace =
        Async.async $ qProc
      where
        qProc = do
          maybeItem ← atomically $ TBQ.readTBQueue evqueue
          case maybeItem of
            Just (LogNamed logname logvalue) → do
                trace' ← Trace.appendName logname trace
                Trace.traceNamedObject trace' logvalue
                qProc
            Nothing → return ()-- stop here
```

Interactive testing EKGView

```
    test :: IO ()
    test = do
      c ← Cardano.BM.Setup.setupTrace (Left "test/config.yaml") "ekg"
      ev ← Cardano.BM.Output ∘ EKGView.realize c

      effectuate ev $ LogNamed "test.questions" (LogValue "answer" 42)
      effectuate ev $ LogNamed "test.monitor023" (LogMessage (LogItem Public Warning "!!!! ALARM !!!!"
```

1.4.25   Cardano.BM.Output.Aggregation

Internal representation

```
    type AggregationMVar = MVar AggregationInternal
    newtype Aggregation = Aggregation
      {getAg :: AggregationMVar}
    data AggregationInternal = AggregationInternal
      {agQueue :: TBQ.TBQueue (Maybe NamedLogItem)
      , agDispatch :: Async.Async ()
      }
```

Relation from context name to aggregated statistics

We keep the aggregated values (Aggregated) for a named context in a *HashMap*.

```
    type AggregationMap = HM.HashMap Text AggregatedExpanded
```

Info for Aggregated operations

Apart from the Aggregated we keep some valuable info regarding to them; such as when was the last time it was sent.

type *Timestamp = Word64*

data AggregatedExpanded = AggregatedExpanded
  {*aeAggregated* :: !Aggregated
  , *aeResetAfter* :: !(*Maybe Word64*)
  , *aeLastSent* :: {-# UNPACK #-} ! *Timestamp*
  }

Aggregation implements *effectuate*

Aggregation is an Accepts a NamedLogItem Enter the log item into the Aggregation queue.

instance IsEffectuator Aggregation **where**
  effectuate agg item = **do**
    ag ← *readMVar* (*getAg agg*)
    nocapacity ← *atomically* $ *TBQ.isFullTBQueue* (*agQueue ag*)
    if *nocapacity*
    then *return* ()
    else *atomically* $! *TBQ.writeTBQueue* (*agQueue ag*) $ *Just item*

Aggregation implements Backend **functions**

Aggregation is an Declaration of a Backend

instance IsBackend Aggregation **where**
  typeof _ = *AggregationBK*
  realize _ = *error* `"Aggregation cannot be instantiated by 'realize'"`
  realizefrom trace0@(*ctx*, _) _ = **do**
    trace ← Trace.subTrace `"#aggregation"` *trace0*
    aggref ← *newEmptyMVar*
    aggregationQueue ← *atomically* $ *TBQ.newTBQueue* 2048
    dispatcher ← *spawnDispatcher* (configuration *ctx*) *HM.empty aggregationQueue trace*
    `-- link the given Async to the current thread, such that if the Async`
    `-- raises an exception, that exception will be re-thrown in the current`
    `-- thread, wrapped in ExceptionInLinkedThread.`
    Async.link dispatcher
    putMVar aggref $ *AggregationInternal aggregationQueue dispatcher*
    return $ Aggregation *aggref*
  unrealize aggregation = **do**
    let *clearMVar* :: *MVar a* → *IO* ()
      clearMVar = *void* ∘ *tryTakeMVar*
    (*dispatcher*, *queue*) ← *withMVar* (*getAg aggregation*) (λ*ag* →
      return (*agDispatch ag*, *agQueue ag*))
    `-- send terminating item to the queue`
    atomically $ *TBQ.writeTBQueue queue Nothing*
    `-- wait for the dispatcher to exit`
    res ← *Async.waitCatch dispatcher*

either throwM return res
(clearMVar ∘ getAg) aggregation

Asynchrouniously reading log items from the queue and their processing

spawnDispatcher :: Configuration
        → *AggregationMap*
        → *TBQ.TBQueue* (*Maybe* NamedLogItem)
        → Trace.Trace *IO*
        → *IO* (*Async.Async* ())
spawnDispatcher conf aggMap aggregationQueue trace = Async.async $ *qProc aggMap*
  where
    qProc aggregatedMap = **do**
      maybeItem ← atomically $ *TBQ.readTBQueue aggregationQueue*
      case *maybeItem* **of**
        Just (LogNamed *logname lo*@(LogObject *lm* _)) → **do**
          (*updatedMap*, *aggregations*) ← *update lo logname aggregatedMap*
          unless (*null aggregations*) $
            sendAggregated (LogObject *lm* (*AggregatedMessage aggregations*)) *logname*
          qProc updatedMap
        Nothing → *return* ()
    createNupdate name value lme agmap = **do**
      case *HM.lookup name agmap* **of**
        Nothing → **do**
          `-- if Aggregated does not exist; initialize it.`
          aggregatedKind ← getAggregatedKind conf name
          case *aggregatedKind* **of**
            StatsAK → *return* $ singletonStats *value*
            EwmaAK aEWMA → **do**
              let *initEWMA = EmptyEWMA aEWMA*
              return $ *AggregatedEWMA* $ ewma *initEWMA value*
        Just a → *return* $ updateAggregation *value* (*aeAggregated a*) *lme* (*aeResetAfter a*)
    update :: LogObject
      → LoggerName
      → *AggregationMap*
      → *IO* (*AggregationMap*, [(*Text*, Aggregated)])
    update (LogObject *lme* (*LogValue iname value*)) *logname agmap* = **do**
      let *fullname = logname* <> " . " <> *iname*
      aggregated ← createNupdate fullname value lme agmap
      now ← getMonotonicTimeNSec
      let *aggregatedX =* AggregatedExpanded {
        aeAggregated = aggregated
        , *aeResetAfter = Nothing*
        , *aeLastSent = now*
        }
        namedAggregated = [(*iname*, *aeAggregated aggregatedX*)]

```haskell
    updatedMap = HM.alter (const $ Just $ aggregatedX) fullname agmap
  return (updatedMap, namedAggregated)
update (LogObject lme (ObserveDiff counterState)) logname agmap =
  updateCounters (csCounters counterState) lme (logname, "diff") agmap []
update (LogObject lme (ObserveOpen counterState)) logname agmap =
  updateCounters (csCounters counterState) lme (logname, "open") agmap []
update (LogObject lme (ObserveClose counterState)) logname agmap =
  updateCounters (csCounters counterState) lme (logname, "close") agmap []

update (LogObject lme (LogMessage msg)) logname agmap = do
  let iname = T.pack $ show (liSeverity msg)
  let fullname = logname <> "." <> iname
  aggregated ← createNupdate fullname (PureI 0) lme agmap
  now ← getMonotonicTimeNSec
  let aggregatedX = AggregatedExpanded {
    aeAggregated = aggregated
    ,aeResetAfter = Nothing
    ,aeLastSent = now
    }
    namedAggregated = [(iname, aeAggregated aggregatedX)]
    updatedMap = HM.alter (const $ Just $ aggregatedX) fullname agmap
  return (updatedMap, namedAggregated)

-- everything else
update _ _ agmap = return (agmap, [])

updateCounters :: [Counter]
               → LOMeta
               → (LoggerName, LoggerName)
               → AggregationMap
               → [(Text, Aggregated)]
               → IO (AggregationMap, [(Text, Aggregated)])
updateCounters [] _ _ aggrMap aggs = return $ (aggrMap, aggs)
updateCounters (counter : cs) lme (logname, msgname) aggrMap aggs = do
  let name = cName counter
    subname = msgname <> "." <> (nameCounter counter) <> "." <> name
    fullname = logname <> "." <> subname
    value = cValue counter
  aggregated ← createNupdate fullname value lme aggrMap
  now ← getMonotonicTimeNSec
  let aggregatedX = AggregatedExpanded {
    aeAggregated = aggregated
    ,aeResetAfter = Nothing
    ,aeLastSent = now
    }
    namedAggregated = (subname, aggregated)
    updatedMap = HM.alter (const $ Just $ aggregatedX) fullname aggrMap
  updateCounters cs lme (logname, msgname) updatedMap (namedAggregated : aggs)

sendAggregated :: LogObject → Text → IO ()
```

sendAggregated aggregatedMsg@(**LogObject** _ (*AggregatedMessage* _)) *logname* = **do**
    `-- enter the aggregated message into the `**`Trace`**
    trace' ← **Trace.appendName** *logname trace*
    liftIO \$ **Trace.traceNamedObject** *trace' aggregatedMsg*
  `-- ingnore every other message`
sendAggregated _ _ = *return* ()

Update aggregation

We distinguish an unitialized from an already initialized aggregation. The latter is properly initialized.

We use Welford's online algorithm to update the estimation of mean and variance of the sample statistics. (see `https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance#Welford's_Online`

updateAggregation :: **Measurable** → **Aggregated** → *LOMeta* → *Maybe Word64* → **Aggregated**
updateAggregation *v* (*AggregatedStats s*) *lme resetAfter* =
    let *count* = *fcount* (*fbasic s*)
       reset = *maybe False* (*count* ⩾) *resetAfter*
    in
    if *reset*
    then
      singletonStats *v*
    else
      AggregatedStats \$! **Stats** {*flast* = *v*
        , *fold* = *mkTimestamp*
        , *fbasic* = *updateBaseStats* (*count* ⩾ 1) *v* (*fbasic s*)
        , *fdelta* = *updateBaseStats* (*count* ⩾ 2) (*v* − *flast s*) (*fdelta s*)
        , *ftimed* = *updateBaseStats* (*count* ⩾ 2) (*mkTimestamp* − *fold s*) (*ftimed s*)
        }
  where
    mkTimestamp = *utc2ns* (*tstamp lme*)
    utc2ns (*UTCTime days secs*) =
      let *yearsecs* :: *Rational*
        yearsecs = 365 ∗ 24 ∗ 3600
        rdays, *rsecs* :: *Rational*
        rdays = *toRational* \$ *toModifiedJulianDay days*
        rsecs = *toRational secs*
        s2ns = 1000000000
      in
      Nanoseconds \$ *round* \$ (*fromRational* \$ *s2ns* ∗ *rsecs* + *rdays* ∗ *yearsecs* :: *Double*)
updateAggregation *v* (*AggregatedEWMA e*) _ _ = *AggregatedEWMA* \$! **ewma** *e v*
updateBaseStats :: *Bool* → **Measurable** → *BaseStats* → *BaseStats*
updateBaseStats False _ *s* = *s* {*fcount* = *fcount s* + 1}
updateBaseStats True v s =
    let *newcount* = *fcount s* + 1
      newvalue = *getDouble v*
      delta = *newvalue* − *fsum_A s*

```
      dincr = (delta / fromIntegral newcount)
      delta2 = newvalue − fsum_A s − dincr
   in
   BaseStats {fmin = min (fmin s) v
      ,fmax         = max v (fmax s)
      ,fcount = newcount
      ,fsum_A = fsum_A s + dincr
      ,fsum_B = fsum_B s + (delta ∗ delta2)
      }
```

Calculation of EWMA

Following https://en.wikipedia.org/wiki/Moving_average#Exponential_moving_average we
calculate the exponential moving average for a series of values Y_t according to:

$$S_t = \begin{cases} Y_1, & t = 1 \\ \alpha \cdot Y_t + (1 - \alpha) \cdot S_{t-1}, & t > 1 \end{cases}$$

The pattern matching below ensures that the EWMA will start with the first value passed in, and
will not change type, once determined.

```
ewma :: EWMA → Measurable → EWMA
ewma (EmptyEWMA a) v = EWMA a v
ewma (EWMA a s@(Microseconds _)) y@(Microseconds _) =
   EWMA a $ Microseconds $ round $ a ∗ (getDouble y) + (1 − a) ∗ (getDouble s)
ewma (EWMA a s@(Seconds _)) y@(Seconds _) =
   EWMA a $ Seconds $ round $ a ∗ (getDouble y) + (1 − a) ∗ (getDouble s)
ewma (EWMA a s@(Bytes _)) y@(Bytes _) =
   EWMA a $ Bytes $ round $ a ∗ (getDouble y) + (1 − a) ∗ (getDouble s)
ewma (EWMA a (PureI s)) (PureI y) =
   EWMA a $ PureI $ round $ a ∗ (fromInteger y) + (1 − a) ∗ (fromInteger s)
ewma (EWMA a (PureD s)) (PureD y) =
   EWMA a $ PureD $ a ∗ y + (1 − a) ∗ s
ewma _ _ = error "Cannot average on values of different type"
```

Index