Cardano.BM - benchmarking and logging

Alexander Diemand   Andreas Triantafyllos

November 2018

Abstract

This is a framework that combines logging, benchmarking and monitoring. Complex evaluations of STM or monadic actions can be observed from outside while reading operating system counters before and after, and calculating their differences, thus relating resource usage to such actions. Through interactive configuration, the runtime behaviour of logging or the measurement of resource usage can be altered. Further reduction in logging can be achieved by redirecting log messages to an aggregation function which will output the running statistics with less frequency than the original message.

Contents

Chapter 1

Cardano BM

1.1 Overview

In figure 1.1 we display the relationships among modules in *Cardano.BM*. The arrows indicate import of a module. The arrows with a triangle at one end would signify "inheritance", but we use it to show that one module replaces the other in the namespace, thus refines its interface.

1.2 Introduction

1.2.1 Logging with *Trace*

1.2.2 Measuring *Observable*s

1.2.3 Monitoring

1.2.4 Information reduction in *Aggregation*

1.2.5 Output selection

1.2.6 Setup procedure

1.3 Examples

1.3.1 Observing evaluation of a STM action

1.3.2 Observing evaluation of a monad action

1.4 Code listings

1.4.1 Cardano.BM.Observer.STM

$$stmWithLog :: STM.STM\,(t, [LogObject]) \rightarrow STM.STM\,(t, [LogObject])$$
$$stmWithLog\ action = action$$
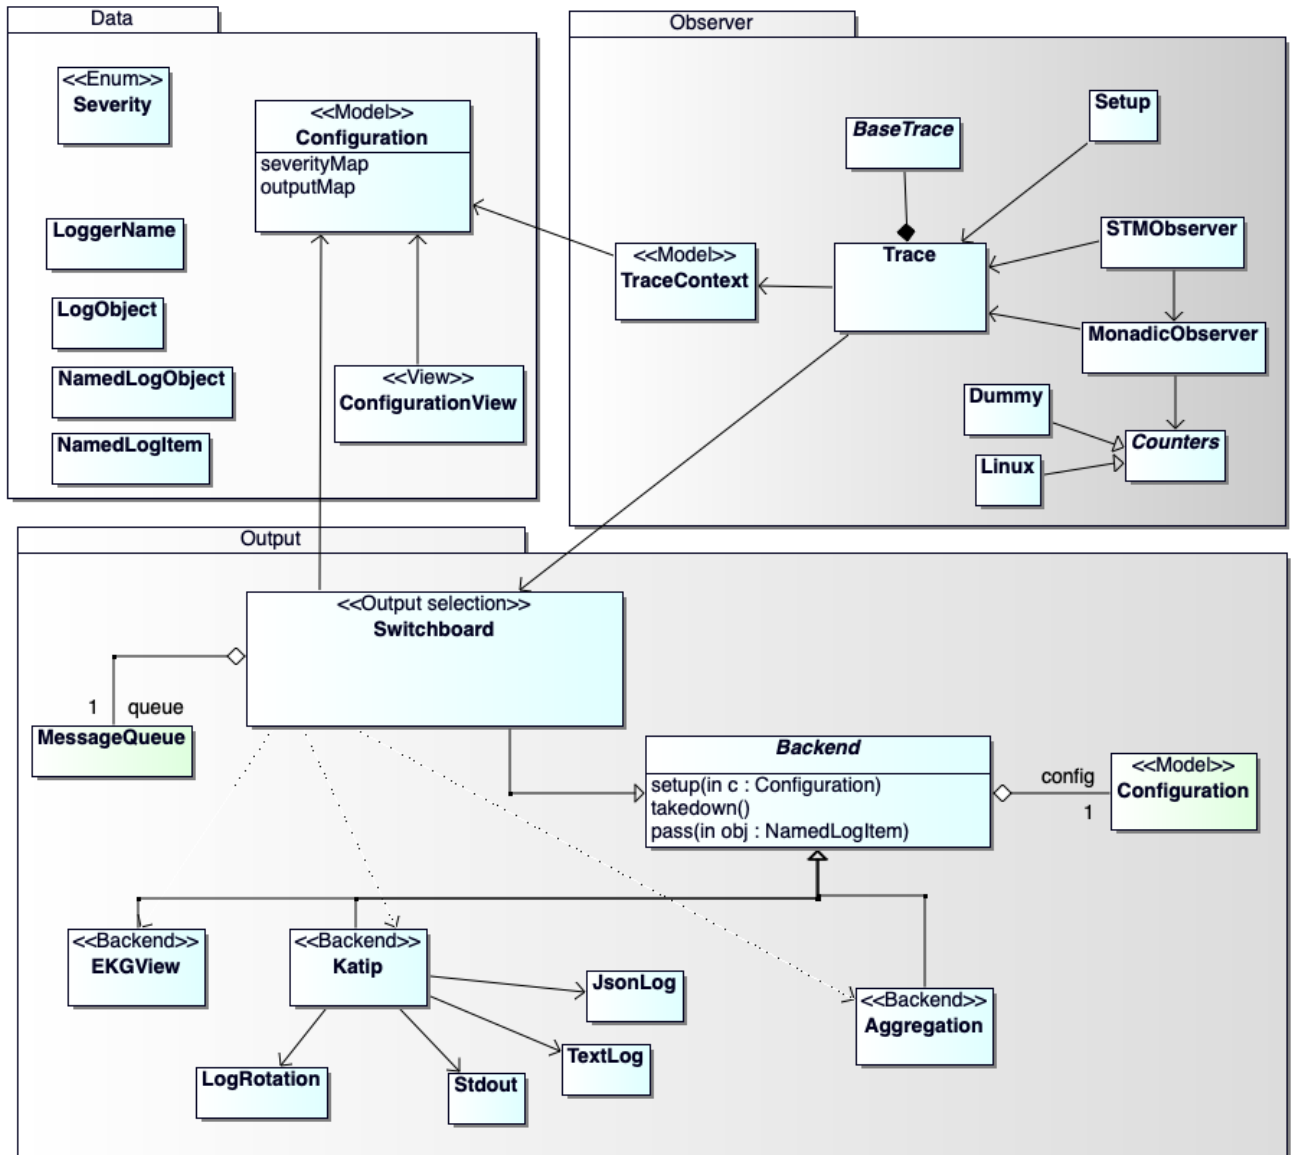
Figure 1.1: Overview of module relationships
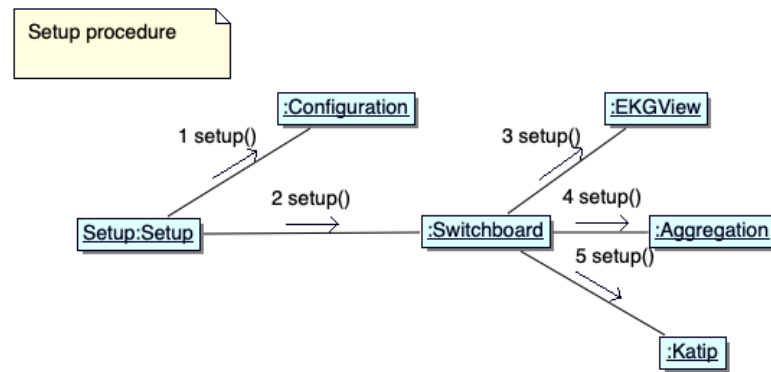
Figure 1.2: Setup procedure

Observe *STM* action in a named context

With given name, create a *SubTrace* according to *Configuration* and run the passed *STM* action on it.

> *bracketObserveIO* :: *Trace IO* → *Text* → *STM.STM t* → *IO t*
> *bracketObserveIO logTrace0 name action* = **do**
>   *logTrace* ← *subTrace name logTrace0*
>   **let** *subtrace* = *typeofTrace logTrace*
>   *bracketObserveIO′ subtrace logTrace action*
>   **where**
>   *bracketObserveIO′* :: *SubTrace* → *Trace IO* → *STM.STM t* → *IO t*
>   *bracketObserveIO′ NoTrace _ act* =
>   *STM.atomically act*
>   *bracketObserveIO′ subtrace logTrace act* = **do**
>   *countersid* ← *observeOpen subtrace logTrace*
>   -- run action, returns result only
>   *t* ← *STM.atomically act*
>   *observeClose subtrace logTrace countersid* [ ]
>   *pure t*

Observe *STM* action in a named context and output captured log items

The *STM* action might output messages, which after "success" will be forwarded to the logging trace. Otherwise, this function behaves the same as Observe *STM* action in a named context.

> *bracketObserveLogIO* :: *Trace IO* → *Text* → *STM.STM (t,* [*LogObject*]) → *IO t*
> *bracketObserveLogIO logTrace0 name action* = **do**
>   *logTrace* ← *subTrace name logTrace0*
>   **let** *subtrace* = *typeofTrace logTrace*
>   *bracketObserveLogIO′ subtrace logTrace action*
>   **where**

bracketObserveLogIO′ :: SubTrace → Trace IO → STM.STM (t, [LogObject]) → IO t
bracketObserveLogIO′ NoTrace _ act = **do**
  (*t*, _) ← *STM.atomically* $ *stmWithLog act*
  pure t
bracketObserveLogIO′ subtrace logTrace act = **do**
  countersid ← observeOpen subtrace logTrace
  `-- run action, return result and log items`
  (*t*, *as*) ← *STM.atomically* $ *stmWithLog act*
  observeClose subtrace logTrace countersid as
  pure t

1.4.2   Cardano.BM.Observer.Monadic

Monadic.bracketObserverIO

Observes an *IO* action and adds a name to the logger name of the passed in *Trace*. An empty *Text* leaves the logger name untouched.

bracketObserveIO :: Trace IO → Text → IO t → IO t
bracketObserveIO logTrace0 name action = **do**
  logTrace ← subTrace name logTrace0
  bracketObserveIO′ (typeofTrace logTrace) logTrace action
  where
  bracketObserveIO′ :: SubTrace → Trace IO → IO t → IO t
  bracketObserveIO′ NoTrace _ act = act
  bracketObserveIO′ subtrace logTrace act = **do**
    countersid ← observeOpen subtrace logTrace
    `-- run action`
    t ← act
    observeClose subtrace logTrace countersid [ ]
    pure t

observerOpen

observeOpen :: SubTrace → Trace IO → IO CounterState
observeOpen subtrace logTrace = **do**
  identifier ← newUnique
  logInfo logTrace $ `"Opening: "` <> *pack* (*show* $ *hashUnique identifier*)
  `-- take measurement`
  counters ← readCounters subtrace
  let *state = CounterState identifier counters*
  `-- send opening message to Trace`
  traceNamedObject logTrace $ *ObserveOpen state*
  return state

observeClose

observeClose :: *SubTrace* → *Trace IO* → *CounterState* → [*LogObject*] → *IO* ()
observeClose subtrace logTrace initState logObjects = **do**
  let *identifier* = *csIdentifier initState*
    initialCounters = *csCounters initState*

  logInfo logTrace $ `"Closing: "` <> *pack* (*show* $ *hashUnique identifier*)
  `-- take measurement`
  counters ← *readCounters subtrace*
  `-- send closing message to Trace`
  traceNamedObject logTrace $ *ObserveClose* (*CounterState identifier counters*)
  `-- send diff message to Trace`
  traceNamedObject logTrace $
    ObserveDiff (*CounterState identifier* (*diffCounters initialCounters counters*))
  `-- trace the messages gathered from inside the action`
  forM_ logObjects $ *traceNamedObject logTrace*

1.4.3 BaseTrace

Contravariant

A covariant is a functor: F A → F B
A contravariant is a functor: F B → F A

  Op a b implements the inverse to 'arrow' "*getOp* :: *b* → *a*", which when applied to a *BaseTrace* of type "*Op* (*m* ()) *s*", yields "*s* → *m* ()". In our case, *Op* accepts an action in a monad *m* with input type *LogNamed LogObject* (see 'Trace').

  newtype *BaseTrace m s* = *BaseTrace* {*runTrace* :: *Op* (*m* ()) *s*}

contramap

A covariant functor defines the function "*fmap* :: (*a* → *b*) → *f a* → *f b*". In case of a contravariant functor, it is the dual function "*contramap* :: (*a* → *b*) → *f b* → *f a*" which is defined.
  In the following instance, *runTrace* extracts type "*Op* (*m* ()) *s*" to which contramap applies *f*, thus "*f s* → *m* ()". The constructor *BaseTrace* restores "*Op* (*m* ()) (*f s*)".

  instance *Contravariant* (*BaseTrace m*) **where**
    contramap f = *BaseTrace* ∘ *contramap f* ∘ *runTrace*

traceWith

Accepts a *Trace* and some payload *s*. First it gets the contravariant from the *Trace* as type "*Op* (*m* ()) *s*" and, after "*getOp* :: *b* → *a*" which translates to "*s* → *m* ()", calls the action on the *LogNamed LogObject*.

  traceWith :: *BaseTrace m s* → *s* → *m* ()
  traceWith = *getOp* ∘ *runTrace*

natTrace

Natural transformation from monad *m* to monad *n*.

> *natTrace* :: (*forall x* ∘ *m x* → *n x*) → *BaseTrace m s* → *BaseTrace n s*
> *natTrace nat* (*BaseTrace* (*Op tr*)) = *BaseTrace* \$ *Op* \$ *nat* ∘ *tr*

noTrace

A *Trace* that discards all inputs.

> *noTrace* :: *Applicative m* ⇒ *BaseTrace m a*
> *noTrace* = *BaseTrace* \$ *Op* \$ *const* (*pure* ())

1.4.4   Cardano.BM.Trace

Utilities

Natural transformation from monad *m* to monad *n*.

> *natTrace* :: (*forall x* ∘ *m x* → *n x*) → *Trace m* → *Trace n*
> *natTrace nat* (*ctx*, *trace*) = (*ctx*, *BaseTrace.natTrace nat trace*)

Access type of *Trace*.

> *typeofTrace* :: *Trace m* → *SubTrace*
> *typeofTrace* (*ctx*, _) = *tracetype ctx*

Update type of *Trace*.

> *updateTracetype* :: *SubTrace* → *Trace m* → *Trace m*
> *updateTracetype subtr* (*ctx*, *tr*) = (*ctx* {*tracetype* = *subtr*}, *tr*)

Enter new named context

The context name is created and checked that its size is below a limit (currently 50 chars). The minimum severity that a log message must be labelled with is looked up in the configuration and recalculated.

> *appendName* :: *MonadIO m* ⇒ *LoggerName* → *Trace m* → *m* (*Trace m*)
> *appendName name* (*ctx*, *trace*) = **do**
>   **let** *prevLoggerName* = *loggerName ctx*
>   *prevMinSeverity* = *minSeverity ctx*
>   *newLoggerName*   = *appendWithDot prevLoggerName name*
>   *globMinSeverity* ← *liftIO* \$ *Config.minSeverity* (*configuration ctx*)
>   *namedSeverity* ← *liftIO* \$ *Config.inspectSeverity* (*configuration ctx*) *newLoggerName*
>   **case** *namedSeverity* **of**
>   *Nothing* → *return* (*ctx* {*loggerName* = *newLoggerName*}, *trace*)
>   *Just sev* → *return* (*ctx* {*loggerName* = *newLoggerName*

```
        , minSeverity = max (max sev prevMinSeverity) globMinSeverity}
        , trace)
appendWithDot :: LoggerName → LoggerName → LoggerName
appendWithDot "" newName = T.take 50 newName
appendWithDot xs "" = xs
appendWithDot xs newName = T.take 50 $ xs <> "." <> newName
```

```
-- return a BaseTrace from a TraceNamed
named :: BaseTrace.BaseTrace m (LogNamed i) → LoggerName → BaseTrace.BaseTrace m i
named trace name = contramap (LogNamed name) trace
```

> TODO remove *locallock*

```
locallock :: MVar ()
locallock = unsafePerformIO $ newMVar ()
```

Trace that forwards to the Switchboard

Every *Trace* ends in the Switchboard which then takes care of dispatching the messages to outputs

```
mainTrace :: Switchboard.Switchboard → TraceNamed IO
mainTrace sb = BaseTrace.BaseTrace $ Op $ λlognamed → do
    Switchboard.effectuate sb lognamed
```

Concrete Trace on stdout

This function returns a trace with an action of type "(*LogNamed LogObject*) → *IO* ()" which will output a text message as text and all others as JSON encoded representation to the console.

```
stdoutTrace :: TraceNamed IO
stdoutTrace = BaseTrace.BaseTrace $ Op $ λlognamed →
    case lnItem lognamed of
      LP (LogMessage logItem) →
        withMVar locallock $ \_ →
          output (lnName lognamed) $ liPayload logItem
      obj →
        withMVar locallock $ \_ →
          output (lnName lognamed) $ toStrict (encodeToLazyText obj)
  where
    output nm msg = TIO.putStrLn $ nm <> " :: " <> msg
```

Concrete Trace into a *TVar*

```
traceInTVar :: STM.TVar [a] → BaseTrace.BaseTrace STM.STM a
traceInTVar tvar = BaseTrace.BaseTrace $ Op $ λa → STM.modifyTVar tvar ((:) a)
```

traceInTVarIO :: *STM.TVar* [*LogObject*] → *TraceNamed IO*
traceInTVarIO tvar = BaseTrace.BaseTrace $ *Op* $ *λln* →
   STM.atomically $ *STM.modifyTVar tvar* ((:) (*lnItem ln*))

traceNamedInTVarIO :: *STM.TVar* [*LogNamed LogObject*] → *TraceNamed IO*
traceNamedInTVarIO tvar = BaseTrace.BaseTrace $ *Op* $ *λln* →
   STM.atomically $ *STM.modifyTVar tvar* ((:) *ln*)

Check a log item's severity against the *Trace*'s minimum severity

> do we need three different *minSeverity* defined?

   We do a lookup of the global *minSeverity* in the configuration. And, a lookup of the *minSeverity* for the current named context. These values might have changed in the meanwhile.
A third filter is the *minSeverity* defined in the current context.

traceConditionally
   :: (*MonadIO m*)
   ⇒ *TraceContext* → *BaseTrace.BaseTrace m LogObject* → *LogObject*
   → *m* ()
traceConditionally ctx logTrace msg@(LP (LogMessage item)) = **do**
   globminsev ← *liftIO* $ *Config.minSeverity* (*configuration ctx*)
   globnamesev ← *liftIO* $ *Config.inspectSeverity* (*configuration ctx*) (*loggerName ctx*)
   let *minsev = max* (*minSeverity ctx*) $ *max globminsev* (*fromMaybe Debug globnamesev*)
      flag = (liSeverity item) ⩾ *minsev*
   when flag $ *BaseTrace.traceWith logTrace msg*
traceConditionally _ logTrace logObject = BaseTrace.traceWith logTrace logObject

Enter message into a trace

The function *traceNamedItem* creates a *LogObject* and threads this through the action defined in the *Trace*.

traceNamedItem
   :: (*MonadIO m*)
   ⇒ *Trace m*
   → *LogSelection*
   → *Severity*
   → *T.Text*
   → *m* ()
traceNamedItem (*ctx, logTrace*) *p s m* =
   let *logmsg = LP* $ *LogMessage* $ *LogItem* {*liSelection = p*
      , *liSeverity = s*
      , *liPayload = m*
      }
   in
   traceConditionally ctx (*named logTrace* (*loggerName ctx*)) $ *logmsg*
logDebug, logInfo, logNotice, logWarning, logError

```
      :: (MonadIO m) ⇒ Trace m → T.Text → m ()
logDebug logTrace = traceNamedItem logTrace Both Debug
logInfo logTrace    = traceNamedItem logTrace Both Info
logNotice logTrace = traceNamedItem logTrace Both Notice
logWarning logTrace = traceNamedItem logTrace Both Warning
logError logTrace   = traceNamedItem logTrace Both Error

logDebugS, logInfoS, logNoticeS, logWarningS, logErrorS
      :: (MonadIO m) ⇒ Trace m → T.Text → m ()
logDebugS logTrace = traceNamedItem logTrace Private Debug
logInfoS logTrace    = traceNamedItem logTrace Private Info
logNoticeS logTrace = traceNamedItem logTrace Private Notice
logWarningS logTrace = traceNamedItem logTrace Private Warning
logErrorS logTrace   = traceNamedItem logTrace Private Error

logDebugP, logInfoP, logNoticeP, logWarningP, logErrorP
      :: (MonadIO m) ⇒ Trace m → T.Text → m ()
logDebugP logTrace = traceNamedItem logTrace Public Debug
logInfoP logTrace    = traceNamedItem logTrace Public Info
logNoticeP logTrace = traceNamedItem logTrace Public Notice
logWarningP logTrace = traceNamedItem logTrace Public Warning
logErrorP logTrace   = traceNamedItem logTrace Public Error

logDebugUnsafeP, logInfoUnsafeP, logNoticeUnsafeP, logWarningUnsafeP, logErrorUnsafeP
      :: (MonadIO m) ⇒ Trace m → T.Text → m ()
logDebugUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Debug
logInfoUnsafeP logTrace    = traceNamedItem logTrace PublicUnsafe Info
logNoticeUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Notice
logWarningUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Warning
logErrorUnsafeP logTrace   = traceNamedItem logTrace PublicUnsafe Error


traceNamedObject
      :: Trace m
      → LogObject
      → m ()
traceNamedObject (ctx, logTrace) = BaseTrace.traceWith (named logTrace (loggerName ctx))
```

subTrace

Transforms the input *Trace* according to the *Configuration* using the logger name of the current *Trace* appended with the new name. If the empty *Text* is passed, then the logger name remains untouched.

```
subTrace :: MonadIO m ⇒ T.Text → Trace m → m (Trace m)
subTrace name tr@(ctx, _) = do
   let newName = appendWithDot (loggerName ctx) name
   subtrace0 ← liftIO $ Config.findSubTrace (configuration ctx) newName
   let subtrace = case subtrace0 of Nothing → Neutral; Just str → str
   case subtrace of
```

Neutral   → **do**
             tr' ← appendName name tr
             return $ *updateTracetype subtrace tr'*
UntimedTrace → **do**
             tr' ← appendName name tr
             return $ *updateTracetype subtrace tr'*
NoTrace   → *return* $ *updateTracetype subtrace* (*ctx, BaseTrace.BaseTrace* $ *Op* $ _ → *pure* ())
DropOpening → *return* $ *updateTracetype subtrace* (*ctx, BaseTrace.BaseTrace* $ *Op* $ λ*lognamed* → **do**
      case *lnItem lognamed* **of**
         ObserveOpen _ → *return* ()
         obj → *traceNamedObject tr obj*)
ObservableTrace _ → **do**
             tr' ← appendName name tr
             return $ *updateTracetype subtrace tr'*

1.4.5   Cardano.BM.Setup
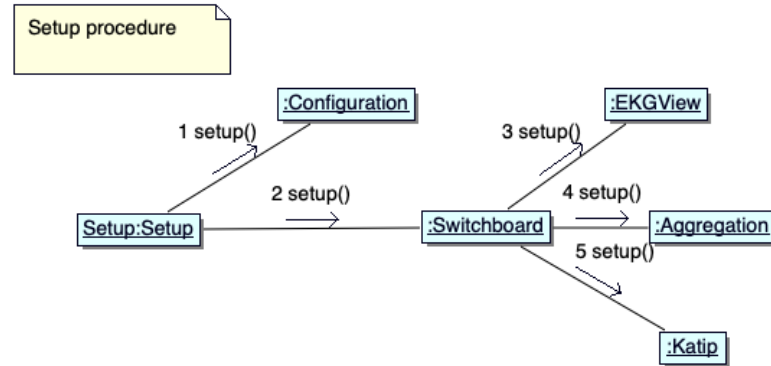
Figure 1.3: Setup procedure

setupTrace

Setup a new *Trace* (Trace) with either a given *Configuration* (Configuration.Model) or a *FilePath* to
a configuration file.

setupTrace :: *MonadIO m* ⇒ *Either FilePath Config.Configuration* → *Text* → *m* (*Trace m*)
setupTrace (*Left cfgFile*) *name* = **do**
   c ← liftIO $ *Config.setup cfgFile*
   setupTrace_ c name
setupTrace (*Right c*) *name* = *setupTrace_ c name*

*setupTrace_* :: *MonadIO m* ⇒ *Config.Configuration* → *Text* → *m* (*Trace m*)
setupTrace_ c name = **do**
   sb ← liftIO $ *Switchboard.realize c*

$sev \leftarrow liftIO \$ Config.minSeverity\ c$
$ctx \leftarrow liftIO \$ newContext\ name\ c\ sev$

let $logTrace = natTrace\ liftIO\ (ctx, mainTrace\ sb)$
$logTrace' \leftarrow subTrace\ "\ "\ logTrace$
$return\ logTrace'$

withTrace

$withTrace :: MonadIO\ m \Rightarrow Config.Configuration \rightarrow Text \rightarrow (Trace\ m \rightarrow m\ t) \rightarrow m\ t$
$withTrace\ cfg\ name\ action =$ **do**
  $logTrace \leftarrow setupTrace\ (Right\ cfg)\ name$
  $action\ logTrace$

newContext

$newContext :: LoggerName \rightarrow Config.Configuration \rightarrow Severity \rightarrow IO\ TraceContext$
$newContext\ name\ cfg\ sev =$ **do**
  $return\ \$ TraceContext\ \{$
    $loggerName = name$
    $, configuration = cfg$
    $, minSeverity = sev$
    $, tracetype = Neutral$
    $\}$

1.4.6 Cardano.BM.Counters

Here the platform is chosen on which we compile this program.
   Currently, we only support *Linux* with its 'proc' filesystem.

```
 {-# LANGUAGE CPP #-}
# if defined (linux_HOST_OS)
# define LINUX
# endif
module Cardano.BM.Counters
  (
    Platform.readCounters
  , diffTimeObserved
  , getMonoClock
  ) where
# ifdef LINUX
import qualified Cardano.BM.Counters.Linux as Platform
# else
import qualified Cardano.BM.Counters.Dummy as Platform
# endif
```

import *Cardano.BM.Counters.Common* (*getMonoClock*)
import *Cardano.BM.Data.Counter*

import *Data.Time.Units* (*Microsecond*)

Calculate difference between clocks

diffTimeObserved :: *CounterState* → *CounterState* → *Microsecond*
diffTimeObserved (*CounterState id0 startCounters*) (*CounterState id1 endCounters*) =
    let
        startTime = *getMonotonicTime startCounters*
        endTime = *getMonotonicTime endCounters*
    in
    if (*id0* ≡ *id1*)
        then *endTime* − *startTime*
        else *error* `"these clocks are not from the same experiment"`
    where
        getMonotonicTime counters = **case** (*filter isMonotonicClockCounter counters*) **of**
            [(*Counter MonotonicClockTime* _ *micros*)] → *fromInteger micros*
            _ → *error* `"A time measurement is missing!"`
        isMonotonicClockCounter :: *Counter* → *Bool*
        isMonotonicClockCounter = (*MonotonicClockTime* ≡) ∘ *cType*

1.4.7   Cardano.BM.Counters.Common

Common functions that serve *readCounters* on all platforms.

nominalTimeToMicroseconds :: *Word64* → *Microsecond*
nominalTimeToMicroseconds = *fromMicroseconds* ∘ *toInteger* ∘ (`div` 1000)

getMonoClock :: *IO* [*Counter*]
getMonoClock = **do**
    t ← *getMonotonicTimeNSec*
    return [*Counter MonotonicClockTime* `"monoclock"` $ *toInteger* $ *nominalTimeToMicroseconds t*]

1.4.8   Cardano.BM.Counters.Dummy

This is a dummy definition of *readCounters* on platforms that do not support the 'proc' filesystem from which we would read the counters.

    The only supported measurement is monotonic clock time for now.

we could well imagine that some day we support all platforms

readCounters :: *SubTrace* → *IO* [*Counter*]
readCounters NoTrace   = *return* [ ]
readCounters Neutral   = *return* [ ]

```
readCounters UntimedTrace = return [ ]
readCounters DropOpening  = return [ ]
readCounters (ObservableTrace tts) = foldrM (λ(sel,fun) a →
    if sel 'member' tts
    then (fun ≫ λxs → return $ a ++ xs)
    else return a)[ ] selectors
  where
    selectors = [(MonotonicClock,getMonoClock)
      -- , (MemoryStats, readProcStatM)
      -- , (ProcessStats, readProcStats)
          -- , (IOStats, readProcIO)
                        ]
```

1.4.9   Cardano.BM.Counters.Linux

> we have to expand the *readMemStats* function
> to read full data from *proc*

```
readCounters :: SubTrace → IO [Counter ]
readCounters NoTrace       = return [ ]
readCounters Neutral       = return [ ]
readCounters UntimedTrace = return [ ]
readCounters DropOpening  = return [ ]
readCounters (ObservableTrace tts) = foldrM (λ(sel,fun) a →
    if sel 'member' tts
    then (fun ≫ λxs → return $ a ++ xs)
    else return a)[ ] selectors
  where
    selectors = [(MonotonicClock,getMonoClock)
      ,(MemoryStats,readProcStatM)
      ,(ProcessStats,readProcStats)
      ,(IOStats,readProcIO)
      ]

pathProc :: FilePath
pathProc = "/proc/"
pathProcStat :: ProcessID → FilePath
pathProcStat pid = pathProc < / > (show pid) < / > "stat"
pathProcStatM :: ProcessID → FilePath
pathProcStatM pid = pathProc < / > (show pid) < / > "statm"
pathProcIO :: ProcessID → FilePath
pathProcIO pid = pathProc < / > (show pid) < / > "io"
```

Reading from a file in /proc/<pid >

```
readProcList :: FilePath → IO [Integer ]
readProcList fp = do
```

$cs \leftarrow readFile\ fp$

$return\ \$\ map\ (\lambda s \rightarrow maybe\ 0\ id\ \$\ (readMaybe\ s :: Maybe\ Integer))\ (words\ cs)$

readProcStatM - /proc/<pid >/statm

```
/proc/[pid]/statm
      Provides information about memory usage, measured in pages.  The columns are:
             size       (1) total program size
                            (same as VmSize in /proc/[pid]/status)
             resident   (2) resident set size
                            (same as VmRSS in /proc/[pid]/status)
             shared     (3) number of resident shared pages (i.e., backed by a file)
                            (same as RssFile+RssShmem in /proc/[pid]/status)
             text       (4) text (code)
             lib        (5) library (unused since Linux 2.6; always 0)
             data       (6) data + stack
             dt         (7) dirty pages (unused since Linux 2.6; always 0)
```

$readProcStatM :: IO\ [\,Counter\,]$

$readProcStatM = \textbf{do}$

$\quad pid \leftarrow getProcessID$

$\quad ps0 \leftarrow readProcList\ (pathProcStatM\ pid)$

$\quad \textbf{let}\ ps = zip\ colnames\ ps0$

$\quad\quad psUseful = filter\ (("unused"\ \not\equiv) \circ fst)\ ps$

$\quad return\ \$\ map\ (\lambda(n,i) \rightarrow Counter\ MemoryCounter\ n\ i)\ psUseful$

$\quad \textbf{where}$

$\quad\quad colnames :: [\,Text\,]$

$\quad\quad colnames = [\,"size","resident","shared","text","unused","data","unused"\,]$

readProcStats - //proc//<pid >//stat

```
/proc/[pid]/stat
      Status  information about the process.  This is used by ps(1).  It is defined in the kernel source file
      fs/proc/array.c.

      The fields, in order, with their proper scanf(3) format specifiers, are listed below.  Whether  or  not
      certain   of   these   fields   display   valid  information  is  governed  by  a  ptrace  access  mode
      PTRACE_MODE_READ_FSCREDS | PTRACE_MODE_NOAUDIT check (refer to ptrace(2)).  If the check denies access,
      then the field value is displayed as 0.  The affected fields are indicated with the marking [PT].

      (1) pid  %d
                 The process ID.

      (2) comm  %s
                 The  filename  of  the  executable,  in parentheses.  This is visible whether or not the exe-
                 cutable is swapped out.

      (3) state  %c
                 One of the following characters, indicating process state:

                 R  Running

                 S  Sleeping in an interruptible wait

                 D  Waiting in uninterruptible disk sleep

                 Z  Zombie
```

```
                    T  Stopped (on a signal) or (before Linux 2.6.33) trace stopped

                    t  Tracing stop (Linux 2.6.33 onward)

                    W  Paging (only before Linux 2.6.0)

                    X  Dead (from Linux 2.6.0 onward)

                    x  Dead (Linux 2.6.33 to 3.13 only)

                    K  Wakekill (Linux 2.6.33 to 3.13 only)

                    W  Waking (Linux 2.6.33 to 3.13 only)

                    P  Parked (Linux 3.9 to 3.13 only)

(4) ppid  %d
            The PID of the parent of this process.

(5) pgrp  %d
            The process group ID of the process.

(6) session  %d
            The session ID of the process.

(7) tty_nr  %d
            The controlling terminal of the process.  (The minor device number is contained in the combi-
            nation of bits 31 to 20 and 7 to 0; the major device number is in bits 15 to 8.)

(8) tpgid  %d
            The ID of the foreground process group of the controlling terminal of the process.

(9) flags  %u
            The  kernel  flags  word of the process.  For bit meanings, see the PF_* defines in the Linux
            kernel source file include/linux/sched.h.  Details depend on the kernel version.

            The format for this field was %lu before Linux 2.6.

(10) minflt  %lu
            The number of minor faults the process has made which have not required loading a memory page
            from disk.

(11) cminflt  %lu
            The number of minor faults that the process's waited-for children have made.

(12) majflt  %lu
            The  number  of  major  faults the process has made which have required loading a memory page
            from disk.

(13) cmajflt  %lu
            The number of major faults that the process's waited-for children have made.

(14) utime  %lu
            Amount of time that this process has been scheduled in user mode,  measured  in  clock  ticks
            (divide by sysconf(_SC_CLK_TCK)).  This includes guest time, guest_time (time spent running a
            virtual CPU, see below), so that applications that are not aware of the guest time  field  do
            not lose that time from their calculations.

(15) stime  %lu
            Amount  of  time that this process has been scheduled in kernel mode, measured in clock ticks
            (divide by sysconf(_SC_CLK_TCK)).

(16) cutime  %ld
            Amount of time that this process's waited-for children have been scheduled in user mode, mea-
            sured  in  clock ticks (divide by sysconf(_SC_CLK_TCK)).  (See also times(2).)  This includes
            guest time, cguest_time (time spent running a virtual CPU, see below).

(17) cstime  %ld
```

Amount of time that this process's waited-for children have been scheduled in kernel mode,
measured in clock ticks (divide by sysconf(_SC_CLK_TCK)).

(18) priority %ld
(Explanation for Linux 2.6) For processes running a real-time scheduling policy (policy
below; see sched_setscheduler(2)), this is the negated scheduling priority, minus one; that
is, a number in the range -2 to -100, corresponding to real-time priorities 1 to 99. For
processes running under a non-real-time scheduling policy, this is the raw nice value (set-
priority(2)) as represented in the kernel. The kernel stores nice values as numbers in the
range 0 (high) to 39 (low), corresponding to the user-visible nice range of -20 to 19.

(19) nice %ld
The nice value (see setpriority(2)), a value in the range 19 (low priority) to -20 (high pri-
ority).

(20) num_threads %ld
Number of threads in this process (since Linux 2.6). Before kernel 2.6, this field was hard
coded to 0 as a placeholder for an earlier removed field.

(21) itrealvalue %ld
The time in jiffies before the next SIGALRM is sent to the process due to an interval timer.
Since kernel 2.6.17, this field is no longer maintained, and is hard coded as 0.

(22) starttime %llu
The time the process started after system boot. In kernels before Linux 2.6, this value was
expressed in jiffies.   Since Linux 2.6, the value is expressed in clock ticks (divide by
sysconf(_SC_CLK_TCK)).

The format for this field was %lu before Linux 2.6.

(23) vsize %lu
Virtual memory size in bytes.

(24) rss %ld
Resident Set Size: number of pages the process has in real memory. This is just the pages
which count toward text, data, or stack space. This does not include pages which have not
been demand-loaded in, or which are swapped out.

(25) rsslim %lu
Current soft limit in bytes on the rss of the process; see the description of RLIMIT_RSS in
getrlimit(2).

(26) startcode %lu [PT]
The address above which program text can run.

(27) endcode %lu [PT]
The address below which program text can run.

(28) startstack %lu [PT]
The address of the start (i.e., bottom) of the stack.

(29) kstkesp %lu [PT]
The current value of ESP (stack pointer), as found in the kernel stack page for the process.

(30) kstkeip %lu [PT]
The current EIP (instruction pointer).

(31) signal %lu
The bitmap of pending signals, displayed as a decimal number. Obsolete, because it does not
provide information on real-time signals; use /proc/[pid]/status instead.

(32) blocked %lu
The bitmap of blocked signals, displayed as a decimal number. Obsolete, because it does not
provide information on real-time signals; use /proc/[pid]/status instead.

(33) sigignore %lu
The bitmap of ignored signals, displayed as a decimal number. Obsolete, because it does not
provide information on real-time signals; use /proc/[pid]/status instead.

```
(34) sigcatch  %lu
              The bitmap of caught signals, displayed as a decimal number.  Obsolete, because it  does  not
              provide information on real-time signals; use /proc/[pid]/status instead.

(35) wchan  %lu  [PT]
              This  is  the  "channel" in which the process is waiting.  It is the address of a location in
              the kernel where the process is sleeping.  The corresponding symbolic name can  be  found  in
              /proc/[pid]/wchan.

(36) nswap  %lu
              Number of pages swapped (not maintained).

(37) cnswap  %lu
              Cumulative nswap for child processes (not maintained).

(38) exit_signal  %d  (since Linux 2.1.22)
              Signal to be sent to parent when we die.

(39) processor  %d  (since Linux 2.2.8)
              CPU number last executed on.

(40) rt_priority  %u  (since Linux 2.5.19)
              Real-time  scheduling priority, a number in the range 1 to 99 for processes scheduled under a
              real-time policy, or 0, for non-real-time processes (see sched_setscheduler(2)).

(41) policy  %u  (since Linux 2.5.19)
              Scheduling policy (see  sched_setscheduler(2)).   Decode  using  the  SCHED_* constants  in
              linux/sched.h.

              The format for this field was %lu before Linux 2.6.22.

(42) delayacct_blkio_ticks  %llu  (since Linux 2.6.18)
              Aggregated block I/O delays, measured in clock ticks (centiseconds).

(43) guest_time  %lu  (since Linux 2.6.24)
              Guest  time  of  the process (time spent running a virtual CPU for a guest operating system),
              measured in clock ticks (divide by sysconf(_SC_CLK_TCK)).

(44) cguest_time  %ld  (since Linux 2.6.24)
              Guest  time  of  the  process's  children,  measured  in  clock  ticks  (divide  by
              sysconf(_SC_CLK_TCK)).

(45) start_data  %lu  (since Linux 3.3)  [PT]
              Address above which program initialized and uninitialized (BSS) data are placed.

(46) end_data  %lu  (since Linux 3.3)  [PT]
              Address below which program initialized and uninitialized (BSS) data are placed.

(47) start_brk  %lu  (since Linux 3.3)  [PT]
              Address above which program heap can be expanded with brk(2).

(48) arg_start  %lu  (since Linux 3.5)  [PT]
              Address above which program command-line arguments (argv) are placed.

(49) arg_end  %lu  (since Linux 3.5)  [PT]
              Address below program command-line arguments (argv) are placed.

(50) env_start  %lu  (since Linux 3.5)  [PT]
              Address above which program environment is placed.

(51) env_end  %lu  (since Linux 3.5)  [PT]
              Address below which program environment is placed.

(52) exit_code  %d  (since Linux 3.5)  [PT]
              The thread's exit status in the form reported by waitpid(2).
```

readProcStats :: *IO* [*Counter*]
readProcStats = **do**
    pid ← *getProcessID*
    ps0 ← *readProcList* (*pathProcStat pid*)
    let *ps* = *zip colnames ps0*
      psUseful = *filter* (("unused" ≢) ∘ *fst*) *ps*
    return $ *map* (λ(*n*,*i*) → *Counter StatInfo n i*) *psUseful*
  where
    colnames :: [*Text*]
    colnames = ["pid","unused","unused","ppid","pgrp","session","ttynr","tpgid","flags","minfl
      ,"cminflt","majflt","cmajflt","utime","stime","cutime","cstime","priority","nice","num
      ,"itrealvalue","starttime","vsize","rss","rsslim","startcode","endcode","startstack","
      ,"signal","blocked","sigignore","sigcatch","wchan","nswap","cnswap","exitsignal","proc
      ,"policy","blkio","guesttime","cguesttime","startdata","enddata","startbrk","argstart"
      ,"envend","exitcode"
      ]

readProcIO - //proc//<pid >//io

```
/proc/[pid]/io (since kernel 2.6.20)
       This file contains I/O statistics for the process, for example:

               # cat /proc/3828/io
               rchar: 323934931
               wchar: 323929600
               syscr: 632687
               syscw: 632675
               read_bytes: 0
               write_bytes: 323932160
               cancelled_write_bytes: 0

       The fields are as follows:

       rchar: characters read
               The number of bytes which this task has caused to be read from storage.  This is simply the  sum
               of bytes which this process passed to read(2) and similar system calls.  It includes things such
               as terminal I/O and is unaffected by whether or not actual physical disk I/O was  required  (the
               read might have been satisfied from pagecache).

       wchar: characters written
               The  number  of bytes which this task has caused, or shall cause to be written to disk.  Similar
               caveats apply here as with rchar.

       syscr: read syscalls
               Attempt to count the number of read I/O operations-that is, system calls  such  as  read(2)  and
               pread(2).

       syscw: write syscalls
               Attempt  to  count the number of write I/O operations-that is, system calls such as write(2) and
               pwrite(2).

       read_bytes: bytes read
               Attempt to count the number of bytes which this process really did cause to be fetched from  the
               storage layer.  This is accurate for block-backed filesystems.

       write_bytes: bytes written
               Attempt to count the number of bytes which this process caused to be sent to the storage layer.

       cancelled_write_bytes:
               The  big  inaccuracy  here  is truncate.  If a process writes 1MB to a file and then deletes the
```

```
file, it will in fact perform no writeout.  But it will have been accounted as having caused 1MB
of  write.   In other words: this field represents the number of bytes which this process caused
to not happen, by truncating pagecache. A task can cause "negative" I/O too.   If  this  task
truncates  some  dirty  pagecache,  some  I/O  which another task has been accounted for (in its
write\_bytes) will not be happening.
```

```
Note: In the current implementation, things are a bit racy  on  32-bit  systems:  if  process  A  reads
process  B's  /proc/[pid]/io  while process B is updating one of these 64-bit counters, process A could
see an intermediate result.
```

```
Permission to access this file is governed by a ptrace access mode PTRACE\_MODE\_READ\_FSCREDS check;  see
ptrace(2).
```

readProcIO :: *IO* [*Counter*]
readProcIO = **do**
    pid ← *getProcessID*
    ps0 ← *readProcList* (*pathProcIO pid*)
    let *ps* = *zip colnames ps0*
    return $ *map* ($\lambda(n,i) \rightarrow$ *Counter IOCounter n i*) *ps*
  where
    colnames :: [*Text*]
    colnames = [`"rchar"`,`"wchar"`,`"syscr"`,`"syscw"`,`"rbytes"`,`"wbytes"`,`"cxwbytes"`]

1.4.10 Cardano.BM.Data.Backend

Accepts a NamedLogItem

class *IsEffectuator t* **where**
  effectuate :: $t \rightarrow$ *NamedLogItem* $\rightarrow$ *IO* ()

Declaration of a *Backend*

class (*IsEffectuator t*) $\Rightarrow$ *IsBackend t* **where**
  typeof :: $t \rightarrow$ *BackendKind*
  realize   :: *Configuration* $\rightarrow$ *IO t*
  realizefrom :: *forall s* ∘ (*IsEffectuator s*) $\Rightarrow$ *Configuration* $\rightarrow$ $s \rightarrow$ *IO t*
  default *realizefrom* :: *forall s* ∘ (*IsEffectuator s*) $\Rightarrow$ *Configuration* $\rightarrow$ $s \rightarrow$ *IO t*
  realizefrom c _ = *realize c*
  unrealize :: $t \rightarrow$ *IO* ()
  effectuatefrom :: *forall s* ∘ (*IsEffectuator s*) $\Rightarrow$ $t \rightarrow$ *NamedLogItem* $\rightarrow$ $s \rightarrow$ *IO* ()
  default *effectuatefrom* :: *forall s* ∘ (*IsEffectuator s*) $\Rightarrow$ $t \rightarrow$ *NamedLogItem* $\rightarrow$ $s \rightarrow$ *IO* ()
  effectuatefrom t nli _ = *effectuate t nli*

Backend

A backend is referenced through the function *bPass* which accepts a NamedLogItem and a terminating function *bTerminate* which is responsible for closing the specific backend.

```
data Backend = MkBackend
  {bPass :: NamedLogItem → IO ()
  ,bTerminate :: IO ()
  }
```

1.4.11   Cardano.BM.Data.Configuration

Data structure to help parsing configuration files.

Representation

```
type Port = Int
data Representation = Representation
  {minSeverity   :: Severity
  ,rotation        :: RotationParameters
  ,setupScribes   :: [ScribeDefinition]
  ,defaultScribes :: [(ScribeKind, Text)]
  ,setupBackends :: [BackendKind]
  ,defaultBackends :: [BackendKind]
  ,hasEKG          :: Maybe Port
  ,hasGUI          :: Maybe Port
  ,options          :: HM.HashMap Text Object
  }
  deriving (Generic, Show, ToJSON, FromJSON)
```

parseRepresentation

```
parseRepresentation :: FilePath → IO Representation
parseRepresentation fp = do
  repr :: Representation ← decodeFileThrow fp
  return $ implicit_fill_representation repr
```

after parsing the configuration representation we implicitly correct it.

```
implicit_fill_representation :: Representation → Representation
implicit_fill_representation =
    remove_ekgview_if_not_defined ∘
    filter_duplicates_from_backends ∘
    filter_duplicates_from_scribes ∘
    union_setup_and_usage_backends ∘
    add_ekgview_if_port_defined ∘
    add_katip_if_any_scribes
  where
    filter_duplicates_from_backends r =
        r {setupBackends = mkUniq $ setupBackends r}
    filter_duplicates_from_scribes r =
```

```
       r {setupScribes = mkUniq $ setupScribes r}
    union_setup_and_usage_backends r =
       r {setupBackends = setupBackends r <> defaultBackends r}
    remove_ekgview_if_not_defined r =
       case hasEKG r of
       Nothing → r {defaultBackends = filter (λbk → bk ≢ EKGViewBK) (defaultBackends r)
          ,setupBackends = filter (λbk → bk ≢ EKGViewBK) (setupBackends r)
          }
       Just _ → r
    add_ekgview_if_port_defined r =
       case hasEKG r of
       Nothing → r
       Just _ → r {setupBackends = setupBackends r <> [EKGViewBK]}
    add_katip_if_any_scribes r =
       if (any ¬ [null $ setupScribes r, null $ defaultScribes r])
       then r {setupBackends = setupBackends r <> [KatipBK]}
       else r
    mkUniq :: Ord a ⇒ [a] → [a]
    mkUniq = Set.toList ∘ Set.fromList
```

1.4.12   Cardano.BM.Data.Counter

Counter

```
data Counter = Counter
          {cType  :: CounterType
          ,cName :: Text
          ,cValue :: Integer
          }
          deriving (Eq, Show, Generic, ToJSON)
data CounterType = MonotonicClockTime
   | MemoryCounter
   | StatInfo
   | IOCounter
   | CpuCounter
     deriving (Eq, Show, Generic, ToJSON)
nameCounter :: Counter → Text
nameCounter (Counter MonotonicClockTime _ _) = "Time"
nameCounter (Counter MemoryCounter _ _) = "Mem"
nameCounter (Counter StatInfo        _ _) = "Stat"
nameCounter (Counter IOCounter       _ _) = "IO"
nameCounter (Counter CpuCounter      _ _) = "Cpu"
instance ToJSON Microsecond where
   toJSON = toJSON ∘ toMicroseconds
   toEncoding = toEncoding ∘ toMicroseconds
```

CounterState

```
data CounterState = CounterState {
  csIdentifier :: Unique
  , csCounters :: [Counter]
  }
  deriving (Generic, ToJSON)
instance ToJSON Unique where
  toJSON = toJSON ∘ hashUnique
  toEncoding = toEncoding ∘ hashUnique
instance Show CounterState where
  show cs = (show ∘ hashUnique) (csIdentifier cs)
      <> " => " <> (show $ csCounters cs)
```

Difference between counters

```
diffCounters :: [Counter] → [Counter] → [Counter]
diffCounters openings closings =
    getCountersDiff openings closings
  where
    getCountersDiff :: [Counter]
                → [Counter]
                → [Counter]
    getCountersDiff as bs =
      let
        getName counter = nameCounter counter <> cName counter

        asNames = map getName as
        aPairs = zip asNames as

        bsNames = map getName bs
        bs' = zip bsNames bs
        bPairs = HM.fromList bs'
      in
        catMaybes $ (flip map) aPairs $ λ(name, Counter _ _ startValue) →
          case HM.lookup name bPairs of
            Nothing     → Nothing
            Just counter → let endValue = cValue counter
                           in Just counter {cValue = endValue − startValue}
```

1.4.13   Cardano.BM.Data.LogItem

LoggerName

```
type LoggerName = Text
```

NamedLogItem

> **type** *NamedLogItem = LogNamed LogObject*

LogItem

> TODO *liPayload* :: *ToObject*

> **data** *LogItem = LogItem*
>   {*liSelection* :: *LogSelection*
>   ,*liSeverity* :: *Severity*
>   ,*liPayload* :: *Text*`-- TODO should become ToObject`
>   } **deriving** (*Show*, *Generic*, *ToJSON*)

> **data** *LogSelection =*
>   *Public* `-- only to public logs.`
> |*PublicUnsafe*`-- only to public logs, not console.`
> |*Private*`-- only to private logs.`
> |*Both* `-- to public and private logs.`
>   **deriving** (*Show*, *Generic*, *ToJSON*, *FromJSON*)

LogObject

> **data** *LogPrims = LogMessage LogItem*
> | *LogValue Text Integer*
>   **deriving** (*Generic*, *Show*, *ToJSON*)
> **data** *LogObject = LP LogPrims*
> | *ObserveOpen CounterState*
> | *ObserveDiff CounterState*
> | *ObserveClose CounterState*
> | *AggregatedMessage Text Aggregated*
> | *KillPill*
>   **deriving** (*Generic*, *Show*, *ToJSON*)

LogNamed

A *LogNamed* contains of a context name and some log item.

> **data** *LogNamed item = LogNamed*
>   {*lnName* :: *LoggerName*
>   ,*lnItem* :: *item*
>   } **deriving** (*Show*)
> **deriving instance** *Generic item* ⇒ *Generic* (*LogNamed item*)
> **deriving instance** (*ToJSON item*, *Generic item*) ⇒ *ToJSON* (*LogNamed item*)

1.4.14   Cardano.BM.Data.Observable

ObservableInstance

```
data ObservableInstance = MonotonicClock
  | MemoryStats
  | ProcessStats
  | IOStats
    deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)
```

1.4.15   Cardano.BM.Data.Output

OutputKind

```
data OutputKind = StdOut
  | TVarList (STM.TVar [LogObject])
  | TVarListNamed (STM.TVar [LogNamed LogObject])
  | Null
  deriving (Eq)
```

ScribeKind

This identifies katip's scribes by type.

```
data ScribeKind = FileTextSK
  | FileJsonSK
  | StdoutSK
  | StderrSK
  deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)
```

ScribeId

A scribe is identified by *ScribeKind x Filename*

```
type ScribeId = Text-- (ScribeKind, Filename)
```

ScribeDefinition

This identifies katip's scribes by type.

```
data ScribeDefinition = ScribeDefinition
  {
    scKind :: ScribeKind
  , scName  :: Text
  , scRotation :: Maybe RotationParameters
  }
  deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)
```

1.4.16   Cardano.BM.Data.Severity

Severity

```
data Severity = Debug | Info | Warning | Notice | Error
  deriving (Show, Eq, Ord, Generic, ToJSON)

instance FromJSON Severity where
  parseJSON = withText "severity" $ λcase
    "Debug"   → pure Debug
    "Info"    → pure Info
    "Notice"  → pure Notice
    "Warning" → pure Warning
    "Error"   → pure Error
    _         → pure Info-- catch all
```

1.4.17   Cardano.BM.Data.SubTrace

SubTrace

```
data SubTrace = Neutral
  | UntimedTrace
  | NoTrace
  | DropOpening
  | ObservableTrace (Set ObservableInstance)
    deriving (Generic, Show, FromJSON, ToJSON)
```

1.4.18   Cardano.BM.Data.Trace

Trace

A *Trace* consists of a TraceContext and a TraceNamed in *m*.

```
type Trace m = (TraceContext, TraceNamed m)
```

TraceNamed

A *TraceNamed* is a specialized Contravariant of type LogNamed with payload LogObject.

```
type TraceNamed m = BaseTrace m (LogNamed LogObject)
```

TraceContext

We keep the context's name and a reference to the *Configuration* in the *TraceContext*.

```
data TraceContext = TraceContext {
  loggerName   :: LoggerName
  , configuration :: Configuration
  , tracetype     :: SubTrace
  , minSeverity :: Severity
  }
```

1.4.19 Cardano.BM.Configuration

see Cardano.BM.Configuration.Model for the implementation.

> $getOptionOrDefault :: CM.Configuration \rightarrow Text \rightarrow Text \rightarrow IO\,(Text)$
> $getOptionOrDefault\ cg\ name\ def =$ **do**
> $opt \leftarrow CM.getOption\ cg\ name$
>   **case** opt **of**
> $Nothing \rightarrow return\ def$
> $Just\ o \rightarrow return\ o$

1.4.20 Cardano.BM.Configuration.Model
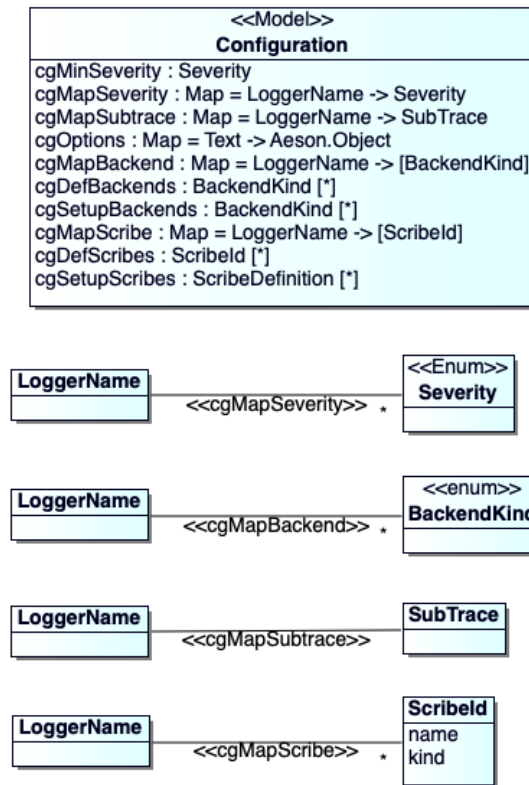
Configuration.Model

Figure 1.4: Configuration model

> **type** $ConfigurationMVar = MVar\ ConfigurationInternal$
> **newtype** $Configuration = Configuration$
> $\{getCG :: ConfigurationMVar\}$
>
> ```
> -- Our internal state; see -"Configuration model"-
> ```
> **data** $ConfigurationInternal = ConfigurationInternal$

```
{cgMinSeverity    :: Severity
,cgMapSeverity    :: HM.HashMap LoggerName Severity
,cgMapSubtrace    :: HM.HashMap LoggerName SubTrace
,cgOptions        :: HM.HashMap Text Object
,cgMapBackend     :: HM.HashMap LoggerName [BackendKind]
,cgDefBackendKs   :: [BackendKind]
,cgSetupBackends  :: [BackendKind]
,cgMapScribe      :: HM.HashMap LoggerName [ScribeId]
,cgDefScribes     :: [ScribeId]
,cgSetupScribes   :: [ScribeDefinition]
,cgPortEKG        :: Int
,cgPortGUI        :: Int
}
```

Backends configured in the *Switchboard*

For a given context name return the list of backends configured, or, in case no such configuration
exists, return the default backends.

```
getBackends :: Configuration → LoggerName → IO [BackendKind]
getBackends configuration name =
    withMVar (getCG configuration) $ λcg → do
      let outs = HM.lookup name (cgMapBackend cg)
      case outs of
        Nothing → do
          return (cgDefBackendKs cg)
        Just os → return $ os-- TODO in (cgDefBackendKs cg)
getDefaultBackends :: Configuration → IO [BackendKind]
getDefaultBackends configuration =
    withMVar (getCG configuration) $ λcg → do
      return (cgDefBackendKs cg)
setDefaultBackends :: Configuration → [BackendKind] → IO ()
setDefaultBackends configuration bes = do
    cg ← takeMVar (getCG configuration)
    putMVar (getCG configuration) $ cg {cgDefBackendKs = bes}
setBackend :: Configuration → LoggerName → Maybe [BackendKind] → IO ()
setBackend configuration name be = do
    cg ← takeMVar (getCG configuration)
    putMVar (getCG configuration) $ cg {cgMapBackend = HM.alter (\_ → be) name (cgMapBackend cg)}
```

Backends to be setup by the *Switchboard*

Defines the list of *Backend*s that need to be setup by the *Switchboard*.

```
setSetupBackends :: Configuration → [BackendKind] → IO ()
setSetupBackends configuration bes = do
```

$cg \leftarrow takeMVar \, (getCG \, configuration)$
$putMVar \, (getCG \, configuration) \, \$ \, cg \, \{cgSetupBackends = bes\}$

Scribes configured in the *Log* backend

For a given context name return the list of scribes to output to, or, in case no such configuration exists, return the default scribes to use.

$getScribes :: Configuration \rightarrow LoggerName \rightarrow IO \, [ScribeId]$
$getScribes \, configuration \, name =$
   $withMVar \, (getCG \, configuration) \, \$ \, \lambda cg \rightarrow \mathbf{do}$
     $\mathbf{let} \, outs = HM.lookup \, name \, (cgMapScribe \, cg)$
     $\mathbf{case} \, outs \, \mathbf{of}$
       $Nothing \rightarrow \mathbf{do}$
         $return \, (cgDefScribes \, cg)$
       $Just \, os \rightarrow return \, \$ \, os$
$setDefaultScribes :: Configuration \rightarrow [ScribeId] \rightarrow IO \, ()$
$setDefaultScribes \, configuration \, scs = \mathbf{do}$
   $cg \leftarrow takeMVar \, (getCG \, configuration)$
   $putMVar \, (getCG \, configuration) \, \$ \, cg \, \{cgDefScribes = scs\}$

Scribes to be setup in the *Log* backend

Defines the list of *Scribe*s that need to be setup in the *Log* backend.

$setSetupScribes :: Configuration \rightarrow [ScribeDefinition] \rightarrow IO \, ()$
$setSetupScribes \, configuration \, sds = \mathbf{do}$
   $cg \leftarrow takeMVar \, (getCG \, configuration)$
   $putMVar \, (getCG \, configuration) \, \$ \, cg \, \{cgSetupScribes = sds\}$

$getSetupScribes :: Configuration \rightarrow IO \, [ScribeDefinition]$
$getSetupScribes \, configuration =$
   $withMVar \, (getCG \, configuration) \, \$ \, \lambda cg \rightarrow \mathbf{do}$
     $return \, \$ \, cgSetupScribes \, cg$

Access port numbers of EKG, GUI

$getEKGport :: Configuration \rightarrow IO \, Int$
$getEKGport \, configuration =$
   $withMVar \, (getCG \, configuration) \, \$ \, \lambda cg \rightarrow \mathbf{do}$
     $return \, \$ \, cgPortEKG \, cg$

$getGUIport :: Configuration \rightarrow IO \, Int$
$getGUIport \, configuration =$
   $withMVar \, (getCG \, configuration) \, \$ \, \lambda cg \rightarrow \mathbf{do}$
     $return \, \$ \, cgPortGUI \, cg$

Options

```
getOption :: Configuration → Text → IO (Maybe Text)
getOption configuration name = do
    withMVar (getCG configuration) $ λcg →
      case HM.lookup name (cgOptions cg) of
        Nothing → return Nothing
        Just o → return $ Just $ pack $ show o
```

Global setting of minimum severity

```
minSeverity :: Configuration → IO Severity
minSeverity configuration = withMVar (getCG configuration) $ λcg →
    return $ cgMinSeverity cg

setMinSeverity :: Configuration → Severity → IO ()
setMinSeverity configuration sev = do
    cg ← takeMVar (getCG configuration)
    putMVar (getCG configuration) $ cg {cgMinSeverity = sev}
```

Relation of context name to minimum severity

```
inspectSeverity :: Configuration → Text → IO (Maybe Severity)
inspectSeverity configuration name = do
    withMVar (getCG configuration) $ λcg →
      return $ HM.lookup name (cgMapSeverity cg)

setSeverity :: Configuration → Text → Maybe Severity → IO ()
setSeverity configuration name sev = do
    cg ← takeMVar (getCG configuration)
    putMVar (getCG configuration) $ cg {cgMapSeverity = HM.alter (\_ → sev) name (cgMapSeverity cg)}
```

Relation of context name to SubTrace

A new context may contain a different type of *Trace*. The function *appendName* (Enter new named context) will look up the *SubTrace* for the context's name.

```
findSubTrace :: Configuration → Text → IO (Maybe SubTrace)
findSubTrace configuration name = do
    withMVar (getCG configuration) $ λcg →
      return $ HM.lookup name (cgMapSubtrace cg)

setSubTrace :: Configuration → Text → Maybe SubTrace → IO ()
setSubTrace configuration name trafo = do
    cg ← takeMVar (getCG configuration)
    putMVar (getCG configuration) $ cg {cgMapSubtrace = HM.alter (\_ → trafo) name (cgMapSubtrace cg)}
```

Parse configuration from file

Parse the configuration into an internal representation first. Then, fill in *Configuration* from it in a second step after refinement.

```
setup :: FilePath → IO Configuration
setup fp = do
    r ← R.parseRepresentation fp
    cgref ← newEmptyMVar
    putMVar cgref $ ConfigurationInternal
        {cgMinSeverity = R.minSeverity r
        ,cgMapSeverity = HM.empty
        ,cgMapSubtrace = HM.empty
        ,cgOptions = R.options r
        ,cgMapBackend = HM.empty
        ,cgDefBackendKs = R.defaultBackends r
        ,cgSetupBackends = R.setupBackends r
        ,cgMapScribe = HM.empty
        ,cgDefScribes = r_defaultScribes r
        ,cgSetupScribes = R.setupScribes r
        ,cgPortEKG = r_hasEKG r
        ,cgPortGUI = r_hasGUI r
        }
    return $ Configuration cgref
  where
    r_hasEKG r = case (R.hasEKG r) of
        Nothing → 0
        Just p → p
    r_hasGUI r = case (R.hasGUI r) of
        Nothing → 0
        Just p → p
    r_defaultScribes r = map (λ(k,n) → pack (show k) <> "::" <> n) (R.defaultScribes r)
```

Setup empty configuration

```
empty :: IO Configuration
empty = do
    cgref ← newEmptyMVar
    putMVar cgref $ ConfigurationInternal Debug HM.empty HM.empty HM.empty HM.empty [ ] [ ] HM.empty [ ]
    return $ Configuration cgref
```

1.4.21   Cardano.BM.Aggregated

```
{-# LANGUAGE DeriveAnyClass #-}
{-# LANGUAGE DeriveGeneric #-}
{-# LANGUAGE StandaloneDeriving #-}
```

```haskell
module Cardano.BM.Aggregated
  (
    Aggregated (..)
  , Stats (..)
  , updateAggregation
  ) where
import GHC.Generics (Generic)
import Data.Aeson (ToJSON)


data Stats = Stats {
  fmin :: Integer,
  fmax :: Integer,
  fcount :: Integer,
  fsum_A :: Integer,
  fsum_B :: Integer
  } deriving (Show, Eq, Generic, ToJSON)
data Aggregated = Aggregated {
  fstats :: Stats,
  flast  :: Integer,
  fdelta :: Stats
  } deriving (Show, Eq, Generic, ToJSON)
```

Update aggregation

We distinguish an unitialized from an already initialized aggregation:

```haskell
updateAggregation :: Integer → Maybe Aggregated → Maybe Aggregated
updateAggregation v Nothing =
  Just $
    Aggregated {fstats = Stats {
        fmin = v, fmax = v, fcount = 1
        , fsum_A = v, fsum_B = v * v}
      , flast = v
      , fdelta = Stats {
        fmin = 0, fmax = 0, fcount = 0
        , fsum_A = 0, fsum_B = 0}
      }
updateAggregation v (Just (Aggregated (Stats _min _max _count _sumA _sumB)
  _last
  (Stats _dmin _dmax _dcount _dsumA _dsumB)
  )) =
  let delta = v − _last
  in
  Just $
    Aggregated {fstats = Stats {
        fmin = (min _min v)
        , fmax = (max _max v)
```

```
            ,fcount = (_count + 1)
            ,fsum_A = (_sumA + v)
            ,fsum_B = (_sumB + v * v)
            }
         ,flast = v
         ,fdelta = Stats {
            fmin = (min _dmin delta)
            ,fmax = (max _dmax delta)
            ,fcount = (_dcount + 1)
            ,fsum_A = (_dsumA + delta)
            ,fsum_B = (_dsumB + delta * delta)
            }
         }
```

1.4.22   Cardano.BM.Output.Switchboard

Switchboard

```
type SwitchboardMVar = MVar SwitchboardInternal
newtype Switchboard = Switchboard
   {getSB :: SwitchboardMVar}
data SwitchboardInternal = SwitchboardInternal
   {sbQueue :: TBQ.TBQueue NamedLogItem
   }
```

Process incoming messages

Incoming messages are put into the queue, and then processed by the dispatcher.
The queue is initialized and the message dispatcher launched.

```
instance IsEffectuator Switchboard where
   effectuate switchboard item = do
      let writequeue :: TBQ.TBQueue NamedLogItem → NamedLogItem → IO ()
         writequeue q i = do
            nocapacity ← atomically $ TBQ.isFullTBQueue q
            if nocapacity
            then return ()
            else atomically $ TBQ.writeTBQueue q i
      withMVar (getSB switchboard) $ λsb →
         writequeue (sbQueue sb) item
```

Switchboard implements Backend functions

Switchboard is an Declaration of a *Backend*

```
instance IsBackend Switchboard where
   typeof _ = SwitchboardBK
```

```
realize cfg = do
  let spawnDispatcher :: IsBackend e ⇒ Configuration → [(BackendKind, e)] → TBQ.TBQueue NamedLogItem
      spawnDispatcher config backends queue =
        let qProc = do
              nli ← atomically $ TBQ.readTBQueue queue
              case lnItem nli of
                AggregatedMessage _ _aggregated → do
                  selectedBackends ← getBackends config (lnName nli)
                  let dropAggrBackends = filter (≢ AggregationBK) selectedBackends
                  forM_ backends (λ(bek, be) →
                    when (bek ∈ dropAggrBackends) (dispatch nli be))
                  qProc
                KillPill →
                  forM_ backends $ λ(_, backend) →
                    unrealize backend
                _ → do
                  selectedBackends ← getBackends config (lnName nli)
                  forM_ backends $ λ(bek, be) →
                    when (bek ∈ selectedBackends) (dispatch nli be)
                  qProc
              -- if Nothing is in the queue then every backend needs to be terminated
              -- and the dispatcher exits, the Switchboard stops.
            dispatch :: IsEffectuator e ⇒ NamedLogItem → e → IO ()
            dispatch nli backend = effectuate backend nli
        in
        Async.async qProc
    q ← atomically $ TBQ.newTBQueue 2048
    sbref ← newEmptyMVar
    putMVar sbref $ SwitchboardInternal q
    let sb = Switchboard sbref

    backends ← getDefaultBackends cfg
    bs ← setupBackends backends cfg sb []

    _ ← spawnDispatcher cfg bs q
    return sb
  unrealize _ = pure ()
```

Realizing the backends according to configuration

```
setupBackends :: IsBackend e
  ⇒ [BackendKind]
  → Configuration
  → e
  → [(BackendKind, e)]
  → IO [(BackendKind, e)]
setupBackends [] _ _ acc = return acc
setupBackends (bk : bes) c sb acc = do
```

   be′ ← setupBackend′ bk c sb
   setupBackends bes c sb ((bk, be′) : acc)
setupBackend′ :: IsBackend e ⇒ BackendKind → Configuration → e → IO e
setupBackend′ EKGViewBK c _ = Cardano.BM.Output ∘ EKGView.realize c
setupBackend′ AggregationBK c sb = Cardano.BM.Output ∘ Aggregation.realizefrom c sb
setupBackend′ KatipBK c _ = Cardano.BM.Output ∘ Log.realize c
setupBackend′ SwitchboardBK _ _ = error `"cannot instantiate a further Switchboard"`

1.4.23   Cardano.BM.Output.Log

Internal representation

   type *LogMVar = MVar LogInternal*
   newtype *Log = Log*
     {*getK :: LogMVar*}

   data *LogInternal = LogInternal*
     {*kLogEnv :: K.LogEnv*
     , *configuration :: Config.Configuration*}

Log implements *effectuate*

   instance *IsEffectuator Log* **where**
     effectuate katip item = **do**
       c ← withMVar (getK katip) $ *λk → return (configuration k)*
       selscribes ← getScribes c (lnName item)
       forM_ selscribes $ *λsc → passN sc katip item*

Log implements backend functions

   instance *IsBackend Log* **where**
     typeof _ = KatipBK

     realize config = **do**
       let *updateEnv :: K.LogEnv → IO UTCTime → K.LogEnv*
         updateEnv le timer =
           le {K._logEnvTimer = timer, K._logEnvHost = `"hostname"`*}*
         register :: [ScribeDefinition] → K.LogEnv → IO K.LogEnv
         register [ ] le = return le
         register (defsc : dscs) le = **do**
           let *kind = scKind defsc*
             name = scName defsc
             name′ = pack (show kind) <> `"::"` *<> name*
           scr ← createScribe kind name
           register dscs ≪ K.registerScribe name′ scr scribeSettings le
         mockVersion :: Version

```
      mockVersion = Version [0,1,0,0] [ ]
      scribeSettings :: KC.ScribeSettings
      scribeSettings =
        let bufferSize = 5000-- size of the queue (in log items)
        in
        KC.ScribeSettings bufferSize
      createScribe FileTextSK name = mkTextFileScribe (FileDescription $ unpack name) False
      createScribe FileJsonSK name = mkJsonFileScribe (FileDescription $ unpack name) False
      createScribe StdoutSK _ = mkStdoutScribe
      createScribe StderrSK _ = mkStderrScribe
    cfoKey ← Config.getOptionOrDefault config (pack "cfokey") (pack "<unknown>")
    le0 ← K.initLogEnv
           (K.Namespace ["iohk"])
           (fromString $ (unpack cfoKey) <> ":" <> showVersion mockVersion)
    -- request a new time 'getCurrentTime' at most 100 times a second
    timer ← mkAutoUpdate defaultUpdateSettings {updateAction = getCurrentTime, updateFreq = 10000}
    let le1 = updateEnv le0 timer
    scribes ← getSetupScribes config
    le ← register scribes le1
    kref ← newEmptyMVar
    putMVar kref $ LogInternal le config
    return $ Log kref
  unrealize katip = do
    le ← withMVar (getK katip) $ λk → return (kLogEnv k)
    void $ K.closeScribes le


example :: IO ()
example = do
  config ← Config.setup "from_some_path.yaml"
  k ← setup config
  passN (pack (show StdoutSK)) k $ LogNamed
    {lnName = "test"
    ,lnItem = LP $ LogMessage $ LogItem
       {liSelection = Both
       ,liSeverity = Info
       ,liPayload = "Hello!"
       }
    }
  passN (pack (show StdoutSK)) k $ LogNamed
    {lnName = "test"
    ,lnItem = LP $ LogValue "cpu-no" 1
    }


-- useful instances for katip
deriving instance K.ToObject LogObject
deriving instance K.ToObject LogItem
```

```
deriving instance K.ToObject (Maybe LogObject)
instance KC.LogItem LogObject where
    payloadKeys _ _ = KC.AllKeys
instance KC.LogItem LogItem where
    payloadKeys _ _ = KC.AllKeys
instance KC.LogItem (Maybe LogObject) where
    payloadKeys _ _ = KC.AllKeys
```

Log.passN

The following function copies the *NamedLogItem* to the queues of all scribes that match on their name. Compare start of name of scribe to (*show backend* <> "::"). This function is non-blocking.

```
passN :: Text → Log → NamedLogItem → IO ()
passN backend katip namedLogItem = do
    env ← withMVar (getK katip) $ λk → return (kLogEnv k)
    forM_ (Map.toList $ K._logEnvScribes env) $
        λ(scName, (KC.ScribeHandle _ shChan)) →
            -- check start of name to match ScribeKind
            if backend 'isPrefixOf' scName
            then do
                let item = lnItem namedLogItem
                let (sev, msg, payload) = case item of
                        (LP (LogMessage logItem)) →
                            (liSeverity logItem, liPayload logItem, Nothing)
                        (AggregatedMessage name aggregated) →
                            (Info, pack (show name ++ ": " ++ show aggregated), Nothing)
                        _ → (Info, "", (Nothing :: Maybe LogObject))
                if (msg ≡ "") ∧ (isNothing payload)
                then return ()
                else do
                    threadIdText ← KC.mkThreadIdText < $ > myThreadId
                    let ns = lnName namedLogItem
                    itemTime ← env ˆ. KC.logEnvTimer
                    let itemKatip = K.Item {
                        _itemApp = env ˆ. KC.logEnvApp
                        , _itemEnv       = env ˆ. KC.logEnvEnv
                        , _itemSeverity = sev2klog sev
                        , _itemThread   = threadIdText
                        , _itemHost       = env ˆ. KC.logEnvHost
                        , _itemProcess   = env ˆ. KC.logEnvPid
                        , _itemPayload  = payload
                        , _itemMessage = K.logStr msg
                        , _itemTime      = itemTime
                        , _itemNamespace = (env ˆ. KC.logEnvApp) <> (K.Namespace [ns])
                        , _itemLoc        = Nothing
                        }
```

```
            void $ atomically $ KC.tryWriteTBQueue shChan (KC.NewItem itemKatip)
        else return ()
```

Scribes

```
mkStdoutScribe :: IO K.Scribe
mkStdoutScribe = mkTextFileScribeH stdout True

mkStderrScribe :: IO K.Scribe
mkStderrScribe = mkTextFileScribeH stderr True

mkTextFileScribeH :: Handle → Bool → IO K.Scribe
mkTextFileScribeH handler color = do
    mkFileScribeH handler formatter color
  where
    formatter h colorize verbosity item =
        TIO.hPutStrLn h $! toLazyText $ formatItem colorize verbosity item

mkFileScribeH
    :: Handle
    → (forall a ∘ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ())
    → Bool
    → IO K.Scribe
mkFileScribeH h formatter colorize = do
    hSetBuffering h LineBuffering
    locklocal ← newMVar ()
    let logger :: forall a ∘ K.LogItem a ⇒ K.Item a → IO ()
        logger item = withMVar locklocal $ \_ →
          formatter h colorize K.V0 item
    pure $ K.Scribe logger (hClose h)

mkTextFileScribe :: FileDescription → Bool → IO K.Scribe
mkTextFileScribe fdesc colorize = do
    mkFileScribe fdesc formatter colorize
  where
    formatter :: Handle → Bool → K.Verbosity → K.Item a → IO ()
    formatter hdl colorize' v' item = do
        let tmsg = toLazyText $ formatItem colorize' v' item
        TIO.hPutStrLn hdl tmsg

mkJsonFileScribe :: FileDescription → Bool → IO K.Scribe
mkJsonFileScribe fdesc colorize = do
    mkFileScribe fdesc formatter colorize
  where
    formatter :: (K.LogItem a) ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ()
    formatter h _ verbosity item = do
        let tmsg = case KC._itemMessage item of
            K.LogStr "" → K.itemJson verbosity item
            K.LogStr msg → K.itemJson verbosity $
              item {KC._itemMessage = K.logStr ("" :: Text)
                 , KC._itemPayload = LogItem Both Info $ toStrict $ toLazyText msg
```

```
              }
          TIO.hPutStrLn h (encodeToLazyText tmsg)
mkFileScribe
    :: FileDescription
    → (forall a ∘ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ())
    → Bool
    → IO K.Scribe
mkFileScribe fdesc formatter colorize = do
    let prefixDir = prefixPath fdesc
    (createDirectoryIfMissing True prefixDir)
        'catchIO' (prtoutException ("cannot log prefix directory: " ++ prefixDir))
    let fpath = filePath fdesc
    h ← catchIO (openFile fpath WriteMode) $
          λe → do
            prtoutException ("error while opening log: " ++ fpath) e
            -- fallback to standard output in case of exception
            return stdout
    hSetBuffering h LineBuffering
    scribestate ← newMVar h
    let finalizer :: IO ()
        finalizer = withMVar scribestate hClose
    let logger :: forall a ∘ K.LogItem a ⇒ K.Item a → IO ()
        logger item =
          withMVar scribestate $ λhandler →
            formatter handler colorize K.V0 item
    return $ K.Scribe logger finalizer


formatItem :: Bool → K.Verbosity → K.Item a → Builder
formatItem withColor _verb K.Item {..} =
    fromText header <>
    fromText " " <>
    brackets (fromText timestamp) <>
    fromText " " <>
    KC.unLogStr _itemMessage
  where
    header = colorBySeverity _itemSeverity $
      "[" <> mconcat namedcontext <> ":" <> severity <> ":" <> threadid <> "]"
    namedcontext = KC.intercalateNs _itemNamespace
    severity = KC.renderSeverity _itemSeverity
    threadid = KC.getThreadIdText _itemThread
    timestamp = pack $ formatTime defaultTimeLocale tsformat _itemTime
    tsformat :: String
    tsformat = "%F %T%2Q %Z"
    colorBySeverity s m = case s of
      K.EmergencyS → red m
      K.AlertS     → red m
      K.CriticalS  → red m
```

```
                K.ErrorS     → red m
                K.NoticeS    → magenta m
                K.WarningS → yellow m
                K.InfoS      → blue m
                _ → m
            red = colorize "31"
            yellow = colorize "33"
            magenta = colorize "35"
            blue = colorize "34"
            colorize c m
                | withColor = "\ESC[" <> c <> "m" <> m <> "\ESC[0m"
                | otherwise = m
-- translate Severity to Log.Severity
sev2klog :: Severity → K.Severity
sev2klog = λcase
        Debug → K.DebugS
        Info     → K.InfoS
        Notice → K.NoticeS
        Warning → K.WarningS
        Error    → K.ErrorS


data FileDescription = FileDescription {
    filePath :: !FilePath}
    deriving (Show)

prefixPath :: FileDescription → FilePath
prefixPath = takeDirectory ∘ filePath

-- display message and stack trace of exception on stdout
prtoutException :: Exception e ⇒ String → e → IO ()
prtoutException msg e = do
    putStrLn msg
    putStrLn ("exception:  " ++ displayException e)
```

1.4.24   Cardano.BM.Output.EKGView

Structure of EKGView

```
type EKGViewMVar = MVar EKGViewInternal
newtype EKGView = EKGView
    {getEV :: EKGViewMVar}
data EKGViewInternal = EKGViewInternal
    {evGauges :: HM.HashMap Text Gauge.Gauge
    ,evLabels   :: HM.HashMap Text Label.Label
    ,evServer :: Server
    }
```

EKG view is an effectuator

```
instance IsEffectuator EKGView where
  effectuate ekgview item =
    let update :: LogObject → LoggerName → EKGViewInternal → IO (Maybe EKGViewInternal)
        update (LP (LogMessage logitem)) logname ekg@(EKGViewInternal _ labels server) =
          case HM.lookup logname labels of
            Nothing → do
              ekghdl ← getLabel logname server
              Label.set ekghdl (liPayload logitem)
              return $ Just $ ekg {evLabels = HM.insert logname ekghdl labels}
            Just ekghdl → do
              Label.set ekghdl (liPayload logitem)
              return Nothing
        update (LP (LogValue iname value)) logname ekg@(EKGViewInternal gauges _ server) =
          let name = logname <> "." <> iname
          in
          case HM.lookup name gauges of
            Nothing → do
              ekghdl ← getGauge name server
              Gauge.set ekghdl (fromInteger value)
              return $ Just $ ekg {evGauges = HM.insert name ekghdl gauges}
            Just ekghdl → do
              Gauge.set ekghdl (fromInteger value)
              return Nothing
        update _ _ _ = return Nothing
    in do
    ekg ← takeMVar (getEV ekgview)
    upd ← update (lnItem item) (lnName item) ekg
    case upd of
      Nothing → putMVar (getEV ekgview) ekg
      Just ekg' → putMVar (getEV ekgview) ekg'
```

EKGView **implements** *Backend* **functions**

EKGView is an Declaration of a *Backend*

```
instance IsBackend EKGView where
  typeof _ = EKGViewBK
  realize config = do
    evref ← newEmptyMVar
    evport ← getEKGport config
    ehdl ← forkServer "127.0.0.1" evport
    ekghdl ← getLabel "iohk-monitoring version" ehdl
    Label.set ekghdl $ pack (showVersion version)
    putMVar evref $ EKGViewInternal
      {evGauges = HM.empty
```

```
      ,evLabels = HM.empty
      ,evServer = ehdl
      }
    return $ EKGView evref
  unrealize ekgview = do
    ekg ← takeMVar $ getEV ekgview
    killThread $ serverThreadId $ evServer ekg
```

Interactive testing *EKGView*

```
test :: IO ()
test = do
  c ← Cardano.BM.Configuration.setup "test/config.yaml"
  ev ← Cardano.BM.Output ∘ EKGView.realize c

  effectuate ev $ LogNamed "test.questions" (LP (LogValue "answer" 42))
  effectuate ev $ LogNamed "test.monitor023" (LP (LogMessage (LogItem Public Warning "!!!! ALARM !!!
```

1.4.25   Cardano.BM.Output.Aggregation

Internal representation

```
type AggregationMVar = MVar AggregationInternal
newtype Aggregation = Aggregation
  {getAg :: AggregationMVar}
data AggregationInternal = AggregationInternal
  {agQueue :: TBQ.TBQueue (Maybe NamedLogItem)
  ,agDispatch :: Async.Async ()
  }
```

Relation from context name to aggregated statistics

We keep the aggregated values (**??**) for a named context in a *HashMap*.

```
type AggregationMap = HM.HashMap Text Aggregated
```

Aggregation **implements** *effectuate*

Aggregation is an **??** Enter the log item into the *Aggregation* queue.

```
instance IsEffectuator Aggregation where
  effectuate agg item = do
    ag ← readMVar (getAg agg)
    atomically $ TBQ.writeTBQueue (agQueue ag) $ Just item
```

Aggregation **implements** *Backend* **functions**

Aggregation is an Declaration of a *Backend*

```
instance IsBackend Aggregation where
  typeof _ = AggregationBK
  realize _ = error "Aggregation cannot be instantiated by 'realize'"
  realizefrom _ switchboard = do
    aggref ← newEmptyMVar
    aggregationQueue ← atomically $ TBQ.newTBQueue 2048
    dispatcher ← spawnDispatcher HM.empty aggregationQueue switchboard
    putMVar aggref $ AggregationInternal aggregationQueue dispatcher
    return $ Aggregation aggref

  unrealize aggregation = do
    let clearMVar :: MVar a → IO ()
        clearMVar = void ∘ tryTakeMVar
    (dispatcher, queue) ← withMVar (getAg aggregation) (λag →
      return (agDispatch ag, agQueue ag))
    -- send terminating item to the queue
    atomically $ TBQ.writeTBQueue queue Nothing
    -- wait for the dispatcher to exit
    res ← Async.waitCatch dispatcher
    either throwM return res
    (clearMVar ∘ getAg) aggregation
```

Asynchrouniously reading log items from the queue and their processing

```
spawnDispatcher :: IsEffectuator e
            ⇒ AggregationMap
            → TBQ.TBQueue (Maybe NamedLogItem)
            → e
            → IO (Async.Async ())
spawnDispatcher aggMap aggregationQueue switchboard = Async.async $ qProc aggMap
  where
    qProc aggregatedMap = do
      maybeItem ← atomically $ TBQ.readTBQueue aggregationQueue
      case maybeItem of
        Just item → do
          let (updatedMap, msgs) =
              update (lnItem item) (lnName item) aggregatedMap
          sendAggregated msgs switchboard (lnName item)
          qProc updatedMap
        Nothing → return ()

    update :: LogObject
      → LoggerName
      → HM.HashMap Text Aggregated
```

```
                  → (HM.HashMap Text Aggregated, [LogObject])
update (LP (LogValue iname value)) logname agmap =
   let name = logname <> "." <> iname
       maybeAggregated = updateAggregation value $ HM.lookup name agmap
       aggregatedMessage = case maybeAggregated of
          Nothing →
             [ ]
          Just aggregated →
             [AggregatedMessage iname aggregated]
   in
   -- use of HM.alter so that in future we can clear the Agrregated
   -- by using as alter's arg a function which returns Nothing.
   (HM.alter (const $ maybeAggregated) name agmap, aggregatedMessage)
update (ObserveDiff counterState) logname agmap =
   let
       counters = csCounters counterState
       (mapNew, msgs) = updateCounter counters logname agmap [ ]
   in
       (mapNew, reverse msgs)
-- TODO for text messages aggregate on delta of timestamps
update _ _ agmap = (agmap, [ ])

updateCounter :: [Counter]
   → LoggerName
   → HM.HashMap Text Aggregated
   → [LogObject]
   → (HM.HashMap Text Aggregated, [LogObject])
updateCounter [ ] _ aggrMap msgs = (aggrMap, msgs)
updateCounter (counter : cs) logname aggrMap msgs =
   let
       name = cName counter
       fullname = logname <> "." <> name
       maybeAggregated = updateAggregation (cValue counter) $ HM.lookup fullname aggrMap
       aggregatedMessage = case maybeAggregated of
          Nothing →
             error "This should not have happened!"
          Just aggregated →
             AggregatedMessage ((nameCounter counter) <> "." <> name) aggregated
       updatedMap = HM.alter (const $ maybeAggregated) fullname aggrMap
   in
       updateCounter cs logname updatedMap (aggregatedMessage : msgs)
sendAggregated :: IsEffectuator e ⇒ [LogObject] → e → Text → IO ()
sendAggregated [ ] _ _ = return ()
sendAggregated (aggregatedMsg@(AggregatedMessage _ _) : ms) sb logname = do
   -- forward the aggregated message to Switchboard
   effectuate sb $
       LogNamed
          {lnName = logname <> ".aggregated"
```

```
                    , lnItem = aggregatedMsg
                    }
        sendAggregated ms sb logname
-- ingnore all other messages that are not of type AggregatedMessage
sendAggregated (_ : ms) sb logname =
        sendAggregated ms sb logname
```