

Cardano.BM - benchmarking and logging

Alexander Diemand

Andreas Triantafyllos

November 2018

Abstract

This is a framework that combines logging, benchmarking and monitoring. Complex evaluations of STM or monadic actions can be observed from outside while reading operating system counters before and after, and calculating their differences, thus relating resource usage to such actions. Through interactive configuration, the runtime behaviour of logging or the measurement of resource usage can be altered. Further reduction in logging can be achieved by redirecting log messages to an aggregation function which will output the running statistics with less frequency than the original message.

Contents

1	Logging, benchmarking and monitoring	3
1.1	Overview	3
1.2	Introduction	3
1.2.1	Logging with Trace	3
1.2.2	Setup procedure	3
1.2.3	Measuring <i>Observables</i>	3
1.2.4	Information reduction in Aggregation	3
1.2.5	Output selection	3
1.2.6	Monitoring	3
1.3	Examples	3
1.3.1	Observing evaluation of a STM action	3
1.3.2	Observing evaluation of a monad action	3
1.3.3	Simple example showing plain logging	3
1.3.4	Complex example showing logging, aggregation of log items, and observing <i>IO</i> actions	5
1.4	Code listings	11
1.4.1	Cardano.BM.Observer.STM	11
1.4.2	Cardano.BM.Observer.Monadlic	12
1.4.3	BaseTrace	15
1.4.4	Cardano.BM.Trace	16
1.4.5	Cardano.BM.Setup	22
1.4.6	Cardano.BM.Counters	23
1.4.7	Cardano.BM.Counters.Common	24
1.4.8	Cardano.BM.Counters.Dummy	25
1.4.9	Cardano.BM.Counters.Linux	25
1.4.10	Cardano.BM.Data.Aggregated	32
1.4.11	Cardano.BM.Data.Backend	37
1.4.12	Cardano.BM.Data.Configuration	37
1.4.13	Cardano.BM.Data.Counter	39
1.4.14	Cardano.BM.Data.LogItem	40
1.4.15	Cardano.BM.Data.Observable	42
1.4.16	Cardano.BM.Data.Output	42
1.4.17	Cardano.BM.Data.Severity	43
1.4.18	Cardano.BM.Data.SubTrace	43
1.4.19	Cardano.BM.Data.Trace	44
1.4.20	Cardano.BM.Configuration	44
1.4.21	Cardano.BM.Configuration.Model	45
1.4.22	Cardano.BM.Output.Switchboard	52

1.4.23	Cardano.BM.Output.Log	55
1.4.24	Cardano.BM.Output.EKGView	61
1.4.25	Cardano.BM.Output.Aggregation	65
2	Testing	71
2.1	Test coverage	71
2.2	Test main entry point	71
2.3	Test case generation	73
2.3.1	instance Arbitrary Aggregated	73
2.4	Tests	74
2.4.1	Testing aggregation	74
2.4.2	STM	76
2.4.3	Trace	76
2.4.4	Testing configuration	86

Chapter 1

Logging, benchmarking and monitoring

1.1 Overview

In figure 1.1 we display the relationships among modules in *Cardano.BM*. The arrows indicate import of a module. The arrows with a triangle at one end would signify "inheritance" in object-oriented programming, but we use it to show that one module replaces the other in the namespace, thus refines its interface.

1.2 Introduction

1.2.1 Logging with **Trace**

1.2.2 Setup procedure

Hierarchy of **Traces**

1.2.3 Measuring *Observables*

1.2.4 Information reduction in **Aggregation**

1.2.5 Output selection

1.2.6 Monitoring

1.3 Examples

1.3.1 Observing evaluation of a STM action

1.3.2 Observing evaluation of a monad action

1.3.3 Simple example showing plain logging

```
{-# LANGUAGE OverloadedStrings #-}
module Main
  (main)
  where
import Control.Concurrent (threadDelay)
import Cardano.BM.Configuration.Static (defaultConfigStdout)
```

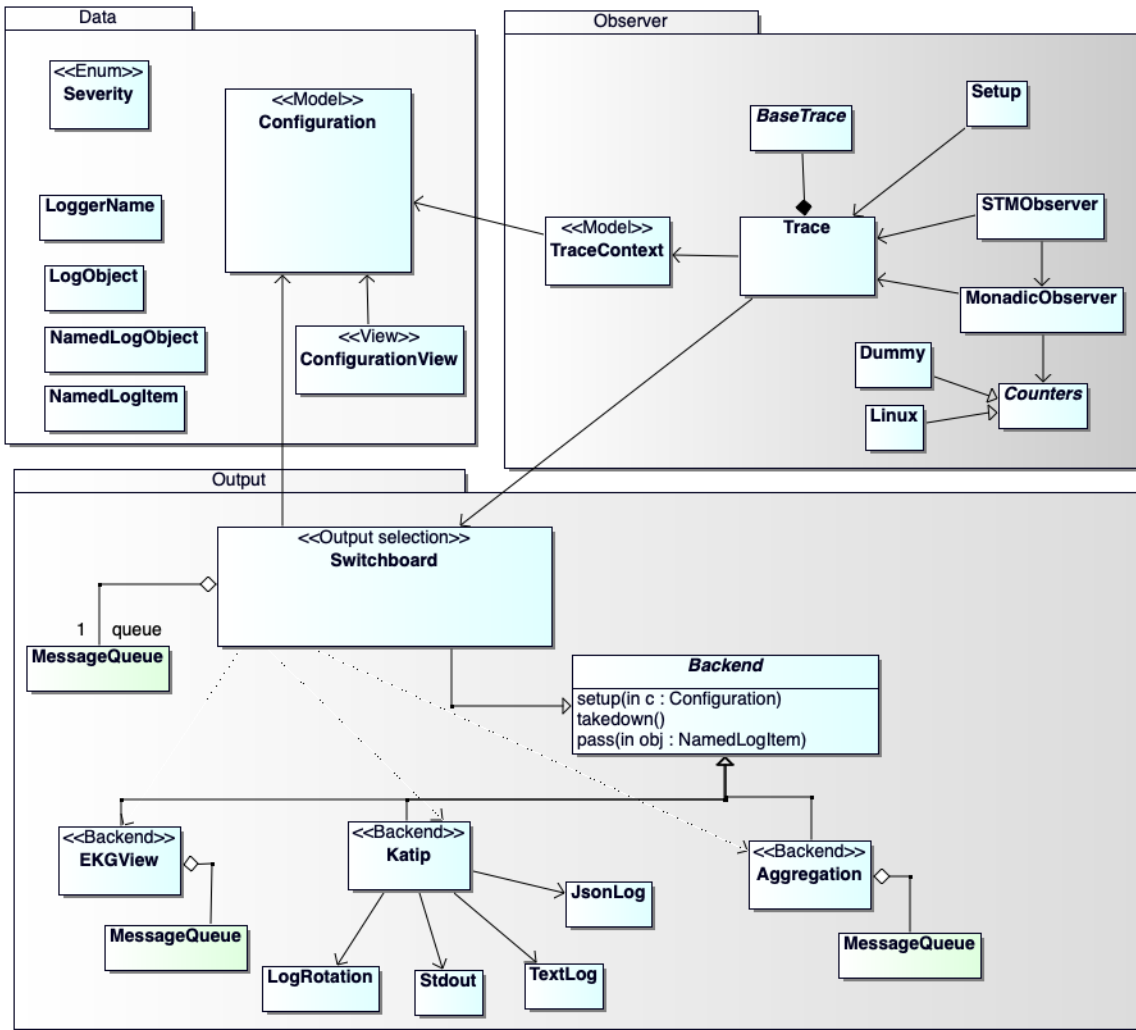


Figure 1.1: Overview of module relationships

```

import Cardano.BM.Setup (setupTrace)
import Cardano.BM.Trace (logDebug, logError, logInfo, logNotice,
                        logWarning)

main :: IO ()
main = do
  c ← defaultConfigStdout
  tr ← setupTrace (Right c) "simple"
  logDebug tr "this is a debug message"
  logInfo tr "this is an information."
  logNotice tr "this is a notice!"
  logWarning tr "this is a warning!"
  logError tr "this is an error!"
  threadDelay 80000

```



Figure 1.2: Setup procedure

```
return ()
```

1.3.4 Complex example showing logging, aggregation of log items, and observing IO actions

Module header and import directives

```

{-# LANGUAGE CPP #-}
{-# LANGUAGE OverloadedStrings #-}
# if defined (linux_HOST_OS)
# define LINUX
# endif
module Main
  (main)
  where
import Control.Concurrent (threadDelay)
import qualified Control.Concurrent.Async as Async
import Control.Monad (forM, forM_)
import GHC.Conc.Sync (STM, TVar, atomically, newTVar, readTVar, writeTVar)
import Data.Text (pack)
# ifdef LINUX
import qualified Data.ByteString.Char8 as BS8
import Network.Download (openURI)
# endif
import System.Random
import qualified Cardano.BM.Configuration.Model as CM
import Cardano.BM.Data.Aggregated (Measurable (..))
import Cardano.BM.Data.AggregatedKind
import Cardano.BM.Data.BackendKind
import Cardano.BM.Data.LogItem

```

```

import Cardano.BM.Data.Observable
import Cardano.BM.Data.Output
import Cardano.BM.Data.Severity
import Cardano.BM.Data.SubTrace
import Cardano.BM.Observer.Monad (bracketObserveIO)
import qualified Cardano.BM.Observer.STM as STM
import Cardano.BM.Setup
import Cardano.BM.Trace

```

Define configuration

The output can be viewed in EKG on <http://localhost:12789>.

```

config :: IO CM.Configuration
config = do
  c ← CM.empty
  CM.setMinSeverity c Debug
  CM.setSetupBackends c [KatipBK, AggregationBK, EKGViewBK]
  CM.setDefaultBackends c [KatipBK]
  CM.setSetupScribes c [ScribeDefinition {
    scName = "stdout"
    ,scKind = StdoutSK
    ,scRotation = Nothing
  }
    ,ScribeDefinition {
    scName = "out.odd.json"
    ,scKind = FileJsonSK
    ,scRotation = Nothing
  }
    ,ScribeDefinition {
    scName = "out.even.json"
    ,scKind = FileJsonSK
    ,scRotation = Nothing
  }
    ,ScribeDefinition {
    scName = "out.txt"
    ,scKind = FileTextSK
    ,scRotation = Nothing
  }
  ]
  CM.setDefaultScribes c ["StdoutSK::stdout"]
  CM.setScribes c "complex.random" (Just ["StdoutSK::stdout", "FileTextSK::out.txt"])
  CM.setScribes c "#aggregated.complex.random" (Just ["StdoutSK::stdout"])
  forM_ [(1 :: Int)..10] $ \x →
    if odd x
    then
      CM.setScribes c ("#aggregation.complex.observeSTM." <> (pack $ show x)) $ Just ["FileJsonSK::ou"]
    else

```



```

    CM.setScribes c ("#aggregation.complex.observeSTM." <> (pack $ show x)) $ Just ["FileJsonSK::ou
# ifdef LINUX
    CM.setSubTrace c "complex.observeDownload" (Just $ ObservableTrace [IOStats])
# endif
    CM.setSubTrace c "complex.random" (Just $ TeeTrace "ewma")
    CM.setSubTrace c "#ekgview"
      (Just $ FilterTrace [(Drop (StartsWith "#ekgview.#aggregation.complex.random"),
        Unhide [(EndsWith ".count"),
          (EndsWith ".avg"),
          (EndsWith ".mean")]),
        (Drop (StartsWith "#ekgview.#aggregation.complex.observeIO"),
          Unhide [(Contains "diff.RTS.cpuNs.timed.")]),
        (Drop (StartsWith "#ekgview.#aggregation.complex.observeSTM"),
          Unhide [(Contains "diff.RTS.gcNum.timed.")]),
        (Drop (StartsWith "#ekgview.#aggregation.complex.message"),
          Unhide [(Contains ".timed.m")])
        ])
    CM.setSubTrace c "complex.observeIO" (Just $ ObservableTrace [GhcRtsStats, MemoryStats])
forM_ [(1 :: Int)..10] $ \x →
    CM.setSubTrace
      c
      ("complex.observeSTM." <> (pack $ show x))
      (Just $ ObservableTrace [GhcRtsStats, MemoryStats])
    CM.setBackends c "complex.message" (Just [AggregationBK, KatipBK])
    CM.setBackends c "complex.random" (Just [AggregationBK, KatipBK])
    CM.setBackends c "complex.random.ewma" (Just [AggregationBK])
    CM.setBackends c "complex.observeIO" (Just [AggregationBK])
forM_ [(1 :: Int)..10] $ \x → do
    CM.setBackends c
      ("complex.observeSTM." <> (pack $ show x))
      (Just [AggregationBK])
    CM.setBackends c
      ("#aggregation.complex.observeSTM." <> (pack $ show x))
      (Just [EKGVViewBK])
    CM.setAggregatedKind c "complex.random.rr" (Just StatsAK)
    CM.setAggregatedKind c "complex.random.ewma.rr" (Just (EwmaAK 0.42))
    CM.setBackends c "#aggregation.complex.message" (Just [EKGVViewBK])
    CM.setBackends c "#aggregation.complex.observeIO" (Just [EKGVViewBK])
    CM.setBackends c "#aggregation.complex.random" (Just [EKGVViewBK])
    CM.setBackends c "#aggregation.complex.random.ewma" (Just [EKGVViewBK])
    CM.setEKGVport c 12789
return c

```

Thread that outputs a random number to a `Trace`

```

randomThr :: Trace IO → IO (Async.Async ())
randomThr trace = do
  logInfo trace "starting random generator"
  trace' ← subTrace "random" trace
  proc ← Async.async (loop trace')
  return proc
where
  loop tr = do
    threadDelay 500000 -- 0.5 second
    num ← randomRIO (42 - 42, 42 + 42) :: IO Double
    lo ← LogObject <$> mkLOMeta <*> pure (LogValue "rr" (PureD num))
    traceNamedObject tr lo
    loop tr

```

Thread that observes an IO action

```

observeIO :: Trace IO → IO (Async.Async ())
observeIO trace = do
  logInfo trace "starting observer"
  proc ← Async.async (loop trace)
  return proc
where
  loop tr = do
    threadDelay 5000000 -- 5 seconds
    _ ← bracketObserveIO tr "observeIO" $ do
      num ← randomRIO (100000, 200000) :: IO Int
      ls ← return $ reverse $ init $ reverse $ 42 : [1..num]
      pure $ const ls ()
    loop tr

```

Thread that observes an IO action which downloads a txt in order to observe the I/O statistics

disabled for now! on Mac OSX this function was blocking all IO.

```

observeDownload :: Trace IO → IO ()
observeDownload trace = loop trace
where
  loop tr = do
    threadDelay 10000000 -- 10 seconds
    tr' ← appendName "observeDownload" tr
    bracketObserveIO tr' "" $ do
      license ← openURI "http://www.gnu.org/licenses/gpl.txt"
      case license of
        Right bs → logNotice tr' $ "downloaded " <> BS8.length bs <> " bytes"

```

```

    Left e → logError tr' e
  threadDelay 500000 -- .5 second
  pure ()
loop tr

```

Threads that observe STM actions on the same TVar

```

observeSTM :: Trace IO → IO [Async.Async ()]
observeSTM trace = do
  logInfo trace "starting STM observer"
  tvar ← atomically $ newTVar ([1..1000] :: [Int])
  -- spawn 10 threads
  proc ← forM [(1 :: Int)..10] $ \x → Async.async (loop trace tvar (pack $ show x))
  return proc
where
  loop tr tvarlist name = do
    threadDelay 10000000 -- 10 seconds
    STM.bracketObserveIO tr ("observeSTM." <> name) (stmAction tvarlist)
    loop tr tvarlist name
stmAction :: TVar [Int] → STM ()
stmAction tvarlist = do
  list ← readTVar tvarlist
  writeTVar tvarlist $ reverse $ init $ reverse $ list
  pure ()

```

Thread that observes an IO action which downloads a txt in order to observe the I/O statistics

```

#ifdef LINUX
observeDownload :: Trace IO → IO (Async.Async ())
observeDownload trace = do
  proc ← Async.async (loop trace)
  return proc
where
  loop tr = do
    threadDelay 1000000 -- 1 second
    tr' ← appendName "observeDownload" tr
    bracketObserveIO tr' "" $ do
      license ← openURI "http://www.gnu.org/licenses/gpl.txt"
      case license of
        Right bs → logNotice tr' $ pack $ BS8.unpack bs
        Left _ → return ()
      threadDelay 50000 -- .05 second
      pure ()
    loop tr
#endif

```

Thread that periodically outputs a message

```

msgThr :: Trace IO → IO (Async.Async ())
msgThr trace = do
  logInfo trace "start messaging .."
  trace' ← subTrace "message" trace
  Async.async (loop trace')
  where
    loop tr = do
      threadDelay 3000000 -- 3 seconds
      logNotice tr "N O T I F I C A T I O N ! ! !"
      logDebug tr "a detailed debug message."
      logError tr "Boooooomm .."
      loop tr

```

Main entry point

```

main :: IO ()
main = do
  -- create configuration
  c ← config
  -- create initial top-level Trace
  tr ← setupTrace (Right c) "complex"
  logNotice tr "starting program; hit CTRL-C to terminate"
  logInfo tr "watch its progress on http://localhost:12789"
  {-start thread sending unbounded sequence of random numbers to a trace which aggregates them into a
  procRandom ← randomThr tr
  -- start thread endlessly reversing lists of random length
  procObsvIO ← observeIO tr
  -- start threads endlessly observing STM actions operating on the same TVar
  procObsvSTMs ← observeSTM tr
  # ifdef LINUX
  -- start thread endlessly which downloads sth in order to check the I/O usage
  procObsvDownload ← observeDownload tr
  # endif
  -- start a thread to output a text messages every n seconds
  procMsg ← msgThr tr
  -- wait for message thread to finish, ignoring any exception
  _ ← Async.waitCatch procMsg
  # ifdef LINUX
  -- wait for download thread to finish, ignoring any exception
  _ ← Async.waitCatch procObsvDownload
  # endif
  -- wait for observer thread to finish, ignoring any exception

```

```

_ ← forM procObsvSTMs Async.waitCatch
-- wait for observer thread to finish, ignoring any exception
_ ← Async.waitCatch procObsvIO
-- wait for random thread to finish, ignoring any exception
_ ← Async.waitCatch procRandom
return ()

```

1.4 Code listings

1.4.1 Cardano.BM.Observer.STM

```

stmWithLog :: STM.STM (t,[LogObject]) → STM.STM (t,[LogObject])
stmWithLog action = action

```

Observe STM action in a named context

With given name, create a **SubTrace** according to **Configuration** and run the passed STM action on it.

```

bracketObserveIO :: Trace IO → Text → STM.STM t → IO t
bracketObserveIO logTrace0 name action = do
  logTrace ← subTrace name logTrace0
  let subtrace = typeofTrace logTrace
  bracketObserveIO' subtrace logTrace action
where
  bracketObserveIO' :: SubTrace → Trace IO → STM.STM t → IO t
  bracketObserveIO' NoTrace _ act =
    STM.atomically act
  bracketObserveIO' subtrace logTrace act = do
    mCountersid ← observeOpen subtrace logTrace
    -- run action; if an exception is caught will be logged and rethrown.
    t ← (STM.atomically act) 'catch' (λ(e :: SomeException) → (logError logTrace (pack (show e)) >> throwM e))
    case mCountersid of
      Left openException →
        -- since observeOpen faced an exception there is no reason to call observeClose
        -- however the result of the action is returned
        logNotice logTrace ("ObserveOpen: " <> pack (show openException))
      Right countersid → do
        res ← observeClose subtrace logTrace countersid [ ]
        case res of
          Left ex → logNotice logTrace ("ObserveClose: " <> pack (show ex))
          _ → pure ()
    pure t

```

Observe STM action in a named context and output captured log items

The STM action might output messages, which after "success" will be forwarded to the logging trace. Otherwise, this function behaves the same as **Observe STM action in a named context**.

```

bracketObserveLogIO :: Trace IO → Text → STM.STM (t, [LogObject]) → IO t
bracketObserveLogIO logTrace0 name action = do
  logTrace ← subTrace name logTrace0
  let subtrace = typeofTrace logTrace
  bracketObserveLogIO' subtrace logTrace action
where
  bracketObserveLogIO' :: SubTrace → Trace IO → STM.STM (t, [LogObject]) → IO t
  bracketObserveLogIO' NoTrace _ act = do
    (t, _) ← STM.atomically $ stmWithLog act
    pure t
  bracketObserveLogIO' subtrace logTrace act = do
    mCountersid ← observeOpen subtrace logTrace
    -- run action, return result and log items; if an exception is
    -- caught will be logged and rethrown.
    (t, as) ← (STM.atomically $ stmWithLog act) 'catch'
      (λ(e :: SomeException) → (logError logTrace (pack (show e)) >> throwM e))
    case mCountersid of
      Left openException →
        -- since observeOpen faced an exception there is no reason to call observeClose
        -- however the result of the action is returned
        logNotice logTrace ("ObserveOpen: " <> pack (show openException))
      Right countersid → do
        res ← observeClose subtrace logTrace countersid as
        case res of
          Left ex → logNotice logTrace ("ObserveClose: " <> pack (show ex))
          _ → pure ()
    pure t

```

1.4.2 Cardano.BM.Observer.Monad

Monadic.bracketObserverIO

Observes an IO action and adds a name to the logger name of the passed in **Trace**. An empty **Text** leaves the logger name untouched.

Microbenchmarking steps:

1. Create a *trace* which will have been configured to observe things besides logging.

```

import qualified Cardano.BM.Configuration.Model as CM
○○○
c ← config
trace@(ctx, _) ← setupTrace (Right c) "demo-playground"

```

```

where
  config :: IO CM.Configuration
  config = do
    c ← CM.empty
    CM.setMinSeverity c Debug
    CM.setSetupBackends c [KatipBK, AggregationBK]
    CM.setDefaultBackends c [KatipBK, AggregationBK]
    CM.setSetupScribes c [ScribeDefinition {
      scName = "stdout"
      , scKind = StdoutSK
      , scRotation = Nothing
    }
    ]
    CM.setDefaultScribes c ["StdoutSK::stdout"]
  return c

```

2. *c* is the **Configuration** of *trace*. In order to enable the collection and processing of measurements (min, max, mean, std-dev) *AggregationBK* is needed.

```
CM.setDefaultBackends c [KatipBK, AggregationBK]
```

in a configuration file (YAML) means

```

defaultBackends:
- KatipBK
- AggregationBK

```

3. Set the measurements that you want to take by changing the configuration of the *trace* using *setSubTrace*, in order to declare the namespace where we want to enable the particular measurements and the list with the kind of measurements.

```

CM.setSubTrace
  (configuration ctx)
  "demo-playground.submit-tx"
  (Just $ ObservableTrace observablesSet)
where
  observablesSet = [MonotonicClock, MemoryStats]

```

4. Find an action to measure. e.g.:

```
runProtocolWithPipe x hdl proto 'catch' (λProtocolStopped → return ())
```

and use **bracketObserveIO**. e.g.:

```

bracketObserveIO trace "submit-tx" $
  runProtocolWithPipe x hdl proto 'catch' (λProtocolStopped → return ())

```

```

bracketObserveIO :: Trace IO → Text → IO t → IO t
bracketObserveIO logTrace0 name action = do

```

```

logTrace ← subTrace name logTrace0
bracketObserveIO' (typeofTrace logTrace) logTrace action
where
bracketObserveIO' :: SubTrace → Trace IO → IO t → IO t
bracketObserveIO' NoTrace _ act = act
bracketObserveIO' subtrace logTrace act = do
  mCountersid ← observeOpen subtrace logTrace
  -- run action; if an exception is caught will be logged and rethrown.
  t ← act 'catch' (λe :: SomeException) → (logError logTrace (pack (show e)) >> throwM e))
  case mCountersid of
    Left openException →
      -- since observeOpen faced an exception there is no reason to call observeClose
      -- however the result of the action is returned
      logNotice logTrace ("ObserveOpen: " <> pack (show openException))
    Right countersid → do
      res ← observeClose subtrace logTrace countersid [ ]
      case res of
        Left ex → logNotice logTrace ("ObserveClose: " <> pack (show ex))
        _ → pure ()
  pure t

```

Monadic.bracketObserverM

Observes a *MonadIO* $m \Rightarrow m$ action and adds a name to the logger name of the passed in **Trace**. An empty *Text* leaves the logger name untouched.

```

bracketObserveM :: (MonadCatch m, MonadIO m) ⇒ Trace IO → Text → m t → m t
bracketObserveM logTrace0 name action = do
  logTrace ← liftIO $ subTrace name logTrace0
  bracketObserveM' (typeofTrace logTrace) logTrace action
where
bracketObserveM' :: (MonadCatch m, MonadIO m) ⇒ SubTrace → Trace IO → m t → m t
bracketObserveM' NoTrace _ act = act
bracketObserveM' subtrace logTrace act = do
  mCountersid ← liftIO $ observeOpen subtrace logTrace
  -- run action; if an exception is caught will be logged and rethrown.
  t ← act 'catch'
    (λe :: SomeException) → (liftIO (logError logTrace (pack (show e)) >> throwM e)))
  case mCountersid of
    Left openException →
      -- since observeOpen faced an exception there is no reason to call observeClose
      -- however the result of the action is returned
      liftIO $ logNotice logTrace ("ObserveOpen: " <> pack (show openException))
    Right countersid → do
      res ← liftIO $ observeClose subtrace logTrace countersid [ ]
      case res of
        Left ex → liftIO (logNotice logTrace ("ObserveClose: " <> pack (show ex)))

```



```

    - → pure ()
  pure t

```

observerOpen

```

observeOpen :: SubTrace → Trace IO → IO (Either SomeException CounterState)
observeOpen subtrace logTrace = (do
  identifier ← newUnique
  -- take measurement
  counters ← readCounters subtrace
  let state = CounterState identifier counters
  if counters == []
  then return ()
  else do
    -- send opening message to Trace
    traceNamedObject logTrace <<
      LogObject <$> mkLOMeta <*> pure (ObserveOpen state)
    return (Right state)) 'catch' (return ∘ Left)

```

observeClose

```

observeClose :: SubTrace → Trace IO → CounterState → [LogObject] → IO (Either SomeException ())
observeClose subtrace logTrace initState logObjects = (do
  let identifier = csIdentifier initState
      initialCounters = csCounters initState
  -- take measurement
  counters ← readCounters subtrace
  if counters == []
  then return ()
  else do
    mle ← mkLOMeta
    -- send closing message to Trace
    traceNamedObject logTrace $
      LogObject mle (ObserveClose (CounterState identifier counters))
    -- send diff message to Trace
    traceNamedObject logTrace $
      LogObject mle (ObserveDiff (CounterState identifier (diffCounters initialCounters counters)))
    -- trace the messages gathered from inside the action
    forM_ logObjects $ traceNamedObject logTrace
    return (Right ())) 'catch' (return ∘ Left)

```

1.4.3 BaseTrace

Contravariant

A covariant is a functor: $F A \rightarrow F B$

A contravariant is a functor: $F B \rightarrow F A$

$Op\ a\ b$ implements the inverse to 'arrow' " $getOp :: b \rightarrow a$ ", which when applied to a **BaseTrace** of type " $Op\ (m\ ())\ s$ ", yields " $s \rightarrow m\ ()$ ". In our case, Op accepts an action in a monad m with input type **LogNamed LogObject** (see 'Trace').

```
newtype BaseTrace m s = BaseTrace {runTrace :: Op (m ()) s}
```

contramap

A covariant functor defines the function " $fmap :: (a \rightarrow b) \rightarrow f\ a \rightarrow f\ b$ ". In case of a contravariant functor, it is the dual function " $contramap :: (a \rightarrow b) \rightarrow f\ b \rightarrow f\ a$ " which is defined.

In the following instance, $runTrace$ extracts type " $Op\ (m\ ())\ s$ " to which $contramap$ applies f , thus " $f\ s \rightarrow m\ ()$ ". The constructor **BaseTrace** restores " $Op\ (m\ ())\ (f\ s)$ ".

```
instance Contravariant (BaseTrace m) where
  contramap f = BaseTrace  $\circ$  contramap f  $\circ$  runTrace
```

traceWith

Accepts a **Trace** and some payload s . First it gets the contravariant from the **Trace** as type " $Op\ (m\ ())\ s$ " and, after " $getOp :: b \rightarrow a$ " which translates to " $s \rightarrow m\ ()$ ", calls the action on the **LogNamed LogObject**.

```
traceWith :: BaseTrace m s  $\rightarrow$  s  $\rightarrow$  m ()
traceWith = getOp  $\circ$  runTrace
```

natTrace

Natural transformation from monad m to monad n .

```
natTrace :: (forall x  $\circ$  m x  $\rightarrow$  n x)  $\rightarrow$  BaseTrace m s  $\rightarrow$  BaseTrace n s
natTrace nat (BaseTrace (Op tr)) = BaseTrace $ Op $ nat  $\circ$  tr
```

noTrace

A **Trace** that discards all inputs.

```
noTrace :: Applicative m  $\Rightarrow$  BaseTrace m a
noTrace = BaseTrace $ Op $ const (pure ())
```

1.4.4 Cardano.BM.Trace

Utilities

Natural transformation from monad m to monad n .

```
natTrace :: (forall x  $\circ$  m x  $\rightarrow$  n x)  $\rightarrow$  Trace m  $\rightarrow$  Trace n
natTrace nat (ctx, trace) = (ctx, BaseTrace.natTrace nat trace)
```

Access type of **Trace**.

```
typeofTrace :: Trace m → SubTrace
typeofTrace (ctx, _) = tracetype ctx
```

Update type of **Trace**.

```
updateTracetype :: SubTrace → Trace m → Trace m
updateTracetype subtr (ctx, tr) = (ctx { tracetype = subtr }, tr)
```

Enter new named context

The context name is created and checked that its size is below a limit (currently 80 chars). The minimum severity that a log message must be labelled with is looked up in the configuration and recalculated.

```
appendName :: MonadIO m ⇒ LoggerName → Trace m → m (Trace m)
appendName name (ctx, trace) = do
  let prevLoggerName = loggerName ctx
      prevMinSeverity = minSeverity ctx
      newLoggerName = appendWithDot prevLoggerName name
      globMinSeverity ← liftIO $ Config.minSeverity (configuration ctx)
      namedSeverity ← liftIO $ Config.inspectSeverity (configuration ctx) newLoggerName
  case namedSeverity of
    Nothing → return (ctx { loggerName = newLoggerName }, trace)
    Just sev → return (ctx { loggerName = newLoggerName
                          , minSeverity = max (max sev prevMinSeverity) globMinSeverity }
                      , trace)

appendWithDot :: LoggerName → LoggerName → LoggerName
appendWithDot "" newName = T.take 80 newName
appendWithDot xs "" = xs
appendWithDot xs newName = T.take 80 $ xs <> " ." <> newName
```

Contramap a trace and produce the naming context

```
named :: BaseTrace.BaseTrace m (LogNamed i) → LoggerName → BaseTrace.BaseTrace m i
named trace name = contramap (LogNamed name) trace
```

Trace a **LogObject** through

```
traceNamedObject
  :: MonadIO m
  ⇒ Trace m
  → LogObject
  → m ()
traceNamedObject trace@(ctx, logTrace) lo@(LogObject _ lc) = do
```

```

let lname = loggerName ctx
doOutput ← case (typeofTrace trace) of
  FilterTrace filters →
    case lc of
      LogValue lname _ →
        return $ evalFilters filters (lname <> "." <> lname)
      _ →
        return $ evalFilters filters lname
  TeeTrace secName → do
    -- create a newly named copy of the LogObject
    BaseTrace.traceWith (named logTrace (lname <> "." <> secName)) lo
    return True
  _ → return True
if doOutput
then BaseTrace.traceWith (named logTrace lname) lo
else return ()

```

Evaluation of FilterTrace

A filter consists of a *DropName* and a list of *UnhideNames*. If the context name matches the *DropName* filter, then at least one of the *UnhideNames* must match the name to have the evaluation of the filters return *True*.

```

evalFilters :: [(DropName, UnhideNames)] → LoggerName → Bool
evalFilters fs nm =
  all (λ(no, yes) → if (dropFilter nm no) then (unhideFilter nm yes) else True) fs
where
  dropFilter :: LoggerName → DropName → Bool
  dropFilter name (Drop sel) = {-not -} (matchName name sel)
  unhideFilter :: LoggerName → UnhideNames → Bool
  unhideFilter _ (Unhide [ ]) = False
  unhideFilter name (Unhide us) = any (λsel → matchName name sel) us
  matchName :: LoggerName → NameSelector → Bool
  matchName name (Exact name') = name ≡ name'
  matchName name (StartsWith prefix) = T.isPrefixOf prefix name
  matchName name (EndsWith postfix) = T.isSuffixOf postfix name
  matchName name (Contains name') = T.isInfixOf name' name

```

Concrete Trace on stdout

This function returns a trace with an action of type "*(LogNamed LogObject) → IO ()*" which will output a text message as text and all others as JSON encoded representation to the console.

TODO remove locallock

```

locallock :: MVar ()
locallock = unsafePerformIO $ newMVar ()

```

```

stdoutTrace :: TraceNamed IO
stdoutTrace = BaseTrace.BaseTrace $ Op $ λ(LogNamed logname (LogObject _ lc)) →
  withMVar locallock $ \_ →
    case lc of
      (LogMessage logItem) →
        output logname $ liPayload logItem
      obj →
        output logname $ toStrict (encodeToLazyText obj)
where
  output nm msg = TIO.putStrLn $ nm <> " :: " <> msg

```

Concrete Trace into a TVar

```

traceInTVar :: STM.TVar [a] → BaseTrace.BaseTrace STM.STM a
traceInTVar tvar = BaseTrace.BaseTrace $ Op $ λa → STM.modifyTVar tvar ((:) a)
traceInTVarIO :: STM.TVar [LogObject] → TraceNamed IO
traceInTVarIO tvar = BaseTrace.BaseTrace $ Op $ λln →
  STM.atomically $ STM.modifyTVar tvar ((:) (lnItem ln))
traceNamedInTVarIO :: STM.TVar [LogNamed LogObject] → TraceNamed IO
traceNamedInTVarIO tvar = BaseTrace.BaseTrace $ Op $ λln →
  STM.atomically $ STM.modifyTVar tvar ((:) ln)

```

Check a log item's severity against the **Trace**'s minimum severity

do we need three different **minSeverity** defined?

We do a lookup of the global **minSeverity** in the configuration. And, a lookup of the **minSeverity** for the current named context. These values might have changed in the meanwhile. A third filter is the **minSeverity** defined in the current context.

```

traceConditionally
  :: MonadIO m
  ⇒ Trace m
  → LogObject
  → m ()
traceConditionally logTrace@(ctx, _) msg@(LogObject _ (LogMessage item)) = do
  globminsev ← liftIO $ Config.minSeverity (configuration ctx)
  globnamesev ← liftIO $ Config.inspectSeverity (configuration ctx) (loggerName ctx)
  let minsev = max (minSeverity ctx) $ max globminsev (fromMaybe Debug globnamesev)
  flag = (liSeverity item) ≥ minsev
  when flag $ traceNamedObject logTrace msg
traceConditionally logTrace logObject =
  traceNamedObject logTrace logObject

```

Enter message into a trace

The function `traceNamedItem` creates a `LogObject` and threads this through the action defined in the `Trace`.

```

traceNamedItem
  :: MonadIO m
  => Trace m
  → LogSelection
  → Severity
  → T.Text
  → m ()
traceNamedItem trace p s m =
  traceConditionally trace <<
    LogObject < $ > liftIO mkLOMeta
      < * > pure (LogMessage LogItem {liSelection = p
        , liSeverity = s
        , liPayload = m
        })

```

Logging functions

```

logDebug, logInfo, logNotice, logWarning, logError, logCritical, logAlert, logEmergency
  :: MonadIO m => Trace m → T.Text → m ()
logDebug  logTrace = traceNamedItem logTrace Both Debug
logInfo    logTrace = traceNamedItem logTrace Both Info
logNotice  logTrace = traceNamedItem logTrace Both Notice
logWarning logTrace = traceNamedItem logTrace Both Warning
logError   logTrace = traceNamedItem logTrace Both Error
logCritical logTrace = traceNamedItem logTrace Both Critical
logAlert   logTrace = traceNamedItem logTrace Both Alert
logEmergency logTrace = traceNamedItem logTrace Both Emergency

logDebugS, logInfoS, logNoticeS, logWarningS, logErrorS, logCriticalS, logAlertS, logEmergencyS
  :: MonadIO m => Trace m → T.Text → m ()
logDebugS  logTrace = traceNamedItem logTrace Private Debug
logInfoS    logTrace = traceNamedItem logTrace Private Info
logNoticeS  logTrace = traceNamedItem logTrace Private Notice
logWarningS logTrace = traceNamedItem logTrace Private Warning
logErrorS   logTrace = traceNamedItem logTrace Private Error
logCriticalS logTrace = traceNamedItem logTrace Private Critical
logAlertS   logTrace = traceNamedItem logTrace Private Alert
logEmergencyS logTrace = traceNamedItem logTrace Private Emergency

logDebugP, logInfoP, logNoticeP, logWarningP, logErrorP, logCriticalP, logAlertP, logEmergencyP
  :: MonadIO m => Trace m → T.Text → m ()
logDebugP  logTrace = traceNamedItem logTrace Public Debug
logInfoP    logTrace = traceNamedItem logTrace Public Info
logNoticeP  logTrace = traceNamedItem logTrace Public Notice

```

```

logWarningP logTrace = traceNamedItem logTrace Public Warning
logErrorP    logTrace = traceNamedItem logTrace Public Error
logCriticalP logTrace = traceNamedItem logTrace Public Critical
logAlertP    logTrace = traceNamedItem logTrace Public Alert
logEmergencyP logTrace = traceNamedItem logTrace Public Emergency
logDebugUnsafeP, logInfoUnsafeP, logNoticeUnsafeP, logWarningUnsafeP, logErrorUnsafeP,
  logCriticalUnsafeP, logAlertUnsafeP, logEmergencyUnsafeP
  :: MonadIO m => Trace m -> T.Text -> m ()
logDebugUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Debug
logInfoUnsafeP   logTrace = traceNamedItem logTrace PublicUnsafe Info
logNoticeUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Notice
logWarningUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Warning
logErrorUnsafeP   logTrace = traceNamedItem logTrace PublicUnsafe Error
logCriticalUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Critical
logAlertUnsafeP   logTrace = traceNamedItem logTrace PublicUnsafe Alert
logEmergencyUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Emergency

```

subTrace

Transforms the input **Trace** according to the **Configuration** using the logger name of the current **Trace** appended with the new name. If the empty **Text** is passed, then the logger name remains untouched.

```

subTrace :: MonadIO m => T.Text -> Trace m -> m (Trace m)
subTrace name tr@(ctx, _) = do
  let newName = appendWithDot (loggerName ctx) name
  subtrace0 <- liftIO $ Config.findSubTrace (configuration ctx) newName
  let subtrace = case subtrace0 of Nothing -> Neutral; Just str -> str
  case subtrace of
    Neutral      -> do
      tr' <- appendName name tr
      return $ updateTracetype subtrace tr'
    UntimedTrace -> do
      tr' <- appendName name tr
      return $ updateTracetype subtrace tr'
    TeeTrace _    -> do
      tr' <- appendName name tr
      return $ updateTracetype subtrace tr'
    FilterTrace _ -> do
      tr' <- appendName name tr
      return $ updateTracetype subtrace tr'
    NoTrace       -> return $ updateTracetype subtrace (ctx, BaseTrace.BaseTrace $ Op $ \_ -> pure ())
    DropOpening   -> return $ updateTracetype subtrace (ctx, BaseTrace.BaseTrace $ Op $
      \ (LogNamed _ lo@(LogObject _ lc)) -> do
        case lc of
          ObserveOpen _ -> return ()
          _ -> traceNamedObject tr lo)

```

```

ObservableTrace _ → do
  tr' ← appendName name tr
  return $ updateTracetype subtrace tr'

```

1.4.5 Cardano.BM.Setup




Figure 1.3: Setup procedure

setupTrace

Setup a new **Trace** (**Trace**) with either a given **Configuration** (**Configuration.Model**) or a **FilePath** to a configuration file. After all tracing operations have ended; **shutdownTrace** must be called.

```

setupTrace :: MonadIO m ⇒ Either FilePath Config.Configuration → Text → m (Trace m)
setupTrace (Left cfgFile) name = do
  c ← liftIO $ Config.setup cfgFile
  setupTrace_ c name
setupTrace (Right c) name = setupTrace_ c name
setupTrace_ :: MonadIO m ⇒ Config.Configuration → Text → m (Trace m)
setupTrace_ c name = do
  sb ← liftIO $ Switchboard.realize c
  sev ← liftIO $ Config.minSeverity c
  ctx ← liftIO $ newContext " " c sev sb
  subTrace name $ natTrace liftIO (ctx, Switchboard.mainTrace sb)

```

shutdownTrace

Shut down a Trace and all the **Traces** related to it.

```

shutdownTrace :: MonadIO m ⇒ Trace m → m ()
shutdownTrace (ctx, _) = liftIO $ shutdown ctx

```


withTrace

```

withTrace :: (MonadIO m, MonadMask m) => Config.Configuration -> Text -> (Trace m -> m t) -> m t
withTrace cfg name action =
    bracket (setupTrace (Right cfg) name) shutdownTrace action

```

newContext

```

newContext :: LoggerName
    -> Config.Configuration
    -> Severity
    -> Switchboard.Switchboard
    -> IO TraceContext
newContext name cfg sev sb = do
    return $ TraceContext {
        loggerName = name
        , configuration = cfg
        , minSeverity = sev
        , tracetype = Neutral
        , shutdown = unrealize sb
    }

```

1.4.6 Cardano.BM.Counters

Here the platform is chosen on which we compile this program.

Currently, we mainly support *Linux* with its 'proc' filesystem.

```

{-# LANGUAGE CPP #-}
# if defined (linux_HOST_OS)
# define LINUX
# endif
module Cardano.BM.Counters
(
    Platform.readCounters
    , diffTimeObserved
    , getMonoClock
) where
# ifdef LINUX
import qualified Cardano.BM.Counters.Linux as Platform
# else
import qualified Cardano.BM.Counters.Dummy as Platform
# endif
import Cardano.BM.Counters.Common (getMonoClock)
import Cardano.BM.Data.Aggregated (Measurable (..))
import Cardano.BM.Data.Counter

```

Calculate difference between clocks

```

diffTimeObserved :: CounterState → CounterState → Measurable
diffTimeObserved (CounterState id0 startCounters) (CounterState id1 endCounters) =
  let
    startTime = getMonotonicTime startCounters
    endTime = getMonotonicTime endCounters
  in
    if (id0 == id1)
    then endTime - startTime
    else error "these clocks are not from the same experiment"
  where
    getMonotonicTime counters = case (filter isMonotonicClockCounter counters) of
      [(Counter MonotonicClockTime _ mus)] → mus
      _ → error "A time measurement is missing!"
    isMonotonicClockCounter :: Counter → Bool
    isMonotonicClockCounter = (MonotonicClockTime ==) ∘ cType

```

1.4.7 Cardano.BM.Counters.Common

Common functions that serve *readCounters* on all platforms.

```

nominalTimeToMicroseconds :: Word64 → Microsecond
nominalTimeToMicroseconds = fromMicroseconds ∘ toInteger ∘ ('div' 1000)

```

Read monotonic clock

```

getMonoClock :: IO [Counter]
getMonoClock = do
  t ← getMonotonicTimeNSec
  return [Counter MonotonicClockTime "monoclock" $ Microseconds (t `div` 1000)]

```

Read GHC RTS statistics

Read counters from GHC's RTS (runtime system). The values returned are as per the last GC (garbage collection) run.

```

readRTSStats :: IO [Counter]
readRTSStats = do
  iscollected ← GhcStats.getRTSStatsEnabled
  if iscollected
  then ghcstats
  else return []
  where
    ghcstats :: IO [Counter]
    ghcstats = do

```

```

-- need to run GC?
rts ← GhcStats.getRTSStats
let getrts = ghcval rts
return [getrts (Bytes ∘ fromIntegral ∘ GhcStats.allocated_bytes, "bytesAllocated")
      ,getrts (Bytes ∘ fromIntegral ∘ GhcStats.max_live_bytes, "liveBytes")
      ,getrts (Bytes ∘ fromIntegral ∘ GhcStats.max_large_objects_bytes, "largeBytes")
      ,getrts (Bytes ∘ fromIntegral ∘ GhcStats.max_compact_bytes, "compactBytes")
      ,getrts (Bytes ∘ fromIntegral ∘ GhcStats.max_slop_bytes, "slopBytes")
      ,getrts (Bytes ∘ fromIntegral ∘ GhcStats.max_mem_in_use_bytes, "usedMemBytes")
      ,getrts (Nanoseconds ∘ fromIntegral ∘ GhcStats.gc_cpu_ns, "gcCpuNs")
      ,getrts (Nanoseconds ∘ fromIntegral ∘ GhcStats.gc_elapsed_ns, "gcElapsedNs")
      ,getrts (Nanoseconds ∘ fromIntegral ∘ GhcStats.cpu_ns, "cpuNs")
      ,getrts (Nanoseconds ∘ fromIntegral ∘ GhcStats.elapsed_ns, "elapsedNs")
      ,getrts (PureI ∘ toInteger ∘ GhcStats.gcs, "gcNum")
      ,getrts (PureI ∘ toInteger ∘ GhcStats.major_gcs, "gcMajorNum")
      ]
ghcval :: GhcStats.RTSStats → ((GhcStats.RTSStats → Measurable), Text) → Counter
ghcval s (f,n) = Counter RTSStats n $ (f s)

```

1.4.8 Cardano.BM.Counters.Dummy

This is a dummy definition of `readCounters` on platforms that do not support the 'proc' filesystem from which we would read the counters.

The only supported measurements are monotonic clock time and RTS statistics for now.

```

readCounters :: SubTrace → IO [Counter]
readCounters NoTrace      = return []
readCounters Neutral      = return []
readCounters (TeeTrace _) = return []
readCounters (FilterTrace _) = return []
readCounters UntimedTrace = return []
readCounters DropOpening  = return []
readCounters (ObservableTrace tts) = foldrM (λ(sel,fun) a →
  if any (≡ sel) tts
  then (fun ≫ λxs → return $ a ++ xs)
  else return a) [] selectors
where
  selectors = [(MonotonicClock, getMonoClock)
    -- , (MemoryStats, readProcStatM)
    -- , (ProcessStats, readProcStats)
    -- , (IOStats, readProcIO)
    , (GhcRtsStats, readRTSStats)
    ]

```

1.4.9 Cardano.BM.Counters.Linux

we have to expand the `readMemStats` function to read full data from `proc`

```

readCounters :: SubTrace → IO [Counter]
readCounters NoTrace      = return []
readCounters Neutral      = return []
readCounters (TeeTrace _) = return []
readCounters (FilterTrace _) = return []
readCounters UntimedTrace = return []
readCounters DropOpening  = return []
readCounters (ObservableTrace tts) = foldrM (λ(sel,fun) a →
  if any (≡ sel) tts
  then (fun >> λxs → return $ a ++ xs)
  else return a) [] selectors
where
  selectors = [ (MonotonicClock, getMonoClock)
               , (MemoryStats, readProcStatM)
               , (ProcessStats, readProcStats)
               , (IOStats, readProcIO)
               ]

```

```

pathProc :: FilePath
pathProc = "/proc/"
pathProcStat :: ProcessID → FilePath
pathProcStat pid = pathProc </> (show pid) </> "stat"
pathProcStatM :: ProcessID → FilePath
pathProcStatM pid = pathProc </> (show pid) </> "statm"
pathProcIO :: ProcessID → FilePath
pathProcIO pid = pathProc </> (show pid) </> "io"

```

Reading from a file in /proc/<pid>

```

readProcList :: FilePath → IO [Integer]
readProcList fp = do
  cs ← readFile fp
  return $ map (λs → maybe 0 id $ (readMaybe s :: Maybe Integer)) (words cs)

```

readProcStatM - /proc/<pid>/statm

/proc/[pid]/statm

Provides information about memory usage, measured in pages. The columns are:

size	(1) total program size (same as VmSize in /proc/[pid]/status)
resident	(2) resident set size (same as VmRSS in /proc/[pid]/status)
shared	(3) number of resident shared pages (i.e., backed by a file) (same as RssFile+RssShmem in /proc/[pid]/status)
text	(4) text (code)
lib	(5) library (unused since Linux 2.6; always 0)
data	(6) data + stack
dt	(7) dirty pages (unused since Linux 2.6; always 0)

```

readProcStatM :: IO [Counter]
readProcStatM = do
    pid ← getProcessID
    ps0 ← readProcList (pathProcStatM pid)
    let ps = zip colnames ps0
        psUseful = filter (("unused" ≠) ∘ fst) ps
    return $ map (λ(n,i) → Counter MemoryCounter n (PureI i)) psUseful
where
    colnames :: [Text]
    colnames = ["size", "resident", "shared", "text", "unused", "data", "unused"]

```

readProcStats - //proc//<pid >//stat

/proc/[pid]/stat

Status information about the process. This is used by ps(1). It is defined in the kernel source file fs/proc/array.c.

The fields, in order, with their proper scanf(3) format specifiers, are listed below. Whether or not certain of these fields display valid information is governed by a ptrace access mode PTRACE_MODE_READ_FSCREDS | PTRACE_MODE_NOAUDIT check (refer to ptrace(2)). If the check denies access, then the field value is displayed as 0. The affected fields are indicated with the marking [PT].

- (1) pid %d
The process ID.
- (2) comm %s
The filename of the executable, in parentheses. This is visible whether or not the executable is swapped out.
- (3) state %c
One of the following characters, indicating process state:
 - R Running
 - S Sleeping in an interruptible wait
 - D Waiting in uninterruptible disk sleep
 - Z Zombie
 - T Stopped (on a signal) or (before Linux 2.6.33) trace stopped
 - t Tracing stop (Linux 2.6.33 onward)
 - W Paging (only before Linux 2.6.0)
 - X Dead (from Linux 2.6.0 onward)
 - x Dead (Linux 2.6.33 to 3.13 only)
 - K Wakekill (Linux 2.6.33 to 3.13 only)
 - W Waking (Linux 2.6.33 to 3.13 only)
 - P Parked (Linux 3.9 to 3.13 only)
- (4) ppid %d
The PID of the parent of this process.
- (5) pgrp %d
The process group ID of the process.

- (6) session %d
The session ID of the process.
- (7) tty_nr %d
The controlling terminal of the process. (The minor device number is contained in the combination of bits 31 to 20 and 7 to 0; the major device number is in bits 15 to 8.)
- (8) tpgid %d
The ID of the foreground process group of the controlling terminal of the process.
- (9) flags %u
The kernel flags word of the process. For bit meanings, see the PF_* defines in the Linux kernel source file include/linux/sched.h. Details depend on the kernel version.

The format for this field was %lu before Linux 2.6.
- (10) minflt %lu
The number of minor faults the process has made which have not required loading a memory page from disk.
- (11) cminflt %lu
The number of minor faults that the process's waited-for children have made.
- (12) majflt %lu
The number of major faults the process has made which have required loading a memory page from disk.
- (13) cmajflt %lu
The number of major faults that the process's waited-for children have made.
- (14) utime %lu
Amount of time that this process has been scheduled in user mode, measured in clock ticks (divide by sysconf(_SC_CLK_TCK)). This includes guest time, guest_time (time spent running a virtual CPU, see below), so that applications that are not aware of the guest time field do not lose that time from their calculations.
- (15) stime %lu
Amount of time that this process has been scheduled in kernel mode, measured in clock ticks (divide by sysconf(_SC_CLK_TCK)).
- (16) cutime %ld
Amount of time that this process's waited-for children have been scheduled in user mode, measured in clock ticks (divide by sysconf(_SC_CLK_TCK)). (See also times(2).) This includes guest time, cguest_time (time spent running a virtual CPU, see below).
- (17) cstime %ld
Amount of time that this process's waited-for children have been scheduled in kernel mode, measured in clock ticks (divide by sysconf(_SC_CLK_TCK)).
- (18) priority %ld
(Explanation for Linux 2.6) For processes running a real-time scheduling policy (policy below; see sched_setscheduler(2)), this is the negated scheduling priority, minus one; that is, a number in the range -2 to -100, corresponding to real-time priorities 1 to 99. For processes running under a non-real-time scheduling policy, this is the raw nice value (setpriority(2)) as represented in the kernel. The kernel stores nice values as numbers in the range 0 (high) to 39 (low), corresponding to the user-visible nice range of -20 to 19.
- (19) nice %ld
The nice value (see setpriority(2)), a value in the range 19 (low priority) to -20 (high priority).
- (20) num_threads %ld
Number of threads in this process (since Linux 2.6). Before kernel 2.6, this field was hard coded to 0 as a placeholder for an earlier removed field.
- (21) itrealvalue %ld
The time in jiffies before the next SIGALRM is sent to the process due to an interval timer. Since kernel 2.6.17, this field is no longer maintained, and is hard coded as 0.

- (22) starttime %llu
The time the process started after system boot. In kernels before Linux 2.6, this value was expressed in jiffies. Since Linux 2.6, the value is expressed in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).

The format for this field was %lu before Linux 2.6.
- (23) vsize %lu
Virtual memory size in bytes.
- (24) rss %ld
Resident Set Size: number of pages the process has in real memory. This is just the pages which count toward text, data, or stack space. This does not include pages which have not been demand-loaded in, or which are swapped out.
- (25) rsslim %lu
Current soft limit in bytes on the rss of the process; see the description of RLIMIT_RSS in `getrlimit(2)`.
- (26) startcode %lu [PT]
The address above which program text can run.
- (27) endcode %lu [PT]
The address below which program text can run.
- (28) startstack %lu [PT]
The address of the start (i.e., bottom) of the stack.
- (29) kstkesp %lu [PT]
The current value of ESP (stack pointer), as found in the kernel stack page for the process.
- (30) kstkeip %lu [PT]
The current EIP (instruction pointer).
- (31) signal %lu
The bitmap of pending signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (32) blocked %lu
The bitmap of blocked signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (33) sigignore %lu
The bitmap of ignored signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (34) sigcatch %lu
The bitmap of caught signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (35) wchan %lu [PT]
This is the "channel" in which the process is waiting. It is the address of a location in the kernel where the process is sleeping. The corresponding symbolic name can be found in `/proc/[pid]/wchan`.
- (36) nswap %lu
Number of pages swapped (not maintained).
- (37) cnswap %lu
Cumulative nswap for child processes (not maintained).
- (38) exit_signal %d (since Linux 2.1.22)
Signal to be sent to parent when we die.
- (39) processor %d (since Linux 2.2.8)
CPU number last executed on.

- (40) `rt_priority` %u (since Linux 2.5.19)
Real-time scheduling priority, a number in the range 1 to 99 for processes scheduled under a real-time policy, or 0, for non-real-time processes (see `sched_setscheduler(2)`).
- (41) `policy` %u (since Linux 2.5.19)
Scheduling policy (see `sched_setscheduler(2)`). Decode using the `SCHED_*` constants in `linux/sched.h`.

The format for this field was %lu before Linux 2.6.22.
- (42) `delayacct_blkio_ticks` %llu (since Linux 2.6.18)
Aggregated block I/O delays, measured in clock ticks (centiseconds).
- (43) `guest_time` %lu (since Linux 2.6.24)
Guest time of the process (time spent running a virtual CPU for a guest operating system), measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (44) `cguest_time` %ld (since Linux 2.6.24)
Guest time of the process's children, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (45) `start_data` %lu (since Linux 3.3) [PT]
Address above which program initialized and uninitialized (BSS) data are placed.
- (46) `end_data` %lu (since Linux 3.3) [PT]
Address below which program initialized and uninitialized (BSS) data are placed.
- (47) `start_brk` %lu (since Linux 3.3) [PT]
Address above which program heap can be expanded with `brk(2)`.
- (48) `arg_start` %lu (since Linux 3.5) [PT]
Address above which program command-line arguments (`argv`) are placed.
- (49) `arg_end` %lu (since Linux 3.5) [PT]
Address below program command-line arguments (`argv`) are placed.
- (50) `env_start` %lu (since Linux 3.5) [PT]
Address above which program environment is placed.
- (51) `env_end` %lu (since Linux 3.5) [PT]
Address below which program environment is placed.
- (52) `exit_code` %d (since Linux 3.5) [PT]
The thread's exit status in the form reported by `waitpid(2)`.

readProcStats :: IO [Counter]

readProcStats = do

pid ← getProcessID

ps0 ← readProcList (pathProcStat pid)

let ps = zip colnames ps0

psUseful = filter ((`"unused"` \neq) \circ fst) ps

*return \$ map ($\lambda(n,i) \rightarrow$ Counter StatInfo *n* (PureI *i*)) psUseful*

where

colnames :: [Text]

colnames = ["pid", "unused", "unused", "ppid", "pgrp", "session", "ttynr", "tpgid", "flags", "minflt", "cmminflt", "majflt", "cmajflt", "utime", "stime", "cutime", "cstime", "priority", "nice", "num", "itrealvalue", "starttime", "vsize", "rss", "rsslim", "startcode", "endcode", "startstack", "signal", "blocked", "sigignore", "sigcatch", "wchan", "nswap", "cnswap", "exitsignal", "proc", "policy", "blkio", "guesttime", "cguesttime", "startdata", "enddata", "startbrk", "argstart", "envend", "exitcode"
]

readProcIO - //proc//<pid >//io

/proc/[pid]/io (since kernel 2.6.20)

This file contains I/O statistics for the process, for example:

```
# cat /proc/3828/io
rchar: 323934931
wchar: 323929600
syscr: 632687
syscw: 632675
read_bytes: 0
write_bytes: 323932160
cancelled_write_bytes: 0
```

The fields are as follows:

rchar: characters read

The number of bytes which this task has caused to be read from storage. This is simply the sum of bytes which this process passed to read(2) and similar system calls. It includes things such as terminal I/O and is unaffected by whether or not actual physical disk I/O was required (the read might have been satisfied from pagecache).

wchar: characters written

The number of bytes which this task has caused, or shall cause to be written to disk. Similar caveats apply here as with rchar.

syscr: read syscalls

Attempt to count the number of read I/O operations—that is, system calls such as read(2) and pread(2).

syscw: write syscalls

Attempt to count the number of write I/O operations—that is, system calls such as write(2) and pwrite(2).

read_bytes: bytes read

Attempt to count the number of bytes which this process really did cause to be fetched from the storage layer. This is accurate for block-backed filesystems.

write_bytes: bytes written

Attempt to count the number of bytes which this process caused to be sent to the storage layer.

cancelled_write_bytes:

The big inaccuracy here is truncate. If a process writes 1MB to a file and then deletes the file, it will in fact perform no writeout. But it will have been accounted as having caused 1MB of write. In other words: this field represents the number of bytes which this process caused to not happen, by truncating pagecache. A task can cause "negative" I/O too. If this task truncates some dirty pagecache, some I/O which another task has been accounted for (in its write_bytes) will not be happening.

Note: In the current implementation, things are a bit racy on 32-bit systems: if process A reads process B's /proc/[pid]/io while process B is updating one of these 64-bit counters, process A could see an intermediate result.

Permission to access this file is governed by a ptrace access mode PTRACE_MODE_READ_FSCREDS check; see ptrace(2).

readProcIO::IO [Counter]

readProcIO = do

pid ← getProcessID

ps0 ← readProcList (pathProcIO pid)

let ps = zip3 colnames ps0 units

return \$ map (λ(n,i,u) → Counter IOCounter n (u i)) ps

where

colnames :: [Text]

```
colnames = [ "rchar", "wchar", "syscr", "syscw", "rbytes", "wbytes", "cxwbytes" ]
units = [ Bytes ◦ fromInteger, Bytes ◦ fromInteger, PureI, PureI, Bytes ◦ fromInteger, Bytes ◦ fromInteger, Bytes ◦ fromInteger ]
```

1.4.10 Cardano.BM.Data.Aggregated

Measurable

A **Measurable** may consist of different types of values. Time measurements are strict, so are *Bytes* which are externally measured. The real or integral numeric values are lazily linked, so we can decide later to drop them.

```
data Measurable = Microseconds {-# UNPACK #-} !Word64
  | Nanoseconds {-# UNPACK #-} !Word64
  | Seconds      {-# UNPACK #-} !Word64
  | Bytes        {-# UNPACK #-} !Word64
  | PureD        Double
  | PureI        Integer
deriving (Eq, Ord, Generic, ToJSON)
```

Measurable can be transformed to an integral value.

```
getInteger :: Measurable → Integer
getInteger (Microseconds a) = toInteger a
getInteger (Nanoseconds a)  = toInteger a
getInteger (Seconds a)      = toInteger a
getInteger (Bytes a)        = toInteger a
getInteger (PureI a)        = a
getInteger (PureD a)        = round a
```

Measurable can be transformed to a rational value.

```
getDouble :: Measurable → Double
getDouble (Microseconds a) = fromIntegral a
getDouble (Nanoseconds a)  = fromIntegral a
getDouble (Seconds a)      = fromIntegral a
getDouble (Bytes a)        = fromIntegral a
getDouble (PureI a)        = fromInteger a
getDouble (PureD a)        = a
```

It is a numerical value, thus supports functions to operate on numbers.

```
instance Num Measurable where
  (+) (Microseconds a) (Microseconds b) = Microseconds (a + b)
  (+) (Nanoseconds a) (Nanoseconds b) = Nanoseconds (a + b)
  (+) (Seconds a) (Seconds b) = Seconds (a + b)
  (+) (Bytes a) (Bytes b) = Bytes (a + b)
  (+) (PureI a) (PureI b) = PureI (a + b)
  (+) (PureD a) (PureD b) = PureD (a + b)
  (+) _ _ = error "Trying to add values with different units"
```

```

(*) (Microseconds a) (Microseconds b) = Microseconds (a * b)
(*) (Nanoseconds a) (Nanoseconds b) = Nanoseconds (a * b)
(*) (Seconds a)      (Seconds b)      = Seconds      (a * b)
(*) (Bytes a)        (Bytes b)        = Bytes        (a * b)
(*) (PureI a)         (PureI b)        = PureI         (a * b)
(*) (PureD a)         (PureD b)        = PureD         (a * b)
(*) -- -- -- -- -- = error "Trying to multiply values with different units"

abs (Microseconds a) = Microseconds (abs a)
abs (Nanoseconds a) = Nanoseconds (abs a)
abs (Seconds a)     = Seconds (abs a)
abs (Bytes a)       = Bytes (abs a)
abs (PureI a)       = PureI (abs a)
abs (PureD a)       = PureD (abs a)

signum (Microseconds a) = Microseconds (signum a)
signum (Nanoseconds a) = Nanoseconds (signum a)
signum (Seconds a)     = Seconds (signum a)
signum (Bytes a)       = Bytes (signum a)
signum (PureI a)       = PureI (signum a)
signum (PureD a)       = PureD (signum a)

negate (Microseconds a) = Microseconds (negate a)
negate (Nanoseconds a) = Nanoseconds (negate a)
negate (Seconds a)     = Seconds (negate a)
negate (Bytes a)       = Bytes (negate a)
negate (PureI a)       = PureI (negate a)
negate (PureD a)       = PureD (negate a)

fromInteger = PureI

```

Pretty printing of **Measurable**.

```

instance Show Measurable where
  show (Microseconds a) = show a
  show (Nanoseconds a)  = show a
  show (Seconds a)      = show a
  show (Bytes a)        = show a
  show (PureI a)        = show a
  show (PureD a)        = show a

showUnits :: Measurable → String
showUnits (Microseconds _) = " μs"
showUnits (Nanoseconds _) = " ns"
showUnits (Seconds _)     = " s"
showUnits (Bytes _)       = " B"
showUnits (PureI _)       = ""
showUnits (PureD _)       = ""

-- show in S.I. units
showSI :: Measurable → String
showSI (Microseconds a) = show (fromFloatDigits ((fromIntegral a) / (1000000 :: Float))) ++
  showUnits (Seconds a)

```

```

showSI (Nanoseconds a) = show (fromFloatDigits ((fromIntegral a) / (1000000000 :: Float))) ++
                           showUnits (Seconds a)
showSI v@(Seconds a)   = show a ++ showUnits v
showSI v@(Bytes a)      = show a ++ showUnits v
showSI v@(PureI a)      = show a ++ showUnits v
showSI v@(PureD a)      = show a ++ showUnits v

```

Stats

A **Stats** statistics is strictly computed.

```

data BaseStats = BaseStats {
  fmin :: !Measurable,
  fmax :: !Measurable,
  fcount :: !Int,
  fsum_A :: !Double,
  fsum_B :: !Double
} deriving (Generic, ToJSON, Show)

instance Eq BaseStats where
  (BaseStats mina maxa counta sumAa sumBa) == (BaseStats minb maxb countb sumAb sumBb) =
    mina == minb & maxa == maxb & counta == countb &
    abs (sumAa - sumAb) < 1.0e-4 &
    abs (sumBa - sumBb) < 1.0e-4

data Stats = Stats {
  flast :: !Measurable,
  fold :: !Measurable,
  fbasic :: !BaseStats,
  fdelta :: !BaseStats,
  ftimed :: !BaseStats
} deriving (Eq, Generic, ToJSON, Show)

```

```

meanOfStats :: BaseStats → Double
meanOfStats = fsum_A

```

```

stdevOfStats :: BaseStats → Double
stdevOfStats s =
  if fcount s < 2
  then 0
  else sqrt $ (fsum_B s) / (fromInteger $ fromIntegral (fcount s) - 1)

```

instance Semigroup Stats disabled for the moment, because not needed.

We use a parallel algorithm to update the estimation of mean and variance from two sample statistics. (see https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance#Parallel_algorithm)

```

instance Semigroup Stats where
  (<>) a b = let counta = fcount a

```

```

countb = fcount b
newcount = counta + countb
delta = fsum_A b - fsum_A a
in
Stats {flast = flast b -- right associative
, fmin    = min (fmin a) (fmin b)
, fmax    = max (fmax a) (fmax b)
, fcount  = newcount
, fsum_A  = fsum_A a + (delta / fromInteger newcount)
, fsum_B  = fsum_B a + fsum_B b + (delta * delta) * (fromInteger (counta * countb) / fromInteger newcount)
}

stats2Text :: Stats → Text
stats2Text (Stats slast _ sbasic sdelta stimed) =
  pack $
    "{ last=" ++ show slast ++
    ", basic-stats=" ++ showStats' (sbasic) ++
    ", delta-stats=" ++ showStats' (sdelta) ++
    ", timed-stats=" ++ showStats' (stimed) ++
    " }"
where
  showStats' :: BaseStats → String
  showStats' s =
    ", { min=" ++ show (fmin s) ++
    ", max=" ++ show (fmax s) ++
    ", mean=" ++ show (meanOfStats s) ++ showUnits (fmin s) ++
    ", std-dev=" ++ show (stdevOfStats s) ++
    ", count=" ++ show (fcount s) ++
    " }"

```

Exponentially Weighted Moving Average (EWMA)

Following https://en.wikipedia.org/wiki/Moving_average#Exponential_moving_average we calculate the exponential moving average for a series of values Y_t according to:

$$S_t = \begin{cases} Y_1, & t = 1 \\ \alpha \cdot Y_t + (1 - \alpha) \cdot S_{t-1}, & t > 1 \end{cases}$$

```

data EWMA = EmptyEWMA {alpha :: Double}
| EWMA {alpha :: Double
, avg :: Measurable
} deriving (Show, Eq, Generic, ToJSON)

```

Aggregated

```
data Aggregated = AggregatedStats Stats
  | AggregatedEWMA EWMA
  deriving (Eq, Generic, ToJSON)
```

instance Semigroup Aggregated disabled for the moment, because not needed.

```
instance Semigroup Aggregated where
  (<>) (AggregatedStats a) (AggregatedStats b) =
    AggregatedStats (a <> b)
  (<>) _ _ = error "Cannot combine different objects"
```

```
singletonStats :: Measurable → Aggregated
```

```
singletonStats a =
  let stats = Stats {flast = a
    ,fold          = 0
    ,fbasic = BaseStats
      {fmin = a
      ,fmax = a
      ,fcount = 1
      ,fsum_A = getDouble a
      ,fsum_B = 0}
    ,fdelta = BaseStats
      {fmin = 0
      ,fmax = 0
      ,fcount = 0
      ,fsum_A = 0
      ,fsum_B = 0}
    ,ftimed = BaseStats
      {fmin = Nanoseconds 999999999999
      ,fmax = Nanoseconds 0
      ,fcount = (-1)
      ,fsum_A = 0
      ,fsum_B = 0}
    }
  in
    AggregatedStats stats
```

```
instance Show Aggregated where
  show (AggregatedStats astats) =
    "{ stats = " ++ show astats ++ " }"
  show (AggregatedEWMA a) = show a
```

1.4.11 Cardano.BM.Data.Backend

Accepts a **NamedLogItem**

Instances of this type class accept a **NamedLogItem** and deal with it.

```
class IsEffectuator t where
  effectuate :: t → NamedLogItem → IO ()
  effectuatefrom :: forall s ◦ (IsEffectuator s) ⇒ t → NamedLogItem → s → IO ()
  default effectuatefrom :: forall s ◦ (IsEffectuator s) ⇒ t → NamedLogItem → s → IO ()
  effectuatefrom t nli _ = effectuate t nli
```

Declaration of a **Backend**

A backend is life-cycle managed, thus can be *realized* and *unrealized*.

```
class (IsEffectuator t) ⇒ IsBackend t where
  typeof    :: t → BackendKind
  realize    :: Configuration → IO t
  realizefrom :: forall s ◦ (IsEffectuator s) ⇒ Trace IO → s → IO t
  default realizefrom :: forall s ◦ (IsEffectuator s) ⇒ Trace IO → s → IO t
  realizefrom (ctx, _) _ = realize (configuration ctx)
  unrealize :: t → IO ()
```

Backend

This data structure for a backend defines its behaviour as an **IsEffectuator** when processing an incoming message, and as an **IsBackend** for unrealizing the backend.

```
data Backend = MkBackend
  { bEffectuate :: NamedLogItem → IO ()
  , bUnrealize :: IO ()
  }
```

1.4.12 Cardano.BM.Data.Configuration

Data structure to help parsing configuration files.

Representation

```
type Port = Int
data Representation = Representation
  { minSeverity    :: Severity
  , rotation       :: RotationParameters
  , setupScribes   :: [ScribeDefinition]
  , defaultScribes :: [(ScribeKind, Text)]
  , setupBackends  :: [BackendKind]
  , defaultBackends :: [BackendKind]
```

```
,hasEKG      :: Maybe Port
,hasGUI      :: Maybe Port
,options     :: HM.HashMap Text Object
}
deriving (Generic, Show, ToJSON, FromJSON)
```

parseRepresentation

```
parseRepresentation :: FilePath → IO Representation
```

```
parseRepresentation fp = do
  repr :: Representation ← decodeFileThrow fp
  return $ implicit_fill_representation repr
```

after parsing the configuration representation we implicitly correct it.

```
implicit_fill_representation :: Representation → Representation
```

```
implicit_fill_representation =
  remove_ekgview_if_not_defined ∘
  filter_duplicates_from_backends ∘
  filter_duplicates_from_scribes ∘
  union_setup_and_usage_backends ∘
  add_ekgview_if_port_defined ∘
  add_katip_if_any_scribes
where
  filter_duplicates_from_backends r =
    r { setupBackends = mkUniq $ setupBackends r }
  filter_duplicates_from_scribes r =
    r { setupScribes = mkUniq $ setupScribes r }
  union_setup_and_usage_backends r =
    r { setupBackends = setupBackends r <> defaultBackends r }
  remove_ekgview_if_not_defined r =
    case hasEKG r of
      Nothing → r { defaultBackends = filter (λbk → bk ≠ EKGViewBK) (defaultBackends r)
                  , setupBackends = filter (λbk → bk ≠ EKGViewBK) (setupBackends r)
                  }
      Just _ → r
  add_ekgview_if_port_defined r =
    case hasEKG r of
      Nothing → r
      Just _ → r { setupBackends = setupBackends r <> [EKGViewBK] }
  add_katip_if_any_scribes r =
    if (any ¬ [null $ setupScribes r, null $ defaultScribes r])
    then r { setupBackends = setupBackends r <> [KatipBK] }
    else r
  mkUniq :: Ord a ⇒ [a] → [a]
  mkUniq = Set.toList ∘ Set.fromList
```


1.4.13 Cardano.BM.Data.Counter

Counter

```

data Counter = Counter
    { cType :: CounterType
    , cName :: Text
    , cValue :: Measurable
    }
    deriving (Eq, Show, Generic, ToJSON)

data CounterType = MonotonicClockTime
    | MemoryCounter
    | StatInfo
    | IOCounter
    | CpuCounter
    | RTSStats
    deriving (Eq, Show, Generic, ToJSON)

instance ToJSON Microsecond where
    toJSON = toJSON ◦ toMicroseconds
    toEncoding = toEncoding ◦ toMicroseconds

```

Names of counters

```

nameCounter :: Counter → Text
nameCounter (Counter MonotonicClockTime _ _) = "Time-interval"
nameCounter (Counter MemoryCounter _ _) = "Mem"
nameCounter (Counter StatInfo _ _) = "Stat"
nameCounter (Counter IOCounter _ _) = "IO"
nameCounter (Counter CpuCounter _ _) = "Cpu"
nameCounter (Counter RTSStats _ _) = "RTS"

```

CounterState

```

data CounterState = CounterState {
    csIdentifier :: Unique
    , csCounters :: [Counter]
    }
    deriving (Generic, ToJSON)

instance ToJSON Unique where
    toJSON = toJSON ◦ hashUnique
    toEncoding = toEncoding ◦ hashUnique

instance Show CounterState where
    show cs = (show ◦ hashUnique) (csIdentifier cs)
        <> " => " <> (show $ csCounters cs)

```

Difference between counters

```

diffCounters :: [Counter] → [Counter] → [Counter]
diffCounters openings closings =
  getCountersDiff openings closings
where
  getCountersDiff :: [Counter]
    → [Counter]
    → [Counter]
  getCountersDiff as bs =
    let
      getName counter = nameCounter counter <> cName counter
      asNames = map getName as
      aPairs = zip asNames as
      bsNames = map getName bs
      bs' = zip bsNames bs
      bPairs = HM.fromList bs'
    in
      catMaybes $ (flip map) aPairs $ \ (name, Counter _ _ startValue) →
        case HM.lookup name bPairs of
          Nothing → Nothing
          Just counter → let endValue = cValue counter
                        in Just counter {cValue = endValue - startValue}

```

1.4.14 Cardano.BM.Data.LogItem**LoggerName**

A **LoggerName** has currently type *Text*.

```
type LoggerName = Text
```

NamedLogItem

```
type NamedLogItem = LogNamed LogObject
```

LogNamed

A **LogNamed** contains of a context name and some log item.

```

data LogNamed item = LogNamed
  { lnName :: LoggerName
  , lnItem :: item
  } deriving (Show)

deriving instance Generic item ⇒ Generic (LogNamed item)
deriving instance (ToJSON item, Generic item) ⇒ ToJSON (LogNamed item)

```

Logging of outcomes with `LogObject`

```
data LogObject = LogObject LOMeta LOContent
  deriving (Generic, Show, ToJSON)
```

Meta data for a `LogObject`:

```
data LOMeta = LOMeta {
  tstamp :: {-# UNPACK #-} !UTCTime
  , tid :: {-# UNPACK #-} !ThreadId
}
deriving (Show)
instance ToJSON LOMeta where
  toJSON (LOMeta _tstamp _tid) =
    object [ "tstamp" . = _tstamp, "tid" . = show _tid ]
mkLOMeta :: IO LOMeta
mkLOMeta =
  LOMeta <$> getCurrentTime
    <*> myThreadId
```

Payload of a `LogObject`:

```
data LOContent = LogMessage LogItem
  | LogValue Text Measurable
  | ObserveOpen CounterState
  | ObserveDiff CounterState
  | ObserveClose CounterState
  | AggregatedMessage [(Text, Aggregated)]
  | KillPill
  deriving (Generic, Show, ToJSON)
```

LogItem

TODO `liPayload` :: *ToObject*

```
data LogItem = LogItem
  { liSelection :: LogSelection
  , liSeverity :: Severity
  , liPayload :: Text -- TODO should become ToObject
  } deriving (Show, Generic, ToJSON)
```

```
data LogSelection =
  Public -- only to public logs.
  | PublicUnsafe -- only to public logs, not console.
  | Private -- only to private logs.
  | Both -- to public and private logs.
deriving (Show, Generic, ToJSON, FromJSON)
```

1.4.15 Cardano.BM.Data.Observable

ObservableInstance

```
data ObservableInstance = MonotonicClock
  | MemoryStats
  | ProcessStats
  | IOStats
  | GhcRtsStats
  deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON, Read)
```

1.4.16 Cardano.BM.Data.Output

OutputKind

```
data OutputKind = TVarList (STM.TVar [LogObject])
  | TVarListNamed (STM.TVar [LogNamed LogObject])
  deriving (Eq)
```

ScribeKind

This identifies katip's scribes by type.

```
data ScribeKind = FileTextSK
  | FileJsonSK
  | StdoutSK
  | StderrSK
  deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)
```

ScribeId

A scribe is identified by **ScribeKind** x *Filename*

```
type ScribeId = Text -- (ScribeKind :: Filename)
```

ScribeDefinition

This identifies katip's scribes by type.

```
data ScribeDefinition = ScribeDefinition
  { scKind :: ScribeKind
  , scName :: Text
  , scRotation :: Maybe RotationParameters
  }
  deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)
```

1.4.17 Cardano.BM.Data.Severity

Severity

The intended meaning of severity codes:

Debug *detailed information about values and decision flow* **Info** general information of events; progressing properly **Notice** *needs attention; something \rightarrow progressing properly* **Warning** may continue into an error condition if continued **Error** *unexpected set of event or condition occurred* **Critical** error condition causing degrade of operation **Alert** *a subsystem is no longer operating correctly, likely requires manual* at this point, the system can never progress without additional intervention

We were informed by the Syslog taxonomy: https://en.wikipedia.org/wiki/Syslog#Severity_level

```
data Severity = Debug
  | Info
  | Notice
  | Warning
  | Error
  | Critical
  | Alert
  | Emergency
  deriving (Show, Eq, Ord, Generic, ToJSON, Read)

instance FromJSON Severity where
  parseJSON = withText "severity" $ \case
    "Debug"    → pure Debug
    "Info"     → pure Info
    "Notice"   → pure Notice
    "Warning"  → pure Warning
    "Error"    → pure Error
    "Critical" → pure Critical
    "Alert"    → pure Alert
    "Emergency" → pure Emergency
    _          → pure Info-- catch all
```

1.4.18 Cardano.BM.Data.SubTrace

SubTrace

```
data NameSelector = Exact Text | StartsWith Text | EndsWith Text | Contains Text
  deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)

data DropName = Drop NameSelector
  deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)

data UnhideNames = Unhide [NameSelector]
  deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)

data SubTrace = Neutral
  | UntimedTrace
  | NoTrace
  | TeeTrace LoggerName
  | FilterTrace [(DropName, UnhideNames)]
```

```

| DropOpening
| ObservableTrace [ ObservableInstance ]
  deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)

```

1.4.19 Cardano.BM.Data.Trace

Trace

A **Trace** consists of a **TraceContext** and a **TraceNamed** in *m*.

```
type Trace m = (TraceContext, TraceNamed m)
```

TraceNamed

A **TraceNamed** is a specialized **Contravariant** of type **NamedLogItem**, a **LogNamed** with payload **LogObject**.

```
type TraceNamed m = BaseTrace m (NamedLogItem)
```

TraceContext

We keep the context's name and a reference to the **Configuration** in the **TraceContext**.

```

data TraceContext = TraceContext
  { loggerName :: LoggerName
  , configuration :: Configuration
  , tracetype   :: SubTrace
  , minSeverity :: Severity
  , shutdown   :: IO ()
  }

```

1.4.20 Cardano.BM.Configuration

see **Cardano.BM.Configuration.Model** for the implementation.

```

getOptionOrDefault :: CM.Configuration → Text → Text → IO (Text)
getOptionOrDefault cg name def = do
  opt ← CM.getOption cg name
  case opt of
    Nothing → return def
    Just o  → return o

```

<<Model>> Configuration	
cgMinSeverity	: Severity
cgMapSeverity	: Map = LoggerName -> Severity
cgMapSubtrace	: Map = LoggerName -> SubTrace
cgOptions	: Map = Text -> Aeson.Object
cgMapBackend	: Map = LoggerName -> [BackendKind]
cgDefBackends	: BackendKind [*]
cgSetupBackends	: BackendKind [*]
cgMapScribe	: Map = LoggerName -> [ScribeId]
cgDefScribes	: ScribeId [*]
cgSetupScribes	: ScribeDefinition [*]
cbMapAggregatedKind	: Map = LoggerName -> AggregatedKind
cgDefAggregatedKind	: AggregatedKind
cgPortEKG	: int
cgPortGUI	: int

Figure 1.4: Configuration model

1.4.21 Cardano.BM.Configuration.Model

Configuration.Model

```

type ConfigurationMVar = MVar ConfigurationInternal
newtype Configuration = Configuration
    {getCG :: ConfigurationMVar}
-- Our internal state; see -"Configuration model"-
data ConfigurationInternal = ConfigurationInternal
    {cgMinSeverity      :: Severity
    -- minimum severity level of every object that will be output
    ,cgMapSeverity      :: HM.HashMap LoggerName Severity
    -- severity filter per loggernaem
    ,cgMapSubtrace      :: HM.HashMap LoggerName SubTrace
    -- type of trace per loggernaem
    ,cgOptions          :: HM.HashMap Text Object
    -- options needed for tracing, logging and monitoring
    ,cgMapBackend       :: HM.HashMap LoggerName [BackendKind]
    -- backends that will be used for the specific loggernaem
    ,cgDefBackendKs     :: [BackendKind]
    -- backends that will be used if a set of backends for the
    -- specific loggernaem is not set
    ,cgSetupBackends    :: [BackendKind]
    -- backends to setup; every backend to be used must have
    -- been declared here
    ,cgMapScribe        :: HM.HashMap LoggerName [ScribeId]
    -- katip scribes that will be used for the specific loggernaem
    ,cgMapScribeCache :: HM.HashMap LoggerName [ScribeId]
    -- map to cache info of the cgMapScribe
    ,cgDefScribes       :: [ScribeId]
    -- katip scribes that will be used if a set of scribes for the

```

```

-- specific loggername is not set
,cgSetupScribes    :: [ScribeDefinition]
-- katip scribes to setup; every scribe to be used must have
-- been declared here
,cgMapAggregatedKind :: HM.HashMap LoggerName AggregatedKind
-- kind of Aggregated that will be used for the specific loggername
,cgDefAggregatedKind :: AggregatedKind
-- kind of Aggregated that will be used if a set of scribes for the
-- specific loggername is not set
,cgPortEKG         :: Int
-- port for EKG server
,cgPortGUI         :: Int
-- port for changes at runtime (NOT IMPLEMENTED YET)
} deriving (Show, Eq)

```

Backends configured in the Switchboard

For a given context name return the list of backends configured, or, in case no such configuration exists, return the default backends.

```

getBackends :: Configuration → LoggerName → IO [BackendKind]
getBackends configuration name =
  withMVar (getCG configuration) $ \cg → do
    let outs = HM.lookup name (cgMapBackend cg)
    case outs of
      Nothing → do
        return (cgDefBackendKs cg)
      Just os → return os

getDefaultBackends :: Configuration → IO [BackendKind]
getDefaultBackends configuration =
  withMVar (getCG configuration) $ \cg → do
    return (cgDefBackendKs cg)

setDefaultBackends :: Configuration → [BackendKind] → IO ()
setDefaultBackends configuration bes = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgDefBackendKs = bes}

setBackends :: Configuration → LoggerName → Maybe [BackendKind] → IO ()
setBackends configuration name be = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMapBackend = HM.alter (\_ → be) name (cgMapBackend cg)}

```

Backends to be setup by the Switchboard

Defines the list of Backends that need to be setup by the Switchboard.

```

setSetupBackends :: Configuration → [BackendKind] → IO ()
setSetupBackends configuration bes = do

```



```

cg ← takeMVar (getCG configuration)
putMVar (getCG configuration) $ cg {cgSetupBackends = bes}
getSetupBackends :: Configuration → IO [BackendKind]
getSetupBackends configuration =
  withMVar (getCG configuration) $ λcg →
    return $ cgSetupBackends cg

```

Scribes configured in the **Log** backend

For a given context name return the list of scribes to output to, or, in case no such configuration exists, return the default scribes to use.

```

getScribes :: Configuration → LoggerName → IO [ScribeId]
getScribes configuration name = do
  (updateCache, scribes) ← withMVar (getCG configuration) $ λcg → do
    let defs = cgDefScribes cg
    let mapScribe = cgMapScribe cg
    let find_s lname = case HM.lookup lname mapScribe of
      Nothing →
        case dropToDot lname of
          Nothing → defs
          Just lname' → find_s lname'
      Just os → os
    let outs = HM.lookup name (cgMapScribeCache cg)
    -- look if scribes are already cached
    return $ case outs of
      -- if no cached scribes found; search the appropriate scribes that
      -- they must inherit and update the cached map
      Nothing → (True, find_s name)
      Just os → (False, os)
  when updateCache $ setCachedScribes configuration name $ Just scribes
  return scribes
where
  dropToDot :: Text → Maybe Text
  dropToDot ts = dropToDot' (breakOnEnd "." ts)
  dropToDot' (_, "") = Nothing
  dropToDot' (name', _) = Just $ dropWhileEnd (≡ ' . ') name'
getCachedScribes :: Configuration → LoggerName → IO (Maybe [ScribeId])
getCachedScribes configuration name =
  withMVar (getCG configuration) $ λcg → do
    return $ HM.lookup name $ cgMapScribeCache cg
setScribes :: Configuration → LoggerName → Maybe [ScribeId] → IO ()
setScribes configuration name scribes = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $
    cg {cgMapScribe = HM.alter (\_ → scribes) name (cgMapScribe cg)}
setCachedScribes :: Configuration → LoggerName → Maybe [ScribeId] → IO ()

```

```

setCachedScribes configuration name scribes = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $
    cg {cgMapScribeCache = HM.alter (\_ → scribes) name (cgMapScribeCache cg)}
setDefaultScribes :: Configuration → [ScribeId] → IO ()
setDefaultScribes configuration scs = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgDefScribes = scs}

```

Scribes to be setup in the **Log** backend

Defines the list of *Scribes* that need to be setup in the **Log** backend.

```

setSetupScribes :: Configuration → [ScribeDefinition] → IO ()
setSetupScribes configuration sds = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgSetupScribes = sds}
getSetupScribes :: Configuration → IO [ScribeDefinition]
getSetupScribes configuration =
  withMVar (getCG configuration) $ λcg → do
    return $ cgSetupScribes cg

```

AggregatedKind to define the type of measurement

For a given context name return its *AggregatedKind* or in case no such configuration exists, return the default *AggregatedKind* to use.

```

getAggregatedKind :: Configuration → LoggerName → IO AggregatedKind
getAggregatedKind configuration name =
  withMVar (getCG configuration) $ λcg → do
    let outs = HM.lookup name (cgMapAggregatedKind cg)
    case outs of
      Nothing → do
        return (cgDefAggregatedKind cg)
      Just os → return $ os
setDefaultAggregatedKind :: Configuration → AggregatedKind → IO ()
setDefaultAggregatedKind configuration defAK = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgDefAggregatedKind = defAK}
setAggregatedKind :: Configuration → LoggerName → Maybe AggregatedKind → IO ()
setAggregatedKind configuration name ak = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMapAggregatedKind = HM.alter (\_ → ak) name (cgMapAggregatedKind cg)}

```

Access port numbers of EKG, GUI

```

getEKGport :: Configuration → IO Int
getEKGport configuration =
  withMVar (getCG configuration) $ \cg → do
    return $ cgPortEKG cg
setEKGport :: Configuration → Int → IO ()
setEKGport configuration port = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgPortEKG = port}
getGUIport :: Configuration → IO Int
getGUIport configuration =
  withMVar (getCG configuration) $ \cg → do
    return $ cgPortGUI cg
setGUIport :: Configuration → Int → IO ()
setGUIport configuration port = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgPortGUI = port}

```

Options

```

getOption :: Configuration → Text → IO (Maybe Text)
getOption configuration name = do
  withMVar (getCG configuration) $ \cg →
    case HM.lookup name (cgOptions cg) of
      Nothing → return Nothing
      Just o → return $ Just $ pack $ show o

```

Global setting of minimum severity

```

minSeverity :: Configuration → IO Severity
minSeverity configuration = withMVar (getCG configuration) $ \cg →
  return $ cgMinSeverity cg
setMinSeverity :: Configuration → Severity → IO ()
setMinSeverity configuration sev = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMinSeverity = sev}

```

Relation of context name to minimum severity

```

inspectSeverity :: Configuration → Text → IO (Maybe Severity)
inspectSeverity configuration name = do
  withMVar (getCG configuration) $ \cg →

```

```

    return $ HM.lookup name (cgMapSeverity cg)
setSeverity :: Configuration → Text → Maybe Severity → IO ()
setSeverity configuration name sev = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMapSeverity = HM.alter (\_ → sev) name (cgMapSeverity cg)}

```

Relation of context name to SubTrace

A new context may contain a different type of **Trace**. The function **appendName** (Enter new named context) will look up the **SubTrace** for the context's name.

```

findSubTrace :: Configuration → Text → IO (Maybe SubTrace)
findSubTrace configuration name = do
  withMVar (getCG configuration) $ λcg →
    return $ HM.lookup name (cgMapSubtrace cg)
setSubTrace :: Configuration → Text → Maybe SubTrace → IO ()
setSubTrace configuration name trafo = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMapSubtrace = HM.alter (\_ → trafo) name (cgMapSubtrace cg)}

```

Parse configuration from file

Parse the configuration into an internal representation first. Then, fill in **Configuration** after refinement.

```

setup :: FilePath → IO Configuration
setup fp = do
  r ← R.parseRepresentation fp
  setupFromRepresentation r
setupFromRepresentation :: R.Representation → IO Configuration
setupFromRepresentation r = do
  cgreg ← newEmptyMVar
  let mapseverity = HM.lookup "mapSeverity" (R.options r)
      mapbackends = HM.lookup "mapBackends" (R.options r)
      mapsubtrace = HM.lookup "mapSubtrace" (R.options r)
      mapscribes = HM.lookup "mapScribes" (R.options r)
      mapAggregatedKinds = HM.lookup "mapAggregatedkinds" (R.options r)
      mapScribe = parseScribeMap mapscribes
  putMVar cgreg $ ConfigurationInternal
    {cgMinSeverity = R.minSeverity r
    ,cgMapSeverity = parseSeverityMap mapseverity
    ,cgMapSubtrace = parseSubtraceMap mapsubtrace
    ,cgOptions = R.options r
    ,cgMapBackend = parseBackendMap mapbackends
    ,cgDefBackendKs = R.defaultBackends r
    ,cgSetupBackends = R.setupBackends r
    ,cgMapScribe = mapScribe}

```

```

    ,cgMapScribeCache = mapScribe
    ,cgDefScribes = r_defaultScribes r
    ,cgSetupScribes = R.setupScribes r
    ,cgMapAggregatedKind = parseAggregatedKindMap mapAggregatedKinds
    ,cgDefAggregatedKind = StatsAK
    ,cgPortEKG = r_hasEKG r
    ,cgPortGUI = r_hasGUI r
  }
  return $ Configuration cgregf
where
  parseSeverityMap :: Maybe (HM.HashMap Text Value) → HM.HashMap Text Severity
  parseSeverityMap Nothing = HM.empty
  parseSeverityMap (Just hmv) = HM.mapMaybe mkSeverity hmv
  mkSeverity (String s) = Just (read (unpack s) :: Severity)
  mkSeverity _ = Nothing
  parseBackendMap Nothing = HM.empty
  parseBackendMap (Just hmv) = HM.map mkBackends hmv
  mkBackends (Array bes) = catMaybes $ map mkBackend $ Vector.toList bes
  mkBackends _ = []
  mkBackend (String s) = Just (read (unpack s) :: BackendKind)
  mkBackend _ = Nothing
  parseScribeMap Nothing = HM.empty
  parseScribeMap (Just hmv) = HM.map mkScribes hmv
  mkScribes (Array scs) = catMaybes $ map mkScribe $ Vector.toList scs
  mkScribes (String s) = [(s :: ScribeId)]
  mkScribes _ = []
  mkScribe (String s) = Just (s :: ScribeId)
  mkScribe _ = Nothing
  parseSubtraceMap :: Maybe (HM.HashMap Text Value) → HM.HashMap Text SubTrace
  parseSubtraceMap Nothing = HM.empty
  parseSubtraceMap (Just hmv) = HM.mapMaybe mkSubtrace hmv
  mkSubtrace (String s) = Just (read (unpack s) :: SubTrace)
  mkSubtrace (Object hm) = mkSubtrace' (HM.lookup "tag" hm) (HM.lookup "contents" hm)
  mkSubtrace _ = Nothing
  mkSubtrace' Nothing _ = Nothing
  mkSubtrace' _ Nothing = Nothing
  mkSubtrace' (Just (String tag)) (Just (Array cs)) =
    if tag == "ObservableTrace"
    then Just $ ObservableTrace $ map (λ(String s) → (read (unpack s) :: ObservableInstance)) $ Vector.toList cs
    else Nothing
  mkSubtrace' _ _ = Nothing
  r_hasEKG repr = case (R.hasEKG repr) of
    Nothing → 0
    Just p → p
  r_hasGUI repr = case (R.hasGUI repr) of
    Nothing → 0
    Just p → p

```

```

r_defaultScribes repr = map (\(k,n) → pack (show k) <> " : " <> n) (R.defaultScribes repr)
parseAggregatedKindMap Nothing = HM.empty
parseAggregatedKindMap (Just hmv) =
  let
    listv = HM.toList hmv
    mapAggregatedKind = HM.fromList $ catMaybes $ map mkAggregatedKind listv
  in
    mapAggregatedKind
mkAggregatedKind (name, String s) = Just (name, read (unpack s) :: AggregatedKind)
mkAggregatedKind _ = Nothing

```

Setup empty configuration

```

empty :: IO Configuration
empty = do
  cgreg ← newEmptyMVar
  putMVar cgreg $ ConfigurationInternal Debug HM.empty HM.empty HM.empty HM.empty [ ] [ ] HM.empty H
  return $ Configuration cgreg

```

1.4.22 Cardano.BM.Output.Switchboard

Switchboard

```

type SwitchboardMVar = MVar SwitchboardInternal
newtype Switchboard = Switchboard
  {getSB :: SwitchboardMVar}
data SwitchboardInternal = SwitchboardInternal
  {sbQueue :: TBQ.TBQueue NamedLogItem
  ,sbDispatch :: Async.Async ()
  }

```

Trace that forwards to the Switchboard

Every **Trace** ends in the **Switchboard** which then takes care of dispatching the messages to outputs

```

mainTrace :: Switchboard → TraceNamed IO
mainTrace sb = BaseTrace.BaseTrace $ Op $ \lognamed → do
  effectuate sb lognamed

```

Process incoming messages

Incoming messages are put into the queue, and then processed by the dispatcher. The queue is initialized and the message dispatcher launched.

```

instance IsEffectuator Switchboard where
  effectuate switchboard item = do

```

```

let writequeue :: TBQ.TBQueue NamedLogItem → NamedLogItem → IO ()
  writequeue q i = do
    nocapacity ← atomically $ TBQ.isFullTBQueue q
    if nocapacity
    then return ()
    else atomically $ TBQ.writeTBQueue q i
  withMVar (getSB switchboard) $ \sb →
    writequeue (sbQueue sb) item

```

Switchboard implements Backend functions

Switchboard is an Declaration of a Backend

```

instance IsBackend Switchboard where
  typeOf _ = SwitchboardBK
  realize cfg =
    let spawnDispatcher
      :: Configuration
      → [(BackendKind, Backend)]
      → TBQ.TBQueue NamedLogItem
      → IO (Async.Async ())
    spawnDispatcher config backends queue =
      let sendMessage nli befilter = do
          selectedBackends ← getBackends config (lnName nli)
          let selBEs = befilter selectedBackends
          forM_ backends $ \ (bek, be) →
            when (bek ∈ selBEs) (bEffectuate be $ nli)
        qProc = do
          nli ← atomically $ TBQ.readTBQueue queue
          case lnItem nli of
            LogObject _ KillPill →
              forM_ backends (λ (_, be) → bUnrealize be)
            LogObject _ (AggregatedMessage _) → do
              sendMessage nli (filter (≠ AggregationBK))
              qProc
            _ → sendMessage nli id >> qProc
      in
        Async.async qProc
    in do
      q ← atomically $ TBQ.newTBQueue 2048
      sbref ← newEmptyMVar
      putMVar sbref $ SwitchboardInternal q $ error "unitialized dispatcher"
      let sb :: Switchboard = Switchboard sbref
      backends ← getSetupBackends cfg
      bs ← setupBackends backends cfg sb []
      dispatcher ← spawnDispatcher cfg bs q
      -- link the given Async to the current thread, such that if the Async

```



```

-- raises an exception, that exception will be re-thrown in the current
-- thread, wrapped in ExceptionInLinkedThread.
Async.link dispatcher
modifyMVar_sbrf $ \sbInternal → return $ sbInternal {sbDispatch = dispatcher}
return sb

unrealize switchboard = do
  let clearMVar :: MVar a → IO ()
      clearMVar = void ∘ tryTakeMVar
  (dispatcher, queue) ← withMVar (getSB switchboard) (\sb → return (sbDispatch sb, sbQueue sb))
  -- send terminating item to the queue
  lo ← LogObject <$> mkLOMeta <*> pure KillPill
  atomically $ TBQ.writeTBQueue queue $ LogNamed "kill.switchboard" lo
  -- wait for the dispatcher to exit
  res ← Async.waitCatch dispatcher
  either throwM return res
  (clearMVar ∘ getSB) switchboard

```

Realizing the backends according to configuration

```

setupBackends :: [BackendKind]
              → Configuration
              → Switchboard
              → [(BackendKind, Backend)]
              → IO [(BackendKind, Backend)]
setupBackends [] _ _ acc = return acc
setupBackends (bk : bes) c sb acc = do
  be' ← setupBackend' bk c sb
  setupBackends bes c sb ((bk, be') : acc)
setupBackend' :: BackendKind → Configuration → Switchboard → IO Backend
setupBackend' SwitchboardBK _ _ = error "cannot instantiate a further Switchboard"
setupBackend' EKGViewBK c _ = do
  be :: Cardano.BM.Output ∘ EKGView.EKGView ← Cardano.BM.Output ∘ EKGView.realize c
  return MkBackend
    { bEffectuate = Cardano.BM.Output ∘ EKGView.effectuate be
    , bUnrealize = Cardano.BM.Output ∘ EKGView.unrealize be
    }
setupBackend' AggregationBK c sb = do
  let trace = mainTrace sb
  ctx = TraceContext { loggerName = ""
                    , configuration = c
                    , minSeverity = Debug
                    , tracetype = Neutral
                    , shutdown = pure ()
                    }
  be :: Cardano.BM.Output ∘ Aggregation.Aggregation ← Cardano.BM.Output ∘ Aggregation.realizefrom (ctx,
  return MkBackend

```



```

    {bEffectuate = Cardano.BM.Output ◦ Aggregation.effectuate be
    ,bUnrealize = Cardano.BM.Output ◦ Aggregation.unrealize be
    }
  setupBackend' KatipBK c _ = do
    be :: Cardano.BM.Output ◦ Log.Log ← Cardano.BM.Output ◦ Log.realize c
    return MkBackend
      {bEffectuate = Cardano.BM.Output ◦ Log.effectuate be
      ,bUnrealize = Cardano.BM.Output ◦ Log.unrealize be
      }

```

1.4.23 Cardano.BM.Output.Log

Internal representation

```

type LogMVar = MVar LogInternal
newtype Log = Log
  {getK :: LogMVar}
data LogInternal = LogInternal
  {kLogEnv :: K.LogEnv
  ,configuration :: Config.Configuration}

```

Log implements effectuate

```

instance IsEffectuator Log where
  effectuate katip item = do
    c ← withMVar (getK katip) $ λk → return (configuration k)
    selscribes ← getScribes c (lnName item)
    forM_ selscribes $ λsc → passN sc katip item

```

Log implements backend functions

```

instance IsBackend Log where
  typeOf _ = KatipBK
  realize config = do
    let updateEnv :: K.LogEnv → IO UTCTime → K.LogEnv
        updateEnv le timer =
          le {K._logEnvTimer = timer, K._logEnvHost = "hostname"}
    register :: [ScribeDefinition] → K.LogEnv → IO K.LogEnv
    register [ ] le = return le
    register (defsc : dscs) le = do
      let kind = scKind defsc
          name = scName defsc
          name' = pack (show kind) <> " : " <> name
          scr ← createScribe kind name

```

```

    register dscs ≡ K.registerScribe name' scr scribeSettings le
mockVersion :: Version
mockVersion = Version [0,1,0,0] []
scribeSettings :: KC.ScribeSettings
scribeSettings =
    let bufferSize = 5000 -- size of the queue (in log items)
    in
        KC.ScribeSettings bufferSize
createScribe FileTextSK name = mkTextFileScribe (FileDescription $ unpack name) False
createScribe FileJsonSK name = mkJsonFileScribe (FileDescription $ unpack name) False
createScribe StdoutSK _ = mkStdoutScribe
createScribe StderrSK _ = mkStderrScribe
cfoKey ← Config.getOptionOrDefault config (pack "cfokey") (pack "<unknown>")
le0 ← K.initLogEnv
    (K.Namespace [ "iohk" ])
    (fromString $ (unpack cfoKey) <> ":" <> showVersion mockVersion)
-- request a new time 'getCurrentTime' at most 100 times a second
timer ← mkAutoUpdate defaultUpdateSettings {updateAction = getCurrentTime, updateFreq = 10000}
let le1 = updateEnv le0 timer
scribes ← getSetupScribes config
le ← register scribes le1
kref ← newEmptyMVar
putMVar kref $ LogInternal le config
return $ Log kref
unrealize katip = do
    le ← withMVar (getK katip) $ λk → return (kLogEnv k)
    void $ K.closeScribes le

example :: IO ()
example = do
    config ← Config.setup "from_some_path.yaml"
    k ← setup config
    passN (pack (show StdoutSK)) k $ LogNamed
        {lnName = "test"
        ,lnItem = LogMessage $ LogItem
            {liSelection = Both
            ,liSeverity = Info
            ,liPayload = "Hello!"
            }
        }
    passN (pack (show StdoutSK)) k $ LogNamed
        {lnName = "test"
        ,lnItem = LogValue "cpu-no" 1
        }

```

Needed instances for *katip*:

```

deriving instance K.ToObject LogObject
deriving instance K.ToObject LogItem
deriving instance K.ToObject (Maybe LOContent)
instance KC.LogItem LogObject where
    payloadKeys _ = KC.AllKeys
instance KC.LogItem LogItem where
    payloadKeys _ = KC.AllKeys
instance KC.LogItem (Maybe LOContent) where
    payloadKeys _ = KC.AllKeys

```

Log.passN

The following function copies the **NamedLogItem** to the queues of all scribes that match on their name. Compare start of name of scribe to (*show backend* <> " : "). This function is non-blocking.

```

passN :: Text → Log → NamedLogItem → IO ()
passN backend katip namedLogItem = do
    env ← withMVar (getK katip) $ λk → return (kLogEnv k)
    forM_ (Map.toList $ K._logEnvScribes env) $
        λ(scName, (KC.ScribeHandle _ shChan)) →
            -- check start of name to match ScribeKind
            if backend 'isPrefixOf' scName
            then do
                let (LogObject lometa loitem) = lnItem namedLogItem
                let (sev, msg, payload) = case loitem of
                    (LogMessage logItem) →
                        (liSeverity logItem, liPayload logItem, Nothing)
                    (ObserveDiff _) →
                        let text = toStrict (encodeToLazyText loitem)
                        in
                            (Info, text, Just loitem)
                    (ObserveOpen _) →
                        let text = toStrict (encodeToLazyText loitem)
                        in
                            (Info, text, Just loitem)
                    (ObserveClose _) →
                        let text = toStrict (encodeToLazyText loitem)
                        in
                            (Info, text, Just loitem)
                    (AggregatedMessage aggregated) →
                        let text = T.concat $ (flip map) aggregated $ λ(name, agg) →
                            "\n" <> name <> " : " <> pack (show agg)
                        in
                            (Info, text, Nothing)
                (LogValue name value) →
                    (Debug, name <> " = " <> pack (showSI value), Nothing)
                KillPill →

```

```

        (Info, "Kill pill received!", Nothing)
    if (msg == "") ^ (isNothing payload)
    then return ()
    else do
        let threadIdText = KC.mkThreadIdText (tid lometa)
        let ns = lnName namedLogItem
        let itemTime = tstamp lometa
        let itemKatip = K.Item {
            _itemApp      = env ^. KC.logEnvApp
            , _itemEnv     = env ^. KC.logEnvEnv
            , _itemSeverity = sev2klog sev
            , _itemThread  = threadIdText
            , _itemHost    = env ^. KC.logEnvHost
            , _itemProcess = env ^. KC.logEnvPid
            , _itemPayload = payload
            , _itemMessage = K.logStr msg
            , _itemTime    = itemTime
            , _itemNamespace = (env ^. KC.logEnvApp) <> (K.Namespace [ ns ])
            , _itemLoc     = Nothing
        }
        void $ atomically $ KC.tryWriteTBQueue shChan (KC.NewItem itemKatip)
    else return ()

```

Scribes

```

mkStdoutScribe :: IO K.Scribe
mkStdoutScribe = do
    -- duplicate stdout so that Katip's closing
    -- action will not close the real stdout
    stdout' ← hDuplicate stdout
    mkTextFileScribeH stdout' True

mkStderrScribe :: IO K.Scribe
mkStderrScribe = do
    -- duplicate stderr so that Katip's closing
    -- action will not close the real stderr
    stderr' ← hDuplicate stderr
    mkTextFileScribeH stderr' True

mkTextFileScribeH :: Handle → Bool → IO K.Scribe
mkTextFileScribeH handler color = do
    mkFileScribeH handler formatter color
    where
        formatter h colorize verbosity item =
            TIO.hPutStrLn h $! toLazyText $ formatItem colorize verbosity item

mkFileScribeH
    :: Handle
    → (forall a ◦ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ())

```

```

    → Bool
    → IO K.Scribe
mkFileScribeH h formatter colorize = do
  hSetBuffering h LineBuffering
  locklocal ← newMVar ()
  let logger :: forall a ◦ K.LogItem a ⇒ K.Item a → IO ()
      logger item = withMVar locklocal $ \_ →
        formatter h colorize K.V0 item
  pure $ K.Scribe logger (hClose h)
mkTextFileScribe :: FileDescription → Bool → IO K.Scribe
mkTextFileScribe fdesc colorize = do
  mkFileScribe fdesc formatter colorize
where
  formatter :: Handle → Bool → K.Verbosity → K.Item a → IO ()
  formatter hdl colorize' v' item =
    case KC._itemMessage item of
      K.LogStr "" →
        -- if message is empty do not output it
        return ()
    _ → do
      let tmsg = toLazyText $ formatItem colorize' v' item
      TIO.hPutStrLn hdl tmsg
mkJsonFileScribe :: FileDescription → Bool → IO K.Scribe
mkJsonFileScribe fdesc colorize = do
  mkFileScribe fdesc formatter colorize
where
  formatter :: (K.LogItem a) ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ()
  formatter h _ verbosity item = do
    let tmsg = case KC._itemMessage item of
      -- if a message is contained in item then only the
      -- message is printed and not the data
      K.LogStr "" → K.itemJson verbosity item
      K.LogStr msg → K.itemJson verbosity $
        item { KC._itemMessage = K.logStr (" " :: Text)
          , KC._itemPayload = LogItem Both Info $ toStrict $ toLazyText msg
        }
    TIO.hPutStrLn h (encodeToLazyText tmsg)
mkFileScribe
  :: FileDescription
  → (forall a ◦ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ())
  → Bool
  → IO K.Scribe
mkFileScribe fdesc formatter colorize = do
  let prefixDir = prefixPath fdesc
  (createDirectoryIfMissing True prefixDir)
  'catchIO' (prtoutException ("cannot log prefix directory: " ++ prefixDir))
  let fpath = filePath fdesc

```

```

h ← catchIO (openFile fpath WriteMode) $
  λe → do
    prtoutException ("error while opening log: " ++ fpath) e
    -- fallback to standard output in case of exception
    return stdout
hSetBuffering h LineBuffering
scribestate ← newMVar h
let finalizer :: IO ()
    finalizer = withMVar scribestate hClose
let logger :: forall a ◦ K.LogItem a ⇒ K.Item a → IO ()
    logger item =
      withMVar scribestate $ λhandler →
        formatter handler colorize K.V0 item
return $ K.Scribe logger finalizer

```

```

formatItem :: Bool → K.Verbosity → K.Item a → Builder
formatItem withColor _verb K.Item {..} =
  fromText header <>
  fromText " " <>
  brackets (fromText timestamp) <>
  fromText " " <>
  KC.unLogStr _itemMessage
where
  header = colorBySeverity _itemSeverity $
    "[ " <> mconcat [namedcontext <> ":" <> severity <> ":" <> threadid <> "]"
  namedcontext = KC.intercalateNs _itemNamespace
  severity = KC.renderSeverity _itemSeverity
  threadid = KC.getThreadIdText _itemThread
  timestamp = pack $ formatTime defaultTimeLocale tsformat _itemTime
  tsformat :: String
  tsformat = "%F %T%2Q %Z"
  colorBySeverity s m = case s of
    K.EmergencyS → red m
    K.AlertS     → red m
    K.CriticalS  → red m
    K.ErrorS     → red m
    K.NoticeS    → magenta m
    K.WarningS   → yellow m
    K.InfoS      → blue m
    _            → m
  red = colorize "31"
  yellow = colorize "33"
  magenta = colorize "35"
  blue = colorize "34"
  colorize c m
    | withColor = "\ESC[" <> c <> "m" <> m <> "\ESC[0m"
    | otherwise = m

```

```

-- translate Severity to Log.Severity
sev2klog :: Severity → K.Severity
sev2klog = λcase
    Debug    → K.DebugS
    Info     → K.InfoS
    Notice   → K.NoticeS
    Warning  → K.WarningS
    Error    → K.ErrorS
    Critical → K.CriticalS
    Alert    → K.AlertS
    Emergency → K.EmergencyS

data FileDescription = FileDescription {
    filePath :: !FilePath}
    deriving (Show)

prefixPath :: FileDescription → FilePath
prefixPath = takeDirectory ∘ filePath

-- display message and stack trace of exception on stdout
prtoutException :: Exception e ⇒ String → e → IO ()
prtoutException msg e = do
    putStrLn msg
    putStrLn ("exception: " ++ displayException e)

```

1.4.24 Cardano.BM.Output.EKGView

Structure of EKGView

```

type EKGViewMVar = MVar EKGViewInternal
newtype EKGView = EKGView
    {getEV :: EKGViewMVar}
data EKGViewInternal = EKGViewInternal
    {evQueue :: TBQ.TBQueue (Maybe NamedLogItem)
    ,evLabels :: EKGViewMap
    ,evServer :: Server
    }

```

Relation from variable name to label handler

We keep the label handlers for later update in a *HashMap*.

```

type EKGViewMap = HM.HashMap Text Label.Label

```

Internal Trace

This is an internal **Trace**, named "#ekgview", which can be used to control the messages that are being displayed by EKG.

```

ekgTrace :: EKGView → Configuration → IO (Trace IO)
ekgTrace ekg c = do
  let trace = ekgTrace' ekg
      ctx = TraceContext { loggerName = " "
                          , configuration = c
                          , minSeverity = Debug
                          , tracetype = Neutral
                          , shutdown = pure ()
                          }
      Trace.subTrace "#ekgview" (ctx, trace)
  where
    ekgTrace' :: EKGView → TraceNamed IO
    ekgTrace' ekgview = BaseTrace.BaseTrace $ Op $ λ(LogNamed lognamed lo) → do
      let setlabel :: Text → Text → EKGViewInternal → IO (Maybe EKGViewInternal)
          setlabel name label ekg_i@(EKGViewInternal _ labels server) =
            case HM.lookup name labels of
              Nothing → do
                ekghdl ← getLabel name server
                Label.set ekghdl label
                return $ Just $ ekg_i { evLabels = HM.insert name ekghdl labels }
              Just ekghdl → do
                Label.set ekghdl label
                return Nothing
      update :: LogObject → LoggerName → EKGViewInternal → IO (Maybe EKGViewInternal)
      update (LogObject _ (LogMessage logitem)) logname ekg_i =
        setlabel logname (liPayload logitem) ekg_i
      update (LogObject _ (LogValue iname value)) logname ekg_i =
        let logname' = logname <> "." <> iname
        in
          setlabel logname' (pack $ show value) ekg_i
      update _ _ _ = return Nothing
    ekgup ← takeMVar (getEV ekgview)
    let -- strip off some prefixes not necessary for display
        lognam1 = case stripPrefix "#ekgview.#aggregation." lognamed of
          Nothing → lognamed
          Just ln' → ln'
        logname = case stripPrefix "#ekgview." lognam1 of
          Nothing → lognam1
          Just ln' → ln'
    upd ← update lo logname ekgup
    case upd of
      Nothing → putMVar (getEV ekgview) ekgup
      Just ekgup' → putMVar (getEV ekgview) ekgup'

```


EKG view is an effectuator

Function *effectuate* is called to pass in a **NamedLogItem** for display in EKG. If the log item is an *AggregatedStats* message, then all its constituents are put into the queue.

instance IsEffectuator EKGView where

```

effectuate ekgview item = do
  ekg ← readMVar (getEV ekgview)
  let queue a = do
    nocapacity ← atomically $ TBQ.isFullTBQueue (evQueue ekg)
    if nocapacity
    then return ()
    else atomically $ TBQ.writeTBQueue (evQueue ekg) (Just a)
  case (lnItem item) of
    (LogObject lometa (AggregatedMessage ags)) → liftIO $ do
      let logname = lnName item
      traceAgg :: [(Text, Aggregated)] → IO ()
      traceAgg [] = return ()
      traceAgg ((n, AggregatedEWMA ewma) : r) = do
        queue $ LogNamed (logname <> "." <> n) $ LogObject lometa (LogValue "avg" $ avg ewma)
        traceAgg r
      traceAgg ((n, AggregatedStats stats) : r) = do
        let statsname = logname <> "." <> n
        qbasestats s' nm = do
          queue $ LogNamed nm $ LogObject lometa (LogValue "mean" (PureD $ meanOfStats s'))
          queue $ LogNamed nm $ LogObject lometa (LogValue "min" $ fmin s')
          queue $ LogNamed nm $ LogObject lometa (LogValue "max" $ fmax s')
          queue $ LogNamed nm $ LogObject lometa (LogValue "count" $ PureI $ fromIntegral $ fcoun
          queue $ LogNamed nm $ LogObject lometa (LogValue "stdev" (PureD $ stdevOfStats s'))
        queue $ LogNamed statsname $ LogObject lometa (LogValue "last" $ flast stats)
        qbasestats (fbasic stats) $ statsname <> ".basic"
        qbasestats (fdelta stats) $ statsname <> ".delta"
        qbasestats (ftimed stats) $ statsname <> ".timed"
        traceAgg r
      traceAgg ags
    (LogObject _ (LogMessage _)) → queue item
    (LogObject _ (LogValue _)) → queue item
    _ → return ()

```

EKGView implements Backend functions

EKGView is an **IsBackend**

instance IsBackend EKGView where

```

typeof _ = EKGViewBK
realize config = do
  evref ← newEmptyMVar
  let ekgview = EKGView evref

```

```

evport ← getEKGport config
ehdl ← forkServer "127.0.0.1" evport
ekghdl ← getLabel "iohk-monitoring version" ehdl
Label.set ekghdl $ pack (showVersion version)
ekgtrace ← ekgTrace ekgview config
queue ← atomically $ TBQ.newTBQueue 512
dispatcher ← spawnDispatcher queue ekgtrace
-- link the given Async to the current thread, such that if the Async
-- raises an exception, that exception will be re-thrown in the current
-- thread, wrapped in ExceptionInLinkedThread.
Async.link dispatcher
putMVar evref $ EKGViewInternal
  {evLabels = HM.empty
  ,evServer = ehdl
  ,evQueue = queue
  }
return ekgview
unrealize ekgview = do
  ekg ← takeMVar $ getEV ekgview
  killThread $ serverThreadId $ evServer ekg

```

Asynchronously reading log items from the queue and their processing

```

spawnDispatcher :: TBQ.TBQueue (Maybe LogNamedItem)
  → Trace.Trace IO
  → IO (Async.Async ())
spawnDispatcher evqueue trace =
  Async.async $ qProc
where
  qProc = do
    maybeItem ← atomically $ TBQ.readTBQueue evqueue
    case maybeItem of
      Just (LogNamed logname logvalue) → do
        trace' ← Trace.appendName logname trace
        Trace.traceNamedObject trace' logvalue
        qProc
      Nothing → return () -- stop here

```

Interactive testing EKGView

```

test :: IO ()
test = do
  c ← Cardano.BM.Setup.setupTrace (Left "test/config.yaml") "ekg"
  ev ← Cardano.BM.Output ◦ EKGView.realize c
  effectuate ev $ LogNamed "test.questions" (LogValue "answer" 42)
  effectuate ev $ LogNamed "test.monitor023" (LogMessage (LogItem Public Warning "!!!! ALARM !!!!!"

```

1.4.25 Cardano.BM.Output.Aggregation

Internal representation

```

type AggregationMVar = MVar AggregationInternal
newtype Aggregation = Aggregation
    {getAg :: AggregationMVar}
data AggregationInternal = AggregationInternal
    {agQueue :: TBQ.TBQueue (Maybe NamedLogItem)
    ,agDispatch :: Async.Async ()
    }

```

Relation from context name to aggregated statistics

We keep the aggregated values (**Aggregated**) for a named context in a *HashMap*.

```

type AggregationMap = HM.HashMap Text AggregatedExpanded

```

Info for Aggregated operations

Apart from the **Aggregated** we keep some valuable info regarding to them; such as when was the last time it was sent.

```

type Timestamp = Word64
data AggregatedExpanded = AggregatedExpanded
    {aeAggregated :: !Aggregated
    ,aeResetAfter :: !(Maybe Int)
    ,aeLastSent :: {-# UNPACK #-} !Timestamp
    }

```

Aggregation implements *effectuate*

Aggregation is an **Accepts a NamedLogItem** Enter the log item into the **Aggregation** queue.

```

instance IsEffectuator Aggregation where
    effectuate agg item = do
        ag ← readMVar (getAg agg)
        nocapacity ← atomically $ TBQ.isFullTBQueue (agQueue ag)
        if nocapacity
        then return ()
        else atomically $! TBQ.writeTBQueue (agQueue ag) $ Just item

```

Aggregation implements Backend functions

Aggregation is an Declaration of a Backend

instance IsBackend Aggregation where

typeof _ = AggregationBK

realize _ = error "Aggregation cannot be instantiated by 'realize'"

realizefrom trace0@(ctx, _) _ = do

trace ← Trace.subTrace "#aggregation" trace0

aggref ← newEmptyMVar

aggregationQueue ← atomically \$ TBQ.newTBQueue 2048

dispatcher ← spawnDispatcher (configuration ctx) HM.empty aggregationQueue trace

-- link the given Async to the current thread, such that if the Async

-- raises an exception, that exception will be re-thrown in the current

-- thread, wrapped in ExceptionInLinkedThread.

Async.link dispatcher

putMVar aggref \$ AggregationInternal aggregationQueue dispatcher

return \$ Aggregation aggref

unrealize aggregation = do

let clearMVar :: MVar a → IO ()

clearMVar = void ∘ tryTakeMVar

(dispatcher, queue) ← withMVar (getAg aggregation) (λag →

return (agDispatch ag, agQueue ag))

-- send terminating item to the queue

atomically \$ TBQ.writeTBQueue queue Nothing

-- wait for the dispatcher to exit

-- TODO add a timeout to waitCatch in order

-- to be sure that it will finish

res ← Async.waitCatch dispatcher

either throwM return res

(clearMVar ∘ getAg) aggregation

Asynchronously reading log items from the queue and their processing

spawnDispatcher :: Configuration

→ AggregationMap

→ TBQ.TBQueue (Maybe NamedLogItem)

→ Trace.Trace IO

→ IO (Async.Async ())

spawnDispatcher conf aggMap aggregationQueue trace = Async.async \$ qProc aggMap

where

qProc aggregatedMap = do

maybeItem ← atomically \$ TBQ.readTBQueue aggregationQueue

case maybeItem of

Just (LogNamed logname lo@(LogObject lm _)) → do

(updatedMap, aggregations) ← update lo logname aggregatedMap

```

    unless (null aggregations) $
        sendAggregated (LogObject lm (AggregatedMessage aggregations)) logname
    qProc updatedMap
    Nothing → return ()
createNupdate name value lme agmap = do
    case HM.lookup name agmap of
        Nothing → do
            -- if Aggregated does not exist; initialize it.
            aggregatedKind ← getAggregatedKind conf name
            case aggregatedKind of
                StatsAK → return $ singletonStats value
                EwmaAK aEWMA → do
                    let initEWMA = EmptyEWMA aEWMA
                    return $ AggregatedEWMA $ ewma initEWMA value
            Just a → return $ updateAggregation value (aeAggregated a) lme (aeResetAfter a)
update :: LogObject
    → LoggerName
    → AggregationMap
    → IO (AggregationMap, [(Text, Aggregated)])
update (LogObject lme (LogValue iname value)) logname agmap = do
    let fullname = logname <> " ." <> iname
    aggregated ← createNupdate fullname value lme agmap
    now ← getMonotonicTimeNSec
    let aggregatedX = AggregatedExpanded {
        aeAggregated = aggregated
        , aeResetAfter = Nothing
        , aeLastSent = now
    }
    namedAggregated = [(iname, aeAggregated aggregatedX)]
    updatedMap = HM.alter (const $ Just $ aggregatedX) fullname agmap
    return (updatedMap, namedAggregated)
update (LogObject lme (ObserveDiff counterState)) logname agmap =
    updateCounters (csCounters counterState) lme (logname, "diff") agmap []
update (LogObject lme (ObserveOpen counterState)) logname agmap =
    updateCounters (csCounters counterState) lme (logname, "open") agmap []
update (LogObject lme (ObserveClose counterState)) logname agmap =
    updateCounters (csCounters counterState) lme (logname, "close") agmap []
update (LogObject lme (LogMessage msg)) logname agmap = do
    let iname = T.pack $ show (liSeverity msg)
    let fullname = logname <> " ." <> iname
    aggregated ← createNupdate fullname (PureI 0) lme agmap
    now ← getMonotonicTimeNSec
    let aggregatedX = AggregatedExpanded {
        aeAggregated = aggregated
        , aeResetAfter = Nothing
        , aeLastSent = now
    }

```

```

    namedAggregated = [(iname, aeAggregated aggregatedX)]
    updatedMap = HM.alter (const $ Just $ aggregatedX) fullname agmap
    return (updatedMap, namedAggregated)
-- everything else
update _ _ agmap = return (agmap, [])
updateCounters :: [Counter]
    → LOMeta
    → (LoggerName, LoggerName)
    → AggregationMap
    → [(Text, Aggregated)]
    → IO (AggregationMap, [(Text, Aggregated)])
updateCounters [] _ _ aggrMap aggs = return $ (aggrMap, aggs)
updateCounters (counter : cs) lme (logname, msgname) aggrMap aggs = do
    let name = cName counter
        subname = msgname <> "." <> (nameCounter counter) <> "." <> name
        fullname = logname <> "." <> subname
        value = cValue counter
    aggregated ← createNupdate fullname value lme aggrMap
    now ← getMonotonicTimeNSec
    let aggregatedX = AggregatedExpanded {
        aeAggregated = aggregated
        , aeResetAfter = Nothing
        , aeLastSent = now
        }
        namedAggregated = (subname, aggregated)
        updatedMap = HM.alter (const $ Just $ aggregatedX) fullname aggrMap
    updateCounters cs lme (logname, msgname) updatedMap (namedAggregated : aggs)
sendAggregated :: LogObject → Text → IO ()
sendAggregated aggregatedMsg@(LogObject _ (AggregatedMessage _)) logname = do
    -- enter the aggregated message into the Trace
    trace' ← Trace.appendName logname trace
    liftIO $ Trace.traceNamedObject trace' aggregatedMsg
-- ignore every other message
sendAggregated _ _ = return ()

```

Update aggregation

We distinguish an uninitialized from an already initialized aggregation. The latter is properly initialized.

We use Welford's online algorithm to update the estimation of mean and variance of the sample statistics. (see https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance#Welford's_Online)

```

updateAggregation :: Measurable → Aggregated → LOMeta → Maybe Int → Aggregated
updateAggregation v (AggregatedStats s) lme resetAfter =
    let count = fcount (fbasic s)
        reset = maybe False (count ≥) resetAfter
    in

```

```

if reset
then
  singletonStats v
else
  AggregatedStats $! Stats {flast = v
    ,fold = mkTimestamp
    ,fbasic = updateBaseStats (count ≥ 1) v (fbasic s)
    ,fdelta = updateBaseStats (count ≥ 2) (v – flast s) (fdelta s)
    ,ftimed = updateBaseStats (count ≥ 2) (mkTimestamp – fold s) (ftimed s)
  }
where
  mkTimestamp = utc2ns (tstamp lme)
  utc2ns (UTCTime days secs) =
    let yearsecs :: Rational
      yearsecs = 365 * 24 * 3600
      rdays, rsecs :: Rational
      rdays = toRational $ toModifiedJulianDay days
      rsecs = toRational secs
      s2ns = 1000000000
    in
      Nanoseconds $ round $ (fromRational $ s2ns * rsecs + rdays * yearsecs :: Double)
  updateAggregation v (AggregatedEWMA e) _ _ = AggregatedEWMA $! ewma e v
  updateBaseStats :: Bool → Measurable → BaseStats → BaseStats
  updateBaseStats False _ s = s {fcount = fcount s + 1}
  updateBaseStats True v s =
    let newcount = fcount s + 1
      newvalue = getDouble v
      delta = newvalue – fsum_A s
      dincr = (delta / fromIntegral newcount)
      delta2 = newvalue – fsum_A s – dincr
    in
      BaseStats {fmin = min (fmin s) v
        ,fmax = max v (fmax s)
        ,fcount = newcount
        ,fsum_A = fsum_A s + dincr
        ,fsum_B = fsum_B s + (delta * delta2)
      }

```

Calculation of EWMA

Following https://en.wikipedia.org/wiki/Moving_average#Exponential_moving_average we calculate the exponential moving average for a series of values Y_t according to:

$$S_t = \begin{cases} Y_1, & t = 1 \\ \alpha \cdot Y_t + (1 - \alpha) \cdot S_{t-1}, & t > 1 \end{cases}$$

The pattern matching below ensures that the EWMA will start with the first value passed in, and

will not change type, once determined.

```

ewma :: EWMA → Measurable → EWMA
ewma (EmptyEWMA a) v = EWMA a v
ewma (EWMA a s@(Microseconds _)) y@(Microseconds _) =
    EWMA a $ Microseconds $ round $ a * (getDouble y) + (1 - a) * (getDouble s)
ewma (EWMA a s@(Seconds _)) y@(Seconds _) =
    EWMA a $ Seconds $ round $ a * (getDouble y) + (1 - a) * (getDouble s)
ewma (EWMA a s@(Bytes _)) y@(Bytes _) =
    EWMA a $ Bytes $ round $ a * (getDouble y) + (1 - a) * (getDouble s)
ewma (EWMA a (PureI s)) (PureI y) =
    EWMA a $ PureI $ round $ a * (fromInteger y) + (1 - a) * (fromInteger s)
ewma (EWMA a (PureD s)) (PureD y) =
    EWMA a $ PureD $ a * y + (1 - a) * s
ewma _ _ = error "Cannot average on values of different type"

```


Chapter 2

Testing

2.1 Test coverage

Test coverage is calculated as the fraction of functions which are called from test routines. This percentage is calculated by the tool *hpc* with a call to

```
cabal new-test
```

Add to a local `cabal.project.local` file these lines:

```
tests:           True
coverage:        True
library-coverage: True
```

2.2 Test main entry point

```
module Main
(
    main
) where
import Test.Tasty
import qualified Cardano.BM.Test.Agregated (tests)
import qualified Cardano.BM.Test.STM (tests)
import qualified Cardano.BM.Test.Trace (tests)
import qualified Cardano.BM.Test.Configuration (tests)
import qualified Cardano.BM.Test.Routing (tests)
main :: IO ()
main = defaultMain tests
tests :: TestTree
tests =
    testGroup "iohk-monitoring"
    [ Cardano.BM.Test.Agregated.tests
    , Cardano.BM.Test.STM.tests
    , Cardano.BM.Test.Trace.tests
    , Cardano.BM.Test.Configuration.tests
```

Cardano.BM.Counters.Dummy	100%
Cardano.BM.Setup	100%
Cardano.BM.Data.Trace	100%
Cardano.BM.Counters.Common	100%
Cardano.BM.Counters	100%
Cardano.BM.Configuration	100%
Cardano.BM.Output.Switchboard	90%
Cardano.BM.Data.Configuration	83%
Cardano.BM.BaseTrace	80%
Cardano.BM.Configuration.Model	79%
Cardano.BM.Observer.Monadic	75%
Cardano.BM.Output.Aggregation	73%
Cardano.BM.Output.Log	66%
Cardano.BM.Data.Aggregated	64%
Cardano.BM.Data.Severity	63%
Cardano.BM.Data.Counter	56%
Cardano.BM.Data.Output	50%
Cardano.BM.Data.BackendKind	50%
Cardano.BM.Data.Backend	50%
Cardano.BM.Configuration.Static	50%
Cardano.BM.Data.LogItem	46%
Cardano.BM.Observer.STM	33%
Cardano.BM.Data.AggregatedKind	33%
Cardano.BM.Trace	31%
Cardano.BM.Data.Observable	20%
Cardano.BM.Data.SubTrace	10%
Cardano.BM.Data.Rotation	10%
Cardano.BM.Output.EKGView	0%
Paths_iohk_monitoring	0%
	52%

Figure 2.1: Test coverage of modules in percent as computed by the tool 'hpc'

```
,Cardano.BM.Test ◦ Routing.tests
]
```

2.3 Test case generation

2.3.1 instance Arbitrary Aggregated

We define an instance of *Arbitrary* for an **Aggregated** which lets *QuickCheck* generate arbitrary instances of **Aggregated**. For this an arbitrary list of *Integer* is generated and this list is aggregated into a structure of **Aggregated**.

```
instance Arbitrary Aggregated where
  arbitrary = do
    vs' ← arbitrary :: Gen [Integer]
    let vs = 42 : 17 : vs'
        ds = map (\(a,b) → a - b) $ zip vs (tail vs)
        (m1,s1) = updateMeanVar $ map fromInteger vs
        (m2,s2) = updateMeanVar $ map fromInteger ds
        mkBasicStats = BaseStats
          (PureI (minimum vs))
          (PureI (maximum vs))
          (fromIntegral $ length vs)
          (m1)
          (s1)
        mkDeltaStats = BaseStats
          (PureI (minimum ds))
          (PureI (maximum ds))
          (fromIntegral $ length ds)
          (m2)
          (s2)
        mkTimedStats = BaseStats
          (Nanoseconds 0)
          (Nanoseconds 0)
          (0)
          (0)
          (0)
    return $ AggregatedStats (Stats
      (PureI (last vs))
      (Nanoseconds 0)
      mkBasicStats
      mkDeltaStats
      mkTimedStats)
```

Estimators for mean and variance must be updated the same way as in the code.

```
updateMeanVar :: [Double] → (Double, Double)
updateMeanVar [] = (0, 0)
updateMeanVar (val : vals) = updateMeanVar' (val, 0) 1 vals
```

```

where
  updateMeanVar' (m,s) _ [] = (m,s)
  updateMeanVar' (m,s) cnt (a:r) =
    let delta = a - m
        newcount = cnt + 1
        m' = m + (delta / newcount)
        s' = s + (delta * (a - m'))
    in
      updateMeanVar' (m',s') newcount r

```

2.4 Tests

2.4.1 Testing aggregation

```

tests :: TestTree
tests = testGroup "aggregation measurements" [
  property_tests
, unit_tests
]

property_tests :: TestTree
property_tests = testGroup "Properties" [
  testProperty "minimal" prop_Aggregation_minimal
, testProperty "commutative" prop_Aggregation_comm
]

unit_tests :: TestTree
unit_tests = testGroup "Unit tests" [
  testCase "initial_minus_1" unit_Aggregation_initial_minus_1
, testCase "initial_plus_1" unit_Aggregation_initial_plus_1
, testCase "initial_0" unit_Aggregation_initial_zero
, testCase "initial_plus_1" unit_Aggregation_initial_plus_1_minus_1
, testCase "stepwise" unit_Aggregation_stepwise
]

prop_Aggregation_minimal :: Bool
prop_Aggregation_minimal = True

lometa :: LOMeta
lometa = unsafePerformIO $ mkLOMeta

prop_Aggregation_comm :: Integer → Integer → Aggregated → Bool
prop_Aggregation_comm v1 v2 ag =
  let AggregatedStats stats1 = updateAggregation (PureI v1) (updateAggregation (PureI v2) ag lometa Not)
      AggregatedStats stats2 = updateAggregation (PureI v2) (updateAggregation (PureI v1) ag lometa Not)
  in
    fbasic stats1 ≡ fbasic stats2 ∧
    (v1 ≡ v2) 'implies' (flast stats1 ≡ flast stats2)

-- implication: if p1 is true, then return p2; otherwise true
implies :: Bool → Bool → Bool

```

```

implies p1 p2 = (¬ p1) ∨ p2
unit_Aggregation_initial_minus_1 :: Assertion
unit_Aggregation_initial_minus_1 = do
  let AggregatedStats stats1 = updateAggregation (-1) firstStateAggregatedStats lometa Nothing
  flast stats1 @? = (-1)
  (fbasic stats1) @? = BaseStats (-1) 0 2 (-0.5) 0.5
  (fdelta stats1) @? = BaseStats 0 0 1 0 0
  -- AggregatedStats (Stats (-1) 0 (BaseStats (-1) 0 2 (-0.5) 0.5) (BaseStats 0 0 1 0 0))
unit_Aggregation_initial_plus_1 :: Assertion
unit_Aggregation_initial_plus_1 = do
  let AggregatedStats stats1 = updateAggregation 1 firstStateAggregatedStats lometa Nothing
  flast stats1 @? = 1
  (fbasic stats1) @? = BaseStats 0 1 2 0.5 0.5
  (fdelta stats1) @? = BaseStats 0 0 1 0 0
  -- AggregatedStats (Stats 1 0 (BaseStats 0 1 2 0.5 0.5) (BaseStats 0 0 1 0 0)) (BaseStats 0 0 1 0 0)
unit_Aggregation_initial_zero :: Assertion
unit_Aggregation_initial_zero = do
  let AggregatedStats stats1 = updateAggregation 0 firstStateAggregatedStats lometa Nothing
  flast stats1 @? = 0
  (fbasic stats1) @? = BaseStats 0 0 2 0 0
  (fdelta stats1) @? = BaseStats 0 0 1 0 0
  -- AggregatedStats (Stats 0 0 (BaseStats 0 0 2 0 0) (BaseStats 0 0 1 0 0)) (BaseStats 0 0 1 0 0)
unit_Aggregation_initial_plus_1_minus_1 :: Assertion
unit_Aggregation_initial_plus_1_minus_1 = do
  let AggregatedStats stats1 = updateAggregation (-1) (updateAggregation 1 firstStateAggregatedStats lometa Nothing)
  (fbasic stats1) @? = BaseStats (-1) 1 3 0.0 2.0
  (fdelta stats1) @? = BaseStats (-2) 0 2 (-1.0) 2.0
unit_Aggregation_stepwise :: Assertion
unit_Aggregation_stepwise = do
  stats0 ← pure $ singletonStats (Bytes 3000)
  putStrLn $ show stats0
  threadDelay 50000 -- 0.05 s
  t1 ← mkLOMeta
  stats1 ← pure $ updateAggregation (Bytes 5000) stats0 t1 Nothing
  putStrLn $ show stats1
  showTimedMean stats1
  threadDelay 50000 -- 0.05 s
  t2 ← mkLOMeta
  stats2 ← pure $ updateAggregation (Bytes 1000) stats1 t2 Nothing
  putStrLn $ show stats2
  showTimedMean stats2
  checkTimedMean stats2
  threadDelay 50000 -- 0.05 s
  t3 ← mkLOMeta
  stats3 ← pure $ updateAggregation (Bytes 3000) stats2 t3 Nothing
  putStrLn $ show stats3
  showTimedMean stats3

```

```

    checkTimedMean stats3
    threadDelay 50000 -- 0.05 s
    t4 ← mkLOMeta
    stats4 ← pure $ updateAggregation (Bytes 1000) stats3 t4 Nothing
    putStrLn $ show stats4
    showTimedMean stats4
    checkTimedMean stats4
  where
    checkTimedMean (AggregatedEWMA _) = return ()
    checkTimedMean (AggregatedStats s) = do
      let mean = meanOfStats (ftimed s)
      assertBool "the mean should be >= the minimum" (mean ≥ getDouble (fmin (ftimed s)))
      assertBool "the mean should be <= the maximum" (mean ≤ getDouble (fmax (ftimed s)))
      showTimedMean (AggregatedEWMA _) = return ()
      showTimedMean (AggregatedStats s) = putStrLn $ "mean = " ++ show (meanOfStats (ftimed s)) ++ showUnit
firstStateAggregatedStats :: Aggregated
firstStateAggregatedStats = AggregatedStats (Stats 0 0 (BaseStats 0 0 1 0 0) (BaseStats 0 0 0 0 0) (BaseStats 0 0 0 0 0))

```

2.4.2 STM

```

module Cardano.BM.Test.STM (
  tests
) where
import Test.Tasty
import Test.Tasty.QuickCheck
tests :: TestTree
tests = testGroup "observing STM actions" [
  testProperty "minimal" prop_STM_observer
]
prop_STM_observer :: Bool
prop_STM_observer = True

```

2.4.3 Trace

```

tests :: TestTree
tests = testGroup "testing Trace" [
  unit_tests
  , testCase "forked traces stress testing" stress_trace_in_fork
  , testCase "stress testing: ObservableTrace vs. NoTrace" timing_Observable_vs_Untimed
  , testCaseInfo "demonstrating nested named context logging" example_with_named_contexts
]
unit_tests :: TestTree
unit_tests = testGroup "Unit tests" [
  testCase "opening messages should not be traced" unit_noOpeningTrace
]

```

```

,testCase "hierarchy of traces" unit_hierarchy
,testCase "forked traces" unit_trace_in_fork
,testCase "hierarchy of traces with NoTrace" $
  unit_hierarchy' [Neutral, NoTrace, (ObservableTrace observablesSet)]
  onlyLevelOneMessage
,testCase "hierarchy of traces with DropOpening" $
  unit_hierarchy' [Neutral, DropOpening, (ObservableTrace observablesSet)]
  notObserveOpen
,testCase "hierarchy of traces with UntimedTrace" $
  unit_hierarchy' [Neutral, UntimedTrace, UntimedTrace]
  observeNoMeasures
,testCase "changing the minimum severity of a trace at runtime"
  unit_trace_min_severity
,testCase "changing the minimum severity of a named context at runtime"
  unit_named_min_severity
,testCase "appending names should not exceed 80 chars" unit_append_name
,testCase "creat subtrace which duplicates messages" unit_trace_duplicate
,testCase "testing name filtering" unit_name_filtering
,testCase "testing throwing of exceptions" unit_exception_throwing
,testCase "NoTrace: check lazy evaluation" unit_test_lazy_evaluation
]
where
  observablesSet = [MonotonicClock, MemoryStats]
  notObserveOpen :: [LogObject] → Bool
  notObserveOpen = all (λcase {LogObject _ (ObserveOpen _) → False; _ → True})
  notObserveClose :: [LogObject] → Bool
  notObserveClose = all (λcase {LogObject _ (ObserveClose _) → False; _ → True})
  notObserveDiff :: [LogObject] → Bool
  notObserveDiff = all (λcase {LogObject _ (ObserveDiff _) → False; _ → True})
  onlyLevelOneMessage :: [LogObject] → Bool
  onlyLevelOneMessage = λcase
    [LogObject _ (LogMessage (LogItem _ _ "Message from level 1."))] → True
    _ → False
  observeNoMeasures :: [LogObject] → Bool
  observeNoMeasures obs = notObserveOpen obs ∧ notObserveClose obs ∧ notObserveDiff obs

```

Helper routines

```

data TraceConfiguration = TraceConfiguration
  { tcOutputKind :: OutputKind
  , tcName       :: LoggerName
  , tcSubTrace   :: SubTrace
  , tcSeverity   :: Severity
  }

setupTrace :: TraceConfiguration → IO (Trace IO)
setupTrace (TraceConfiguration outk name subTr sev) = do

```

```

c ← liftIO $ Cardano.BM.Configuration ◦ Model.empty
mockSwitchboard ← newMVar $ error "Switchboard uninitialized."
ctx ← liftIO $ newContext name c sev $ Switchboard mockSwitchboard
let logTrace0 = case outk of
  TVarList tvar → BaseTrace.natTrace liftIO $ traceInTVarIO tvar
  TVarListNamed tvar → BaseTrace.natTrace liftIO $ traceNamedInTVarIO tvar
setSubTrace (configuration ctx) name (Just subTr)
logTrace' ← subTrace "" (ctx, logTrace0)
return logTrace'

setTransformer_ :: Trace IO → LoggerName → Maybe SubTrace → IO ()
setTransformer_ (ctx, _) name subtr = do
  let c = configuration ctx
  n = (loggerName ctx) <> " ." <> name
  setSubTrace c n subtr

```

Example of using named contexts with **Trace**

```

example_with_named_contexts :: IO String
example_with_named_contexts = do
  cfg ← defaultConfigTesting
  logTrace ← Setup.setupTrace (Right cfg) "test"
  putStrLn "\n"
  logInfo logTrace "entering"
  logTrace0 ← appendName "simple-work-0" logTrace
  work0 ← complexWork0 logTrace0 "0"
  logTrace1 ← appendName "complex-work-1" logTrace
  work1 ← complexWork1 logTrace1 "42"
  Async.wait work0
  Async.wait work1
  -- the named context will include "complex" in the logged message
  logInfo logTrace "done."
  threadDelay 1000
  Setup.shutdownTrace logTrace
  return ""
where
  complexWork0 tr msg = Async.async $ logInfo tr ("let's see (0): " 'append' msg)
  complexWork1 tr msg = Async.async $ do
    logInfo tr ("let's see (1): " 'append' msg)
    trInner@(ctx, _) ← appendName "inner-work-1" tr
    let observablesSet = [MonotonicClock]
    setSubTrace (configuration ctx) "test.complex-work-1.inner-work-1.STM-action" $
      Just $ ObservableTrace observablesSet
    _ ← STMObserver.bracketObserveIO trInner "STM-action" setVar_
    logInfo trInner "let's see: done."
    -- logInfo logTrace' "let's see: done."

```


Show effect of turning off observables

```

run_timed_action :: Trace IO → Int → IO Measurable
run_timed_action logTrace reps = do
  runid ← newUnique
  t0 ← getMonoClock
  forM_ [1 :: Int]..reps $ const $ observeAction logTrace
  t1 ← getMonoClock
  return $ diffTimeObserved (CounterState runid t0) (CounterState runid t1)
where
  observeAction trace = do
    _ ← MonadicObserver.bracketObserveIO trace "" action
    return ()
  action = return $ forM [1 :: Int..100] $ \x → [x] ++ (init $ reverse [1 :: Int..10000])

timing_Observable_vs_Untimed :: Assertion
timing_Observable_vs_Untimed = do
  msgs1 ← STM.newTVarIO []
  traceObservable ← setupTrace $ TraceConfiguration
    (TVarList msgs1)
    "observables"
    (ObservableTrace observablesSet)
    Debug
  msgs2 ← STM.newTVarIO []
  traceUntimed ← setupTrace $ TraceConfiguration
    (TVarList msgs2)
    "no timing"
    UntimedTrace
    Debug
  msgs3 ← STM.newTVarIO []
  traceNoTrace ← setupTrace $ TraceConfiguration
    (TVarList msgs3)
    "no trace"
    NoTrace
    Debug
  t_observable ← run_timed_action traceObservable 100
  t_untimed ← run_timed_action traceUntimed 100
  t_notrace ← run_timed_action traceNoTrace 100
  assertBool
    ("Untimed consumed more time than ObservableTrace " ++ (show [t_untimed, t_observable]))
    (t_untimed < t_observable)
  assertBool
    ("NoTrace consumed more time than ObservableTrace " ++ (show [t_notrace, t_observable]))
    (t_notrace < t_observable)
  assertBool
    ("NoTrace consumed more time than Untimed " ++ (show [t_notrace, t_untimed]))
    True

```

where

```
observablesSet = [MonotonicClock, GhcRtsStats, MemoryStats]
```

Control tracing in a hierarchy of **Traces**

We can lay out traces in a hierarchical manner, that the children forward traced items to the parent **Trace**. A **NoTrace** introduced in this hierarchy will cut off a branch from messaging to the root.

```
unit_hierarchy :: Assertion
unit_hierarchy = do
  msgs ← STM.newTVarIO []
  trace0 ← setupTrace $ TraceConfiguration (TVarList msgs) "test" Neutral Debug
  logInfo trace0 "This should have been displayed!"
  -- subtrace of trace which traces nothing
  setTransformer_trace0 "inner" (Just NoTrace)
  trace1 ← subTrace "inner" trace0
  logInfo trace1 "This should NOT have been displayed!"
  setTransformer_trace1 "innermost" (Just Neutral)
  trace2 ← subTrace "innermost" trace1
  logInfo trace2 "This should NOT have been displayed also due to the trace one level above"
  -- acquire the traced objects
  res ← STM.readTVarIO msgs
  -- only the first message should have been traced
  assertBool
    ("Found more or less messages than expected: " ++ show res)
    (length res == 1)
```

Change a trace's minimum severity

A trace is configured with a minimum severity and filters out messages that are labelled with a lower severity. This minimum severity of the current trace can be changed.

```
unit_trace_min_severity :: Assertion
unit_trace_min_severity = do
  msgs ← STM.newTVarIO []
  trace@(ctx, _) ← setupTrace $ TraceConfiguration (TVarList msgs) "test min severity" Neutral Debug
  logInfo trace "Message #1"
  -- raise the minimum severity to Warning
  setMinSeverity (configuration ctx) Warning
  msev ← Cardano.BM.Configuration.minSeverity (configuration ctx)
  assertBool ("min severity should be Warning, but is " ++ (show msev))
    (msev == Warning)
  -- this message will not be traced
  logInfo trace "Message #2"
  -- lower the minimum severity to Info
```

```

setMinSeverity (configuration ctx) Info
-- this message is traced
logInfo trace "Message #3"
-- acquire the traced objects
res ← STM.readTVarIO msgs
-- only the first and last messages should have been traced
assertBool
  ("Found more or less messages than expected: " ++ show res)
  (length res ≡ 2)
assertBool
  ("Found Info message when Warning was minimum severity: " ++ show res)
  (all (λcase {LogObject _ (LogMessage (LogItem _ Info "Message #2")) → False; _ → True}) res)

```

Define a subtrace's behaviour to duplicate all messages

The **SubTrace** will duplicate all messages that pass through it. Each message will be in its own named context.

```

unit_trace_duplicate :: Assertion
unit_trace_duplicate = do
  msgs ← STM.newTVarIO []
  trace0@(ctx, _) ← setupTrace $ TraceConfiguration (TVarList msgs) "test duplicate" Neutral Debug
  logInfo trace0 "Message #1"
  -- create a subtrace which duplicates all messages
  setSubTrace (configuration ctx) "test duplicate.orig" $ Just (TeeTrace "dup")
  trace ← subTrace "orig" trace0
  -- this message will be duplicated
  logInfo trace "You will see me twice!"
  -- acquire the traced objects
  res ← STM.readTVarIO msgs
  -- only the first and last messages should have been traced
  assertBool
    ("Found more or less messages than expected: " ++ show res)
    (length res ≡ 3)

```

Change the minimum severity of a named context

A trace of a named context can be configured with a minimum severity, such that the trace will filter out messages that are labelled with a lower severity.

```

unit_named_min_severity :: Assertion
unit_named_min_severity = do
  msgs ← STM.newTVarIO []
  trace0 ← setupTrace $ TraceConfiguration (TVarList msgs) "test named severity" Neutral Debug
  trace@(ctx, _) ← appendName "sev-change" trace0
  logInfo trace "Message #1"

```

```

-- raise the minimum severity to Warning
setSeverity (configuration ctx) (loggerName ctx) (Just Warning)
msev ← Cardano.BM.Configuration.inspectSeverity (configuration ctx) (loggerName ctx)
assertBool ("min severity should be Warning, but is " ++ (show msev))
  (msev ≡ Just Warning)
-- this message will not be traced
logInfo trace "Message #2"
-- lower the minimum severity to Info
setSeverity (configuration ctx) (loggerName ctx) (Just Info)
-- this message is traced
logInfo trace "Message #3"
-- acquire the traced objects
res ← STM.readTVarIO msgs
-- only the first and last messages should have been traced
assertBool
  ("Found more or less messages than expected: " ++ show res)
  (length res ≡ 2)
assertBool
  ("Found Info message when Warning was minimum severity: " ++ show res)
  (all (λcase {LogObject _ (LogMessage (LogItem _ Info "Message #2")) → False; _ → True}) res)

```

```

unit_hierarchy' :: [SubTrace] → ([LogObject] → Bool) → Assertion
unit_hierarchy' subtraces f = do
  let (t1 : t2 : t3 : _) = cycle subtraces
  msgs ← STM.newTVarIO []
  -- create trace of type 1
  trace1 ← setupTrace $ TraceConfiguration (TVarList msgs) "test" t1 Debug
  logInfo trace1 "Message from level 1."
  -- subtrace of type 2
  setTransformer_trace1 "inner" (Just t2)
  trace2 ← subTrace "inner" trace1
  logInfo trace2 "Message from level 2."
  -- subsubtrace of type 3
  setTransformer_trace2 "innermost" (Just t3)
  _ ← STMObserver.bracketObserveIO trace2 "innermost" setVar_
  logInfo trace2 "Message from level 3."
  -- acquire the traced objects
  res ← STM.readTVarIO msgs
  -- only the first message should have been traced
  assertBool
    ("Found more or less messages than expected: " ++ show res)
    (f res)

```

Logging in parallel

```

unit_trace_in_fork :: Assertion
unit_trace_in_fork = do
  msgs ← STM.newTVarIO [ ]
  trace ← setupTrace $ TraceConfiguration (TVarListNamed msgs) "test" Neutral Debug
  trace0 ← appendName "work0" trace
  trace1 ← appendName "work1" trace
  work0 ← work trace0
  threadDelay 5000
  work1 ← work trace1
  Async.wait $ work0
  Async.wait $ work1
  res ← STM.readTVarIO msgs
  let names@(_:namesTail) = map lnName res
  -- each trace should have its own name and log right after the other
  assertBool
    ("Consecutive loggernames are not different: " ++ show names)
    (and $ zipWith (≠) names namesTail)
where
  work :: Trace IO → IO (Async.Async ())
  work trace = Async.async $ do
    logInfoDelay trace "1"
    logInfoDelay trace "2"
    logInfoDelay trace "3"
  logInfoDelay :: Trace IO → Text → IO ()
  logInfoDelay trace msg =
    logInfo trace msg >>
    threadDelay 10000

```

Stress testing parallel logging

```

stress_trace_in_fork :: Assertion
stress_trace_in_fork = do
  msgs ← STM.newTVarIO [ ]
  trace ← setupTrace $ TraceConfiguration (TVarListNamed msgs) "test" Neutral Debug
  let names = map (λa → ("work-" <> pack (show a))) [1..(10::Int)]
  ts ← forM names $ λname → do
    trace' ← appendName name trace
    work trace'
  forM ts Async.wait
  res ← STM.readTVarIO msgs
  let resNames = map lnName res
  let frequencyMap = fromListWith (+) [(x,1) | x ← resNames]
  -- each trace should have traced 'totalMessages' messages
  assertBool

```

```

    ("Frequencies of logged messages according to logname: " + show frequencyMap)
    (all ( $\lambda$ name  $\rightarrow$  (lookup ("test." <> name) frequencyMap)  $\equiv$  Just totalMessages) names)
  where
    work :: Trace IO  $\rightarrow$  IO (Async.Async ())
    work trace = Async.async $ forM_ [1..totalMessages] $ (logInfo trace)  $\circ$  pack  $\circ$  show
    totalMessages :: Int
    totalMessages = 10

```

Dropping ObserveOpen messages in a subtrace

```

unit_noOpening_Trace :: Assertion
unit_noOpening_Trace = do
  msgs  $\leftarrow$  STM.newTVarIO []
  logTrace  $\leftarrow$  setupTrace $ TraceConfiguration (TVarList msgs) "test" DropOpening Debug
  _  $\leftarrow$  STMObserver.bracketObserveIO logTrace "setTVar" setVar_
  res  $\leftarrow$  STM.readTVarIO msgs
  assertBool
    ("Found non-expected ObserveOpen message: " + show res)
    (all ( $\lambda$ case {LogObject _ (ObserveOpen _)  $\rightarrow$  False; _  $\rightarrow$  True}) res)

```

Assert maximum length of log context name

The name of the log context cannot grow beyond a maximum number of characters, currently the limit is set to 80.

```

unit_append_name :: Assertion
unit_append_name = do
  cfg  $\leftarrow$  defaultConfigTesting
  Setup.withTrace cfg "test" $  $\lambda$ trace0  $\rightarrow$  do
    trace1  $\leftarrow$  appendName bigName trace0
    (ctx2, _)  $\leftarrow$  appendName bigName trace1
    assertBool
      ("Found logger name with more than 80 chars: " + show (loggerName ctx2))
      (T.length (loggerName ctx2)  $\leq$  80)
  where
    bigName = T.replicate 30 "abcdefghijklmnopqrstuvwxy"

```

```

setVar_ :: STM.STM Integer
setVar_ = do
  t  $\leftarrow$  STM.newTVar 0
  STM.writeTVar t 42
  res  $\leftarrow$  STM.readTVar t
  return res

```

Testing log context name filters

```

unit_name_filtering::Assertion
unit_name_filtering = do
  let contextName = "test.sub.1"
  let loname = "sum"-- would be part of a "LogValue loname 42"
  let filter1 = [(Drop (Exact "test.sub.1"), Unhide [ ])]
  assertBool ("Dropping a specific name should filter it out and thus return False")
    (False == evalFilters filter1 contextName)
  let filter2 = [(Drop (EndsWith ".1"), Unhide [ ])]
  assertBool ("Dropping a name ending with a specific text should filter out the context name")
    (False == evalFilters filter2 contextName)
  let filter3 = [(Drop (StartsWith "test."), Unhide [ ])]
  assertBool ("Dropping a name starting with a specific text should filter out the context name")
    (False == evalFilters filter3 contextName)
  let filter4 = [(Drop (Contains ".sub."), Unhide [ ])]
  assertBool ("Dropping a name starting containing a specific text should filter out the context name")
    (False == evalFilters filter4 contextName)
  let filter5 = [(Drop (StartsWith "test."),
    Unhide [(Exact "test.sub.1")])]
  assertBool ("Dropping all and un hiding a specific name should the context name allow passing")
    (True == evalFilters filter5 contextName)
  let filter6 = [(Drop (StartsWith "test."),
    Unhide [(EndsWith ".sum"),
    (EndsWith ".other")])]
  assertBool ("Dropping all and un hiding some names, the LogObject should pass the filter")
    (True == evalFilters filter6 (contextName <> "." <> loname))
  let filter7 = [(Drop (StartsWith "test."),
    Unhide [(EndsWith ".product")])]
  assertBool ("Dropping all and un hiding an inexistant named value, the LogObject should not pass")
    (False == evalFilters filter7 (contextName <> "." <> loname))
  let filter8 = [(Drop (StartsWith "test."),
    Unhide [(Exact "test.sub.1")]),
    (Drop (StartsWith "something.else."),
    Unhide [(EndsWith ".this")])]
  assertBool ("Disjunction of filters that should pass")
    (True == evalFilters filter8 contextName)
  let filter9 = [(Drop (StartsWith "test."),
    Unhide [(Exact ".that")]),
    (Drop (StartsWith "something.else."),
    Unhide [(EndsWith ".this")])]
  assertBool ("Disjunction of filters that should not pass")
    (False == evalFilters filter9 contextName)

```

Exception throwing

Exceptions encountered should be thrown.

```

unit_exception_throwing :: Assertion
unit_exception_throwing = do
    action ← work msg
    res ← Async.waitCatch action
    assertBool
        ("Exception should have been rethrown")
        (isLeft res)
where
    msg :: Text
    msg = error "faulty message"
    work :: Text → IO (Async.Async ())
    work message = Async.async $ do
        cfg ← defaultConfigTesting
        trace ← Setup.setupTrace (Right cfg) "test"
        logInfo trace message
        Setup.shutdownTrace trace

```

Check lazy evaluation of trace

Exception should not be thrown when type of **Trace** is **NoTrace**.

```

unit_test_lazy_evaluation :: Assertion
unit_test_lazy_evaluation = do
    action ← work msg
    res ← Async.waitCatch action
    assertBool
        ("Exception should not have been rethrown when type of Trace is NoTrace")
        (isRight res)
where
    msg :: Text
    msg = error "faulty message"
    work :: Text → IO (Async.Async ())
    work message = Async.async $ do
        cfg ← defaultConfigTesting
        trace0@(ctx, _) ← Setup.setupTrace (Right cfg) "test"
        setSubTrace (configuration ctx) "test.work" (Just NoTrace)
        trace ← subTrace "work" trace0
        logInfo trace message
        Setup.shutdownTrace trace

```

2.4.4 Testing configuration

Test declarations

```

tests :: TestTree
tests = testGroup "config tests" [

```



```

    property_tests
  , unit_tests
]

property_tests :: TestTree
property_tests = testGroup "Properties" [
  testProperty "minimal" prop_Configuration_minimal
]

unit_tests :: TestTree
unit_tests = testGroup "Unit tests" [
  testCase "static_representation" unit_Configuration_static_representation
  , testCase "parsed_representation" unit_Configuration_parsed_representation
  , testCase "parsed_configuration" unit_Configuration_parsed
  , testCase "include_EKG_if_defined" unit_Configuration_check_EKG_positive
  , testCase "not_include_EKG_if_ndef" unit_Configuration_check_EKG_negative
  , testCase "check_scribe_caching" unit_Configuration_check_scribe_cache
]

```

Property tests

```

prop_Configuration_minimal :: Bool
prop_Configuration_minimal = True

```

Unit tests

The configuration file only indicates that EKG is listening on port nnnnn. Infer that *EKGViewBK* needs to be started as a backend.

```

unit_Configuration_check_EKG_positive :: Assertion
unit_Configuration_check_EKG_positive = do
  let c = ["rotation:"
    , "  rpLogLimitBytes: 5000000"
    , "  rpKeepFilesNum: 10"
    , "  rpMaxAgeHours: 24"
    , "minSeverity: Info"
    , "defaultBackends:"
    , "  - KatipBK"
    , "setupBackends:"
    , "  - KatipBK"
    , "defaultScribes:"
    , "  - StdoutSK"
    , "  - stdout"
    , "setupScribes:"
    , "  - scName: stdout"
    , "  scRotation: null"
    , "  scKind: StdoutSK"
    , "hasEKG: 18321"

```

```

    , "options:"
    , "  test:"
    , "    value: nothing"
  ]
  fp = "/tmp/test_ekgv_config.yaml"
  writeFile fp $unlines c
  repr ← parseRepresentation fp
  assertBool "expecting EKGViewBK to be setup" $
    EKGViewBK ∈ (setupBackends repr)

```

If there is no port defined for EKG, then do not start it even if present in the config.

```

unit_Configuration_check_EKG_negative :: Assertion
unit_Configuration_check_EKG_negative = do
  let c = [ "rotation:"
    , "  rpLogLimitBytes: 5000000"
    , "  rpKeepFilesNum: 10"
    , "  rpMaxAgeHours: 24"
    , "minSeverity: Info"
    , "defaultBackends:"
    , "  - KatipBK"
    , "  - EKGViewBK"
    , "setupBackends:"
    , "  - KatipBK"
    , "  - EKGViewBK"
    , "defaultScribes:"
    , "- - StdoutSK"
    , "  - stdout"
    , "setupScribes:"
    , "- scName: stdout"
    , "  scRotation: null"
    , "  scKind: StdoutSK"
    , "###hasEKG: 18321"
    , "options:"
    , "  test:"
    , "    value: nothing"
  ]
  fp = "/tmp/test_ekgv_config.yaml"
  writeFile fp $unlines c
  repr ← parseRepresentation fp
  assertBool "EKGViewBK shall not be setup" $
    ¬ $ EKGViewBK ∈ (setupBackends repr)
  assertBool "EKGViewBK shall not receive messages" $
    ¬ $ EKGViewBK ∈ (defaultBackends repr)

unit_Configuration_static_representation :: Assertion
unit_Configuration_static_representation =

```

```

let r = Representation
  {minSeverity = Info
  ,rotation = RotationParameters 5000000 24 10
  ,setupScribes =
    [ScribeDefinition {scName = "stdout"
                      ,scKind = StdoutSK
                      ,scRotation = Nothing}
    ]
  ,defaultScribes = [(StdoutSK, "stdout")]
  ,setupBackends = [EKGVViewBK, KatipBK]
  ,defaultBackends = [KatipBK]
  ,hasGUI = Just 12789
  ,hasEKG = Just 18321
  ,options =
    HM.fromList [("test1", (HM.singleton "value" "object1"))
                ,("test2", (HM.singleton "value" "object2")))]
}

in
  encode r @? = "
rotation:\n
  rpLogLimitBytes: 5000000\n
  rpKeepFilesNum: 10\n
  rpMaxAgeHours: 24\n
defaultBackends:\n
- KatipBK\n
setupBackends:\n
- EKGVViewBK\n
- KatipBK\n
hasGUI: 12789\n
defaultScribes:\n
- - StdoutSK\n
  - stdout\n
options:\n
  test2:\n
    value: object2\n
  test1:\n
    value: object1\n
setupScribes:\n
- scName: stdout\n
  scRotation: null\n
  scKind: StdoutSK\n
hasEKG: 18321\n
minSeverity: Info\n
unit_Configuration_parsed_representation :: Assertion
unit_Configuration_parsed_representation = do
  repr ← parseRepresentation "test/config.yaml"
  encode repr @? = "

```

```

"rotation:\n"
"  rpLogLimitBytes: 5000000\n"
"  rpKeepFilesNum: 10\n"
"  rpMaxAgeHours: 24\n"
"defaultBackends:\n"
"- KatipBK\n"
"setupBackends:\n"
"- AggregationBK\n"
"- EKGViewBK\n"
"- KatipBK\n"
"hasGUI: null\n"
"defaultScribes:\n"
"- - StdoutSK\n"
"  - stdout\n"
"options:\n"
"  mapSubtrace:\n"
"    iohk.benchmarking:\n"
"      tag: ObservableTrace\n"
"      contents:\n"
"        - GhcRtsStats\n"
"        - MonotonicClock\n"
"    iohk.deadend: NoTrace\n"
"  mapSeverity:\n"
"    iohk.startup: Debug\n"
"    iohk.background.process: Error\n"
"    iohk.testing.uncritical: Warning\n"
"  mapAggregatedkinds:\n"
"    iohk.interesting.value: EwmaAK {alpha = 0.75}\n"
"    iohk.background.process: StatsAK\n"
"  cfokey:\n"
"    value: Release-1.0.0\n"
"  mapScribes:\n"
"    iohk.interesting.value:\n"
"      - StdoutSK::stdout\n"
"      - FileTextSK::testlog\n"
"    iohk.background.process: FileTextSK::testlog\n"
"  mapBackends:\n"
"    iohk.interesting.value:\n"
"      - EKGViewBK\n"
"      - AggregationBK\n"
"setupScribes:\n"
"- scName: testlog\n"
"  scRotation:\n"
"    rpLogLimitBytes: 25000000\n"
"    rpKeepFilesNum: 3\n"
"    rpMaxAgeHours: 24\n"
"  scKind: FileTextSK\n"
"- scName: stdout\n"

```

```

"  scRotation: null\n"
"  scKind: StdoutSK\n"
"hasEKG: 12789\n"
"minSeverity: Info\n"
unit_Configuration_parsed :: Assertion
unit_Configuration_parsed = do
  cfg ← setup "test/config.yaml"
  cfgInternal ← readMVar $ getCG cfg
  cfgInternal @? = ConfigurationInternal
    {cgMinSeverity      = Info
    ,cgMapSeverity      = HM.fromList [("iohk.startup", Debug)
                                       ,("iohk.background.process", Error)
                                       ,("iohk.testing.uncritical", Warning)
                                       ]
    ,cgMapSubtrace      = HM.fromList [("iohk.benchmarking",
                                       ObservableTrace [GhcRtsStats, MonotonicClock])
                                       ,("iohk.deadend", NoTrace)
                                       ]
    ,cgOptions          = HM.fromList
      [ ("mapSubtrace",
        HM.fromList [("iohk.benchmarking",
                      Object (HM.fromList [("tag", String "ObservableTrace")
                                           ,("contents", Array $ V.fromList
                                             [String "GhcRtsStats"
                                              ,String "MonotonicClock" ])))
                      ,("iohk.deadend", String "NoTrace")])
        , ("mapSeverity", HM.fromList [("iohk.startup", String "Debug")
                                       ,("iohk.background.process", String "Error")
                                       ,("iohk.testing.uncritical", String "Warning")])
        , ("mapAggregatedkinds", HM.fromList [("iohk.interesting.value",
                                                String "EwmaAK {alpha = 0.75}")
                                                ,("iohk.background.process",
                                                String "StatsAK")])
        , ("cfokey", HM.fromList [("value", String "Release-1.0.0")])
        , ("mapScribes", HM.fromList [("iohk.interesting.value",
                                        Array $ V.fromList [String "StdoutSK::stdout"
                                                            ,String "FileTextSK::testlog" ])
                                        ,("iohk.background.process", String "FileTextSK::testlog")])
        , ("mapBackends", HM.fromList [("iohk.interesting.value",
                                        Array $ V.fromList [String "EKGViewBK"
                                                            ,String "AggregationBK" ]))
        ]
    ,cgMapBackend        = HM.fromList [("iohk.interesting.value", [EKGViewBK, AggregationBK])
    ,cgDefBackendKs      = [KatipBK]
    ,cgSetupBackends     = [AggregationBK, EKGViewBK, KatipBK]
    ,cgMapScribe         = HM.fromList [("iohk.interesting.value",
                                       ["StdoutSK::stdout", "FileTextSK::testlog"])

```

```

        ,("iohk.background.process",["FileTextSK::testlog"])
      ]
    ,cgMapScribeCache = HM.fromList [("iohk.interesting.value",
      ["StdoutSK::stdout","FileTextSK::testlog"])
      ,("iohk.background.process",["FileTextSK::testlog"])
    ]
    ,cgDefScribes      = ["StdoutSK::stdout"]
    ,cgSetupScribes    = [ScribeDefinition
      {scKind = FileTextSK
      ,scName = "testlog"
      ,scRotation = Just $ RotationParameters
        {rpLogLimitBytes = 25000000
        ,rpMaxAgeHours = 24
        ,rpKeepFilesNum = 3
        }
      }
      ,ScribeDefinition
      {scKind = StdoutSK
      ,scName = "stdout"
      ,scRotation = Nothing
      }
    ]
    ,cgMapAggregatedKind = HM.fromList [("iohk.interesting.value",EwmaAK {alpha = 0.75})
      ,("iohk.background.process",StatsAK)
    ]
    ,cgDefAggregatedKind = StatsAK
    ,cgPortEKG            = 12789
    ,cgPortGUI            = 0
  }

```

Test caching and inheritance of Scribes.

```

unit_Configuration_check_scribe_cache :: Assertion
unit_Configuration_check_scribe_cache = do
  configuration ← empty
  let defScribes = ["FileTextSK::node.log"]
  setDefaultScribes configuration defScribes
  let scribes12 = ["StdoutSK::stdout","FileTextSK::out.txt"]
  setScribes configuration "name1.name2" $ Just scribes12
  scribes1234 ← getScribes configuration "name1.name2.name3.name4"
  scribes1 ← getScribes configuration "name1"
  scribes1234cached ← getCacheScribes configuration "name1.name2.name3.name4"
  scribesXcached ← getCacheScribes configuration "nameX"
  assertBool "Scribes for name1.name2.name3.name4 must be the same as name1.name2" $
    scribes1234 == scribes12
  assertBool "Scribes for name1 must be the default ones" $
    scribes1 == defScribes

```

```
assertBool "Scribes for name1.name2.name3.name4 must have been cached" $  
  scribes1234cached ≡ Just scribes1234  
assertBool "Scribes for nameX must not have been cached since getScribes was not called" $  
  scribesXcached ≡ Nothing
```

Index

Aggregated, 36
 instance of Semigroup, 36
 instance of Show, 36
AggregatedExpanded, 65
Aggregation, 65
appendName, 17

Backend, 37
BaseTrace, 16
 instance of Contravariant, 16

Counter, 39
Counters
 Dummy
 readCounters, 25
 Linux
 readCounters, 25
CounterState, 39
CounterType, 39

diffCounters, 40
diffTimeObserved, 24

evalFilters, 18
EWMA, 35
ewma, 69

getMonoClock, 24
getOptionOrDefault, 44

IsBackend, 37
IsEffectuator, 37

logAlert, 20
logAlertP, 20
logAlertS, 20
logCritical, 20
logCriticalP, 20
logCriticalS, 20
logDebug, 20
logDebugP, 20
logDebugS, 20

logEmergency, 20
logEmergencyP, 20
logEmergencyS, 20
logError, 20
logErrorP, 20
logErrorS, 20
LoggerName, 40
logInfo, 20
logInfoP, 20
logInfoS, 20
LogItem, 41
 liPayload, 41
 liSelection, 41
 liSeverity, 41
LogNamed, 40
logNotice, 20
logNoticeP, 20
logNoticeS, 20
LogObject, 41
LogSelection, 41
 Both, 41
 Private, 41
 Public, 41
 PublicUnsafe, 41
logWarning, 20
logWarningP, 20
logWarningS, 20

mainTrace, 52
Measurable, 32
 instance of Num, 32
 instance of Show, 33

nameCounter, 39
NamedLogItem, 40
natTrace, 16
newContext, 23
nominalTimeToMicroseconds, 24
noTrace, 16

ObservableInstance, 42

- GhcRtsStats, 42
- IOStats, 42
- MemoryStats, 42
- MonotonicClock, 42
- ProcessStats, 42
- OutputKind, 42
- TVarList, 42
- TVarListNamed, 42
- parseRepresentation, 38
- Port, 37
- readRTSStats, 24
- Representation, 37
- ScribeDefinition, 42
 - scKind, 42
 - scName, 42
 - scRotation, 42
- ScribeId, 42
- ScribeKind
 - FileJsonSK, 42
 - FileTextSK, 42
 - StderrSK, 42
 - StdoutSK, 42
- setupTrace, 22
- Severity, 43
 - Alert, 43
 - Critical, 43
 - Debug, 43
 - Emergency, 43
 - Error, 43
 - Info, 43
 - instance of FromJSON, 43
 - Notice, 43
 - Warning, 43
- shutdownTrace, 22
- singletonStats, 36
- Stats, 34
 - instance of Semigroup, 34
- stats2Text, 35
- stdoutTrace, 18
- SubTrace, 43
 - DropOpening, 43
 - FilterTrace, 43
 - NameOperator, 43
 - NameSelector, 43
 - Neutral, 43
 - NoTrace, 43
 - ObservableTrace, 43
 - TeeTrace, 43
 - UntimedTrace, 43
- subTrace, 21
- Switchboard, 52
 - instance of IsBackend, 53
 - instance of IsEffectuator, 52
 - setupBackends, 54
- Trace, 44
- traceConditionally, 19
- TraceContext, 44
 - configuration, 44
 - loggerName, 44
 - minSeverity, 44
 - shutdown, 44
 - tracetype, 44
- traceInTVar, 19
- traceInTVarIO, 19
- TraceNamed, 44
- traceNamedInTVarIO, 19
- traceNamedItem, 20
- traceNamedObject, 17
- traceWith, 16
- typeofTrace, 17
- updateAggregation, 68
- updateTracetype, 17
- withTrace, 23