

Cardano.BM - logging, benchmarking and monitoring

Alexander Diemand

Andreas Triantafyllos

November 2018

Abstract

This is a framework that combines logging, benchmarking and monitoring. Complex evaluations of STM or monadic actions can be observed from outside while reading operating system counters before and after, and calculating their differences, thus relating resource usage to such actions. Through interactive configuration, the runtime behaviour of logging or the measurement of resource usage can be altered. Further reduction in logging can be achieved by redirecting log messages to an aggregation function which will output the running statistics with less frequency than the original message.

Contents

1	Logging, benchmarking and monitoring	3
1.1	Overview	3
1.2	Requirements	3
1.2.1	Observables	3
1.2.2	Traces	5
1.2.3	Aggregation	6
1.2.4	Monitoring	6
1.2.5	Reporting	6
1.2.6	Visualisation	6
1.3	Description	7
1.3.1	Logging with Trace	7
1.3.2	Micro-benchmarks record observables	7
1.3.3	Configuration	8
1.3.4	Information reduction in Aggregation	10
1.3.5	Output selection	10
1.3.6	Monitoring	10
1.4	Examples	10
1.4.1	Simple example showing plain logging	10
1.4.2	Complex example showing logging, aggregation, and observing <i>IO</i> actions	10
1.5	Code listings - basic-tracer package	17
1.5.1	Cardano.BM.Tracer	17
1.5.2	Cardano.BM.Tracer.Class	18
1.5.3	Cardano.BM.Tracer.Output	18
1.5.4	Cardano.BM.Tracer.Transformers	19
1.6	Code listings - iohk-monitoring package	19
1.6.1	Cardano.BM.Observer.STM	19
1.6.2	Cardano.BM.Observer.Monad	21
1.6.3	Cardano.BM.Trace	24
1.6.4	Cardano.BM.Setup	27
1.6.5	Cardano.BM.Counters	28
1.6.6	Cardano.BM.Counters.Common	29
1.6.7	Cardano.BM.Counters.Dummy	30
1.6.8	Cardano.BM.Counters.Linux	30
1.6.9	Cardano.BM.Data.Aggregated	38
1.6.10	Cardano.BM.Data.AggregatedKind	43
1.6.11	Cardano.BM.Data.Backend	43
1.6.12	Cardano.BM.Data.BackendKind	44
1.6.13	Cardano.BM.Data.Configuration	44
1.6.14	Cardano.BM.Data.Counter	45
1.6.15	Cardano.BM.Data.LogItem	47
1.6.16	Cardano.BM.Data.Observable	48

1.6.17	Cardano.BM.Data.Output	48
1.6.18	Cardano.BM.Data.Rotation	50
1.6.19	Cardano.BM.Data.Severity	50
1.6.20	Cardano.BM.Data.SubTrace	50
1.6.21	Cardano.BM.Data.Trace	51
1.6.22	Cardano.BM.Data.Tracer	51
1.6.23	Cardano.BM.Configuration	54
1.6.24	Cardano.BM.Configuration.Model	55
1.6.25	Cardano.BM.Configuration.Static	63
1.6.26	Cardano.BM.Configuration.Editor	64
1.6.27	Cardano.BM.Output.Switchboard	68
1.6.28	Cardano.BM.Output.Log	73
1.6.29	Cardano.BM.Output.EKGView	80
1.6.30	Cardano.BM.Output.Aggregation	84
1.6.31	Cardano.BM.Output.Monitoring	89
2	Testing	93
2.1	Test coverage	93
2.2	Test main entry point	93
2.3	Test case generation	95
2.3.1	instance Arbitrary Aggregated	95
2.4	Tests	96
2.4.1	Testing aggregation	96
2.4.2	Cardano.BM.Test.STM	100
2.4.3	Cardano.BM.Test.Trace	100
2.4.4	Testing configuration	111
2.4.5	Rotator	119

Chapter 1

Logging, benchmarking and monitoring

1.1 Overview

In figure 1.1 we display the relationships among modules in *Cardano.BM*. Central is the **Switchboard** (see `Cardano.BM.Output.Switchboard`) that will redirect incoming log messages to selected backends according the **Configuration** (see `Cardano.BM.Configuration.Model`). The log items are created in the application's context and passed via a hierarchy of **Traces** (see `Cardano.BM.Trace`). Such a hierarchy can be built with the function `setSubTrace`. The newly added child **Trace** will add its name to the logging context and behave as configured. Among the different kinds of **Traces** implemented are **NoTrace** which suppresses all log items, **FilterTrace** which filters the log items passing through it, and **ObservableTrace** which allows capturing of operating system counters (see `Cardano.BM.Data.SubTrace`). The backend **EKGView** (see `Cardano.BM.Output.EKGView`) displays selected values in a browser. The **Log** backend is based on *katip* and outputs log items in files or the console. The format can be chosen to be textual or JSON representation. And finally, the **Aggregation** backend computes simple statistics over incoming log items (e.g. last, min, max, mean, etc.) (see `Cardano.BM.Data.Aggregated`).

Output selection determines which log items of a named context are routed to which backend. In the case of the **Log** output, this includes a configured output sink (e.g. which file). Items that are aggregated lead to the creation of an output of their current statistics. To prevent a potential infinite loop these aggregation statistics cannot be routed again back into the **Aggregation**.

With *Monitoring* we aim to shortcut the logging-analysis cycle and immediately evaluate monitors on logged values when they become available. In case a monitor is triggered a number of actions can be run: either internal actions that can alter the **Configuration**, or actions that can lead to alerting in external systems.

It is not the intention that this framework should (as part of normal use) record sufficient information so as to make the sequence of events reproducible, i.e. it is not an audit or transaction log.

1.2 Requirements

1.2.1 Observables

We can observe the passage of the flow of execution through particular points in the code (really the points at which the graph is reduced). Typically observables would be part of an outcome (which has a start and an end). Where the environment permits these outcomes could also gather additional environmental context (e.g read system counters, know the time). The

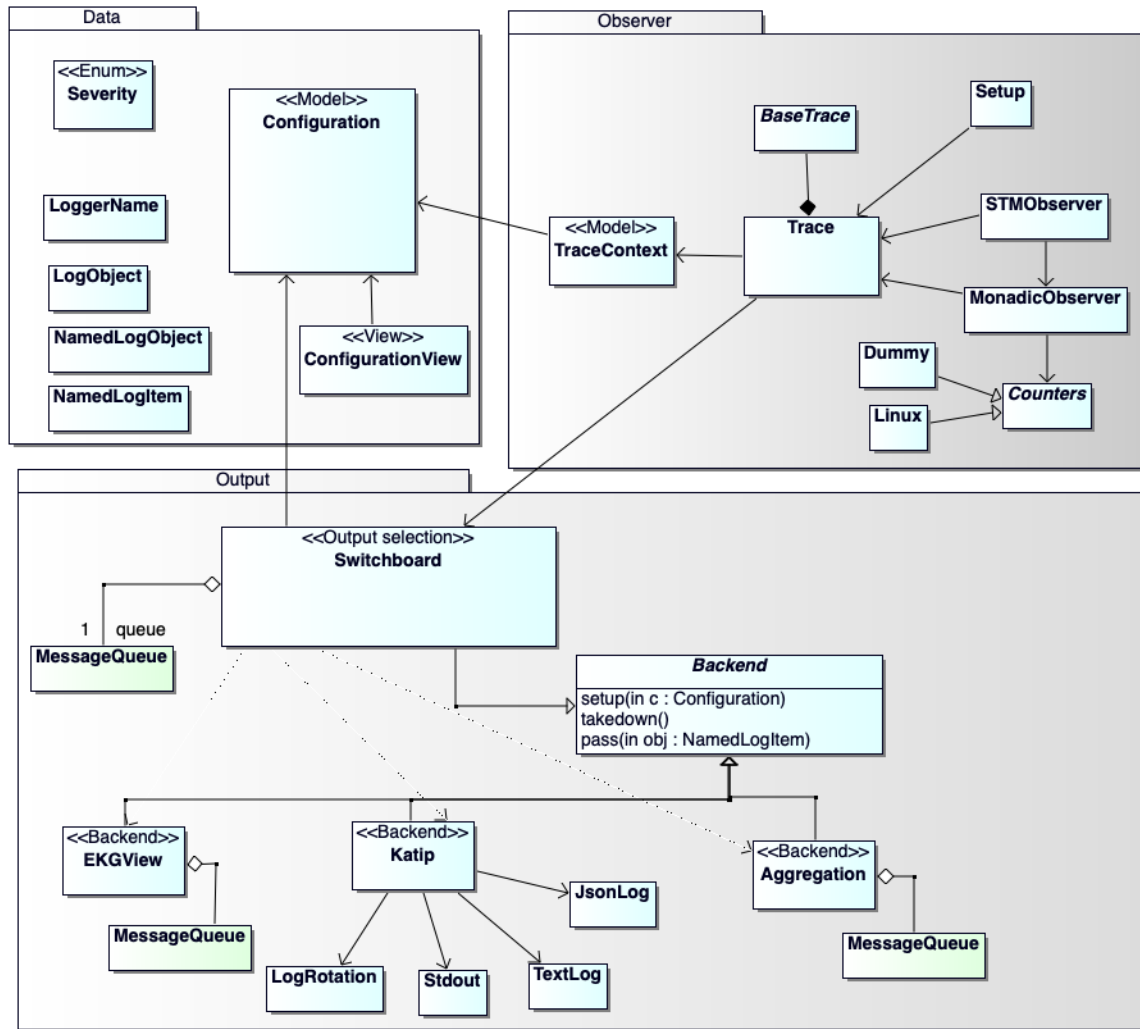


Figure 1.1: Overview of module relationships. The arrows indicate import of a module. The arrows with a triangle at one end would signify "inheritance" in object-oriented programming, but we use it to show that one module replaces the other in the namespace, thus specializes its interface.

proposed framework would be able to aggregate, filter such outcome measures so as to calculation things (where appropriate) such as:

- min/max/mean/variance of the resource costs of achieving an outcome
- elapsed wall-clock time
- CPU cycles
- memory allocations, etc
- exponentially weighted moving average of outcomes, events
- min/max/mean/variance of inter-arrival times of demand for service (the arrival pattern)

- measuring offered load against the system (e.g rate/distribution of requests against the wallet by an exchange, transactions being forwarded between nodes)

STM evaluation

We treat STM evaluation as a black box and register measurables (counters) before entering, and report the difference at exit together with the result. Logging in an STM will keep a list of log items which at the exit of the evaluation will be passed to the logging subsystem. Since we do not know the exact time an event occurred in the STM action, we annotate the event afterwards with the time interval of the STM action.

Function evaluation

We treat a function call as a black box and register measurables (counters) before entering, and report the difference at exit together with the result. The function is expected to accept a 'Trace' argument which receives the events.

QuickCheck properties *tentatively*

The function

```
quickCheckResult :: Testable prop => prop -> IO Result
```

will return a *Result* data structure from which we can extract the number of tests performed. Recording the start and end times allows us to derive the time spent for a single test. (although this measurement is wrong as it includes the time spent in QuickCheck setting up the test case (and shrinking?))

1.2.2 Traces

Log items are sent as streams of events to the logging system for processing (aggregation, ..) before output. Functions that need to log events must accept a *Trace* argument. There is no monad related to logging in the monad stack, thus this can work in any monadic environment.

Trace Context

A Trace maintains a named context stack. A new name can be put onto it, and all subsequent log messages are labeled with this named context. This is also true to all downstream functions which receive the modified Trace. We thus can see the call tree and how the evaluation entered the context where a logging function was called. The context also maintains a mapping from name to Severity: this way a logging function call can early end and not produce a log item when the minimum severity is not reached.

SubTrace

A Trace is created in *IO* within `setupTrace` with the intent to pass the traced items to a downstream logging framework for outputting to various destinations in different formats. Apart from adding a name to the naming stack we can also alter the behaviour of the Trace. The newly created Trace with a specific function to process the recorded items will forward these to the upstream Trace. This way we can, for example, locally turn on aggregation of observables and only report a summary to the logs.

1.2.3 Aggregation

Log items contain a named context, severity and a payload (message, structured value). Thinking of a relation

`(name, severity) -> value`

, folding a summarizing function over it outputs

`(name, severity) -> Summary`

. Depending on the type of *value*, the summary could provide for example:

- `*` : first, last, count, the time between events (mean, sigma)
- `Num` : min, max, median, quartiles, mean, sigma, the delta between events (mean, sigma)

Other possible aggregations:

- exponentially weighted moving average
- histograms

1.2.4 Monitoring

- Enable (or disable) measuring events and performance at runtime (e.g. measure how block holding time has changed).
- Send alarms when observables give evidence for abnormalities
- Observe actions in progress, i.e. have started and not yet finished
- Bridge to *Datadog*?

1.2.5 Reporting

We might want to buffer events in case an exception is detected. This FIFO queue could then be output to the log for post-factum inspection.

1.2.6 Visualisation

EKG

<https://hackage.haskell.org/package/ekg>

This library allows live monitor a running instance over HTTP. There is a way we can add our own metrics to it and update them.

Log files

The output of observables immediately or aggregated to log files. The format is chosen to be JSON for easier post-processing.

Web app

Could combine EKG, log files and parameterization into one GUI.
(e.g. <https://github.com/HeinrichApfelmus/threepenny-gui>)

1.3 Description

1.3.1 Logging with **Trace**

Setup procedure




Figure 1.2: Setup procedure

Hierarchy of **Traces**

1.3.2 Micro-benchmarks record observables

Micro-benchmarks are recording observables that measure resource usage of the whole program for a specific time. These measurements are then associated with the subsystem that was observed at that time. Caveat: if the executable under observation runs on a multiprocessor computer where more than one parallel thread executes at the same time, it becomes difficult to associate resource usage to a single function. Even more so, as Haskell's thread do not map directly to operating system threads. So the expressiveness of our approach is only valid statistically when a large number of observables have been captured.

Counters

The framework provides access to the following O/S counters (defined in **ObservableInstance**) on *Linux*:

- monotonic clock (see **MonotonicClock**)
- CPU or total time (`/proc/<pid>/stat`) (see **ProcessStats**)
- memory allocation (`/proc/<pid>/statm`) (see **MemoryStats**)
- network bytes received/sent (`/proc/<pid>/net/netstat`) (see **NetStats**)
- disk input/output (`/proc/<pid>/io`) (see **IOStats**)

On all platforms, access is provided to the RTS counters (see **GhcRtsStats**).

Implementing micro-benchmarks

In a micro-benchmark we capture operating system counters over an STM evaluation or a function, before and afterwards. Then, we compute the difference between the two and report all three measurements via a *Trace* to the logging system. Here we refer to the example that can be found in [complex example](#).

The capturing of STM actions is defined in [Cardano.BM.Observer.STM](#)

and the function `STM.bracketObserveIO` has type:

```
bracketObserveIO :: Trace IO -> Text -> STM.STM t -> IO t
```

It accepts a *Trace* to which it logs, adds

a name to the context name and enters this with a *SubTrace*, and finally the STM action which will be evaluated. Because this evaluation can be retried, we cannot pass to it a *Trace* to which it could log directly. A variant of this function [bracketObserveLogIO](#) also captures log items in its result, which then are threaded through the *Trace*.

Capturing observables for a function evaluation in *IO*, the type of `bracketObserveIO` (defined in [Cardano.BM.Observer.MonadIO](#)) is:

```
bracketObserveIO :: Trace IO -> Text -> IO t -> IO t
```

It accepts a *Trace* to which it logs items, adds a name to the context name and enters this with a *SubTrace*, and then the IO action which will be evaluated.

Counters are eval-

uated before the evaluation and afterwards. We trace these as log items

[ObserveOpen](#) and

[ObserveClose](#), as well

as the difference

with type [ObserveDiff](#).

```
bracketObserveIO trace "observeDownload" $ do
  license <- openURI "http://www.gnu.org/licenses/gpl.txt"
  case license of
    Right bs -> logInfo trace $ pack $ BS8.unpack bs
    Left e    -> logError trace $ "failed to download; error: " ++ (show e)
  threadDelay 50000 -- .05 second
  pure ()
```

Configuration of mu-benchmarks

Observed STM actions or functions enter a new named context with a *SubTrace*. Thus, they need a configuration of the behaviour of this *SubTrace* in the new context. We can define this in the configuration for our example:

```
CM.setSubTrace c "complex.observeDownload" (Just $ ObservableTrace [NetStats,IOStats])
```

This enables the capturing of network and I/O stats from the operating system. Other Observables are implemented in [Cardano.BM.Data.Observable](#).

Captured observables need to be routed to backends. In our example we configure:

```
CM.setBackends c "complex.observeIO" (Just [AggregationBK])
```

to direct observables from named context *complex.observeIO* to the Aggregation backend.

1.3.3 Configuration

Format

The configuration is parsed from a file in *Yaml* format (see <https://en.wikipedia.org/wiki/YAML>) on startup. In a first parsing step the file is loaded into an internal *Representation*. This structure is then further processed and validated before copied into the runtime [Configuration](#).

Configuration editor

The configuration editor (figure 1.3) provides a minimalistic GUI accessible through a browser that directly modifies the runtime configuration of the logging system. Most importantly, the global minimum severity filter can be set. This will suppress all log messages that have a severity assigned that is lower than this setting. Moreover, the following behaviours of the logging system can be changed through the GUI:

- *Backends*: relates the named logging context to a **BackendKind**
- *Scribes*: if the backend is *KatipBK*, defines to which outputs the messages are directed (see **ScribeId**)
- *Severities* a local minimum severity filter for just the named context (see **Severity**)
- *SubTrace* entering a new named context can create a new **Trace** with a specific behaviour (see **SubTrace**)
- *Aggregation* if the backend is *AggregationBK*, defines which aggregation method to use (see **AggregatedKind**)

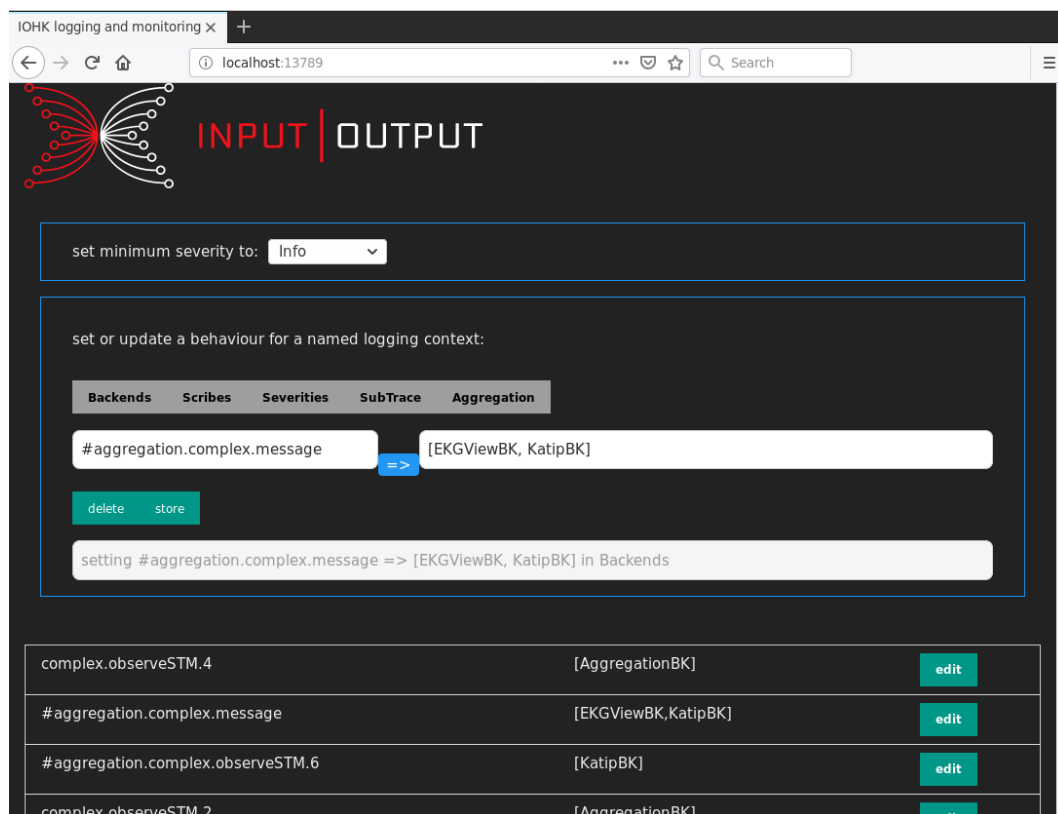


Figure 1.3: The configuration editor is listening on *localhost* and can be accessed through a browser. At the top is the setting for the global minimum severity filter, that drops all messages that have a severity lower than this setting. Below are the settings for various behaviours of the logging system.

1.3.4 Information reduction in **Aggregation**

Statistics

Configuration

1.3.5 Output selection

Configuration

1.3.6 Monitoring

Configuration

Evaluation of monitors

Actions fired

1.4 Examples

1.4.1 Simple example showing plain logging

```
{-# LANGUAGE ScopedTypeVariables #-}
module Main
  (main)
  where
import Control.Concurrent (threadDelay)
import Cardano.BM.Configuration.Static (defaultConfigStdout)
import Cardano.BM.Setup (setupTrace)
import Cardano.BM.Trace (Trace, logDebug, logError, logInfo,
                          logNotice, logWarning)

main :: IO ()
main = do
  c ← defaultConfigStdout
  tr :: Trace IO String ← setupTrace (Right c) "simple"
  logDebug tr "this is a debug message"
  logInfo tr "this is an information."
  logNotice tr "this is a notice!"
  logWarning tr "this is a warning!"
  logError tr "this is an error!"
  threadDelay 80000
  return ()
```

1.4.2 Complex example showing logging, aggregation, and observing IO actions

Module header and import directives

```
{-# LANGUAGE CPP #-}
{-# LANGUAGE ScopedTypeVariables #-}
# if defined (linux_HOST_OS)
# define LINUX
# endif
```

```

{-define the parallel procedures that create messages -}
# define RUN_ProcMessageOutput
# define RUN_ProcObserveIO
# define RUN_ProcObserveSTM
# define RUN_ProcObserveDownload
# define RUN_ProcRandom

module Main
  (main)
  where

import Control.Concurrent (threadDelay)
import qualified Control.Concurrent.Async as Async
import Control.Monad (forM_)
# ifdef ENABLE_OBSERVABLES
import Control.Monad (forM)
import GHC.Conc.Sync (atomically, STM, TVar, newTVar, readTVar, writeTVar)
# ifdef LINUX
import qualified Data.ByteString.Char8 as BS8
import Network.Download (openURI)
# endif
# endif
import Data.Text (Text, pack)
import System.Random
# ifdef ENABLE_GUI
import qualified Cardano.BM.Configuration.Editor as CME
# endif
import qualified Cardano.BM.Configuration.Model as CM
import Cardano.BM.Data.Aggregated (Measurable (...))
import Cardano.BM.Data.AggregatedKind
import Cardano.BM.Data.BackendKind
import Cardano.BM.Configuration (Configuration)
import Cardano.BM.Data.LogItem
import Cardano.BM.Data.Output
import Cardano.BM.Data.Rotation
import Cardano.BM.Data.Severity
import Cardano.BM.Data.SubTrace
# ifdef ENABLE_OBSERVABLES
import Cardano.BM.Data.Observable
import Cardano.BM.Observer.Monad (bracketObserveIO)
import qualified Cardano.BM.Observer.STM as STM
# endif
import Cardano.BM.Setup
import Cardano.BM.Trace

```

Define configuration

Selected values can be viewed in EKG on <http://localhost:12789>.
 The configuration editor listens on <http://localhost:13789>.

```

prepare_configuration :: IO CM.Configuration
prepare_configuration = do
  c ← CM.empty

```

```

CM.setMinSeverity c Debug
CM.setSetupBackends c [KatipBK
# ifdef ENABLE_AGGREGATION
    ,AggregationBK
# endif
# ifdef ENABLE_EKG
    ,EKGVViewBK
# endif
]
CM.setDefaultBackends c [KatipBK]
CM.setSetupScribes c [ ScribeDefinition {
    scName = "stdout"
    ,scKind = StdoutSK
    ,scPrivacy = ScPublic
    ,scRotation = Nothing
}
, ScribeDefinition {
    scName = "logs/out.odd.json"
    ,scKind = FileJsonSK
    ,scPrivacy = ScPublic
    ,scRotation = Nothing
}
, ScribeDefinition {
    scName = "logs/out.even.json"
    ,scKind = FileJsonSK
    ,scPrivacy = ScPublic
    ,scRotation = Nothing
}
, ScribeDefinition {
    scName = "logs/downloading.json"
    ,scKind = FileJsonSK
    ,scPrivacy = ScPublic
    ,scRotation = Nothing
}
, ScribeDefinition {
    scName = "logs/out.txt"
    ,scKind = FileTextSK
    ,scPrivacy = ScPublic
    ,scRotation = Just $ RotationParameters
        {rpLogLimitBytes = 5000-- 5kB
        ,rpMaxAgeHours = 24
        ,rpKeepFilesNum = 3
        }
}
]
CM.setDefaultScribes c ["StdoutSK::stdout"]
CM.setScribes c "complex.random" (Just ["StdoutSK::stdout", "FileTextSK::logs/out.txt"])
CM.setScribes c "#aggregated.complex.random" (Just ["StdoutSK::stdout"])
forM_ [(1::Int)..10] $ \x →
    if odd x
    then

```

```

    CM.setScribes c ("#aggregation.complex.observeSTM." <> (pack $ show x)) $ Just [ "FileJsonSK:"
else
    CM.setScribes c ("#aggregation.complex.observeSTM." <> (pack $ show x)) $ Just [ "FileJsonSK:"
# ifdef LINUX
# ifdef ENABLE_OBSERVABLES
    CM.setSubTrace c "complex.observeDownload" (Just $ ObservableTrace [IOStats,NetStats])
# endif
    CM.setBackends c "complex.observeDownload" (Just [KatipBK])
    CM.setScribes c "complex.observeDownload" (Just [ "StdoutSK::stdout", "FileJsonSK::logs/download" ])
# endif
    CM.setSubTrace c "complex.random" (Just $ TeeTrace "ewma")
    CM.setSubTrace c "#ekgview"
        (Just $ FilterTrace [(Drop (StartsWith "#ekgview.#aggregation.complex.random"),
            Unhide [(EndsWith ".count"),
                (EndsWith ".avg"),
                (EndsWith ".mean")]),
            (Drop (StartsWith "#ekgview.#aggregation.complex.observeIO"),
                Unhide [(Contains "diff.RTS.cpuNs.timed.")]),
            (Drop (StartsWith "#ekgview.#aggregation.complex.observeSTM"),
                Unhide [(Contains "diff.RTS.gcNum.timed.")]),
            (Drop (StartsWith "#ekgview.#aggregation.complex.message"),
                Unhide [(Contains ".timed.m")])
        ])
# ifdef ENABLE_OBSERVABLES
    CM.setSubTrace c "complex.observeIO" (Just $ ObservableTrace [GhcRtsStats,MemoryStats])
    forM_ [(1 :: Int)..10] $ \x →
        CM.setSubTrace
            c
            ("complex.observeSTM." <> (pack $ show x))
            (Just $ ObservableTrace [GhcRtsStats,MemoryStats])
# endif
# ifdef ENABLE_AGGREGATION
    CM.setBackends c "complex.message" (Just [AggregationBK,KatipBK])
    CM.setBackends c "complex.random" (Just [AggregationBK,KatipBK])
    CM.setBackends c "complex.random.ewma" (Just [AggregationBK])
    CM.setBackends c "complex.observeIO" (Just [AggregationBK])
# endif
    forM_ [(1 :: Int)..10] $ \x → do
# ifdef ENABLE_AGGREGATION
        CM.setBackends c
            ("complex.observeSTM." <> (pack $ show x))
            (Just [AggregationBK])
# endif
        CM.setBackends c
            ("#aggregation.complex.observeSTM." <> (pack $ show x))
            (Just [KatipBK])
        CM.setAggregatedKind c "complex.random.rr" (Just StatsAK)
        CM.setAggregatedKind c "complex.random.ewma.rr" (Just (EwmaAK 0.42))
# ifdef ENABLE_EKG
    CM.setBackends c "#aggregation.complex.message" (Just [EKGVViewBK])

```

```

CM.setBackends c "#aggregation.complex.observeIO" (Just [EKGViewBK])
CM.setBackends c "#aggregation.complex.random" (Just [EKGViewBK])
CM.setBackends c "#aggregation.complex.random.ewma" (Just [EKGViewBK])
CM.setEKGport c 12789
# endif
CM.setGUIport c 13789
return c

```

Thread that outputs a random number to a **Trace**

```

randomThr :: Trace IO Text → IO (Async.Async ())
randomThr trace = do
  logInfo trace "starting random generator"
  trace' ← appendName "random" trace
  proc ← Async.async (loop trace')
  return proc
where
  loop tr = do
    threadDelay 500000 -- 0.5 second
    num ← randomRIO (42 - 42, 42 + 42) :: IO Double
    lo ← (,) < $ > (mkLOMeta Debug Public) < * > pure (LogValue "rr" (PureD num))
    traceNamedObject tr lo
    loop tr

```

Thread that observes an IO action

```

# ifdef ENABLE_OBSERVABLES
observeIO :: Configuration → Trace IO Text → IO (Async.Async ())
observeIO config trace = do
  logInfo trace "starting observer"
  proc ← Async.async (loop trace)
  return proc
where
  loop tr = do
    threadDelay 5000000 -- 5 seconds
    _ ← bracketObserveIO config tr Debug "observeIO" $ do
      num ← randomRIO (100000, 200000) :: IO Int
      ls ← return $ reverse $ init $ reverse $ 42 : [1..num]
      pure $ const ls ()
    loop tr
# endif

```

Threads that observe STM actions on the same TVar

```

# ifdef ENABLE_OBSERVABLES
observeSTM :: Configuration → Trace IO Text → IO [Async.Async ()]
observeSTM config trace = do
  logInfo trace "starting STM observer"

```



```

tvar ← atomically $ newTVar ([1..1000] :: [Int])
-- spawn 10 threads
proc ← forM [(1 :: Int)..10] $ \x → Async.async (loop trace tvar (pack $ show x))
return proc
where
  loop tr tvarlist name = do
    threadDelay 10000000 -- 10 seconds
    STM.bracketObserveIO config tr Debug ("observeSTM." <> name) (stmAction tvarlist)
    loop tr tvarlist name
stmAction :: TVar [Int] → STM ()
stmAction tvarlist = do
  list ← readTVar tvarlist
  writeTVar tvarlist $ reverse $ init $ reverse $ list
  pure ()
# endif

```

Thread that observes an IO action which downloads a text in order to observe the I/O statistics

```

# ifdef LINUX
# ifdef ENABLE_OBSERVABLES
observeDownload :: Configuration → Trace IO Text → IO (Async.Async ())
observeDownload config trace = do
  proc ← Async.async (loop trace)
  return proc
where
  loop tr = do
    threadDelay 1000000 -- 1 second
    tr' ← appendName "observeDownload" tr
    bracketObserveIO config tr' Debug "" $ do
      license ← openURI "http://www.gnu.org/licenses/gpl.txt"
      case license of
        Right bs → logNotice tr' $ pack $ BS8.unpack bs
        Left _ → return ()
    threadDelay 50000 -- .05 second
    pure ()
  loop tr
# endif
# endif

```

Thread that periodically outputs a message

```

msgThr :: Trace IO Text → IO (Async.Async ())
msgThr trace = do
  logInfo trace "start messaging .."
  trace' ← appendName "message" trace
  Async.async (loop trace')
where
  loop tr = do

```

```

threadDelay 3000000 -- 3 seconds
logNotice tr "N O T I F I C A T I O N ! ! !"
logDebug tr "a detailed debug message."
logError tr "Boooooomm .."
loop tr

```

Main entry point

```

main :: IO ()
main = do
  -- create configuration
  c ← prepare_configuration
# ifdef ENABLE_GUI
  -- start configuration editor
  CME.startup c
# endif

  -- create initial top-level Trace
  tr :: Trace IO Text ← setupTrace (Right c) "complex"
  logNotice tr "starting program; hit CTRL-C to terminate"
-- user can watch the progress only if EKG is enabled.
# ifdef ENABLE_EKG
  logInfo tr "watch its progress on http://localhost:12789"
# endif

# ifdef RUN_ProcRandom
  {-start thread sending unbounded sequence of random numbers to a trace which aggregates them in
  procRandom ← randomThr tr
# endif
# ifdef RUN_ProcObserveIO
  -- start thread endlessly reversing lists of random length
# ifdef ENABLE_OBSERVABLES
  procObsvIO ← observeIO c tr
# endif
# endif
# ifdef RUN_ProcObserveSTM
  -- start threads endlessly observing STM actions operating on the same TVar
# ifdef ENABLE_OBSERVABLES
  procObsvSTMs ← observeSTM c tr
# endif
# endif
# ifdef LINUX
# ifdef RUN_ProcObserveDownload
  -- start thread endlessly which downloads sth in order to check the I/O usage
# ifdef ENABLE_OBSERVABLES
  procObsvDownload ← observeDownload c tr
# endif
# endif
# endif
# ifdef RUN_ProcMessageOutput
  -- start a thread to output a text messages every n seconds

```

```

    procMsg ← msgThr tr
    -- wait for message thread to finish, ignoring any exception
    _ ← Async.waitCatch procMsg
# endif
# ifdef LINUX
# ifdef RUN_ProcObserveDownload
    -- wait for download thread to finish, ignoring any exception
# ifdef ENABLE_OBSERVABLES
    _ ← Async.waitCatch procObsvDownload
# endif
# endif
# endif
# ifdef RUN_ProcObserverSTM
    -- wait for observer thread to finish, ignoring any exception
# ifdef ENABLE_OBSERVABLES
    _ ← forM procObsvSTMs Async.waitCatch
# endif
# endif
# ifdef RUN_ProcObserveIO
    -- wait for observer thread to finish, ignoring any exception
# ifdef ENABLE_OBSERVABLES
    _ ← Async.waitCatch procObsvIO
# endif
# endif
# ifdef RUN_ProcRandom
    -- wait for random thread to finish, ignoring any exception
    _ ← Async.waitCatch procRandom
# endif
    return ()

```

1.5 Code listings - basic-tracer package

1.5.1 Cardano.BM.Tracer

Tracing using the contravariant **Tracer** naturally reads:

```

let logTrace = tracing $ showTracing $ stdoutTracer
in logTrace "hello world"

example1 :: IO ()
example1 = do
    let logTrace a = tracingWith (showTracing (contramap ("Debug: " ++) stdoutTracer)) a
    void $ callFun1 logTrace
callFun1 :: (String → IO ()) → IO Int
callFun1 logTrace = do
    logTrace "in function 1"
    return 42

```

1.5.2 Cardano.BM.Tracer.Class

Contravariant Tracer

The notion of a **Tracer** is an action that can be used to observe information of interest during evaluation. **Tracers** can capture (and annotate) such observations with additional information from their execution context.

```
newtype Tracer m s = Tracer {tracing :: Op (m ()) s}
```

A **Tracer** is an instance of *Contravariant*, which permits new **Tracers** to be constructed that feed into the existing Tracer by use of *contramap*.

```
instance Contravariant (Tracer m) where  
  contramap f = Tracer ∘ contramap f ∘ tracing
```

Although a **Tracer** is invoked in a monadic context (which may be *Identity*), the construction of a new **Tracer** is a pure function.

This brings with it the constraint that the derived **Tracers** form a hierarchy which has its root at the top level tracer.

nullTracer

The simplest tracer - one that suppresses all output.

```
nullTracer :: (Applicative m) ⇒ Tracer m a  
nullTracer = Tracer $ Op $ \_ → pure ()
```

tracingWith

Accepts a **Tracer** and some payload *s*. First it gets the contravariant from the **Tracer** as type "*Op (m ()) s*" and, after "*getOp :: b → a*" which translates to "*s → m ()*".

```
tracingWith :: Tracer m a → a → m ()  
tracingWith = getOp ∘ tracing
```

1.5.3 Cardano.BM.Tracer.Output

Directing a **Tracer**'s output to stdout

The Tracer that prints a string (as a line) to stdout (usual caveats about interleaving should be heeded).

```
stdoutTracer :: (MonadIO m) ⇒ Tracer m String  
stdoutTracer = Tracer $ Op $ liftIO ∘ putStrLn
```

Outputting a **Tracer** with *Debug.Trace*

A Tracer that uses *TraceM* (from **Debug.Trace**) as its output mechanism.

```
debugTracer :: (Applicative m) ⇒ Tracer m String  
debugTracer = Tracer $ Op traceM
```

1.5.4 Cardano.BM.Tracer.Transformers

Applying *show* on a **Tracer**'s messages

The Tracer transformer exploiting Show.

```
showTracing :: (Show a) => Tracer m String -> Tracer m a
showTracing = contramap show
```

Conditional tracing - statically defined

The Tracer transformer that allows for on/off control of tracing at trace creation time.

```
condTracing :: (Monad m) => (a -> Bool) -> Tracer m a -> Tracer m a
condTracing active tr = Tracer $ Op $ \s -> do
  when (active s) (tracingWith tr s)
```

Conditional tracing - dynamically evaluated

The tracer transformer that can exercise dynamic control over tracing, the dynamic decision being made using the context accessible in the monadic context.

```
condTracingM :: (Monad m) => m (a -> Bool) -> Tracer m a -> Tracer m a
condTracingM activeP tr = Tracer $ Op $ \s -> do
  active <- activeP
  when (active s) (tracingWith tr s)
```

natTrace

Natural transformation from monad *m* to monad *n*.

```
natTrace :: (forall x o m x -> n x) -> Tracer m s -> Tracer n s
natTrace nat (Tracer (Op tr)) = Tracer $ Op $ nat o tr
```

1.6 Code listings - iohk-monitoring package

1.6.1 Cardano.BM.Observer.STM

```
stmWithLog :: STM.STM (t, [(LOMeta, LOContent a)]) -> STM.STM (t, [(LOMeta, LOContent a)])
stmWithLog action = action
```

Observe STM action in a named context

With given name, create a **SubTrace** according to **Configuration** and run the passed STM action on it.

```
bracketObserveIO :: Config.Configuration -> Trace IO a -> Severity -> Text -> STM.STM t -> IO t
bracketObserveIO config trace severity name action = do
  subTrace <- fromMaybe Neutral <$> Config.findSubTrace config name
  bracketObserveIO' subTrace severity trace action
```

where

```

bracketObserveIO' :: SubTrace → Severity → Trace IO a → STM.STM t → IO t
bracketObserveIO' NoTrace _ _ act =
  STM.atomically act
bracketObserveIO' subtrace sev logTrace act = do
  mCountersid ← observeOpen subtrace sev logTrace
  -- run action; if an exception is caught, then it will be logged and rethrown.
  t ← (STM.atomically act) 'catch' (λ(e :: SomeException) → (TIO.hPutStrLn stderr (pack (show e)) >> throwM e))
  case mCountersid of
    Left openException →
      -- since observeOpen faced an exception there is no reason to call observeClose
      -- however the result of the action is returned
      TIO.hPutStrLn stderr ("ObserveOpen: " <> pack (show openException))
    Right countersid → do
      res ← observeClose subtrace sev logTrace countersid [ ]
      case res of
        Left ex → TIO.hPutStrLn stderr ("ObserveClose: " <> pack (show ex))
        _ → pure ()
  pure t

```

Observe STM action in a named context and output captured log items

The STM action might output messages, which after "success" will be forwarded to the logging trace. Otherwise, this function behaves the same as **Observe STM action in a named context**.

```

bracketObserveLogIO :: Config.Configuration → Trace IO a → Severity → Text → STM.STM (t, [(LOMeta, LOContent a)])
bracketObserveLogIO config trace severity name action = do
  subTrace ← fromMaybe Neutral < $ > Config.findSubTrace config name
  bracketObserveLogIO' subTrace severity trace action
where
  bracketObserveLogIO' :: SubTrace → Severity → Trace IO a → STM.STM (t, [(LOMeta, LOContent a)])
  bracketObserveLogIO' NoTrace _ _ act = do
    (t, _) ← STM.atomically $ stmWithLog act
    pure t
  bracketObserveLogIO' subtrace sev logTrace act = do
    mCountersid ← observeOpen subtrace sev logTrace
    -- run action, return result and log items; if an exception is
    -- caught, then it will be logged and rethrown.
    (t, as) ← (STM.atomically $ stmWithLog act) 'catch'
      (λ(e :: SomeException) → (TIO.hPutStrLn stderr (pack (show e)) >> throwM e))
    case mCountersid of
      Left openException →
        -- since observeOpen faced an exception there is no reason to call observeClose
        -- however the result of the action is returned
        TIO.hPutStrLn stderr ("ObserveOpen: " <> pack (show openException))
      Right countersid → do
        res ← observeClose subtrace sev logTrace countersid as
        case res of
          Left ex → TIO.hPutStrLn stderr ("ObserveClose: " <> pack (show ex))
          _ → pure ()
    pure t

```

1.6.2 Cardano.BM.Observer.Monad

Monad.bracketObserverIO

Observes an *IO* action and adds a name to the logger name of the passed in **Trace**. An empty *Text* leaves the logger name untouched.

Microbenchmarking steps:

1. Create a *trace* which will have been configured to observe things besides logging.

```
import qualified Cardano.BM.Configuration.Model as CM
ooo
c ← config
trace ← setupTrace (Right c) "demo-playground"
where
  config :: IO CM.Configuration
  config = do
    c ← CM.empty
    CM.setMinSeverity c Debug
    CM.setSetupBackends c [KatipBK, AggregationBK]
    CM.setDefaultBackends c [KatipBK, AggregationBK]
    CM.setSetupScribes c [ScribeDefinition {
      scName = "stdout"
    , scKind = StdoutSK
    , scRotation = Nothing
    }
    ]
    CM.setDefaultScribes c ["StdoutSK::stdout"]
  return c
```

2. *c* is the **Configuration** of *trace*. In order to enable the collection and processing of measurements (min, max, mean, std-dev) *AggregationBK* is needed.

```
CM.setDefaultBackends c [KatipBK, AggregationBK]
```

in a configuration file (YAML) means

```
defaultBackends:
  - KatipBK
  - AggregationBK
```

3. Set the measurements that you want to take by changing the configuration of the *trace* using **setSubTrace**, in order to declare the namespace where we want to enable the particular measurements and the list with the kind of measurements.

```
CM.setSubTrace
  config
  "demo-playground.submit-tx"
  (Just $ ObservableTrace observablesSet)
where
  observablesSet = [MonotonicClock, MemoryStats]
```

4. Find an action to measure. e.g.:

```
runProtocolWithPipe x hdl proto 'catch' (λProtocolStopped → return ())
```

and use **bracketObserveIO**. e.g.:

```
bracketObserveIO trace "submit-tx" $
  runProtocolWithPipe x hdl proto 'catch' (λProtocolStopped → return ())
```

```
bracketObserveIO :: Config.Configuration → Trace IO a → Severity → Text → IO t → IO t
```

```
bracketObserveIO config trace severity name action = do
```

```
  subTrace ← fromMaybe Neutral < $ > Config.findSubTrace config name
```

```
  bracketObserveIO' subTrace severity trace action
```

```
where
```

```
bracketObserveIO' :: SubTrace → Severity → Trace IO a → IO t → IO t
```

```
bracketObserveIO' NoTrace _ _ act = act
```

```
bracketObserveIO' subtrace sev logTrace act = do
```

```
  mCountersid ← observeOpen subtrace sev logTrace
```

```
  -- run action; if an exception is caught it will be logged and rethrown.
```

```
  t ← act 'catch' (λ(e :: SomeException) → (TIO.hPutStrLn stderr (pack (show e)) >> throwM e))
```

```
  case mCountersid of
```

```
    Left openException →
```

```
      -- since observeOpen faced an exception there is no reason to call observeClose
```

```
      -- however the result of the action is returned
```

```
      TIO.hPutStrLn stderr ("ObserveOpen: " <> pack (show openException))
```

```
    Right countersid → do
```

```
      res ← observeClose subtrace sev logTrace countersid [ ]
```

```
      case res of
```

```
        Left ex → TIO.hPutStrLn stderr ("ObserveClose: " <> pack (show ex))
```

```
        _ → pure ()
```

```
  pure t
```

Monadic.bracketObserverM

Observes a `MonadIO m ⇒ m` action and adds a name to the logger name of the passed in **Trace**. An empty `Text` leaves the logger name untouched.

```
bracketObserveM :: (MonadCatch m, MonadIO m) ⇒ Config.Configuration → Trace IO a → Severity → Text → m t → m t
```

```
bracketObserveM config trace severity name action = do
```

```
  subTrace ← liftIO $ fromMaybe Neutral < $ > Config.findSubTrace config name
```

```
  bracketObserveM' subTrace severity trace action
```

```
where
```

```
bracketObserveM' :: (MonadCatch m, MonadIO m) ⇒ SubTrace → Severity → Trace IO a → m t → m t
```

```
bracketObserveM' NoTrace _ _ act = act
```

```
bracketObserveM' subtrace sev logTrace act = do
```

```
  mCountersid ← liftIO $ observeOpen subtrace sev logTrace
```

```
  -- run action; if an exception is caught it will be logged and rethrown.
```

```
  t ← act 'catch' (λ(e :: SomeException) → liftIO (TIO.hPutStrLn stderr (pack (show e)) >> throwM e))
```

```
  case mCountersid of
```

```
    Left openException →
```

```
      -- since observeOpen faced an exception there is no reason to call observeClose
```

```
      -- however the result of the action is returned
```



```

    liftIO $ TIO.hPutStrLn stderr ("ObserveOpen: " <> pack (show openException))
  Right countersid → do
    res ← liftIO $ observeClose subtrace sev logTrace countersid [ ]
    case res of
      Left ex → liftIO (TIO.hPutStrLn stderr ("ObserveClose: " <> pack (show ex)))
      _ → pure ()
  pure t

```

observerOpen

```

observeOpen :: SubTrace → Severity → Trace IO a → IO (Either SomeException CounterState)
observeOpen subtrace severity logTrace = (do
  identifier ← newUnique
  -- take measurement
  counters ← readCounters subtrace
  let state = CounterState identifier counters
  if counters ≡ [ ]
  then return ()
  else do
    -- send opening message to Trace
    traceNamedObject logTrace ≪
      (,) < $ > (mkLOMeta severity Confidential) < * > pure (ObserveOpen state)
  return (Right state)) 'catch' (return ∘ Left)

```

observeClose

```

observeClose
  :: SubTrace
  → Severity
  → Trace IO a
  → CounterState
  → [(LOMeta, LOContent a)]
  → IO (Either SomeException ())
observeClose subtrace sev logTrace initState logObjects = (do
  let identifier = csIdentifier initState
      initialCounters = csCounters initState
  -- take measurement
  counters ← readCounters subtrace
  if counters ≡ [ ]
  then return ()
  else do
    mle ← mkLOMeta sev Confidential
    -- send closing message to Trace
    traceNamedObject logTrace $
      (mle, ObserveClose (CounterState identifier counters))
    -- send diff message to Trace
    traceNamedObject logTrace $
      (mle, ObserveDiff (CounterState identifier (diffCounters initialCounters counters)))
    -- trace the messages gathered from inside the action

```

```
forM_ logObjects $ traceNamedObject logTrace
return (Right ())) 'catch' (return ◦ Left)
```

1.6.3 Cardano.BM.Trace

Utilities

Natural transformation from monad m to monad n .

```
natTrace :: (forall x ◦ m x → n x) → Trace m a → Trace n a
natTrace nat trace = TracerT.natTrace nat trace
```

Enter new named context

A new context name is added.

```
appendName :: MonadIO m ⇒ LoggerName → Trace m a → m (Trace m a)
appendName name =
  modifyName (λprevLoggerName → appendWithDot name prevLoggerName)
appendWithDot :: LoggerName → LoggerName → LoggerName
appendWithDot "" newName = newName
appendWithDot xs "" = xs
appendWithDot xs newName = xs <> "." <> newName
```

Change named context

The context name is overwritten.

```
modifyName :: MonadIO m ⇒ (LoggerName → LoggerName) → Trace m a → m (Trace m a)
modifyName f trace0 = return $ modifyNameBase f trace0
modifyNameBase
  :: (LoggerName → LoggerName)
  → Tracer m (LogObject a)
  → Tracer m (LogObject a)
modifyNameBase k = contramap f
where
  f (LogObject name meta item) = LogObject (k name) meta item
```

Contramap a trace and produce the naming context

```
named :: Tracer m (LogObject a) → Tracer m (LOMeta, LOContent a)
named = contramap $ uncurry (LogObject mempty)
```

Trace a LogObject through

```
traceNamedObject
  :: MonadIO m
  ⇒ Trace m a
  → (LOMeta, LOContent a)
  → m ()
traceNamedObject logTrace lo =
  tracingWith (named logTrace) lo
```

Evaluation of **FilterTrace**

A filter consists of a *DropName* and a list of *UnhideNames*. If the context name matches the *DropName* filter, then at least one of the *UnhideNames* must match the name to have the evaluation of the filters return *True*.

```
evalFilters :: [(DropName, UnhideNames)] → LoggerName → Bool
evalFilters fs nm =
  all (λ(no, yes) → if (dropFilter nm no) then (unhideFilter nm yes) else True) fs
where
  dropFilter :: LoggerName → DropName → Bool
  dropFilter name (Drop sel) = {-not -} (matchName name sel)
  unhideFilter :: LoggerName → UnhideNames → Bool
  unhideFilter _ (Unhide []) = False
  unhideFilter name (Unhide us) = any (λsel → matchName name sel) us
  matchName :: LoggerName → NameSelector → Bool
  matchName name (Exact name') = name == name'
  matchName name (StartsWith prefix) = T.isPrefixOf prefix name
  matchName name (EndsWith postfix) = T.isSuffixOf postfix name
  matchName name (Contains name') = T.isInfixOf name' name
```

Concrete Trace on stdout

This function returns a trace with an action of type "**LogObject** *a* → IO ()" which will output a text message as text and all others as JSON encoded representation to the console.

TODO remove *locallock*

```
locallock :: MVar ()
locallock = unsafePerformIO $ newMVar ()

stdoutTrace :: Tracer IO (LogObject T.Text)
stdoutTrace = Tracer $ Op $ λ(LogObject logname _ lc) →
  withMVar locallock $ \_ →
    case lc of
      (LogMessage logItem) →
        output logname $ logItem
      obj →
        output logname $ toStrict (encodeToLazyText obj)
where
  output nm msg = TIO.putStrLn $ nm <> " :: " <> msg
```

Concrete Trace into a TVar

```
traceInTVar :: STM.TVar [a] → Tracer STM.STM a
traceInTVar tvar = Tracer $ Op $ λa → STM.modifyTVar tvar ((:) a)
traceInTVarIO :: STM.TVar [LogObject a] → Tracer IO (LogObject a)
traceInTVarIO tvar = Tracer $ Op $ λln →
  STM.atomically $ STM.modifyTVar tvar ((:) ln)
traceInTVarIOConditionally :: STM.TVar [LogObject a] → Config.Configuration → Tracer IO (LogObject a)
```

```

traceInTVarIOConditionally tvar config =
  Tracer $ Op $ \item@(LogObject loggename meta _) → do
    globminsev ← Config.minSeverity config
    globnamesev ← Config.inspectSeverity config loggename
    let minsev = max globminsev $ fromMaybe Debug globnamesev
    subTrace ← fromMaybe Neutral < $ > Config.findSubTrace config loggename
    let doOutput = subtraceOutput subTrace item
    when ((severity meta) ≥ minsev ∧ doOutput) $
      STM.atomically $ STM.modifyTVar tvar ((:) item)
    case subTrace of
      TeeTrace secName →
        STM.atomically $ STM.modifyTVar tvar ((:) item {loName = secName})
      _ → return ()
subtraceOutput :: SubTrace → LogObject a → Bool
subtraceOutput subTrace (LogObject lname _ loitem) =
  case subTrace of
    FilterTrace filters →
      case loitem of
        LogValue name _ →
          evalFilters filters (lname <> "." <> name)
        _ →
          evalFilters filters lname
    DropOpening → case loitem of
      ObserveOpen _ → False
      _ → True
    NoTrace → False
    _ → True

```

Enter message into a trace

The function `traceNamedItem` creates a `LogObject` and threads this through the action defined in the `Trace`.

```

traceNamedItem
  :: MonadIO m
  ⇒ Trace m a
  → PrivacyAnnotation
  → Severity
  → a
  → m ()
traceNamedItem trace p s m =
  traceNamedObject trace ≪≪
    (,) < $ > liftIO (mkLOMeta s p)
    < * > pure (LogMessage m)

```

Logging functions

```

logDebug, logInfo, logNotice, logWarning, logError, logCritical, logAlert, logEmergency
  :: MonadIO m ⇒ Trace m a → a → m ()

```

```

logDebug  logTrace = traceNamedItem logTrace Public Debug
logInfo   logTrace = traceNamedItem logTrace Public Info
logNotice logTrace = traceNamedItem logTrace Public Notice
logWarning logTrace = traceNamedItem logTrace Public Warning
logError  logTrace = traceNamedItem logTrace Public Error
logCritical logTrace = traceNamedItem logTrace Public Critical
logAlert  logTrace = traceNamedItem logTrace Public Alert
logEmergency logTrace = traceNamedItem logTrace Public Emergency
logDebugS, logInfoS, logNoticeS, logWarningS, logErrorS, logCriticalS, logAlertS, logEmergencyS
  :: MonadIO m => Trace m a -> a -> m ()
logDebugS  logTrace = traceNamedItem logTrace Confidential Debug
logInfoS    logTrace = traceNamedItem logTrace Confidential Info
logNoticeS  logTrace = traceNamedItem logTrace Confidential Notice
logWarningS logTrace = traceNamedItem logTrace Confidential Warning
logErrorS   logTrace = traceNamedItem logTrace Confidential Error
logCriticalS logTrace = traceNamedItem logTrace Confidential Critical
logAlertS   logTrace = traceNamedItem logTrace Confidential Alert
logEmergencyS logTrace = traceNamedItem logTrace Confidential Emergency

```

1.6.4 Cardano.BM.Setup




Figure 1.4: Setup procedure

setupTrace

Setup a new **Trace** (**Trace**) with either a given **Configuration** (**Configuration.Model**) or a **FilePath** to a configuration file. After all tracing operations have ended; **shutdownTrace** must be called.

```

setupTrace :: (MonadIO m, Show a, ToJSON a) => Either FilePath Config.Configuration -> Text -> m (Trace m a)
setupTrace (Left cfgFile) name = do
  c <- liftIO $ Config.setup cfgFile
  fst < $ > setupTrace_ c name
setupTrace (Right c) name = fst < $ > setupTrace_ c name
setupTrace_ :: (MonadIO m, Show a, ToJSON a) => Config.Configuration -> Text -> m (Trace m a, Switchboard)
setupTrace_ c name = do
  sb <- liftIO $ Switchboard.realize c

```

```
tr ← appendName name $ natTrace liftIO (Switchboard.mainTraceConditionally c sb)
return (tr, sb)
```

shutdown

Shut down the Switchboard and all the **Traces** related to it.

```
shutdown :: (Show a, ToJSON a) ⇒ Switchboard.Switchboard a → IO ()
shutdown = Switchboard.unrealize
```

withTrace

Setup a **Trace** from **Configuration** and pass it to the action. At the end, shutdown all the components and close the trace.

```
withTrace :: (MonadIO m, MonadMask m, Show a, ToJSON a) ⇒ Config.Configuration → Text → (Trace m a)
withTrace cfg name action =
  bracket
    (setupTrace_cfg name)           -- acquire
    (λ(., sb) → liftIO $ shutdown sb) -- release
    (λ(tr, _) → action tr)         -- action
```

1.6.5 Cardano.BM.Counters

Here the platform is chosen on which we compile this program.

Currently, we mainly support *Linux* with its 'proc' filesystem.

```
{-# LANGUAGE CPP #-}
# if defined (linux_HOST_OS)
# define LINUX
# endif
module Cardano.BM.Counters
(
  Platform.readCounters
, diffTimeObserved
, getMonoClock
) where
# ifdef LINUX
import qualified Cardano.BM.Counters.Linux as Platform
# else
import qualified Cardano.BM.Counters.Dummy as Platform
# endif
import Cardano.BM.Counters.Common (getMonoClock)
import Cardano.BM.Data.Aggregated (Measurable (..))
import Cardano.BM.Data.Counter
```

Calculate difference between clocks

```

diffTimeObserved :: CounterState → CounterState → Measurable
diffTimeObserved (CounterState id0 startCounters) (CounterState id1 endCounters) =
  let
    startTime = getMonotonicTime startCounters
    endTime = getMonotonicTime endCounters
  in
    if (id0 ≡ id1)
    then endTime - startTime
    else error "these clocks are not from the same experiment"
where
  getMonotonicTime counters = case (filter isMonotonicClockCounter counters) of
    [(Counter MonotonicClockTime _ mus)] → mus
    _ → error "A time measurement is missing!"
  isMonotonicClockCounter :: Counter → Bool
  isMonotonicClockCounter = (MonotonicClockTime ≡) ∘ cType

```

1.6.6 Cardano.BM.Counters.Common

Common functions that serve *readCounters* on all platforms.

```

nominalTimeToMicroseconds :: Word64 → Microsecond
nominalTimeToMicroseconds = fromMicroseconds ∘ toInteger ∘ ('div' 1000)

```

Read monotonic clock

```

getMonoClock :: IO [Counter]
getMonoClock = do
  t ← getMonotonicTimeNSec
  return [Counter MonotonicClockTime "monoclock" $ Microseconds (t `div' 1000)]

```

Read GHC RTS statistics

Read counters from GHC's *RTS* (runtime system). The values returned are as per the last GC (garbage collection) run.

```

readRTSStats :: IO [Counter]
readRTSStats = do
  iscollected ← GhcStats.getRTSStatsEnabled
  if iscollected
  then ghcstats
  else return []
where
  ghcstats :: IO [Counter]
  ghcstats = do
    -- need to run GC?
    rts ← GhcStats.getRTSStats
    let getrts = ghcval rts

```

```

return [getrts (Bytes ◦ fromIntegral ◦ GHCStats.allocated_bytes, "bytesAllocated")
      ,getrts (Bytes ◦ fromIntegral ◦ GHCStats.max_live_bytes, "liveBytes")
      ,getrts (Bytes ◦ fromIntegral ◦ GHCStats.max_large_objects_bytes, "largeBytes")
      ,getrts (Bytes ◦ fromIntegral ◦ GHCStats.max_compact_bytes, "compactBytes")
      ,getrts (Bytes ◦ fromIntegral ◦ GHCStats.max_slop_bytes, "slopBytes")
      ,getrts (Bytes ◦ fromIntegral ◦ GHCStats.max_mem_in_use_bytes, "usedMemBytes")
      ,getrts (Nanoseconds ◦ fromIntegral ◦ GHCStats.gc_cpu_ns, "gcCpuNs")
      ,getrts (Nanoseconds ◦ fromIntegral ◦ GHCStats.gc_elapsed_ns, "gcElapsedNs")
      ,getrts (Nanoseconds ◦ fromIntegral ◦ GHCStats.cpu_ns, "cpuNs")
      ,getrts (Nanoseconds ◦ fromIntegral ◦ GHCStats.elapsed_ns, "elapsedNs")
      ,getrts (PureI ◦ toInteger ◦ GHCStats.gcs, "gcNum")
      ,getrts (PureI ◦ toInteger ◦ GHCStats.major_gcs, "gcMajorNum")
      ]
ghcval :: GHCStats.RTSStats → ((GHCStats.RTSStats → Measurable), Text) → Counter
ghcval s (f, n) = Counter RTSStats n $ (f s)

```

1.6.7 Cardano.BM.Counters.Dummy

This is a dummy definition of `readCounters` on platforms that do not support the 'proc' filesystem from which we would read the counters.

The only supported measurements are monotonic clock time and RTS statistics for now.

```

readCounters :: SubTrace → IO [Counter]
readCounters NoTrace      = return []
readCounters Neutral      = return []
readCounters (TeeTrace _) = return []
readCounters (FilterTrace _) = return []
readCounters UntimedTrace = return []
readCounters DropOpening  = return []
# ifdef ENABLE_OBSERVABLES
readCounters (ObservableTrace tts) = foldrM (\sel, fun) a →
  if any (≡ sel) tts
  then (fun ≫ λxs → return $ a ++ xs)
  else return a [ ] selectors
  where
    selectors = [ (MonotonicClock, getMonoClock)
                  , (GhcRtsStats, readRTSStats)
                  ]
# else
readCounters (ObservableTrace _) = return []
# endif

```

1.6.8 Cardano.BM.Counters.Linux

we have to expand the `readMemStats` function to read full data from `proc`

```

readCounters :: SubTrace → IO [Counter]
readCounters NoTrace      = return []
readCounters Neutral      = return []
readCounters (TeeTrace _) = return []

```



```

readCounters (FilterTrace _) = return []
readCounters UntimedTrace = return []
readCounters DropOpening = return []
# ifdef ENABLE_OBSERVABLES
readCounters (ObservableTrace tts) = do
    pid ← getProcessID
    foldrM (λ(sel,fun) a →
        if any (≡ sel) tts
        then (fun ≧ λxs → return $ a ++ xs)
        else return a) [] (selectors pid)
    where
        selectors pid = [ (MonotonicClock, getMonoClock)
            , (MemoryStats, readProcStatM pid)
            , (ProcessStats, readProcStats pid)
            , (NetStats, readProcNet pid)
            , (IOStats, readProcIO pid)
            , (GhcRtsStats, readRTSStats)
            ]
# else
readCounters (ObservableTrace _) = return []
# endif

# ifdef ENABLE_OBSERVABLES
pathProc :: FilePath
pathProc = "/proc/"
pathProcStat :: ProcessID → FilePath
pathProcStat pid = pathProc </> (show pid) </> "stat"
pathProcStatM :: ProcessID → FilePath
pathProcStatM pid = pathProc </> (show pid) </> "statm"
pathProcIO :: ProcessID → FilePath
pathProcIO pid = pathProc </> (show pid) </> "io"
pathProcNet :: ProcessID → FilePath
pathProcNet pid = pathProc </> (show pid) </> "net" </> "netstat"
# endif

```

Reading from a file in /proc/<pid>

```

# ifdef ENABLE_OBSERVABLES
readProcList :: FilePath → IO [Integer]
readProcList fp = do
    cs ← readFile fp
    return $ map (λs → maybe 0 id $ (readMaybe s :: Maybe Integer)) (words cs)
# endif

```

readProcStatM - /proc/<pid>/statm

```

/proc/[pid]/statm
    Provides information about memory usage, measured in pages. The columns are:
        size      (1) total program size
                  (same as VmSize in /proc/[pid]/status)
        resident  (2) resident set size

```

```

        (same as VmRSS in /proc/[pid]/status)
shared    (3) number of resident shared pages (i.e., backed by a file)
        (same as RssFile+RssShmem in /proc/[pid]/status)
text      (4) text (code)
lib       (5) library (unused since Linux 2.6; always 0)
data      (6) data + stack
dt        (7) dirty pages (unused since Linux 2.6; always 0)

```

```

# ifdef ENABLE_OBSERVABLES
readProcStatM :: ProcessID → IO [Counter]
readProcStatM pid = do
    ps0 ← readProcList (pathProcStatM pid)
    let ps = zip colnames ps0
    psUseful = filter (("unused" ≠) ∘ fst) ps
    return $ map (λ(n,i) → Counter MemoryCounter n (PureI i)) psUseful
where
    colnames :: [Text]
    colnames = ["size", "resident", "shared", "text", "unused", "data", "unused"]
# endif

```

readProcStats - //proc//<pid>//stat

```
/proc/[pid]/stat
```

Status information about the process. This is used by ps(1). It is defined in the kernel source file fs/proc/array.c.

The fields, in order, with their proper scanf(3) format specifiers, are listed below. Whether or not certain of these fields display valid information is governed by a ptrace access mode PTRACE_MODE_READ_FSCREDS | PTRACE_MODE_NOAUDIT check (refer to ptrace(2)). If the check denies access, then the field value is displayed as 0. The affected fields are indicated with the marking [PT].

```

(1) pid %d
    The process ID.

(2) comm %s
    The filename of the executable, in parentheses. This is visible whether or not the executable is swapped out.

(3) state %c
    One of the following characters, indicating process state:

    R  Running

    S  Sleeping in an interruptible wait

    D  Waiting in uninterruptible disk sleep

    Z  Zombie

    T  Stopped (on a signal) or (before Linux 2.6.33) trace stopped

    t  Tracing stop (Linux 2.6.33 onward)

    W  Paging (only before Linux 2.6.0)

    X  Dead (from Linux 2.6.0 onward)

    x  Dead (Linux 2.6.33 to 3.13 only)

    K  Wakekill (Linux 2.6.33 to 3.13 only)

    W  Waking (Linux 2.6.33 to 3.13 only)

    P  Parked (Linux 3.9 to 3.13 only)

```

- (4) `ppid %d`
The PID of the parent of this process.
- (5) `pgrp %d`
The process group ID of the process.
- (6) `session %d`
The session ID of the process.
- (7) `tty_nr %d`
The controlling terminal of the process. (The minor device number is contained in the combination of bits 31 to 20 and 7 to 0; the major device number is in bits 15 to 8.)
- (8) `tpgid %d`
The ID of the foreground process group of the controlling terminal of the process.
- (9) `flags %u`
The kernel flags word of the process. For bit meanings, see the `PF_*` defines in the Linux kernel source file `include/linux/sched.h`. Details depend on the kernel version.

The format for this field was `%lu` before Linux 2.6.
- (10) `minflt %lu`
The number of minor faults the process has made which have not required loading a memory page from disk.
- (11) `cminflt %lu`
The number of minor faults that the process's waited-for children have made.
- (12) `majflt %lu`
The number of major faults the process has made which have required loading a memory page from disk.
- (13) `cmajflt %lu`
The number of major faults that the process's waited-for children have made.
- (14) `utime %lu`
Amount of time that this process has been scheduled in user mode, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`). This includes guest time, `guest_time` (time spent running a virtual CPU, see below), so that applications that are not aware of the guest time field do not lose that time from their calculations.
- (15) `stime %lu`
Amount of time that this process has been scheduled in kernel mode, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (16) `cutime %ld`
Amount of time that this process's waited-for children have been scheduled in user mode, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`). (See also `times(2)`.) This includes guest time, `cguest_time` (time spent running a virtual CPU, see below).
- (17) `cstime %ld`
Amount of time that this process's waited-for children have been scheduled in kernel mode, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (18) `priority %ld`
(Explanation for Linux 2.6) For processes running a real-time scheduling policy (policy below; see `sched_setscheduler(2)`), this is the negated scheduling priority, minus one; that is, a number in the range -2 to -100, corresponding to real-time priorities 1 to 99. For processes running under a non-real-time scheduling policy, this is the raw nice value (`setpriority(2)`) as represented in the kernel. The kernel stores nice values as numbers in the range 0 (high) to 39 (low), corresponding to the user-visible nice range of -20 to 19.
- (19) `nice %ld`
The nice value (see `setpriority(2)`), a value in the range 19 (low priority) to -20 (high priority).
- (20) `num_threads %ld`
Number of threads in this process (since Linux 2.6). Before kernel 2.6, this field was hard coded to 0 as a placeholder for an earlier removed field.
- (21) `itrealvalue %ld`
The time in jiffies before the next `SIGALRM` is sent to the process due to an interval timer.

Since kernel 2.6.17, this field is no longer maintained, and is hard coded as 0.

- (22) starttime %llu
The time the process started after system boot. In kernels before Linux 2.6, this value was expressed in jiffies. Since Linux 2.6, the value is expressed in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).

The format for this field was %lu before Linux 2.6.
- (23) vsize %lu
Virtual memory size in bytes.
- (24) rss %ld
Resident Set Size: number of pages the process has in real memory. This is just the pages which count toward text, data, or stack space. This does not include pages which have not been demand-loaded in, or which are swapped out.
- (25) rsslim %lu
Current soft limit in bytes on the rss of the process; see the description of `RLIMIT_RSS` in `getrlimit(2)`.
- (26) startcode %lu [PT]
The address above which program text can run.
- (27) endcode %lu [PT]
The address below which program text can run.
- (28) startstack %lu [PT]
The address of the start (i.e., bottom) of the stack.
- (29) kstkesp %lu [PT]
The current value of ESP (stack pointer), as found in the kernel stack page for the process.
- (30) kstkeip %lu [PT]
The current EIP (instruction pointer).
- (31) signal %lu
The bitmap of pending signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (32) blocked %lu
The bitmap of blocked signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (33) sigignore %lu
The bitmap of ignored signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (34) sigcatch %lu
The bitmap of caught signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (35) wchan %lu [PT]
This is the "channel" in which the process is waiting. It is the address of a location in the kernel where the process is sleeping. The corresponding symbolic name can be found in `/proc/[pid]/wchan`.
- (36) nswap %lu
Number of pages swapped (not maintained).
- (37) cnswap %lu
Cumulative nswap for child processes (not maintained).
- (38) exit_signal %d (since Linux 2.1.22)
Signal to be sent to parent when we die.
- (39) processor %d (since Linux 2.2.8)
CPU number last executed on.
- (40) rt_priority %u (since Linux 2.5.19)
Real-time scheduling priority, a number in the range 1 to 99 for processes scheduled under a real-time policy, or 0, for non-real-time processes (see `sched_setscheduler(2)`).

- (41) `policy %u` (since Linux 2.5.19)
Scheduling policy (see `sched_setscheduler(2)`). Decode using the `SCHED_*` constants in `linux/sched.h`.

The format for this field was `%lu` before Linux 2.6.22.
- (42) `delayacct_blkio_ticks %llu` (since Linux 2.6.18)
Aggregated block I/O delays, measured in clock ticks (centiseconds).
- (43) `guest_time %lu` (since Linux 2.6.24)
Guest time of the process (time spent running a virtual CPU for a guest operating system), measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (44) `cguest_time %ld` (since Linux 2.6.24)
Guest time of the process's children, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (45) `start_data %lu` (since Linux 3.3) [PT]
Address above which program initialized and uninitialized (BSS) data are placed.
- (46) `end_data %lu` (since Linux 3.3) [PT]
Address below which program initialized and uninitialized (BSS) data are placed.
- (47) `start_brk %lu` (since Linux 3.3) [PT]
Address above which program heap can be expanded with `brk(2)`.
- (48) `arg_start %lu` (since Linux 3.5) [PT]
Address above which program command-line arguments (`argv`) are placed.
- (49) `arg_end %lu` (since Linux 3.5) [PT]
Address below program command-line arguments (`argv`) are placed.
- (50) `env_start %lu` (since Linux 3.5) [PT]
Address above which program environment is placed.
- (51) `env_end %lu` (since Linux 3.5) [PT]
Address below which program environment is placed.
- (52) `exit_code %d` (since Linux 3.5) [PT]
The thread's exit status in the form reported by `waitpid(2)`.

```
# ifdef ENABLE_OBSERVABLES
```

```
readProcStats :: ProcessID → IO [Counter]
```

```
readProcStats pid = do
```

```
  ps0 ← readProcList (pathProcStat pid)
```

```
  let ps = zip colnames ps0
```

```
      psUseful = filter (("unused" ≠) ∘ fst) ps
```

```
  return $ map (λ(n,i) → Counter StatInfo n (PureI i)) psUseful
```

```
where
```

```
  colnames :: [Text]
```

```
  colnames = [ "pid", "unused", "unused", "ppid", "ppgrp", "session", "ttynr", "tpgid", "flags", "minflt", "cminflt", "majflt", "cmajflt", "utime", "stime", "cutime", "cstime", "priority", "nice", "rtime", "rtmin", "rtmax", "itrealvalue", "starttime", "vsize", "rss", "rsslim", "startcode", "endcode", "startstack", "signal", "blocked", "sigignore", "sigcatch", "wchan", "nswap", "cnsnap", "exitcode", "policy", "blkio", "guesttime", "cguesttime", "startdata", "enddata", "startbrk", "argstart", "envend", "exitcode" ]
```

```
# endif
```

```
readProcIO - //proc//<pid >//io
```

```
/proc/[pid]/io (since kernel 2.6.20)
```

This file contains I/O statistics for the process, for example:

```
# cat /proc/3828/io
rchar: 323934931
wchar: 323929600
syscr: 632687
syscw: 632675
read_bytes: 0
write_bytes: 323932160
cancelled_write_bytes: 0
```

The fields are as follows:

rchar: characters read

The number of bytes which this task has caused to be read from storage. This is simply the sum of bytes which this process passed to read(2) and similar system calls. It includes things such as terminal I/O and is unaffected by whether or not actual physical disk I/O was required (the read might have been satisfied from pagecache).

wchar: characters written

The number of bytes which this task has caused, or shall cause to be written to disk. Similar caveats apply here as with rchar.

syscr: read syscalls

Attempt to count the number of read I/O operations—that is, system calls such as read(2) and pread(2).

syscw: write syscalls

Attempt to count the number of write I/O operations—that is, system calls such as write(2) and pwrite(2).

read_bytes: bytes read

Attempt to count the number of bytes which this process really did cause to be fetched from the storage layer. This is accurate for block-backed filesystems.

write_bytes: bytes written

Attempt to count the number of bytes which this process caused to be sent to the storage layer.

cancelled_write_bytes:

The big inaccuracy here is truncate. If a process writes 1MB to a file and then deletes the file, it will in fact perform no writeout. But it will have been accounted as having caused 1MB of write. In other words: this field represents the number of bytes which this process caused to not happen, by truncating pagecache. A task can cause "negative" I/O too. If this task truncates some dirty pagecache, some I/O which another task has been accounted for (in its write_bytes) will not be happening.

Note: In the current implementation, things are a bit racy on 32-bit systems: if process A reads process B's /proc/[pid]/io while process B is updating one of these 64-bit counters, process A could see an intermediate result.

Permission to access this file is governed by a ptrace access mode PTRACE_MODE_READ_FSCREDS check; see ptrace(2).

```
# ifdef ENABLE_OBSERVABLES
```

```
readProcIO :: ProcessID → IO [Counter]
```

```
readProcIO pid = do
```

```
    ps0 ← readProcList (pathProcIO pid)
```

```
    let ps = zip3 colnames ps0 units
```

```
    return $ map (λ(n,i,u) → Counter IOCounter n (u i)) ps
```

```
where
```

```
    colnames :: [Text]
```

```
    colnames = [ "rchar", "wchar", "syscr", "syscw", "rbytes", "wbytes", "cxwbytes" ]
```

```
    units = [ Bytes ∘ fromInteger, Bytes ∘ fromInteger, PureI, PureI, Bytes ∘ fromInteger, Bytes ∘ fromInteger, Bytes ∘ fromInteger ]
```

```
# endif
```

Network TCP/IP counters

example:

```

\\
cat /proc/<pid>/net/netstat
\\
TcpExt: SyncookiesSent SyncookiesRecv SyncookiesFailed EmbryonicRsts PruneCalled RcvPruned OfoPruned OutOfWindowIcmps Lo!
!ckDroppedIcmps ArpFilter TW TWRcycled TWKilled PAWSActive PAWSEstab DelayedACKs DelayedACKLocked DelayedACKLost Listen0!
!verflows ListenDrops TCPHPHits TCPPureAcks TCPHPAcks TCPReNoRecovery TCPSackRecovery TCPSACKReneging TCPSACKReorder TCPR!
!enoReorder TCPTSReorder TCPFullUndo TCPPartialUndo TCPDSACKUndo TCPLossUndo TCPLostRetransmit TCPReNoFailures TCPSackFai!
!lures TCPLossFailures TCPFastRetrans TCPSlowStartRetrans TCPTimeouts TCPLossProbes TCPLossProbeRecovery TCPReNoRecoveryF!
!ail TCPSackRecoveryFail TCPRecvCollapsed TCPDSACKOldSent TCPDSACKOfoSent TCPDSACKRecv TCPDSACKOfoRecv TCPAbortOnData TCPA!
!bortOnClose TCPAbortOnMemory TCPAbortOnTimeout TCPAbortOnLinger TCPAbortFailed TCPMemoryPressures TCPMemoryPressuresChro!
!no TCPSACKDiscard TCPDSACKIgnoredOld TCPDSACKIgnoredNoUndo TCPSpuriousRTOs TCPMD5NotFound TCPMD5Unexpected TCPMD5Failure!
! TCPSackShifted TCPSackMerged TCPSackShiftFallback TCPBacklogDrop PFMemallocDrop TCPMinTTLDrop TCPDeferAcceptDrop IPReve!
!rsePathFilter TCPTIMEwaitOverflow TCPReqQFullDoCookies TCPReqQFullDrop TCPReTransFail TCPRecvCoalesce TCPOFOQueue TCPOFOD!
!rop TCPOFOMerge TCPChallengeACK TCPSYNChallenge TCPFastOpenActive TCPFastOpenActiveFailTCPFastOpenPassive TCPFastOpenPas!
!siveFail TCPFastOpenListenOverflow TCPFastOpenCookieReqd TCPFastOpenBlackhole TCPSpuriousRtxHostQueues BusyPollRxPackets!
! TCPAutoCorking TCPFromZeroWindowAdv TCPToZeroWindowAdv TCPWantZeroWindowAdv TCPSynRetrans TCPOrigDataSent TCPHystartTra!
!inDetect TCPHystartTrainCwnd TCPHystartDelayDetect TCPHystartDelayCwnd TCPACKSkippedSynRecv TCPACKSkippedPAWS TCPACKSkip!
!pedSeq TCPACKSkippedFinWait2 TCPACKSkippedTimeWait TCPACKSkippedChallenge TCPWinProbe TCPKeepAlive TCPMTUPFail TCPMTUPSu!
!ccess TCPDelivered TCPDeliveredCE TCPAckCompressed
TcpExt: 0 0 0 0 28 0 0 0 0 1670 1 0 0 6 6029 1 1766 0 0 384612 66799 105553 0 21 0 638 0 1 7 1 1 32 128 0 1 0 22 0 116!
! 383 19 0 0 0 1788 224 178 0 435 224 0 13 0 0 0 0 0 67 0 0 0 0 3 1 668 0 0 0 4 0 0 0 0 91870 4468 0 224 22 23 0 0 0 !
!0 0 0 0 6 0 21492 0 0 11 188 188680 6 145 13 425 0 3 4 0 0 1 117 22984 0 0 192495 0 4500
IpExt: InNoRoutes InTruncatedPkts InMcastPkts OutMcastPkts InBcastPkts OutBcastPkts InOctets OutOctets InMcastOctets Out!
!McastOctets InBcastOctets OutBcastOctets InCsumErrors InNoECTPkts InECT1Pkts InECT0Pkts InCEPkts
IpExt: 0 0 20053 8977 2437 23 3163525943 196480057 2426648 1491754 394285 5523 0 3513269 0 217426 0

```

```

# ifdef ENABLE_OBSERVABLES
readProcNet :: ProcessID → IO [Counter]
readProcNet pid = do
  ls0 ← lines < $ > readFile (pathProcNet pid)
  let ps0 = readinfo ls0
  let ps1 = map (λ(n,c) → (n,readMaybe c :: Maybe Integer)) ps0
  return $ mapCounters $ filter selcolumns ps1
where
  construct "IpExt:OutOctets" i = Bytes $ fromInteger i
  construct "IpExt:InOctets" i = Bytes $ fromInteger i
  construct _i = PureI i
  -- only a few selected columns will be returned
  selcolumns (n,_) = n ∈ [ "IpExt:OutOctets", "IpExt:InOctets" ]
  mapCounters [] = []
  mapCounters ((n,c):r) = case c of
    Nothing → mapCounters r
    Just i → mapCounters r <> [Counter NetCounter (pack n) (construct n i)]
  readinfo :: [String] → [(String,String)]
  readinfo [] = []
  readinfo (_:[]) = []
  readinfo (l1:l2:r) =
    let col0 = words l1
        cols = tail col0
        vals = tail $ words l2
        pref = head col0
    in
      readinfo r <> zip (map (λn → pref ++ n) cols) vals
# endif

```

1.6.9 Cardano.BM.Data.Aggregated

Measurable

A **Measurable** may consist of different types of values. Time measurements are strict, so are *Bytes* which are externally measured. The real or integral numeric values are lazily linked, so we can decide later to drop them.

```
data Measurable = Microseconds {-# UNPACK #-} ! Word64
  | Nanoseconds {-# UNPACK #-} ! Word64
  | Seconds      {-# UNPACK #-} ! Word64
  | Bytes        {-# UNPACK #-} ! Word64
  | PureD        Double
  | PureI        Integer
  | Severity     S.Severity
  deriving (Eq, Read, Generic, ToJSON)
```

Measurable can be transformed to an integral value.

instance Ord **Measurable** where

```
compare (Seconds a) (Seconds b)           = compare a b
compare (Microseconds a) (Microseconds b) = compare a b
compare (Nanoseconds a) (Nanoseconds b)   = compare a b
compare (Seconds a) (Microseconds b)      = compare (a * 1000000) b
compare (Nanoseconds a) (Microseconds b)  = compare a (b * 1000)
compare (Seconds a) (Nanoseconds b)       = compare (a * 1000000000) b
compare (Microseconds a) (Nanoseconds b)  = compare (a * 1000) b
compare (Microseconds a) (Seconds b)       = compare a (b * 1000000)
compare (Nanoseconds a) (Seconds b)       = compare a (b * 1000000000)
compare (Bytes a) (Bytes b)               = compare a b
compare (PureD a) (PureD b)               = compare a b
compare (PureI a) (PureI b)               = compare a b
compare (Severity a) (Severity b)         = compare a b
compare (PureI a) (Seconds b) | a ≥ 0 = compare a (toInteger b)
compare (PureI a) (Microseconds b) | a ≥ 0 = compare a (toInteger b)
compare (PureI a) (Nanoseconds b) | a ≥ 0 = compare a (toInteger b)
compare (PureI a) (Bytes b) | a ≥ 0 = compare a (toInteger b)
compare (Seconds a) (PureI b) | b ≥ 0 = compare (toInteger a) b
compare (Microseconds a) (PureI b) | b ≥ 0 = compare (toInteger a) b
compare (Nanoseconds a) (PureI b) | b ≥ 0 = compare (toInteger a) b
compare (Bytes a) (PureI b) | b ≥ 0 = compare (toInteger a) b
compare a@(PureD _) (PureI b)             = compare (getInteger a) b
compare (PureI a) b@(PureD _)             = compare a (getInteger b)
compare a b                               = error $ "cannot compare " ++ (showType a) ++ " " ++ (showType b)
```

Measurable can be transformed to an integral value.

```
getInteger :: Measurable → Integer
getInteger (Microseconds a) = toInteger a
getInteger (Nanoseconds a) = toInteger a
getInteger (Seconds a)     = toInteger a
getInteger (Bytes a)       = toInteger a
getInteger (PureI a)       = a
```



```

getInteger (PureD a)      = round a
getInteger (Severity a)   = toInteger (fromEnum a)

```

Measurable can be transformed to a rational value.

```

getDouble :: Measurable → Double
getDouble (Microseconds a) = fromIntegral a
getDouble (Nanoseconds a)  = fromIntegral a
getDouble (Seconds a)      = fromIntegral a
getDouble (Bytes a)        = fromIntegral a
getDouble (PureI a)        = fromIntegral a
getDouble (PureD a)        = a
getDouble (Severity a)     = fromIntegral (fromEnum a)

```

It is a numerical value, thus supports functions to operate on numbers.

instance Num Measurable where

```

(+) (Microseconds a) (Microseconds b) = Microseconds (a + b)
(+) (Nanoseconds a) (Nanoseconds b) = Nanoseconds (a + b)
(+) (Seconds a)      (Seconds b)      = Seconds      (a + b)
(+) (Bytes a)        (Bytes b)        = Bytes        (a + b)
(+) (PureI a)        (PureI b)        = PureI        (a + b)
(+) (PureD a)        (PureD b)        = PureD        (a + b)
(+) --              = error "Trying to add values with different units"

(*) (Microseconds a) (Microseconds b) = Microseconds (a * b)
(*) (Nanoseconds a) (Nanoseconds b) = Nanoseconds (a * b)
(*) (Seconds a)      (Seconds b)      = Seconds      (a * b)
(*) (Bytes a)        (Bytes b)        = Bytes        (a * b)
(*) (PureI a)        (PureI b)        = PureI        (a * b)
(*) (PureD a)        (PureD b)        = PureD        (a * b)
(*) --              = error "Trying to multiply values with different units"

abs (Microseconds a) = Microseconds (abs a)
abs (Nanoseconds a) = Nanoseconds (abs a)
abs (Seconds a)      = Seconds      (abs a)
abs (Bytes a)        = Bytes        (abs a)
abs (PureI a)        = PureI        (abs a)
abs (PureD a)        = PureD        (abs a)
abs (Severity _)     = error "cannot compute absolute value for Severity"

signum (Microseconds a) = Microseconds (signum a)
signum (Nanoseconds a) = Nanoseconds (signum a)
signum (Seconds a)      = Seconds      (signum a)
signum (Bytes a)        = Bytes        (signum a)
signum (PureI a)        = PureI        (signum a)
signum (PureD a)        = PureD        (signum a)
signum (Severity _)     = error "cannot compute sign of Severity"

negate (Microseconds a) = Microseconds (negate a)
negate (Nanoseconds a) = Nanoseconds (negate a)
negate (Seconds a)      = Seconds      (negate a)
negate (Bytes a)        = Bytes        (negate a)
negate (PureI a)        = PureI        (negate a)
negate (PureD a)        = PureD        (negate a)
negate (Severity _)     = error "cannot negate Severity"

```

fromInteger = PureI

Pretty printing of **Measurable**.

instance Show Measurable where

```
show (Microseconds a) = show a
show (Nanoseconds a)  = show a
show (Seconds a)      = show a
show (Bytes a)         = show a
show (PureI a)         = show a
show (PureD a)         = show a
show (Severity a)      = show a
```

```
showUnits :: Measurable → String
showUnits (Microseconds _) = " s"
showUnits (Nanoseconds _) = " ns"
showUnits (Seconds _)     = " s"
showUnits (Bytes _)       = " B"
showUnits (PureI _)       = ""
showUnits (PureD _)       = ""
showUnits (Severity _)    = ""
```

```
showType :: Measurable → String
showType (Microseconds _) = "Microseconds"
showType (Nanoseconds _) = "Nanoseconds"
showType (Seconds _)     = "Seconds"
showType (Bytes _)       = "Bytes"
showType (PureI _)       = "PureI"
showType (PureD _)       = "PureD"
showType (Severity _)    = "Severity"
```

-- show in S.I. units

```
showSI :: Measurable → String
showSI (Microseconds a) = show (fromFloatDigits ((fromIntegral a) / (1000000 :: Float))) ++
    showUnits (Seconds a)
showSI (Nanoseconds a) = show (fromFloatDigits ((fromIntegral a) / (1000000000 :: Float))) ++
    showUnits (Seconds a)
showSI v@(Seconds a)   = show a ++ showUnits v
showSI v@(Bytes a)     = show a ++ showUnits v
showSI v@(PureI a)     = show a ++ showUnits v
showSI v@(PureD a)     = show a ++ showUnits v
showSI v@(Severity a)  = show a ++ showUnits v
```

Stats

A **Stats** statistics is strictly computed.

```
data BaseStats = BaseStats {
    fmin :: !Measurable,
    fmax :: !Measurable,
    fcount :: {-# UNPACK #-} !Int,
    fsum_A :: {-# UNPACK #-} !Double,
    fsum_B :: {-# UNPACK #-} !Double
} deriving (Generic, ToJSON, Show)
```

```

instance Eq BaseStats where
  (BaseStats mina maxa counta sumAa sumBa)  $\equiv$  (BaseStats minb maxb countb sumAb sumBb) =
    mina  $\equiv$  minb  $\wedge$  maxa  $\equiv$  maxb  $\wedge$  counta  $\equiv$  countb  $\wedge$ 
    abs (sumAa - sumAb) < 1.0e-4  $\wedge$ 
    abs (sumBa - sumBb) < 1.0e-4

```

```

data Stats = Stats {
  flast :: !Measurable,
  fold  :: !Measurable,
  fbasic :: !BaseStats,
  fdelta :: !BaseStats,
  ftimed :: !BaseStats
} deriving (Eq, Generic, ToJSON, Show)

```

```

meanOfStats :: BaseStats  $\rightarrow$  Double
meanOfStats = fsum_A

```

```

stdevOfStats :: BaseStats  $\rightarrow$  Double
stdevOfStats s =
  if fcount s < 2
  then 0
  else sqrt $ (fsum_B s) / (fromInteger $ fromIntegral (fcount s) - 1)

```

instance Semigroup Stats disabled for the moment, because not needed.

We use a parallel algorithm to update the estimation of mean and variance from two sample statistics. (see https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance#Parallel_algorithm)

```

instance Semigroup Stats where
  (<>) a b = let counta = fcount a
    countb = fcount b
    newcount = counta + countb
    delta = fsum_A b - fsum_A a
  in
    Stats {flast = flast b -- right associative
    ,fmin   = min (fmin a) (fmin b)
    ,fmax   = max (fmax a) (fmax b)
    ,fcount = newcount
    ,fsum_A = fsum_A a + (delta / fromInteger newcount)
    ,fsum_B = fsum_B a + fsum_B b + (delta * delta) * (fromInteger (counta * countb) / fromInteger newcount)
    }

```

```

stats2Text :: Stats  $\rightarrow$  Text
stats2Text (Stats slast _ sbasic sdelta stimed) =
  pack $
    "{ last=" ++ show slast ++
    ", basic-stats=" ++ showStats' (sbasic) ++
    ", delta-stats=" ++ showStats' (sdelta) ++
    ", timed-stats=" ++ showStats' (stimed) ++
    " }"
where

```

```

showStats' :: BaseStats → String
showStats' s =
  ", { min=" ++ show (fmin s) ++
  ", max=" ++ show (fmax s) ++
  ", mean=" ++ show (meanOfStats s) ++ showUnits (fmin s) ++
  ", std-dev=" ++ show (stdevOfStats s) ++
  ", count=" ++ show (fcount s) ++
  " }"

```

Exponentially Weighted Moving Average (EWMA)

Following https://en.wikipedia.org/wiki/Moving_average#Exponential_moving_average we calculate the exponential moving average for a series of values Y_t according to:

$$S_t = \begin{cases} Y_1, & t = 1 \\ \alpha \cdot Y_t + (1 - \alpha) \cdot S_{t-1}, & t > 1 \end{cases}$$

```

data EWMA = EmptyEWMA {alpha :: Double}
  | EWMA {alpha :: Double
  , avg :: Measurable
  } deriving (Show, Eq, Generic, ToJSON)

```

Aggregated

```

data Aggregated = AggregatedStats Stats
  | AggregatedEWMA EWMA
  deriving (Eq, Generic, ToJSON)

```

instance Semigroup Aggregated disabled for the moment, because not needed.

```

instance Semigroup Aggregated where
  (<>) (AggregatedStats a) (AggregatedStats b) =
    AggregatedStats (a <> b)
  (<>) _ _ = error "Cannot combine different objects"

```

```

singletonStats :: Measurable → Aggregated
singletonStats a =

```

```

  let stats = Stats {flast = a
    , fold      = Nanoseconds 0
    , fbasic    = BaseStats
      { fmin = a
      , fmax = a
      , fcount = 1
      , fsum_A = getDouble a
      , fsum_B = 0 }
    , fdelta    = BaseStats
      { fmin = 0
      , fmax = 0

```

```

    ,fcount = 0
    ,fsum_A = 0
    ,fsum_B = 0}
    ,ftimed = BaseStats
    {fmin = Nanoseconds 999999999999
    ,fmax = Nanoseconds 0
    ,fcount = (-1)
    ,fsum_A = 0
    ,fsum_B = 0}
  }
in
  AggregatedStats stats

```

```

instance Show Aggregated where
  show (AggregatedStats astats) =
    "{ stats = " ++ show astats ++ " }"
  show (AggregatedEWMA a) = show a

```

1.6.10 Cardano.BM.Data.AggregatedKind

AggregatedKind

This identifies the type of Aggregated.

```

data AggregatedKind = StatsAK
  | EwmaAK {alpha :: Double}
  deriving (Generic, Eq, Show, FromJSON, ToJSON, Read)

```

1.6.11 Cardano.BM.Data.Backend

Accepts a Logging of outcomes with LogObject

Instances of this type class accept a LogObject and deal with it.

```

class IsEffectuator t a where
  effectuate :: t a → LogObject a → IO ()
  effectuatefrom :: forall s o (IsEffectuator s a) ⇒ t a → LogObject a → s a → IO ()
  default effectuatefrom :: forall s o (IsEffectuator s a) ⇒ t a → LogObject a → s a → IO ()
  effectuatefrom t nli _ = effectuate t nli
  handleOverflow :: t a → IO ()

```

Declaration of a Backend

A backend is life-cycle managed, thus can be realized and unrealized.

```

class IsEffectuator t a ⇒ IsBackend t a where
  typeOf :: t a → BackendKind
  realize :: Configuration → IO (t a)
  realizefrom :: forall s o (IsEffectuator s a) ⇒ Configuration → Trace IO a → s a → IO (t a)
  default realizefrom :: forall s o (IsEffectuator s a) ⇒ Configuration → Trace IO a → s a → IO (t a)
  realizefrom cfg _ = realize cfg
  unrealize :: t a → IO ()

```

Backend

This data structure for a backend defines its behaviour as an **IsEffectuator** when processing an incoming message, and as an **IsBackend** for unrealizing the backend.

```
data Backend a = MkBackend
  { bEffectuate :: LogObject a → IO ()
  , bUnrealize :: IO ()
  }
```

1.6.12 Cardano.BM.Data.BackendKind

BackendKind

This identifies the backends that can be attached to the **Switchboard**.

```
data BackendKind =
  AggregationBK
  | EKGViewBK
  | KatipBK
  | MonitoringBK
  | SwitchboardBK
  deriving (Generic, Eq, Ord, Show, ToJSON, FromJSON, Read)
```

1.6.13 Cardano.BM.Data.Configuration

Data structure to help parsing configuration files.

Representation

```
type Port = Int
data Representation = Representation
  { minSeverity    :: Severity
  , rotation      :: Maybe RotationParameters
  , setupScribes  :: [ScribeDefinition]
  , defaultScribes :: [(ScribeKind, Text)]
  , setupBackends :: [BackendKind]
  , defaultBackends :: [BackendKind]
  , hasEKG        :: Maybe Port
  , hasGUI        :: Maybe Port
  , options       :: HM.HashMap Text Object
  }
  deriving (Generic, Show, ToJSON, FromJSON)
```

parseRepresentation

```
parseRepresentation :: FilePath → IO Representation
parseRepresentation fp = do
  repr :: Representation ← decodeFileThrow fp
  return $ implicit_fill_representation repr
```

after parsing the configuration representation we implicitly correct it.

```

implicit_fill_representation :: Representation → Representation
implicit_fill_representation =
  remove_ekgview_if_not_defined ∘
  filter_duplicates_from_backends ∘
  filter_duplicates_from_scribes ∘
  union_setup_and_usage_backends ∘
  add_ekgview_if_port_defined ∘
  add_katip_if_any_scribes
where
  filter_duplicates_from_backends r =
    r {setupBackends = mkUniq $ setupBackends r}
  filter_duplicates_from_scribes r =
    r {setupScribes = mkUniq $ setupScribes r}
  union_setup_and_usage_backends r =
    r {setupBackends = setupBackends r <> defaultBackends r}
# ifdef ENABLE_EKG
  remove_ekgview_if_not_defined r =
    case hasEKG r of
      Nothing → r {defaultBackends = filter (λbk → bk ≠ EKGViewBK) (defaultBackends r)
                  , setupBackends = filter (λbk → bk ≠ EKGViewBK) (setupBackends r)
                  }
      Just _ → r
  add_ekgview_if_port_defined r =
    case hasEKG r of
      Nothing → r
      Just _ → r {setupBackends = setupBackends r <> [EKGViewBK]}
# else
  remove_ekgview_if_not_defined = id
  add_ekgview_if_port_defined = id
# endif
  add_katip_if_any_scribes r =
    if (any ¬ [null $ setupScribes r, null $ defaultScribes r])
    then r {setupBackends = setupBackends r <> [KatipBK]}
    else r
mkUniq :: Ord a ⇒ [a] → [a]
mkUniq = Set.toList ∘ Set.fromList

```

1.6.14 Cardano.BM.Data.Counter

Counter

```

data Counter = Counter
  { cType :: CounterType
  , cName :: Text
  , cValue :: Measurable
  }
  deriving (Eq, Show, Generic, ToJSON)

data CounterType = MonotonicClockTime
  | MemoryCounter

```

```

| StatInfo
| IOCounter
| NetCounter
| CpuCounter
| RTSSStats
  deriving (Eq, Show, Generic, ToJSON)
instance ToJSON Microsecond where
  toJSON = toJSON ◦ toMicroseconds
  toEncoding = toEncoding ◦ toMicroseconds

```

Names of counters

```

nameCounter :: Counter → Text
nameCounter (Counter MonotonicClockTime _) = "Time-interval"
nameCounter (Counter MemoryCounter _) = "Mem"
nameCounter (Counter StatInfo _) = "Stat"
nameCounter (Counter IOCounter _) = "IO"
nameCounter (Counter NetCounter _) = "Net"
nameCounter (Counter CpuCounter _) = "Cpu"
nameCounter (Counter RTSSStats _) = "RTS"

```

CounterState

```

data CounterState = CounterState {
  csIdentifier :: Unique
  , csCounters :: [Counter]
}
  deriving (Generic, ToJSON)
instance ToJSON Unique where
  toJSON = toJSON ◦ hashUnique
  toEncoding = toEncoding ◦ hashUnique
instance Show CounterState where
  show cs = (show ◦ hashUnique) (csIdentifier cs)
    <> " => " <> (show $ csCounters cs)

```

Difference between counters

```

diffCounters :: [Counter] → [Counter] → [Counter]
diffCounters openings closings =
  getCountersDiff openings closings
where
  getCountersDiff :: [Counter]
    → [Counter]
    → [Counter]
  getCountersDiff as bs =
    let
      getName counter = nameCounter counter <> cName counter

```



```

asNames = map getName as
aPairs = zip asNames as
bsNames = map getName bs
bs' = zip bsNames bs
bPairs = HM.fromList bs'
in
catMaybes $ (flip map) aPairs $ \ (name, Counter _ _ startValue) →
  case HM.lookup name bPairs of
    Nothing → Nothing
    Just counter → let endValue = cValue counter
                    in Just counter {cValue = endValue - startValue}

```

1.6.15 Cardano.BM.Data.LogItem

LoggerName

A **LoggerName** has currently type *Text*.

```
type LoggerName = Text
```

Logging of outcomes with LogObject

```

data LogObject a = LogObject
  { loName :: LoggerName
  , loMeta :: !LOMeta
  , loContent :: (LOContent a)
  }
deriving (Generic, Show, ToJSON)

```

Meta data for a **LogObject**. Text was selected over ThreadId in order to be able to use the logging system under SimM of ourboros-network because ThreadId from Control.Concurrent lacks a Read instance.

```

data LOMeta = LOMeta {
  tstamp :: {-# UNPACK #-} !UTCTime
  , tid   :: {-# UNPACK #-} !Text
  , severity :: !Severity
  , privacy :: !PrivacyAnnotation
  }
deriving (Show)
instance ToJSON LOMeta where
  toJSON (LOMeta _tstamp _tid _sev _priv) =
    object [ "tstamp" . = _tstamp
            , "tid"     . = show _tid
            , "severity" . = show _sev
            , "privacy" . = show _priv
            ]
mkLOMeta :: Severity → PrivacyAnnotation → IO LOMeta
mkLOMeta sev priv =
  LOMeta <$> getCurrentTime

```

```

<*> (pack ◦ show <$> myThreadId)
<*> pure sev
<*> pure priv

```

Payload of a **LogObject**:

```

data LOContent a = LogMessage a
  | LogValue Text Measurable
  | ObserveOpen CounterState
  | ObserveDiff CounterState
  | ObserveClose CounterState
  | AggregatedMessage [(Text, Aggregated)]
  | MonitoringEffect (LogObject a)
  | KillPill
  deriving (Generic, Show, ToJSON)

data PrivacyAnnotation =
  Confidential -- confidential information - handle with care
  | Public -- indifferent - can be public.
  deriving (Show, Generic, ToJSON, FromJSON)

```

Data structure for annotating the severity and privacy of an object.

```

data PrivacyAndSeverityAnnotated a
  = PSA
  { psaSeverity :: !Severity
  , psaPrivacy :: !PrivacyAnnotation
  , psaPayload :: a
  }
  deriving (Show)

```

1.6.16 Cardano.BM.Data.Observable

ObservableInstance

```

data ObservableInstance = MonotonicClock
  | MemoryStats
  | ProcessStats
  | NetStats
  | IOStats
  | GhcRtsStats
  deriving (Generic, Eq, Show, FromJSON, ToJSON, Read)

```

1.6.17 Cardano.BM.Data.Output

OutputKind

```

data OutputKind a = TVarList (STM.TVar [LogObject a])
  deriving (Eq)

```

ScribeKind

This identifies katip's scribes by type.

```

data ScribeKind = FileTextSK
  | FileJsonSK
  | StdoutSK
  | StderrSK
  deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)

```

ScribeId

A scribe is identified by **ScribeKind** \times *Filename*

```

type ScribeId = Text -- (ScribeKind :: Filename)

```

ScribePrivacy

This declares if a scribe will be public (and must not contain sensitive data) or private.

```

data ScribePrivacy = ScPublic | ScPrivate
  deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)

```

ScribeDefinition

This identifies katip's scribes by type.

```

data ScribeDefinition = ScribeDefinition
  { scKind    :: ScribeKind
  , scName    :: Text
  , scPrivacy :: ScribePrivacy
  , scRotation :: Maybe RotationParameters
  }
  deriving (Generic, Eq, Ord, Show, ToJSON)
instance FromJSON ScribeDefinition where
  parseJSON (Object o) = do
    kind    ← o .: "scKind"
    name    ← o .: "scName"
    mayPrivacy ← o .:? "scPrivacy"
    rotation ← o .:? "scRotation"
    return $ ScribeDefinition
      { scKind    = kind
      , scName    = name
      , scPrivacy = fromMaybe ScPublic mayPrivacy
      , scRotation = rotation
      }
  parseJSON invalid = typeMismatch "ScribeDefinition" invalid

```

1.6.18 Cardano.BM.Data.Rotation

RotationParameters

```
data RotationParameters = RotationParameters
  { rpLogLimitBytes :: !Word64 -- max size of file in bytes
  , rpMaxAgeHours   :: !Word   -- hours
  , rpKeepFilesNum  :: !Word   -- number of files to keep
  } deriving (Generic, Show, Eq, Ord, FromJSON, ToJSON)
```

1.6.19 Cardano.BM.Data.Severity

Severity

The intended meaning of severity codes:

Debug *detailed information about values and decision flow* **Info** general information of events; progressing properly **Notice** *needs attention; something – progressing properly* **Warning** may continue into an error condition if continued **Error** *unexpected set of event or condition occurred* **Critical** error condition causing degrade of operation **Alert** *a subsystem is no longer operating correctly, likely requires man* at this point, the system can never progress without additional intervention

We were informed by the Syslog taxonomy: https://en.wikipedia.org/wiki/Syslog#Severity_level

```
data Severity = Debug
  | Info
  | Notice
  | Warning
  | Error
  | Critical
  | Alert
  | Emergency
  deriving (Show, Eq, Ord, Enum, Generic, ToJSON, Read)

instance FromJSON Severity where
  parseJSON = withText "severity" $ \case
    "Debug"    → pure Debug
    "Info"     → pure Info
    "Notice"   → pure Notice
    "Warning"  → pure Warning
    "Error"    → pure Error
    "Critical" → pure Critical
    "Alert"    → pure Alert
    "Emergency" → pure Emergency
    _          → pure Info -- catch all
```

1.6.20 Cardano.BM.Data.SubTrace

SubTrace

```
data NameSelector = Exact Text | StartsWith Text | EndsWith Text | Contains Text
  deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)
data DropName     = Drop NameSelector
  deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)
```

```

data UnhideNames = Unhide [NameSelector]
                        deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)

data SubTrace = Neutral
  | UntimedTrace
  | NoTrace
  | TeeTrace LoggerName
  | FilterTrace [(DropName, UnhideNames)]
  | DropOpening
  | ObservableTrace [ObservableInstance]
  deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)

```

1.6.21 Cardano.BM.Data.Trace

Trace

A **Trace** *m a* is a **Tracer** *m* (**LogObject** *a*).

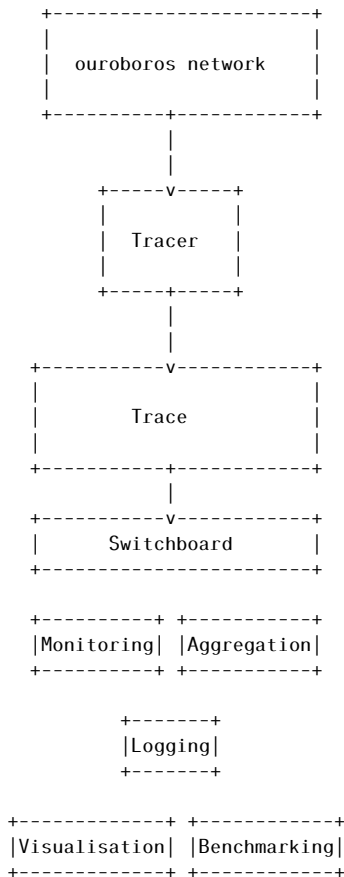
```

type Trace m a = Tracer m (LogObject a)

```

1.6.22 Cardano.BM.Data.Tracer

This module extends the basic **Tracer** with one that keeps a list of context names to create the basis for **Trace** which accepts messages from a Tracer and ends in the **Switchboard** for further processing of the messages.



LogNamed

A **LogNamed** contains of a context name and some log item.

```
data LogNamed item = LogNamed
  { lnName :: LoggerName
  , lnItem :: item
  } deriving (Show)
```

```
renderNamedItemTracing :: Show a => Tracer m String -> Tracer m (LogNamed a)
renderNamedItemTracing = contramap $ \item ->
  unpack (lnName item) ++ ": " ++ show (lnItem item)
renderNamedItemTracing' :: Show a => Tracer m String -> Tracer m (LogObject a)
renderNamedItemTracing' = contramap $ \item ->
  unpack (loName item) ++ ": " ++ show (loContent item) ++ ", (meta): " ++ show (loMeta item)
```

```
named :: Tracer m (LogNamed a) -> Tracer m a
named = contramap (LogNamed mempty)
```

Add a new name to the logging context

```
appendNamed :: LoggerName -> Tracer m (LogNamed a) -> Tracer m (LogNamed a)
appendNamed name = contramap $ (\(LogNamed oldName item) ->
  LogNamed (name <> "." <> oldName) item)
```

Add a new name to the logging context

```
appendNamed' :: LoggerName -> Tracer m (LogObject a) -> Tracer m (LogObject a)
appendNamed' name = contramap $ (\(LogObject oldName meta item) ->
  if oldName == ""
  then LogObject name meta item
  else LogObject (name <> "." <> oldName) meta item)
```

The function `toLogObject` can be specialized for various environments

```
class Monad m => ToLogObject m where
  toLogObject :: Tracer m (LogObject a) -> Tracer m a
instance ToLogObject IO where
  toLogObject :: Tracer IO (LogObject a) -> Tracer IO a
  toLogObject tr = Tracer $ Op $ \a -> do
    lo <- LogObject <$> pure ""
    <*> (mkLOMeta Debug Public)
    <*> pure (LogMessage a)
    tracingWith tr lo
```

To be placed in `ouroboros-network` ○

```
instance (MonadFork m, MonadTimer m) => ToLogObject m where
  toLogObject (Tracer (Op tr)) = Tracer $ Op $ \a -> do
    lo <- LogObject <$> pure ""
    <*> (LOMeta <$> getMonotonicTime-- must be evaluated at the calling site
    <*> (pack o show <$> myThreadId)
```

```

    < * > pure Debug
    < * > pure Public)
  < * > pure (LogMessage a)
tr lo

tracingNamed :: Show a => Tracer IO (LogObject a) -> Tracer IO a
tracingNamed = toLogObject
example2 :: IO ()
example2 = do
  let logTrace = appendNamed' "example2" (renderNamedItemTracing' stdoutTracer)
  void $ callFun2 logTrace
callFun2 :: Tracer IO (LogObject Text) -> IO Int
callFun2 logTrace = do
  let logTrace' = appendNamed' "fun2" logTrace
  tracingWith (tracingNamed logTrace') "in function 2"
  callFun3 logTrace'
callFun3 :: Tracer IO (LogObject Text) -> IO Int
callFun3 logTrace = do
  tracingWith (tracingNamed (appendNamed' "fun3" logTrace)) "in function 3"
  return 42

```

A **Tracer** transformer creating a **LogObject** from *PrivacyAndSeverityAnnotated*.

```

logObjectFromAnnotated :: Show a
=> Tracer IO (LogObject a)
-> Tracer IO (PrivacyAndSeverityAnnotated a)
logObjectFromAnnotated (Tracer (Op tr)) = Tracer $ Op $ \ (PSA sev priv a) -> do
  lometa <- mkLOMeta sev priv
  tr $ LogObject "" lometa (LogMessage a)

example3 :: IO ()
example3 = do
  let logTrace =
    logObjectFromAnnotated $ appendNamed' "example3" $ renderNamedItemTracing' stdoutTracer
  tracingWith logTrace $ PSA Info Confidential ("Hello" :: String)
  tracingWith logTrace $ PSA Warning Public "World"

filterAppendNameTracing :: Monad m
=> m (LogNamed a -> Bool)
-> LoggerName
-> Tracer m (LogNamed a)
-> Tracer m (LogNamed a)
filterAppendNameTracing test name = (appendNamed name) o (condTracingM test)
example4 :: IO ()
example4 = do
  let appendF = filterAppendNameTracing oracle
  logTrace = appendF "example4" (renderNamedItemTracing stdoutTracer)
  tracingWith (named logTrace) ("Hello" :: String)
  let logTrace' = appendF "inner" logTrace

```

```

tracingWith (named logTrace') "World"
let logTrace' = appendF "innest" logTrace'
tracingWith (named logTrace'') "!!"
where
  oracle :: Monad m => m (LogNamed a -> Bool)
  oracle = return $ ((≠) "example4.inner.") ∘ lnName

-- severity anotated
example5 :: IO ()
example5 = do
  let logTrace =
    condTracingM oracle $
      logObjectFromAnnotated $
        appendNamed' "test5" $ renderNamedItemTracing' stdoutTracer
  tracingWith logTrace $ PSA Debug Confidential ("Hello" :: String)
  tracingWith logTrace $ PSA Warning Public "World"
where
  oracle :: Monad m => m (PrivacyAndSeverityAnnotated a -> Bool)
  oracle = return $ λ(PSA sev _priv _) -> (sev > Debug)

-- test for combined name and severity
example6 :: IO ()
example6 = do
  let logTrace0 = -- the basis, will output using the local renderer to stdout
    appendNamed' "test6" $ renderNamedItemTracing' stdoutTracer
  logTrace1 = -- the trace from Privacy...Annotated to LogObject
    condTracingM oracleSev $ logObjectFromAnnotated $ logTrace0
  logTrace2 =
    appendNamed' "row" $ condTracingM oracleName $ logTrace0
  logTrace3 = -- oracle should eliminate messages from this trace
    appendNamed' "raw" $ condTracingM oracleName $ logTrace0
  tracingWith logTrace1 $ PSA Debug Confidential ("Hello" :: String)
  tracingWith logTrace1 $ PSA Warning Public "World"
  lometa ← mkLOMeta Info Public
  tracingWith logTrace2 $ LogObject "" lometa (LogMessage " ", Row!)
  tracingWith logTrace3 $ LogObject "" lometa (LogMessage " ", Row!)
where
  oracleSev :: Monad m => m (PrivacyAndSeverityAnnotated a -> Bool)
  oracleSev = return $ λ(PSA sev _priv _) -> (sev > Debug)
  oracleName :: Monad m => m (LogObject a -> Bool)
  oracleName = return $ λ(LogObject name _ _) -> (name ≡ "row") -- we only see the names from

```

1.6.23 Cardano.BM.Configuration

see `Cardano.BM.Configuration.Model` for the implementation.

```

getOptionOrDefault :: CM.Configuration -> Text -> Text -> IO (Text)
getOptionOrDefault cg name def = do
  opt ← CM.getOption cg name
  case opt of

```


Nothing → return *def*

Just o → return *o*

1.6.24 Cardano.BM.Configuration.Model

Configuration.Model

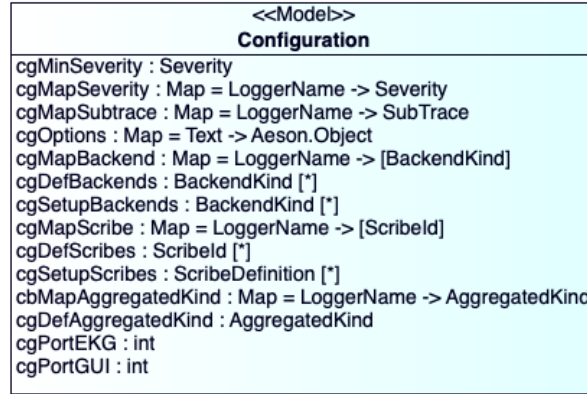


Figure 1.5: Configuration model

```

type ConfigurationMVar = MVar ConfigurationInternal
newtype Configuration = Configuration
  {getCG :: ConfigurationMVar}
-- Our internal state; see -"Configuration model"-
data ConfigurationInternal = ConfigurationInternal
  {cgMinSeverity      :: Severity
  -- minimum severity level of every object that will be output
  ,cgMapSeverity      :: HM.HashMap LoggerName Severity
  -- severity filter per loggernaem
  ,cgMapSubtrace      :: HM.HashMap LoggerName SubTrace
  -- type of trace per loggernaem
  ,cgOptions          :: HM.HashMap Text Object
  -- options needed for tracing, logging and monitoring
  ,cgMapBackend       :: HM.HashMap LoggerName [BackendKind]
  -- backends that will be used for the specific loggernaem
  ,cgDefBackendKs     :: [BackendKind]
  -- backends that will be used if a set of backends for the
  -- specific loggernaem is not set
  ,cgSetupBackends    :: [BackendKind]
  -- backends to setup; every backend to be used must have
  -- been declared here
  ,cgMapScribe        :: HM.HashMap LoggerName [ScribeId]
  -- katip scribes that will be used for the specific loggernaem
  ,cgMapScribeCache   :: HM.HashMap LoggerName [ScribeId]
  -- map to cache info of the cgMapScribe
  ,cgDefScribes       :: [ScribeId]
  -- katip scribes that will be used if a set of scribes for the

```

```

-- specific loggername is not set
,cgSetupScribes    :: [ScribeDefinition]
-- katip scribes to setup; every scribe to be used must have
-- been declared here
,cgMapAggregatedKind :: HM.HashMap LoggerName AggregatedKind
-- kind of Aggregated that will be used for the specific loggername
,cgDefAggregatedKind :: AggregatedKind
-- kind of Aggregated that will be used if a set of scribes for the
-- specific loggername is not set
,cgMonitors        :: HM.HashMap LoggerName (MEvExpr, [MEvAction])
,cgPortEKG          :: Int
-- port for EKG server
,cgPortGUI          :: Int
-- port for changes at runtime (NOT IMPLEMENTED YET)
} deriving (Show, Eq)

```

Backends configured in the Switchboard

For a given context name return the list of backends configured, or, in case no such configuration exists, return the default backends.

```

getBackends :: Configuration → LoggerName → IO [BackendKind]
getBackends configuration name = do
  cg ← readMVar $ getCG configuration
  let outs = HM.lookup name (cgMapBackend cg)
  case outs of
    Nothing → return (cgDefBackendKs cg)
    Just os → return os

getDefaultBackends :: Configuration → IO [BackendKind]
getDefaultBackends configuration =
  cgDefBackendKs < $ > (readMVar $ getCG configuration)

setDefaultBackends :: Configuration → [BackendKind] → IO ()
setDefaultBackends configuration bes =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgDefBackendKs = bes}

setBackends :: Configuration → LoggerName → Maybe [BackendKind] → IO ()
setBackends configuration name be =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgMapBackend = HM.alter (\_ → be) name (cgMapBackend cg)}

```

Backends to be setup by the Switchboard

Defines the list of Backends that need to be setup by the Switchboard.

```

setSetupBackends :: Configuration → [BackendKind] → IO ()
setSetupBackends configuration bes =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgSetupBackends = bes}

getSetupBackends :: Configuration → IO [BackendKind]
getSetupBackends configuration =
  cgSetupBackends < $ > (readMVar $ getCG configuration)

```

Scribes configured in the **Log** backend

For a given context name return the list of scribes to output to, or, in case no such configuration exists, return the default scribes to use.

```

getScribes :: Configuration → LoggerName → IO [ScribeId]
getScribes configuration name = do
  cg ← readMVar (getCG configuration)
  (updateCache, scribes) ← do
    let defs = cgDefScribes cg
    let mapscribes = cgMapScribe cg
    let find_s lname = case HM.lookup lname mapscribes of
      Nothing →
        case dropToDot lname of
          Nothing → defs
          Just lname' → find_s lname'
    Just os → os
  let outs = HM.lookup name (cgMapScribeCache cg)
  -- look if scribes are already cached
  return $ case outs of
    -- if no cached scribes found; search the appropriate scribes that
    -- they must inherit and update the cached map
    Nothing → (True, find_s name)
    Just os → (False, os)

  when updateCache $ setCachedScribes configuration name $ Just scribes
  return scribes

dropToDot :: Text → Maybe Text
dropToDot ts = dropToDot' (breakOnEnd "." ts)
  where
    dropToDot' (_, "") = Nothing
    dropToDot' (name', _) = Just $ dropWhileEnd (≡ '.' ) name'

getCachedScribes :: Configuration → LoggerName → IO (Maybe [ScribeId])
getCachedScribes configuration name = do
  cg ← readMVar $ getCG configuration
  return $ HM.lookup name $ cgMapScribeCache cg

setScribes :: Configuration → LoggerName → Maybe [ScribeId] → IO ()
setScribes configuration name scribes =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgMapScribe = HM.alter (\_ → scribes) name (cgMapScribe cg)}

setCachedScribes :: Configuration → LoggerName → Maybe [ScribeId] → IO ()
setCachedScribes configuration name scribes =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgMapScribeCache = HM.alter (\_ → scribes) name (cgMapScribeCache cg)}

setDefaultScribes :: Configuration → [ScribeId] → IO ()
setDefaultScribes configuration scs =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgDefScribes = scs}

```

Scribes to be setup in the **Log** backend

Defines the list of Scribes that need to be setup in the **Log** backend.

```

setSetupScribes :: Configuration → [ScribeDefinition] → IO ()
setSetupScribes configuration sds =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgSetupScribes = sds}
getSetupScribes :: Configuration → IO [ScribeDefinition]
getSetupScribes configuration =
  cgSetupScribes < $ > readMVar (getCG configuration)

```

AggregatedKind to define the type of measurement

For a given context name return its **AggregatedKind** or in case no such configuration exists, return the default **AggregatedKind** to use.

```

getAggregatedKind :: Configuration → LoggerName → IO AggregatedKind
getAggregatedKind configuration name = do
  cg ← readMVar $ getCG configuration
  let outs = HM.lookup name (cgMapAggregatedKind cg)
  case outs of
    Nothing → return $ cgDefAggregatedKind cg
    Just os → return $ os

setDefaultAggregatedKind :: Configuration → AggregatedKind → IO ()
setDefaultAggregatedKind configuration defAK =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgDefAggregatedKind = defAK}

setAggregatedKind :: Configuration → LoggerName → Maybe AggregatedKind → IO ()
setAggregatedKind configuration name ak =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgMapAggregatedKind = HM.alter (\_ → ak) name (cgMapAggregatedKind cg)}

```

Access port numbers of EKG, GUI

```

getEKGport :: Configuration → IO Int
getEKGport configuration =
  cgPortEKG < $ > (readMVar $ getCG configuration)

setEKGport :: Configuration → Int → IO ()
setEKGport configuration port =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgPortEKG = port}

getGUIport :: Configuration → IO Int
getGUIport configuration =
  cgPortGUI < $ > (readMVar $ getCG configuration)

setGUIport :: Configuration → Int → IO ()
setGUIport configuration port =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgPortGUI = port}

```

Options

```
getOption :: Configuration → Text → IO (Maybe Text)
getOption configuration name = do
  cg ← readMVar $ getCG configuration
  case HM.lookup name (cgOptions cg) of
    Nothing → return Nothing
    Just o → return $ Just $ pack $ show o
```

Global setting of minimum severity

```
minSeverity :: Configuration → IO Severity
minSeverity configuration =
  cgMinSeverity < $ > (readMVar $ getCG configuration)
setMinSeverity :: Configuration → Severity → IO ()
setMinSeverity configuration sev =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgMinSeverity = sev}
```

Relation of context name to minimum severity

```
inspectSeverity :: Configuration → Text → IO (Maybe Severity)
inspectSeverity configuration name = do
  cg ← readMVar $ getCG configuration
  return $ HM.lookup name (cgMapSeverity cg)
setSeverity :: Configuration → Text → Maybe Severity → IO ()
setSeverity configuration name sev =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgMapSeverity = HM.alter (\_ → sev) name (cgMapSeverity cg)}
```

Relation of context name to SubTrace

A new context may contain a different type of **Trace**. The function **appendName** (Enter new named context) will look up the **SubTrace** for the context's name.

```
findSubTrace :: Configuration → Text → IO (Maybe SubTrace)
findSubTrace configuration name = do
  cg ← readMVar $ getCG configuration
  return $ HM.lookup name (cgMapSubtrace cg)
setSubTrace :: Configuration → Text → Maybe SubTrace → IO ()
setSubTrace configuration name trafo =
  modifyMVar_ (getCG configuration) $ \cg →
    return cg {cgMapSubtrace = HM.alter (\_ → trafo) name (cgMapSubtrace cg)}
```

Monitors

```
Just (
  fromList [
    ("chain.creation.block", Array [
      Object (fromList [ ("monitor", String "((time > (23 s)) Or (time < (17 s)))")),
      Object (fromList [ ("actions", Array [
        String "AlterMinSeverity \"chain.creation\" Debug" ]))]
    ), ("#aggregation.critproc.observable", Array [
      Object (fromList [ ("monitor", String "(mean >= (42))")),
      Object (fromList [ ("actions", Array [
        String "CreateMessage \"exceeded\" \"the observable has been too long too high!\",
        String "AlterGlobalMinSeverity Info" ]))]
    ])]
```

```
getMonitors :: Configuration → IO (HM.HashMap LoggerName (MEvExpr, [MEvAction]))
getMonitors configuration = do
  cg ← readMVar $ getCG configuration
  return (cgMonitors cg)
```

Parse configuration from file

Parse the configuration into an internal representation first. Then, fill in **Configuration** after refinement.

```
setup :: FilePath → IO Configuration
setup fp = do
  r ← R.parseRepresentation fp
  setupFromRepresentation r

parseMonitors :: Maybe (HM.HashMap Text Value) → HM.HashMap LoggerName (MEvExpr, [MEvAction])
parseMonitors Nothing = HM.empty
parseMonitors (Just hmv) = HM.mapMaybe mkMonitor hmv
  where
    mkMonitor (Array a) =
      if Vector.length a == 2
      then do
        e ← mkExpression $ aVector.!0
        as ← mkActions $ aVector.!1
        return (e, as)
      else Nothing
    mkMonitor _ = Nothing
    mkExpression :: Value → Maybe MEvExpr
    mkExpression (Object o1) =
      case HM.lookup "monitor" o1 of
        Nothing → Nothing
        Just (String expr) → MEv.parseMaybe expr
        Just _ → Nothing
    mkExpression _ = Nothing
    mkActions :: Value → Maybe [MEvAction]
    mkActions (Object o2) =
      case HM.lookup "actions" o2 of
```

```

    Nothing → Nothing
    Just (Array as) → Just $ map (λ(String s) → s) $ Vector.toList as
    Just _ → Nothing

mkActions _ = Nothing

setupFromRepresentation :: R.Representation → IO Configuration
setupFromRepresentation r = do
    let mapseverities0 = HM.lookup "mapSeverity" (R.options r)
        mapbackends = HM.lookup "mapBackends" (R.options r)
        mapsubtrace = HM.lookup "mapSubtrace" (R.options r)
        mapscribes0 = HM.lookup "mapScribes" (R.options r)
        mapaggregatedkinds = HM.lookup "mapAggregatedkinds" (R.options r)
        mapmonitors = HM.lookup "mapMonitors" (R.options r)
        mapseverities = parseSeverityMap mapseverities0
        mapscribes = parseScribeMap mapscribes0
    cgreg ← newMVar $ ConfigurationInternal
        { cgMinSeverity = R.minSeverity r
        , cgMapSeverity = mapseverities
        , cgMapSubtrace = parseSubtraceMap mapsubtrace
        , cgOptions = R.options r
        , cgMapBackend = parseBackendMap mapbackends
        , cgDefBackendKs = R.defaultBackends r
        , cgSetupBackends = R.setupBackends r
        , cgMapScribe = mapscribes
        , cgMapScribeCache = mapscribes
        , cgDefScribes = r_defaultScribes r
        , cgSetupScribes = fillRotationParams (R.rotation r) (R.setupScribes r)
        , cgMapAggregatedKind = parseAggregatedKindMap mapaggregatedkinds
        , cgDefAggregatedKind = StatsAK
        , cgMonitors = parseMonitors mapmonitors
        , cgPortEKG = r_hasEKG r
        , cgPortGUI = r_hasGUI r
        }
    return $ Configuration cgreg
where
    parseSeverityMap :: Maybe (HM.HashMap Text Value) → HM.HashMap Text Severity
    parseSeverityMap Nothing = HM.empty
    parseSeverityMap (Just hmv) = HM.mapMaybe mkSeverity hmv
    where
        mkSeverity (String s) = Just (read (unpack s) :: Severity)
        mkSeverity _ = Nothing

    fillRotationParams :: Maybe RotationParameters → [ScribeDefinition] → [ScribeDefinition]
    fillRotationParams defaultRotation = map $ λsd →
        if (scKind sd ≠ StdoutSK) ∧ (scKind sd ≠ StderrSK)
        then
            sd { scRotation = maybe defaultRotation Just (scRotation sd) }
        else
            -- stdout and stderr cannot be rotated
            sd { scRotation = Nothing }

    parseBackendMap Nothing = HM.empty
    parseBackendMap (Just hmv) = HM.map mkBackends hmv

```



```

where
  mkBackends (Array bes) = catMaybes $ map mkBackend $ Vector.toList bes
  mkBackends _ = [ ]
  mkBackend (String s) = Just (read (unpack s) :: BackendKind)
  mkBackend _ = Nothing

parseScribeMap Nothing = HM.empty
parseScribeMap (Just hmv) = HM.map mkScribes hmv

where
  mkScribes (Array scs) = catMaybes $ map mkScribe $ Vector.toList scs
  mkScribes (String s) = [ (s :: ScribeId) ]
  mkScribes _ = [ ]
  mkScribe (String s) = Just (s :: ScribeId)
  mkScribe _ = Nothing

parseSubtraceMap :: Maybe (HM.HashMap Text Value) → HM.HashMap Text SubTrace
parseSubtraceMap Nothing = HM.empty
parseSubtraceMap (Just hmv) = HM.mapMaybe mkSubtrace hmv

where
  mkSubtrace (String s) = Just (read (unpack s) :: SubTrace)
  mkSubtrace (Object hm) = mkSubtrace' (HM.lookup "tag" hm) (HM.lookup "contents" hm)
  mkSubtrace _ = Nothing
  mkSubtrace' Nothing _ = Nothing
  mkSubtrace' _ Nothing = Nothing
  mkSubtrace' (Just (String tag)) (Just (Array cs)) =
    if tag ≡ "ObservableTrace"
    then Just $ ObservableTrace $ map (λ(String s) → (read (unpack s) :: ObservableInstance)) $ Vec
    else Nothing
  mkSubtrace' _ _ = Nothing

r_hasEKG repr = case (R.hasEKG repr) of
  Nothing → 0
  Just p → p
r_hasGUI repr = case (R.hasGUI repr) of
  Nothing → 0
  Just p → p
r_defaultScribes repr = map (λ(k,n) → pack (show k) <> " :: " <> n) (R.defaultScribes repr)
parseAggregatedKindMap Nothing = HM.empty
parseAggregatedKindMap (Just hmv) =
  let
    listv = HM.toList hmv
    mapAggregatedKind = HM.fromList $ catMaybes $ map mkAggregatedKind listv
  in
    mapAggregatedKind

where
  mkAggregatedKind (name, String s) = Just (name, read (unpack s) :: AggregatedKind)
  mkAggregatedKind _ = Nothing

```

Setup empty configuration

```

empty :: IO Configuration
empty = do

```



```

cgreg ← newMVar $ ConfigurationInternal
  {cgMinSeverity      = Debug
  ,cgMapSeverity      = HM.empty
  ,cgMapSubtrace      = HM.empty
  ,cgOptions          = HM.empty
  ,cgMapBackend       = HM.empty
  ,cgDefBackendKs     = []
  ,cgSetupBackends    = []
  ,cgMapScribe        = HM.empty
  ,cgMapScribeCache   = HM.empty
  ,cgDefScribes       = []
  ,cgSetupScribes     = []
  ,cgMapAggregatedKind = HM.empty
  ,cgDefAggregatedKind = StatsAK
  ,cgMonitors         = HM.empty
  ,cgPortEKG          = 0
  ,cgPortGUI          = 0
  }
return $ Configuration cgreg

```

1.6.25 Cardano.BM.Configuration.Static

Default configuration outputting on *stdout*

```

defaultConfigStdout :: IO CM.Configuration
defaultConfigStdout = do
  c ← CM.empty
  CM.setMinSeverity c Debug
  CM.setSetupBackends c [KatipBK]
  CM.setDefaultBackends c [KatipBK]
  CM.setSetupScribes c [ScribeDefinition {
    scName = "stdout"
  ,scKind = StdoutSK
  ,scPrivacy = ScPublic
  ,scRotation = Nothing
  }
  ]
  CM.setDefaultScribes c ["StdoutSK::stdout"]
  return c

```

Default configuration for testing

```

defaultConfigTesting :: IO CM.Configuration
defaultConfigTesting = do
  c ← CM.empty
  CM.setMinSeverity c Debug
  # ifdef ENABLE_AGGREGATION
  CM.setSetupBackends c [KatipBK, AggregationBK]
  CM.setDefaultBackends c [KatipBK, AggregationBK]
  # else

```

```

CM.setSetupBackends c [KatipBK]
CM.setDefaultBackends c [KatipBK]
# endif
CM.setSetupScribes c [ScribeDefinition {
  scName = "stdout"
  ,scKind = StdoutSK
  ,scPrivacy = ScPublic
  ,scRotation = Nothing
}]
CM.setDefaultScribes c ["StdoutSK::stdout"]
return c

```

1.6.26 Cardano.BM.Configuration.Editor

This simple configuration editor is accessible through a browser on <http://127.0.0.1:13789>, or whatever port has been set in the configuration.

A number of maps that relate logging context name to behaviour can be changed. And, most importantly, the global minimum severity that defines the filtering of log messages.

links

The GUI is built on top of *Threepenny-GUI* (<http://hackage.haskell.org/package/threepenny-gui>). The appearance is due to *w3-css* (<https://www.w3schools.com/w3css>).

```

startup :: Configuration → IO ()
startup config = do
  port ← getGUIport config
  when (port > 0) $ do
    thd ← Async.async $
      startGUI defaultConfig {jsPort = Just port
                             ,jsAddr   = Just "127.0.0.1"
                             ,jsStatic = Just "iohk-monitoring/static"
                             ,jsCustomHTML = Just "configuration-editor.html"
                             } $ prepare config
    Async.link thd

```

```

data Cmd = Backends | Scribes | Severities | SubTrace | Aggregation
  deriving (Enum, Eq, Show, Read)

```

```

prepare :: Configuration → Window → UI ()
prepare config window = void $ do
  let commands = [Backends..Aggregation]
  inputKey ← UI.input #. "w3-input w3-border" # set UI.size "34"
  inputValue ← UI.input #. "w3-input w3-border" # set UI.size "60"
  outputMsg ← UI.input #. "w3-input w3-border"
  currentCmd ← UI.p #. "current-cmd"
  let performActionOnId anId action =
      getElementById window anId >>= λcase
        Nothing      → return ()
        Just anElement → action anElement

```

```

let turn    anElement toState    = void $ element anElement # set UI.enabled toState
let setValueOf anElement aValue = void $ element anElement # set UI.value aValue
let setClasses classes anElement = void $ element anElement # set UI.class_ classes
let setError m = setValueOf outputMsg ("ERROR: " ++ m)
let setMessage m = setValueOf outputMsg m
let enable anElement = turn anElement True
let disable anElement = turn anElement False
let clean    anElement = setValueOf anElement ""
let cleanAndDisable anElement = clean anElement >> disable anElement
let rememberCurrent cmd = setValueOf currentCmd $ show cmd
let removeItem Backends k = CM.setBackends  config k Nothing
    removeItem Severities k = CM.setSeverity  config k Nothing
    removeItem Scribes    k = CM.setScribes    config k Nothing
    removeItem SubTrace   k = CM.setSubTrace   config k Nothing
    removeItem Aggregation k = CM.setAggregatedKind config k Nothing
let updateItem Backends k v = case (readMay v :: Maybe [BackendKind]) of
    Nothing → setError "parse error on backend list"
    Just v' → liftIO $ CM.setBackends config k $ Just v'
updateItem Severities k v = case (readMay v :: Maybe Severity) of
    Nothing → setError "parse error on severity"
    Just v' → liftIO $ CM.setSeverity config k $ Just v'
updateItem Scribes k v = case (readMay v :: Maybe [ScribeId]) of
    Nothing → setError "parse error on scribe list"
    Just v' → liftIO $ CM.setScribes config k $ Just v'
updateItem SubTrace k v = case (readMay v :: Maybe SubTrace) of
    Nothing → setError "parse error on subtrace"
    Just v' → liftIO $ CM.setSubTrace config k $ Just v'
updateItem Aggregation k v = case (readMay v :: Maybe AggregatedKind) of
    Nothing → setError "parse error on aggregated kind"
    Just v' → liftIO $ CM.setAggregatedKind config k $ Just v'

disable inputKey
disable inputValue
disable outputMsg

let saveItemButtonId      = "save-item-button"
let cancelSaveItemButtonId = "cancel-save-item-button"
let addItemButtonId      = "add-item-button"
let outputTableId        = "output-table"
let saveItemButton       = performActionOnId saveItemButtonId
let cancelSaveItemButton = performActionOnId cancelSaveItemButtonId
let mkTableRow :: Show t ⇒ Cmd → LoggerName → t → UI Element
mkTableRow cmd n v = UI.tr #. "itemrow" #+
    [ UI.td #+ [string (unpack n)]
    , UI.td #+ [string (show v)]
    , UI.td #+
        [ do
            b ← UI.button #. "w3-small w3-btn w3-ripple w3-orange edit-item-button"
              #+ [ UI.bold #+ [string "Edit"] ]
            on UI.click b $ const $ do
                saveItemButton enable

```



```

switchToTab c@Aggregation = showCorrespondingItems c CM.cgMapAggregatedKind
let mkEditInputs =
  row [element inputKey
    ,UI.span#. "key-value-separator" #+ [string ":"]
    ,element inputValue
    ,UI.span#. "key-value-separator" #+ [string ""]
    ,do
      b ← UI.button#. "w3-btn w3-ripple w3-green save-item-button"
        #set (UI.attr "id") addItemButtonId
        #set UI.enabled False
        #+ [UI.bold #+ [string "New" ]]
      on UI.click b $ const $ do
        enable inputKey
        enable inputValue
        saveItemButton enable
        cancelSaveItemButton enable
      return b
    ,UI.span#. "key-value-separator" #+ [string ""]
    ,do
      b ← UI.button#. "w3-btn w3-ripple w3-lime save-item-button"
        #set (UI.attr "id") saveItemButtonId
        #set UI.enabled False
        #+ [UI.bold #+ [string "Save" ]]
      on UI.click b $ const $ do
        k ← inputKey # get UI.value
        v ← inputValue # get UI.value
        m ← currentCmd # get UI.value
        case (readMay m :: Maybe Cmd) of
          Nothing → setError "parse error on cmd"
          Just c → do
            cleanAndDisable inputKey
            cleanAndDisable inputValue
            saveItemButton disable
            cancelSaveItemButton disable
            setMessage$ "Setting '" ++ k ++ "' to '" ++ v ++ "' in " ++ m
            updateItem c (pack k) v
            switchToTab c
        return b
    ,UI.span#. "key-value-separator" #+ [string ""]
    ,do
      b ← UI.button#. "w3-btn w3-ripple w3-white"
        #set (UI.attr "id") cancelSaveItemButtonId
        #set UI.enabled False
        #+ [UI.bold #+ [string "Cancel" ]]
      on UI.click b $ const $ do
        cleanAndDisable inputKey
        cleanAndDisable inputValue
        saveItemButton disable
        cancelSaveItemButton disable
      return b
    ]

```

```

let minimumSeveritySelection = do
  confMinSev ← liftIO $ minSeverity config
  let setMinSev _el Nothing = pure ()
      setMinSev _el (Just sev) = liftIO $
        setMinSeverity config (toEnum sev :: Severity)
  mkSevOption sev = UI.option # set UI.text (show sev)
    # set UI.value (show sev)
    # if (confMinSev == sev) then set UI.selected True else id
  minsev ← UI.select #. "minsevfield" #+
    map mkSevOption (enumFrom Debug)
  on UI.selectionChange minsev $ setMinSev minsev
  row [string "Set minimum severity to:"
    , UI.span # set html "&nbsp;"
    , UI.span #. "severity-dropdown big" #+ [element minsev]
    ]
let commandTabs =
  row $ flip map commands $ \cmd → do
    b ← UI.button #. "w3-bar-item w3-button w3-grey"
      # set (UI.attr "id") (show cmd)
      #+ [UI.bold #+ [string (show cmd)]]
    on UI.click b $ const $ do
      cleanAndDisable inputKey
      cleanAndDisable inputValue
      clean outputMsg
      switchToTab cmd
    return b
getElementById window "main-section" >>= \case
  Nothing → pure ()
  Just mainSection → void $ element mainSection #+
    [UI.div #. "w3-panel" #+
      [UI.div #. "w3-border w3-border-dark-grey" #+
        [UI.div #. "w3-panel" #+ [minimumSeveritySelection]]
      , UI.div #. "w3-panel" #+ []
      , UI.div #. "w3-border w3-border-dark-grey" #+
        [UI.div #. "w3-bar w3-grey" #+ [commandTabs]
        , UI.div #. "w3-panel" #+ [mkEditInputs]
        , UI.div #. "w3-panel" #+ [element outputMsg]
        ]
      ]
    ]
]

```

1.6.27 Cardano.BM.Output.Switchboard

Switchboard

We are using an *MVar* because we spawn a set of backends that may try to send messages to the switchboard before it is completely setup.

```

type SwitchboardMVar a = MVar (SwitchboardInternal a)
newtype Switchboard a = Switchboard

```

```

    {getSB :: SwitchboardMVar a}
data SwitchboardInternal a = SwitchboardInternal
    {sbQueue :: TBQ.TBQueue (LogObject a)
    ,sbDispatch :: Async.Async ()
    }

```

Trace that forwards to the Switchboard

Every **Trace** ends in the **Switchboard** which then takes care of dispatching the messages to the selected backends.

This **Tracer** will forward all messages unconditionally to the **Switchboard**. (currently disabled)

```

mainTrace :: Switchboard a → Tracer IO (LogObject a)
mainTrace sb = Tracer.Tracer $ Op $ effectuate sb

```

This function will apply to every message the severity filter as defined in the **Configuration**.

```

mainTraceConditionally :: Configuration → Switchboard a → Tracer IO (LogObject a)
mainTraceConditionally config sb = Tracer.Tracer $ Op $ \item@(LogObject loggename meta _) → do
    passSevFilter ← testSeverity loggename meta
    passSubTrace ← testSubTrace loggename item
    if passSevFilter ∧ passSubTrace
    then effectuate sb item
    else return ()
where
    testSeverity :: LoggerName → LOMeta → IO Bool
    testSeverity loggename meta = do
        globminsev ← liftIO $ Config.minSeverity config
        globnamesev ← liftIO $ Config.inspectSeverity config loggename
        let minsev = max globminsev $ fromMaybe Debug globnamesev
        return $ (severity meta) ≥ minsev
    testSubTrace :: LoggerName → LogObject a → IO Bool
    testSubTrace loggename lo = do
        subtrace ← fromMaybe Neutral < $ > Config.findSubTrace config loggename
        return $ testSubTrace' lo subtrace
    testSubTrace' :: LogObject a → SubTrace → Bool
    testSubTrace' _ NoTrace = False
    testSubTrace' (LogObject _ _ (ObserveOpen _)) DropOpening = False
    testSubTrace' (LogObject loname _ (LogValue vname _)) (FilterTrace filters) = evalFilters filters (loname vname)
    testSubTrace' (LogObject loname _ _) (FilterTrace filters) = evalFilters filters loname
    testSubTrace' _ _ = True -- fallback: all pass

```

Process incoming messages

Incoming messages are put into the queue, and then processed by the dispatcher. The switchboard will never block when processing incoming messages ("eager receiver").

The queue is initialized and the message dispatcher launched.

```

instance IsEffectuator Switchboard a where
    effectuate switchboard item = do
        let writequeue :: TBQ.TBQueue (LogObject a) → LogObject a → IO ()

```



```

writequeue q i = do
  nocapacity ← atomically $ TBQ.isFullTBQueue q
  if nocapacity
    then handleOverflow switchboard
    else atomically $ TBQ.writeTBQueue q i
sb ← readMVar (getSB switchboard)
writequeue (sbQueue sb) item
handleOverflow _ = TIO.hPutStrLn stderr "Error: Switchboard's queue full, dropping log item"

instead of 'writequeue ...':
  evalMonitoringAction config item >>=
    mapM_ (writequeue (sbQueue sb))
evalMonitoringAction :: Configuration → LogObject a → m [LogObject a]
evalMonitoringAction c item = return [item]
-- let action = LogObject loName=(loName item) <> ".action", loContent=LogMessage .
-- return (action : item)

```

Switchboard implements Backend functions

Switchboard is an Declaration of a Backend

```

instance (Show a, ToJSON a) ⇒ IsBackend Switchboard a where
  typeof _ = SwitchboardBK
  realize cfg =
    let spawnDispatcher
      :: (Show a)
      ⇒ [(BackendKind, Backend a)]
      → TBQ.TBQueue (LogObject a)
      → IO (Async.Async ())
    spawnDispatcher backends queue = do
      now ← getCurrentTime
      let messageCounters = resetCounters now
      countersMVar ← newMVar messageCounters
      let traceInQueue q =
          Tracer.Tracer $ Op $ \lognamed → do
            nocapacity ← atomically $ TBQ.isFullTBQueue q
            if nocapacity
              then putStrLn "Error: Switchboard's queue full, dropping log items!"
              else atomically $ TBQ.writeTBQueue q lognamed
      _timer ← Async.async $ sendAndResetAfter
        (traceInQueue queue)
        "#messagecounters.switchboard"
      countersMVar
      60000 -- 60000 ms = 1 min
      Warning -- Debug
    let sendMessage nli befilter = do
        selectedBackends ← getBackends cfg (loName nli)
        let selBEs = befilter selectedBackends
        forM_ backends $ \ (bek, be) →
          when (bek ∈ selBEs) (bEffectuate be nli)

```



```

qProc counters = do
  -- read complete queue at once and process items
  nlis ← atomically $ do
    r ← TBQ.flushTBQueue queue
    when (null r) retry
    return r
  let processItem nli@(LogObject loname _ loitem) = do
    when (loname ≠ "#messagecounters.switchboard") $
      modifyMVar_ counters $
        λcnt → return $ updateMessageCounters cnt nli
  Config.findSubTrace cfg loname ≫ λcase
    Just (TeeTrace secName) →
      atomically $ TBQ.writeTBQueue queue $ nli {loName = secName}
    _ → return ()
  case loitem of
    KillPill → do
      forM_ backends (λ(., be) → bUnrealize be)
      return False
    (AggregatedMessage _) → do
      sendMessage nli (filter (≠ AggregationBK))
      return True
    (MonitoringEffect inner) → do
      sendMessage (inner {loName = loname}) (filter (≠ MonitoringBK))
      return True
    _ → do
      sendMessage nli id
      return True
  res ← mapM processItem nlis
  when (and res) $ qProc counters
  Async.async $ qProc countersMVar
in do
  q ← atomically $ TBQ.newTBQueue 2048
  sbref ← newEmptyMVar
  let sb :: Switchboard a = Switchboard sbref
  backends ← getSetupBackends cfg
  bs ← setupBackends backends cfg sb []
  dispatcher ← spawnDispatcher bs q
  -- link the given Async to the current thread, such that if the Async
  -- raises an exception, that exception will be re-thrown in the current
  -- thread, wrapped in ExceptionInLinkedThread.
  Async.link dispatcher
  putMVar sbref $ SwitchboardInternal {sbQueue = q, sbDispatch = dispatcher}
  return sb
unrealize switchboard = do
  let clearMVar :: MVar some → IO ()
    clearMVar = void ∘ tryTakeMVar
  (dispatcher, queue) ← withMVar (getSB switchboard) (λsb → return (sbDispatch sb, sbQueue sb))
  -- send terminating item to the queue
  lo ← LogObject <$> pure "kill.switchboard"

```

```

    < * > (mkLOMeta Warning Confidential)
    < * > pure KillPill
    atomically $ TBQ.writeTBQueue queue lo
    -- wait for the dispatcher to exit
    res ← Async.waitCatch dispatcher
    either throwM return res
    (clearMVar ◦ getSB) switchboard

```

Realizing the backends according to configuration

```

setupBackends :: (Show a, ToJSON a)
    ⇒ [ BackendKind ]
    → Configuration
    → Switchboard a
    → [(BackendKind, Backend a)]
    → IO [(BackendKind, Backend a)]
setupBackends [ ] _ _ acc = return acc
setupBackends (bk : bes) c sb acc = do
    be' ← setupBackend' bk c sb
    setupBackends bes c sb ((bk, be') : acc)
setupBackend' :: (Show a, ToJSON a) ⇒ BackendKind → Configuration → Switchboard a → IO (Backend a)
setupBackend' SwitchboardBK _ _ = error "cannot instantiate a further Switchboard"
# ifdef ENABLE_MONITORING
setupBackend' MonitoringBK c sb = do
    let trace = mainTraceConditionally c sb
    be :: Cardano.BM.Output ◦ Monitoring.Monitor a ← Cardano.BM.Output ◦ Monitoring.realizefrom c trace s
    return MkBackend
        { bEffectuate = Cardano.BM.Output ◦ Monitoring.effectuate be
        , bUnrealize = Cardano.BM.Output ◦ Monitoring.unrealize be
        }
# else
    -- We need it anyway, to avoid "Non-exhaustive patterns" warning.
setupBackend' MonitoringBK _ _ =
    TIO.hPutStrLn stderr "disabled! will not setup backend 'Monitoring'"
# endif
# ifdef ENABLE_EKG
setupBackend' EKGViewBK c sb = do
    let trace = mainTraceConditionally c sb
    be :: Cardano.BM.Output ◦ EKGView.EKGView a ← Cardano.BM.Output ◦ EKGView.realizefrom c trace s
    return MkBackend
        { bEffectuate = Cardano.BM.Output ◦ EKGView.effectuate be
        , bUnrealize = Cardano.BM.Output ◦ EKGView.unrealize be
        }
# else
    -- We need it anyway, to avoid "Non-exhaustive patterns" warning.
setupBackend' EKGViewBK _ _ =
    TIO.hPutStrLn stderr "disabled! will not setup backend 'EKGView'"
# endif
# ifdef ENABLE_AGGREGATION
setupBackend' AggregationBK c sb = do

```

```

let trace = mainTraceConditionally c sb
be :: Cardano.BM.Output ◦ Aggregation.Aggregation a ← Cardano.BM.Output ◦ Aggregation.realizefrom
return MkBackend
  { bEffectuate = Cardano.BM.Output ◦ Aggregation.effectuate be
  , bUnrealize = Cardano.BM.Output ◦ Aggregation.unrealize be
  }
# else
-- We need it anyway, to avoid "Non-exhaustive patterns" warning.
setupBackend' AggregationBK _ =
  TIO.hPutStrLn stderr "disabled! will not setup backend 'Aggregation'"
# endif
setupBackend' KatipBK c _ = do
  be :: Cardano.BM.Output ◦ Log.Log a ← Cardano.BM.Output ◦ Log.realize c
  return MkBackend
    { bEffectuate = Cardano.BM.Output ◦ Log.effectuate be
    , bUnrealize = Cardano.BM.Output ◦ Log.unrealize be
    }

```

1.6.28 Cardano.BM.Output.Log

Internal representation

```

type LogMVar = MVar LogInternal
newtype Log a = Log
  { getK :: LogMVar }
data LogInternal = LogInternal
  { kLogEnv      :: K.LogEnv
  , msgCounters :: MessageCounter
  , configuration :: Config.Configuration }

```

Log implements effectuate

```

instance (ToJSON a, Show a) ⇒ IsEffectuator Log a where
  effectuate katip item = do
    let logMVar = getK katip
    c ← configuration < $ > readMVar logMVar
    setupScribes ← getSetupScribes c
    selscribes ← getScribes c (loName item)
    let selscribesFiltered =
      case item of
        LogObject _ (LOMeta _ _ Confidential) (LogMessage _)
          → removePublicScribes setupScribes selscribes
        _ → selscribes
    forM_ selscribesFiltered $ λsc → passN sc katip item
    -- increase the counter for the specific severity and message type
    modifyMVar_ logMVar $ λli → return $
      li { msgCounters = updateMessageCounters (msgCounters li) item }
    -- reset message counters afer 60 sec = 1 min
    resetMessageCounters logMVar 60 Warning selscribesFiltered

```

```

where
  removePublicScribes allScribes = filter $ \sc →
    let (_, nameD) = T.breakOn ":" sc
    -- drop ":" from the start of name
    name = T.drop 2 nameD
  in
    case find (λx → scName x ≡ name) allScribes of
      Nothing → False
      Just scribe → scPrivacy scribe ≡ ScPrivate
resetMessageCounters logMVar interval sev scribes = do
  counters ← msgCounters < $ > readMVar logMVar
  let start = mcStart counters
  now = case item of
    LogObject _ meta _ → tstamp meta
  diffTime = round $ diffUTCTime now start
  when (diffTime > interval) $ do
    countersObjects ← forM (HM.toList $ mcCountersMap counters) $ λ(key, count) →
      LogObject
        < $ > pure "#messagecounters.katip"
        < * > (mkLOMeta sev Confidential)
        < * > pure (LogValue (pack key) (PureI $ toInteger count))
    intervalObject ←
      LogObject
        < $ > pure "#messagecounters.katip"
        < * > (mkLOMeta sev Confidential)
        < * > pure (LogValue "time_interval_(s)" (PureI diffTime))
  let namedCounters = countersObjects ++ [intervalObject]
  forM_ scribes $ \sc →
    forM_ namedCounters $ λnamedCounter →
      passN sc katip namedCounter
    modifyMVar_ logMVar $ λli → return $
      li {msgCounters = resetCounters now}
  handleOverflow _ = TIO.hPutStrLn stderr "Notice: Katip's queue full, dropping log items!"

```

Log implements backend functions

```

instance (ToJSON a, Show a) ⇒ IsBackend Log a where
  typeof _ = KatipBK
  realize config = do
    let updateEnv :: K.LogEnv → IO UTCTime → K.LogEnv
    updateEnv le timer =
      le {K._logEnvTimer = timer, K._logEnvHost = "hostname"}
    register :: [ScribeDefinition] → K.LogEnv → IO K.LogEnv
    register [] le = return le
    register (defsc : dscs) le = do
      let kind = scKind defsc
      name = scName defsc
      rotParams = scRotation defsc
      name' = pack (show kind) <> ":" <> name
      scr ← createScribe kind name rotParams

```

```

    register dscs ≡ K.registerScribe name' scr scribeSettings le
mockVersion :: Version
mockVersion = Version [0,1,0,0] []
scribeSettings :: KC.ScribeSettings
scribeSettings =
    let bufferSize = 5000 -- size of the queue (in log items)
    in
        KC.ScribeSettings bufferSize
createScribe FileTextSK name rotParams = mkTextFileScribe
    rotParams
    (FileDescription $ unpack name)
    False
createScribe FileJsonSK name rotParams = mkJsonFileScribe
    rotParams
    (FileDescription $ unpack name)
    False
createScribe StdoutSK _ _ = mkStdoutScribe
createScribe StderrSK _ _ = mkStderrScribe
cfoKey ← Config.getOptionOrDefault config (pack "cfokey") (pack "<unknown>")
le0 ← K.initLogEnv
    (K.Namespace [ "iohk" ])
    (fromString $ (unpack cfoKey) <> ":" <> showVersion mockVersion)
-- request a new time 'getCurrentTime' at most 100 times a second
timer ← mkAutoUpdate defaultUpdateSettings {updateAction = getCurrentTime, updateFreq = 10000}
let le1 = updateEnv le0 timer
scribes ← getSetupScribes config
le ← register scribes le1
messageCounters ← resetCounters <$ > getCurrentTime
kref ← newMVar $ LogInternal le messageCounters config
return $ Log kref
unrealize katip = do
    le ← withMVar (getK katip) $ λk → return (kLogEnv k)
    void $ K.closeScribes le

```

```

example :: IO ()
example = do
    config ← Config.setup "from_some_path.yaml"
    k ← setup config
    meta ← mkLOMeta Info Public
    passN (pack (show StdoutSK)) k $ LogObject
        { loName = "test"
        , loMeta = meta
        , loContent = LogMessage "Hello!"
        }
    meta' ← mkLOMeta Info Public
    passN (pack (show StdoutSK)) k $ LogObject
        { loName = "test"
        , loMeta = meta'
        , loContent = LogValue "cpu-no" 1
        }

```

Needed instances for *katip*:

```

deriving instance ToJSON a ⇒ K.ToObject (LogObject a)
deriving instance K.ToObject Text
deriving instance ToJSON a ⇒ K.ToObject (Maybe (LOContent a))
instance ToJSON a ⇒ KC.LogItem (LogObject a) where
    payloadKeys _ = KC.AllKeys
instance KC.LogItem Text where
    payloadKeys _ = KC.AllKeys
instance ToJSON a ⇒ KC.LogItem (Maybe (LOContent a)) where
    payloadKeys _ = KC.AllKeys

```

Log.passN

The following function copies the **LogObject** to the queues of all scribes that match on their name. Compare start of name of scribe to (*show backend* <> "::*name*"). This function is non-blocking.

```

passN :: (ToJSON a, Show a) ⇒ ScribeId → Log a → LogObject a → IO ()
passN backend katip (LogObject lname lmeta loitem) = do
    env ← kLogEnv < $ > readMVar (getK katip)
    forM_ (Map.toList $ K._logEnvScribes env) $
        λ(scName, (KC.ScribeHandle _ shChan)) →
            -- check start of name to match ScribeKind
            if backend 'isPrefixOf' scName
            then do
                let (sev, msg, payload) = case loitem of
                    (LogMessage logItem) →
                        (severity lmeta, pack $ show logItem, Nothing)
                    (ObserveDiff _) →
                        let text = TL.toStrict (encodeToLazyText loitem)
                        in
                        (severity lmeta, text, Just loitem)
                    (ObserveOpen _) →
                        let text = TL.toStrict (encodeToLazyText loitem)
                        in
                        (severity lmeta, text, Just loitem)
                    (ObserveClose _) →
                        let text = TL.toStrict (encodeToLazyText loitem)
                        in
                        (severity lmeta, text, Just loitem)
                    (AggregatedMessage aggregated) →
                        let text = T.concat $ (flip map) aggregated $ λ(name, agg) →
                            "\n" <> name <> " : " <> pack (show agg)
                        in
                        (severity lmeta, text, Nothing)
                    (LogValue name value) →
                        (severity lmeta, name <> " = " <> pack (showSI value), Nothing)
                (MonitoringEffect logitem) →
                    let text = TL.toStrict (encodeToLazyText logitem)
                    in

```

```

        (severity lometa, text, Just loitem)
    KillPill →
        (severity lometa, "Kill pill received!", Nothing)
    if (msg == "") ^ (isNothing payload)
    then return ()
    else do
        let threadIdText = KC.ThreadIdText $ tid lometa
        let itemTime = tstamp lometa
        let itemKatip = K.Item {
            _itemApp      = env ^. KC.logEnvApp
            , _itemEnv     = env ^. KC.logEnvEnv
            , _itemSeverity = sev2klog sev
            , _itemThread  = threadIdText
            , _itemHost    = env ^. KC.logEnvHost
            , _itemProcess = env ^. KC.logEnvPid
            , _itemPayload = payload
            , _itemMessage = K.logStr msg
            , _itemTime    = itemTime
            , _itemNamespace = (env ^. KC.logEnvApp) <> (K.Namespace [loname])
            , _itemLoc     = Nothing
        }
        void $ atomically $ KC.tryWriteTBQueue shChan (KC.NewItem itemKatip)
    else return ()

```

Scribes

```

mkStdoutScribe :: IO K.Scribe
mkStdoutScribe = do
    -- duplicate stdout so that Katip's closing
    -- action will not close the real stdout
    stdout' ← hDuplicate stdout
    mkTextFileScribeH stdout' True

mkStderrScribe :: IO K.Scribe
mkStderrScribe = do
    -- duplicate stderr so that Katip's closing
    -- action will not close the real stderr
    stderr' ← hDuplicate stderr
    mkTextFileScribeH stderr' True

mkTextFileScribeH :: Handle → Bool → IO K.Scribe
mkTextFileScribeH handler color = do
    mkFileScribeH handler formatter color
    where
        formatter h colorize verbosity item =
            TIO.hPutStrLn h $! toLazyText $ formatItem colorize verbosity item

mkFileScribeH
    :: Handle
    → (forall a ◦ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ())
    → Bool
    → IO K.Scribe
mkFileScribeH h formatter colorize = do

```



```

hSetBuffering h LineBuffering
locklocal ← newMVar ()
let logger :: forall a ◦ K.LogItem a ⇒ K.Item a → IO ()
    logger item = withMVar locklocal $ \_ →
        formatter h colorize K.V0 item
pure $ K.Scribe logger (hClose h)

mkTextFileScribe :: Maybe RotationParameters → FileDescription → Bool → IO K.Scribe
mkTextFileScribe rotParams fdesc colorize = do
    mkFileScribe rotParams fdesc formatter colorize
where
    formatter :: Handle → Bool → K.Verbosity → K.Item a → IO Int
    formatter hdl colorize' v' item =
        case KC._itemMessage item of
            K.LogStr "" →
                -- if message is empty do not output it
                return 0
            _ → do
                let tmsg = toLazyText $ formatItem colorize' v' item
                TIO.hPutStrLn hdl tmsg
                return $ fromIntegral $ TL.length tmsg

mkJsonFileScribe :: Maybe RotationParameters → FileDescription → Bool → IO K.Scribe
mkJsonFileScribe rotParams fdesc colorize = do
    mkFileScribe rotParams fdesc formatter colorize
where
    formatter :: (K.LogItem a) ⇒ Handle → Bool → K.Verbosity → K.Item a → IO Int
    formatter h _ verbosity item = do
        let jmsg = case KC._itemMessage item of
            -- if a message is contained in item then only the
            -- message is printed and not the data
            K.LogStr "" → K.itemJson verbosity item
            K.LogStr msg → K.itemJson verbosity $
                item { KC._itemMessage = K.logStr (" " :: Text)
                    , KC._itemPayload = TL.toStrict $ toLazyText msg
                    -- do we need the severity from meta?
                    }
        tmsg = encodeToLazyText jmsg
        TIO.hPutStrLn h tmsg
        return $ fromIntegral $ TL.length tmsg

mkFileScribe
    :: Maybe RotationParameters
    → FileDescription
    → (forall a ◦ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO Int)
    → Bool
    → IO K.Scribe
mkFileScribe (Just rotParams) fdesc formatter colorize = do
    let prefixDir = prefixPath fdesc
    (createDirectoryIfMissing True prefixDir)
    'catchIO' (prtoutException ("cannot log prefix directory: " ++ prefixDir))
    let fpath = filePath fdesc
    trp ← initializeRotator rotParams fpath
    scribestate ← newMVar trp -- triple of (handle), (bytes remaining), (rotate time)

```



```

-- sporadically remove old log files - every 10 seconds
cleanup ← mkAutoUpdate defaultUpdateSettings {
    updateAction = cleanupRotator rotParams fpath
    ,updateFreq = 10000000
}
let finalizer :: IO ()
    finalizer = withMVar scribestate $
        λ(h,_,_) → hClose h
let logger :: forall a ◦ K.LogItem a ⇒ K.Item a → IO ()
    logger item =
        modifyMVar_ scribestate $ λ(h,bytes,rotime) → do
            byteswritten ← formatter h colorize K.V0 item
            -- remove old files
            cleanup
            -- detect log file rotation
            let bytes' = bytes - (toInteger $ byteswritten)
            let tdiff' = round $ diffUTCTime rotime (K._itemTime item)
            if bytes' < 0 ∨ tdiff' < (0 :: Integer)
            then do -- log file rotation
                hClose h
                (h2,bytes2,rotime2) ← evalRotator rotParams fpath
                return (h2,bytes2,rotime2)
            else
                return (h,bytes',rotime)
        return $ K.Scribe logger finalizer
-- log rotation disabled.
mkFileScribe Nothing fdesc formatter colorize = do
    let prefixDir = prefixPath fdesc
    (createDirectoryIfMissing True prefixDir)
    'catchIO' (prtoutException ("cannot log prefix directory: " ++ prefixDir))
    let fpath = filePath fdesc
    h ← catchIO (openFile fpath WriteMode) $
        λe → do
            prtoutException ("error while opening log: " ++ fpath) e
            -- fallback to standard output in case of exception
            return stdout
    hSetBuffering h LineBuffering
    scribestate ← newMVar h
    let finalizer :: IO ()
        finalizer = withMVar scribestate hClose
    let logger :: forall a ◦ K.LogItem a ⇒ K.Item a → IO ()
        logger item =
            withMVar scribestate $ λhandler →
                void $ formatter handler colorize K.V0 item
    return $ K.Scribe logger finalizer

formatItem :: Bool → K.Verbosity → K.Item a → Builder
formatItem withColor _verb K.Item {..} =
    fromText header <>
    fromText " " <>
    brackets (fromText timestamp) <>

```

```

fromText " " <>
KC.unLogStr _itemMessage
where
  header = colorBySeverity _itemSeverity $
    "[ " <> mconcat namedcontext <> ":" <> severity <> ":" <> threadid <> "]"
  namedcontext = KC.intercalateNs _itemNamespace
  severity = KC.renderSeverity _itemSeverity
  threadid = KC.getThreadIdText _itemThread
  timestamp = pack $ formatTime defaultTimeLocale tsformat _itemTime
  tsformat :: String
  tsformat = "%F %T%2Q %Z"
  colorBySeverity s m = case s of
    K.EmergencyS → red m
    K.AlertS     → red m
    K.CriticalS  → red m
    K.ErrorS     → red m
    K.NoticeS    → magenta m
    K.WarningS   → yellow m
    K.InfoS      → blue m
    _           → m
  red = colorize "31"
  yellow = colorize "33"
  magenta = colorize "35"
  blue = colorize "34"
  colorize c m
    | withColor = "\ESC[" <> c <> "m" <> m <> "\ESC[0m"
    | otherwise = m
-- translate Severity to Log.Severity
sev2klog :: Severity → K.Severity
sev2klog = λcase
  Debug    → K.DebugS
  Info     → K.InfoS
  Notice   → K.NoticeS
  Warning  → K.WarningS
  Error    → K.ErrorS
  Critical → K.CriticalS
  Alert    → K.AlertS
  Emergency → K.EmergencyS

data FileDescription = FileDescription {
  filePath :: !FilePath
  deriving (Show)
  prefixPath :: FileDescription → FilePath
  prefixPath = takeDirectory ∘ filePath

```

1.6.29 Cardano.BM.Output.EKGView

Structure of EKGView

```

type EKGViewMVar a = MVar (EKGViewInternal a)
newtype EKGView a = EKGView

```

```

    {getEV :: EKGViewMVar a}
data EKGViewInternal a = EKGViewInternal
    {evQueue :: TBQ.TBQueue (Maybe (LogObject a))
    ,evLabels :: EKGViewMap
    ,evServer :: Server
    }

```

Relation from variable name to label handler

We keep the label handlers for later update in a *HashMap*.

```

type EKGViewMap = HM.HashMap Text Label.Label

```

Internal Trace

This is an internal **Trace**, named "#ekgview", which can be used to control the messages that are being displayed by EKG.

```

ekgTrace :: Show a => EKGView a -> Configuration -> IO (Trace IO a)
ekgTrace ekg _c = do
    let trace = ekgTrace' ekg
    Trace.appendName "#ekgview" trace
where
    ekgTrace' :: Show a => EKGView a -> Tracer IO (LogObject a)
    ekgTrace' ekgview = Tracer $ Op $ \lo@ (LogObject lname _) -> do
        let setlabel :: Text -> Text -> EKGViewInternal a -> IO (Maybe (EKGViewInternal a))
        setlabel name label ekg_i@(EKGViewInternal _ labels server) =
            case HM.lookup name labels of
                Nothing -> do
                    ekghdl ← getLabel name server
                    Label.set ekghdl label
                    return $ Just $ ekg_i {evLabels = HM.insert name ekghdl labels}
                Just ekghdl -> do
                    Label.set ekghdl label
                    return Nothing
    update :: Show a => LogObject a -> EKGViewInternal a -> IO (Maybe (EKGViewInternal a))
    update (LogObject lname _ (LogMessage logitem)) ekg_i =
        setlabel lname (pack $ show logitem) ekg_i
    update (LogObject lname _ (LogValue lname value)) ekg_i =
        let lname' = lname <> "." <> lname
        in
        setlabel lname' (pack $ show value) ekg_i
    update _ _ = return Nothing
    modifyMVar_ (getEV ekgview) $ \ekgup -> do
        let -- strip off some prefixes not necessary for display
        lognam1 = case stripPrefix "#ekgview.#aggregation." lname of
            Nothing -> lname
            Just ln' -> ln'
        lname = case stripPrefix "#ekgview." lognam1 of
            Nothing -> lognam1

```

```

    Just ln' → ln'
  upd ← update lo {loName = logname} ekgup
  case upd of
    Nothing → return ekgup
    Just ekgup' → return ekgup'

```

EKG view is an effectuator

Function *effectuate* is called to pass in a **LogObject** for display in EKG. If the log item is an *AggregatedStats* message, then all its constituents are put into the queue. In case the queue is full, all new items are dropped.

```

instance IsEffectuator EKGView a where
  effectuate ekgview item = do
    ekg ← readMVar (getEV ekgview)
    let enqueue a = do
      nocapacity ← atomically $ TBQ.isFullTBQueue (evQueue ekg)
      if nocapacity
      then handleOverflow ekgview
      else atomically $ TBQ.writeTBQueue (evQueue ekg) (Just a)
    case item of
      (LogObject logname lometa (AggregatedMessage ags)) → liftIO $ do
        let traceAgg :: [(Text, Aggregated)] → IO ()
        traceAgg [] = return ()
        traceAgg ((n, AggregatedEWMA ewma) : r) = do
          enqueue $ LogObject (logname <> "." <> n) lometa (LogValue "avg" $ avg ewma)
          traceAgg r
        traceAgg ((n, AggregatedStats stats) : r) = do
          let statsname = logname <> "." <> n
          qbasestats s' nm = do
            enqueue $ LogObject nm lometa (LogValue "mean" (PureD $ meanOfStats s'))
            enqueue $ LogObject nm lometa (LogValue "min" $ fmin s')
            enqueue $ LogObject nm lometa (LogValue "max" $ fmax s')
            enqueue $ LogObject nm lometa (LogValue "count" $ PureI $ fromIntegral $ fcount s')
            enqueue $ LogObject nm lometa (LogValue "stdev" (PureD $ stdevOfStats s'))
          enqueue $ LogObject statsname lometa (LogValue "last" $ flast stats)
          qbasestats (fbasic stats) $ statsname <> ".basic"
          qbasestats (fdelta stats) $ statsname <> ".delta"
          qbasestats (ftimed stats) $ statsname <> ".timed"
          traceAgg r
        traceAgg ags
      (LogObject _ _ (LogMessage _)) → enqueue item
      (LogObject _ _ (LogValue _)) → enqueue item
      _ → return ()
  handleOverflow _ = TIO.hPutStrLn stderr "Notice: EKGViews's queue full, dropping log items"

```

EKGView implements Backend functions

EKGView is an IsBackend

```

instance Show a ⇒ IsBackend EKGView a where
  typeOf _ = EKGViewBK

```

```

realize _ = error "EKGView cannot be instantiated by 'realize'"
realizefrom config sbtrace _ = do
  evref ← newEmptyMVar
  let ekview = EKGView evref
  evport ← getEKGport config
  ehdl ← forkServer "127.0.0.1" evport
  ekghdl ← getLabel "iohk-monitoring version" ehdl
  Label.set ekghdl $ pack (showVersion version)
  ekgtrace ← ekgTrace ekview config
  queue ← atomically $ TBQ.newTBQueue 512
  dispatcher ← spawnDispatcher queue sbtrace ekgtrace
  -- link the given Async to the current thread, such that if the Async
  -- raises an exception, that exception will be re-thrown in the current
  -- thread, wrapped in ExceptionInLinkedThread.
  Async.link dispatcher
  putMVar evref $ EKGViewInternal
    { evLabels = HM.empty
    , evServer = ehdl
    , evQueue = queue
    }
  return ekview
unrealize ekview =
  withMVar (getEV ekview) $ \ekg →
    killThread $ serverThreadId $ evServer ekg

```

Asynchronously reading log items from the queue and their processing

```

spawnDispatcher :: (Show a)
  ⇒ TBQ.TBQueue (Maybe (LogObject a))
  → Trace.Trace IO a
  → Trace.Trace IO a
  → IO (Async.Async ())
spawnDispatcher evqueue sbtrace trace = do
  now ← getCurrentTime
  let messageCounters = resetCounters now
  countersMVar ← newMVar messageCounters
  _timer ← Async.async $ sendAndResetAfter
    sbtrace
    "#messagecounters.ekgview"
    countersMVar
    60000 -- 60000 ms = 1 min
    Warning -- Debug
  Async.async $ qProc countersMVar
where
  qProc counters = do
    maybeItem ← atomically $ TBQ.readTBQueue evqueue
    case maybeItem of
      Just obj@(LogObject logname meta content) → do
        trace' ← Trace.appendName logname trace
        Trace.traceNamedObject trace' (meta, content)

```

```

-- increase the counter for the type of message
modifyMVar_ counters $ \cnt → return $ updateMessageCounters cnt obj
qProc counters
Nothing → return ()-- stop here

```

Interactive testing **EKGView**

```

test :: IO ()
test = do
  c ← Cardano.BM.Setup.setupTrace (Left "test/config.yaml") "ekg"
  ev ← Cardano.BM.Output ◦ EKGView.realize c
  meta ← mkLOMeta Info Public
  effectuate ev $ LogObject "test.questions" meta (LogValue "answer" 42)
  effectuate ev $ LogObject "test.monitor023" meta (LogMessage (LogItem Public Warning "!!!! ALAR

```

1.6.30 Cardano.BM.Output.Aggregation

Internal representation

```

type AggregationMVar a = MVar (AggregationInternal a)
newtype Aggregation a = Aggregation
  {getAg :: AggregationMVar a}
data AggregationInternal a = AggregationInternal
  {agQueue :: TBQ.TBQueue (Maybe (LogObject a))
  ,agDispatch :: Async.Async ()
  }

```

Relation from context name to aggregated statistics

We keep the aggregated values (**Aggregated**) for a named context in a *HashMap*.

```

type AggregationMap = HM.HashMap Text AggregatedExpanded

```

Info for Aggregated operations

Apart from the **Aggregated** we keep some valuable info regarding to them; such as when was the last time it was sent.

```

type Timestamp = Word64
data AggregatedExpanded = AggregatedExpanded
  {aeAggregated :: !Aggregated
  ,aeResetAfter :: !(Maybe Int)
  ,aeLastSent :: {-# UNPACK #-} !Timestamp
  }

```

Aggregation implements effectuate

Aggregation is an **IsEffectuator** Enter the log item into the **Aggregation** queue.

```
instance IsEffectuator Aggregation a where
  effectuate agg item = do
    ag ← readMVar (getAg agg)
    nocapacity ← atomically $ TBQ.isFullTBQueue (agQueue ag)
    if nocapacity
    then handleOverflow agg
    else atomically $ TBQ.writeTBQueue (agQueue ag) $! Just item
  handleOverflow _ = TIO.hPutStrLn stderr "Notice: Aggregation's queue full, dropping log item"
```

Aggregation implements Backend functions

Aggregation is an **IsBackend**

```
instance Show a ⇒ IsBackend Aggregation a where
  typeOf _ = AggregationBK
  realize _ = error "Aggregation cannot be instantiated by 'realize'"
  realizeFrom config trace _ = do
    aggref ← newEmptyMVar
    aggregationQueue ← atomically $ TBQ.newTBQueue 2048
    dispatcher ← spawnDispatcher config HM.empty aggregationQueue trace
    -- link the given Async to the current thread, such that if the Async
    -- raises an exception, that exception will be re-thrown in the current
    -- thread, wrapped in ExceptionInLinkedThread.
    Async.link dispatcher
    putMVar aggref $ AggregationInternal aggregationQueue dispatcher
    return $ Aggregation aggref
  unrealize aggregation = do
    let clearMVar :: MVar a → IO ()
        clearMVar = void ∘ tryTakeMVar
    (dispatcher, queue) ← withMVar (getAg aggregation) (λag →
      return (agDispatch ag, agQueue ag))
    -- send terminating item to the queue
    atomically $ TBQ.writeTBQueue queue Nothing
    -- wait for the dispatcher to exit
    -- TODO add a timeout to waitCatch in order
    -- to be sure that it will finish
    res ← Async.waitCatch dispatcher
    either throwM return res
    (clearMVar ∘ getAg) aggregation
```

Asynchronously reading log items from the queue and their processing

```
spawnDispatcher :: (Show a)
  ⇒ Configuration
  → AggregationMap
  → TBQ.TBQueue (Maybe (LogObject a))
```

```

    → Trace.Trace IO a
    → IO (Async.Async ())
spawnDispatcher conf aggMap aggregationQueue trace0 = do
  now ← getCurrentTime
  trace ← Trace.appendName "#aggregation" trace0
  let messageCounters = resetCounters now
  countersMVar ← newMVar messageCounters
  _timer ← Async.async $ sendAndResetAfter
    trace0
    "#messagecounters.aggregation"
    countersMVar
  60000 -- 60000 ms = 1 min
  Warning-- Debug
  Async.async $ qProc trace countersMVar aggMap
where
  qProc trace counters aggregatedMap = do
    maybeItem ← atomically $ TBQ.readTBQueue aggregationQueue
    case maybeItem of
      Just (lo@(LogObject logname lm _)) → do
        (updatedMap,aggregations) ← update lo aggregatedMap
        unless (null aggregations) $
          sendAggregated trace (LogObject logname lm (AggregatedMessage aggregations))
        -- increase the counter for the specific severity and message type
        modifyMVar_ counters $ \cnt → return $ updateMessageCounters cnt lo
        qProc trace counters updatedMap
      Nothing → return ()
  createNupdate name value lme agmap = do
    case HM.lookup name agmap of
      Nothing → do
        -- if Aggregated does not exist; initialize it.
        aggregatedKind ← getAggregatedKind conf name
        case aggregatedKind of
          StatsAK → return $ singletonStats value
          EwmaAK aEWMA → do
            let initEWMA = EmptyEWMA aEWMA
            return $ AggregatedEWMA $ ewma initEWMA value
        Just a → return $ updateAggregation value (aeAggregated a) lme (aeResetAfter a)
  update :: LogObject a
    → AggregationMap
    → IO (AggregationMap, [(Text, Aggregated)])
  update (LogObject logname lme (LogValue iname value)) agmap = do
    let fullname = logname <> " ." <> iname
    aggregated ← createNupdate fullname value lme agmap
    now ← getMonotonicTimeNSec
    let aggregatedX = AggregatedExpanded {
      aeAggregated = aggregated
      ,aeResetAfter = Nothing
      ,aeLastSent = now
    }
    namedAggregated = [(iname, aeAggregated aggregatedX)]
    updatedMap = HM.alter (const $ Just $ aggregatedX) fullname agmap

```



```

    return (updatedMap, namedAggregated)
update (LogObject logname lme (ObserveDiff counterState)) agmap =
  updateCounters (csCounters counterState) lme (logname, "diff") agmap [ ]
update (LogObject logname lme (ObserveOpen counterState)) agmap =
  updateCounters (csCounters counterState) lme (logname, "open") agmap [ ]
update (LogObject logname lme (ObserveClose counterState)) agmap =
  updateCounters (csCounters counterState) lme (logname, "close") agmap [ ]
update (LogObject logname lme (LogMessage _)) agmap = do
  let iname = pack $ show (severity lme)
  let fullname = logname <> "." <> iname
  aggregated ← createNupdate fullname (PureI 0) lme agmap
  now ← getMonotonicTimeNSec
  let aggregatedX = AggregatedExpanded {
    aeAggregated = aggregated
    ,aeResetAfter = Nothing
    ,aeLastSent = now
  }
  namedAggregated = [(iname, aeAggregated aggregatedX)]
  updatedMap = HM.alter (const $ Just $ aggregatedX) fullname agmap
  return (updatedMap, namedAggregated)
-- everything else
update _ agmap = return (agmap, [ ])
updateCounters :: [Counter]
  → LOMeta
  → (LoggerName, LoggerName)
  → AggregationMap
  → [(Text, Aggregated)]
  → IO (AggregationMap, [(Text, Aggregated)])
updateCounters [ ] _ _ aggrMap aggs = return $ (aggrMap, aggs)
updateCounters (counter : cs) lme (logname, msgname) aggrMap aggs = do
  let name = cName counter
  subname = msgname <> "." <> (nameCounter counter) <> "." <> name
  fullname = logname <> "." <> subname
  value = cValue counter
  aggregated ← createNupdate fullname value lme aggrMap
  now ← getMonotonicTimeNSec
  let aggregatedX = AggregatedExpanded {
    aeAggregated = aggregated
    ,aeResetAfter = Nothing
    ,aeLastSent = now
  }
  namedAggregated = (subname, aggregated)
  updatedMap = HM.alter (const $ Just $ aggregatedX) fullname aggrMap
  updateCounters cs lme (logname, msgname) updatedMap (namedAggregated : aggs)
sendAggregated :: Trace.Trace IO a → LogObject a → IO ()
sendAggregated trace (LogObject logname meta v@(AggregatedMessage _)) = do
  -- enter the aggregated message into the Trace
  trace' ← Trace.appendName logname trace
  liftIO $ Trace.traceNamedObject trace' (meta, v)
-- ignore every other message

```

```
sendAggregated _ _ = return ()
```

Update aggregation

We distinguish an uninitialized from an already initialized aggregation. The latter is properly initialized.

We use Welford's online algorithm to update the estimation of mean and variance of the sample statistics. (see [https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance#Welford's](https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance#Welford's_algorithm)

```
updateAggregation :: Measurable → Aggregated → LOMeta → Maybe Int → Aggregated
```

```
updateAggregation v (AggregatedStats s) lme resetAfter =
```

```
  let count = fcount (fbasic s)
```

```
      reset = maybe False (count ≥) resetAfter
```

```
  in
```

```
  if reset
```

```
  then
```

```
    singletonStats v
```

```
  else
```

```
    AggregatedStats $! Stats {flast = v
```

```
      ,fold = mkTimestamp
```

```
      ,fbasic = updateBaseStats (count ≥ 1) v (fbasic s)
```

```
      ,fdelta = updateBaseStats (count ≥ 2) (v - flast s) (fdelta s)
```

```
      ,ftimed = updateBaseStats (count ≥ 2) (mkTimestamp - fold s) (ftimed s)
```

```
    }
```

```
  where
```

```
    mkTimestamp = utc2ns (tstamp lme)
```

```
    utc2ns (UTCTime days secs) =
```

```
      let yearsecs :: Rational
```

```
          yearsecs = 365 * 24 * 3600
```

```
          rdays, rsecs :: Rational
```

```
          rdays = toRational $ toModifiedJulianDay days
```

```
          rsecs = toRational secs
```

```
          s2ns = 10000000000
```

```
      in
```

```
        Nanoseconds $ round $ (fromRational $ s2ns * rsecs + rdays * yearsecs :: Double)
```

```
updateAggregation v (AggregatedEWMA e) _ _ = AggregatedEWMA $! ewma e v
```

```
updateBaseStats :: Bool → Measurable → BaseStats → BaseStats
```

```
updateBaseStats False _ s = s {fcount = fcount s + 1}
```

```
updateBaseStats True v s =
```

```
  let newcount = fcount s + 1
```

```
      newvalue = getDouble v
```

```
      delta = newvalue - fsum_A s
```

```
      dincr = (delta / fromIntegral newcount)
```

```
      delta2 = newvalue - fsum_A s - dincr
```

```
  in
```

```
    BaseStats {fmin = min (fmin s) v
```

```
      ,fmax = max v (fmax s)
```

```
      ,fcount = newcount
```

```
      ,fsum_A = fsum_A s + dincr
```

```
      ,fsum_B = fsum_B s + (delta * delta2)
```

```
    }
```

Calculation of EWMA

Following https://en.wikipedia.org/wiki/Moving_average#Exponential_moving_average we calculate the exponential moving average for a series of values Y_t according to:

$$S_t = \begin{cases} Y_1, & t = 1 \\ \alpha \cdot Y_t + (1 - \alpha) \cdot S_{t-1}, & t > 1 \end{cases}$$

The pattern matching below ensures that the **EWMA** will start with the first value passed in, and will not change type, once determined.

```
ewma :: EWMA → Measurable → EWMA
ewma (EmptyEWMA a) v = EWMA a v
ewma (EWMA a s@(Microseconds _)) y@(Microseconds _) =
    EWMA a $ Microseconds $ round $ a * (getDouble y) + (1 - a) * (getDouble s)
ewma (EWMA a s@(Seconds _)) y@(Seconds _) =
    EWMA a $ Seconds $ round $ a * (getDouble y) + (1 - a) * (getDouble s)
ewma (EWMA a s@(Bytes _)) y@(Bytes _) =
    EWMA a $ Bytes $ round $ a * (getDouble y) + (1 - a) * (getDouble s)
ewma (EWMA a (PureI s)) (PureI y) =
    EWMA a $ PureI $ round $ a * (fromInteger y) + (1 - a) * (fromInteger s)
ewma (EWMA a (PureD s)) (PureD y) =
    EWMA a $ PureD $ a * y + (1 - a) * s
ewma _ _ = error "Cannot average on values of different type"
```

1.6.31 Cardano.BM.Output.Monitoring

Structure of Monitoring

```
type MonitorMVar a = MVar (MonitorInternal a)
newtype Monitor a = Monitor
    { getMon :: MonitorMVar a }
data MonitorInternal a = MonitorInternal
    { monQueue :: TBQ.TBQueue (Maybe (LogObject a))
    }
```

Relation from context name to monitoring state

We remember the state of each monitored context name.

```
data MonitorState = MonitorState {
    _expression :: MEvExpr
    , _actions   :: [MEvAction]
    , _environment :: Environment
    }
type MonitorMap = HM.HashMap LoggerName MonitorState
```

Monitor view is an effectuator

Function *effectuate* is called to pass in a **LogObject** for monitoring.

```
instance IsEffectuator Monitor a where
  effectuate monitor item = do
    mon ← readMVar (getMon monitor)
    nocapacity ← atomically $ TBQ.isFullTBQueue (monQueue mon)
    if nocapacity
    then handleOverflow monitor
    else atomically $ TBQ.writeTBQueue (monQueue mon) $ Just item
  handleOverflow _ = TIO.hPutStrLn stderr "Notice: Monitor's queue full, dropping log items!"
```

Monitor implements Backend functions

Monitor is an **IsBackend**

```
instance Show a ⇒ IsBackend Monitor a where
  typeof _ = MonitoringBK
  realize _ = error "Monitoring cannot be instantiated by 'realize'"
  realizefrom config sbtrace _ = do
    monref ← newEmptyMVar
    let monitor = Monitor monref
    queue ← atomically $ TBQ.newTBQueue 512
    dispatcher ← spawnDispatcher queue config sbtrace
    -- link the given Async to the current thread, such that if the Async
    -- raises an exception, that exception will be re-thrown in the current
    -- thread, wrapped in ExceptionInLinkedThread.
    Async.link dispatcher
    putMVar monref $ MonitorInternal
      { monQueue = queue
      -- , monState = mempty
      }
    return monitor
  unrealize _ = return ()
```

Asynchronously reading log items from the queue and their processing

```
spawnDispatcher :: (Show a)
  ⇒ TBQ.TBQueue (Maybe (LogObject a))
  → Configuration
  → Trace.Trace IO a
  → IO (Async.Async ())
spawnDispatcher mqueue config sbtrace = do
  now ← getCurrentTime
  let messageCounters = resetCounters now
  countersMVar ← newMVar messageCounters
  timer ← Async.async $ sendAndResetAfter
    sbtrace
    "#messagecounters.monitoring"
```

```

countersMVar
60000 -- 60000 ms = 1 min
Warning-- Debug
Async.async (initMap >>= qProc countersMVar)
where
qProc counters state = do
  maybeItem ← atomically $ TBQ.readTBQueue mqueue
  case maybeItem of
    Just (logvalue@(LogObject _ _)) → do
      state' ← evalMonitoringAction state logvalue
      -- increase the counter for the type of message
      modifyMVar_ counters $ \cnt → return $ updateMessageCounters cnt logvalue
      qProc counters state'
    Nothing → return () -- stop here
  initMap = do
    ls ← getMonitors config
    return $ HM.fromList $ map (\(n, (e, as)) → (n, MonitorState e as HM.empty)) $ HM.toList ls

```

Evaluation of monitoring action

Inspect the log message and match it against configured thresholds. If positive, then run the action on the current state and return the updated state.

```

evalMonitoringAction :: MonitorMap → LogObject a → IO MonitorMap
evalMonitoringAction mmap logObj@(LogObject logname _) =
  case HM.lookup logname mmap of
    Nothing → return mmap
    Just mon@(MonitorState expr acts env0) → do
      let env' = updateEnv env0 logObj
      if evaluate env' expr
      then do
        now ← getMonotonicTimeNSec
        let env'' = HM.insert "lastalert" (Nanoseconds now) env'
        TIO.putStrLn $ "alert! " <> logname <> " " <> (pack $ show acts) <> " " <> (pack $ show env')
        return $ HM.insert logname mon { _environment = env'' } mmap
      else return mmap
where
  utc2ns (UTCTime days secs) =
    let yearsecs :: Rational
        yearsecs = 365 * 24 * 3600
        rdays, rsecs :: Rational
        rdays = toRational $ toModifiedJulianDay days
        rsecs = toRational secs
        s2ns = 1000000000
    in
      Nanoseconds $ round $ (fromRational $ s2ns * rsecs + rdays * yearsecs :: Double)
  updateEnv env (LogObject _ _ (ObserveOpen _)) = env
  updateEnv env (LogObject _ _ (ObserveDiff _)) = env
  updateEnv env (LogObject _ _ (ObserveClose _)) = env
  updateEnv env (LogObject _ lmeta (LogValue vn val)) =
    let addenv = HM.fromList [(vn, val)]

```

```

        ,("timestamp",utc2ns (tstamp lometa))
      ]
    in
      HM.union addenv env
updateEnv env (LogObject _ lometa (LogMessage _logitem)) =
  let addenv = HM.fromList [ ("severity", (Severity (severity lometa)))
    -- , ("selection", (liSelection logitem))
    -- , ("message", (liPayload logitem))
    ,("timestamp",utc2ns (tstamp lometa))
  ]
  in
    HM.union addenv env
updateEnv env (LogObject _ lometa (AggregatedMessage vals)) =
  let addenv = ("timestamp",utc2ns (tstamp lometa)): aggs2measurables vals [ ]
  in
    HM.union (HM.fromList addenv) env
where
  aggs2measurables [ ] acc = acc
  aggs2measurables ((n, AggregatedEWMA ewma): r) acc = aggs2measurables r $ (n <> ".avg", avg ewma)
  aggs2measurables ((n, AggregatedStats s): r) acc = aggs2measurables r $
    (n <> ".mean", PureD ◦ meanOfStats $ fbasic s)
    : (n <> ".flast", flast s)
    : (n <> ".fcount", PureI ◦ fromIntegral ◦ fcount $ fbasic s)
    : acc
-- catch all
updateEnv env _ = env

```

Chapter 2

Testing

2.1 Test coverage

Test coverage is calculated as the fraction of functions which are called from test routines. This percentage is calculated by the tool *hpc* with a call to

```
cabal new-test
```

Add to a local `cabal.project.local` file these lines:

```
tests:           True
coverage:        True
library-coverage: True
```

2.2 Test main entry point

```
{-# LANGUAGE CPP #-}
module Main
(
    main
) where
import Test.Tasty
# ifdef ENABLE_AGGREGATION
import qualified Cardano.BM.Test.Agregated (tests)
# endif
import qualified Cardano.BM.Test.STM (tests)
import qualified Cardano.BM.Test.Trace (tests)
import qualified Cardano.BM.Test.Configuration (tests)
import qualified Cardano.BM.Test.Rotator (tests)
import qualified Cardano.BM.Test.Routing (tests)
# ifdef ENABLE_MONITORING
import qualified Cardano.BM.Test.Monitoring (tests)
# endif
main :: IO ()
main = defaultMain tests
tests :: TestTree
tests =
    testGroup "iohk-monitoring"
    [
```

Cardano.BM.Counters.Dummy	100%
Cardano.BM.Setup	100%
Cardano.BM.Data.Trace	100%
Cardano.BM.Configuration.Static	100%
Cardano.BM.Configuration	100%
Cardano.BM.Configuration.Model	94%
Cardano.BM.Data.MessageCounter	85%
Cardano.BM.Data.Configuration	83%
Cardano.BM.Data.MonitoringEval	81%
Cardano.BM.BaseTrace	80%
Cardano.BM.Observer.Monad	75%
Cardano.BM.Output.Switchboard	75%
Cardano.BM.Output.Aggregation	68%
Cardano.BM.Output.Log	68%
Cardano.BM.Trace	63%
Cardano.BM.Data.Aggregated	63%
Cardano.BM.Rotator	50%
Cardano.BM.Data.BackendKind	50%
Cardano.BM.Data.Backend	50%
Cardano.BM.Counters.Common	50%
Cardano.BM.Data.Output	48%
Cardano.BM.Data.Severity	47%
Cardano.BM.Data.LogItem	41%
Cardano.BM.Data.Observable	40%
Cardano.BM.Observer.STM	33%
Cardano.BM.Data.AggregatedKind	33%
Cardano.BM.Data.Rotation	20%
Cardano.BM.Data.SubTrace	10%
Cardano.BM.Output.Monitoring	0%
Cardano.BM.Output.EKGView	0%
Cardano.BM.Data.Counter	0%
Cardano.BM.Counters	0%
Paths_iohk_monitoring	0%
	54%

Figure 2.1: Test coverage of modules in percent as computed by the tool 'hpc'


```

# ifdef ENABLE_AGGREGATION
    Cardano.BM.Test ◦ Aggregated.tests
    ,
# endif
    Cardano.BM.Test ◦ STM.tests
    , Cardano.BM.Test ◦ Trace.tests
    , Cardano.BM.Test ◦ Configuration.tests
    , Cardano.BM.Test ◦ Rotator.tests
    , Cardano.BM.Test ◦ Routing.tests
# ifdef ENABLE_MONITORING
    , Cardano.BM.Test ◦ Monitoring.tests
# endif
]

```

2.3 Test case generation

2.3.1 instance Arbitrary Aggregated

We define an instance of *Arbitrary* for an **Aggregated** which lets *QuickCheck* generate arbitrary instances of **Aggregated**. For this an arbitrary list of *Integer* is generated and this list is aggregated into a structure of **Aggregated**.

```

instance Arbitrary Aggregated where
    arbitrary = do
        vs' ← arbitrary :: Gen [Integer]
        let vs = 42 : 17 : vs'
        ds = map (λ(a,b) → a - b) $ zip vs (tail vs)
        (m1,s1) = updateMeanVar $ map fromInteger vs
        (m2,s2) = updateMeanVar $ map fromInteger ds
        mkBasicStats = BaseStats
            (PureI (minimum vs))
            (PureI (maximum vs))
            (fromIntegral $ length vs)
            (m1)
            (s1)
        mkDeltaStats = BaseStats
            (PureI (minimum ds))
            (PureI (maximum ds))
            (fromIntegral $ length ds)
            (m2)
            (s2)
        mkTimedStats = BaseStats
            (Nanoseconds 0)
            (Nanoseconds 0)
            (0)
            (0)
            (0)
        return $ AggregatedStats (Stats
            (PureI (last vs))
            (Nanoseconds 0)
            mkBasicStats

```

```
mkDeltaStats
mkTimedStats)
```

Estimators for mean and variance must be updated the same way as in the code.

```
updateMeanVar :: [Double] → (Double, Double)
updateMeanVar [] = (0, 0)
updateMeanVar (val : vals) = updateMeanVar' (val, 0) 1 vals
  where
    updateMeanVar' (m, s) _ [] = (m, s)
    updateMeanVar' (m, s) cnt (a : r) =
      let delta = a - m
          newcount = cnt + 1
          m' = m + (delta / newcount)
          s' = s + (delta * (a - m'))
      in
        updateMeanVar' (m', s') newcount r
```

2.4 Tests

2.4.1 Testing aggregation

```
tests :: TestTree
tests = testGroup "Aggregation measurements" [
  propertyTests
  , unitTests1
  , unitTests2
]
propertyTests :: TestTree
propertyTests = testGroup "Properties" [
  testProperty "minimal" prop_Aggregation_minimal
  , testProperty "commutative" prop_Aggregation_comm
]
unitTests1 :: TestTree
unitTests1 = testGroup "Unit tests for Aggregated" [
  testCase "compare equal >" unitAggregatedEqualGT
  , testCase "compare equal <" unitAggregatedEqualLT
  , testCase "compare different >" unitAggregatedDiffGT
  , testCase "compare different <" unitAggregatedDiffLT
]
unitTests2 :: TestTree
unitTests2 = testGroup "Unit tests for Aggregation" [
  testCase "initial -1" unitAggregationInitialMinus1
  , testCase "initial +1" unitAggregationInitialPlus1
  , testCase "initial +0" unitAggregationInitialZero
  , testCase "initial +1, -1" unitAggregationInitialPlus1Minus1
  , testCase "stepwise" unitAggregationStepwise
]
```

Property tests

```

prop_Aggregation_minimal :: Bool
prop_Aggregation_minimal = True

lometa :: LOMeta
lometa = unsafePerformIO $ mkLOMeta Debug Public

prop_Aggregation_comm :: Integer → Integer → Aggregated → Property
prop_Aggregation_comm v1 v2 ag =
  let AggregatedStats stats1 = updateAggregation (PureI v1) (updateAggregation (PureI v2) ag lometa)
      AggregatedStats stats2 = updateAggregation (PureI v2) (updateAggregation (PureI v1) ag lometa)
  in
    fbasic stats1 == fbasic stats2 .&&.
    (v1 ≡ v2) 'implies' (fbase stats1 == fbase stats2)
-- implication: if p1 is true, then return p2; otherwise true
implies :: Bool → Property → Property
implies p1 p2 = property (¬ p1) .|. p2

```

Unit tests for Aggregation

```

unitAggregationInitialMinus1 :: Assertion
unitAggregationInitialMinus1 = do
  let AggregatedStats stats1 = updateAggregation (-1) firstStateAggregatedStats lometa Nothing
      fbase stats1 @? = (-1)
      (fbasic stats1) @? = BaseStats (-1) 0 2 (-0.5) 0.5
      (fdelta stats1) @? = BaseStats 0 0 1 0 0
  -- AggregatedStats (Stats (-1) 0 (BaseStats (-1) 0 2 (-0.5) 0.5) (BaseStats 0 0
unitAggregationInitialPlus1 :: Assertion
unitAggregationInitialPlus1 = do
  let AggregatedStats stats1 = updateAggregation 1 firstStateAggregatedStats lometa Nothing
      fbase stats1 @? = 1
      (fbasic stats1) @? = BaseStats 0 1 2 0.5 0.5
      (fdelta stats1) @? = BaseStats 0 0 1 0 0
  -- AggregatedStats (Stats 1 0 (BaseStats 0 1 2 0.5 0.5) (BaseStats 0 0 1 0 0) (B
unitAggregationInitialZero :: Assertion
unitAggregationInitialZero = do
  let AggregatedStats stats1 = updateAggregation 0 firstStateAggregatedStats lometa Nothing
      fbase stats1 @? = 0
      (fbasic stats1) @? = BaseStats 0 0 2 0 0
      (fdelta stats1) @? = BaseStats 0 0 1 0 0
  -- AggregatedStats (Stats 0 0 (BaseStats 0 0 2 0 0) (BaseStats 0 0 1 0 0) (BaseS
unitAggregationInitialPlus1Minus1 :: Assertion
unitAggregationInitialPlus1Minus1 = do
  let AggregatedStats stats1 = updateAggregation (PureI (-1)) (updateAggregation (PureI 1) firstState
      (fbasic stats1) @? = BaseStats (PureI (-1)) (PureI 1) 3 0.0 2.0
      (fdelta stats1) @? = BaseStats (PureI (-2)) (PureI 0) 2 (-1.0) 2.0
unitAggregationStepwise :: Assertion
unitAggregationStepwise = do
  stats0 ← pure $ singletonStats (Bytes 3000)
  -- putStrLn (show stats0)

```

```

threadDelay 50000-- 0.05 s
t1 ← mkLOMeta Debug Public
stats1 ← pure $ updateAggregation (Bytes 5000) stats0 t1 Nothing
-- putStrLn (show stats1)
-- showTimedMean stats1
threadDelay 50000-- 0.05 s
t2 ← mkLOMeta Debug Public
stats2 ← pure $ updateAggregation (Bytes 1000) stats1 t2 Nothing
-- putStrLn (show stats2)
-- showTimedMean stats2
checkTimedMean stats2
threadDelay 50000-- 0.05 s
t3 ← mkLOMeta Debug Public
stats3 ← pure $ updateAggregation (Bytes 3000) stats2 t3 Nothing
-- putStrLn (show stats3)
-- showTimedMean stats3
checkTimedMean stats3
threadDelay 50000-- 0.05 s
t4 ← mkLOMeta Debug Public
stats4 ← pure $ updateAggregation (Bytes 1000) stats3 t4 Nothing
-- putStrLn (show stats4)
-- showTimedMean stats4
checkTimedMean stats4
where
  checkTimedMean (AggregatedEWMA _) = return ()
  checkTimedMean (AggregatedStats s) = do
    let mean = meanOfStats (ftimed s)
    assertBool "the mean should be >= the minimum" (mean ≥ getDouble (fmin (ftimed s)))
    assertBool "the mean should be <= the maximum" (mean ≤ getDouble (fmax (ftimed s)))

```

commented out:

```

showTimedMean (AggregatedEWMA _) = return ()
showTimedMean (AggregatedStats s) = putStrLn $ "mean = " ++ show (meanOfStats (ftimed s)) ++ showUnits

```

```

firstStateAggregatedStats :: Aggregated
firstStateAggregatedStats = AggregatedStats (Stats z z (BaseStats z z 1 0 0) (BaseStats z z 0 0 0) (BaseStats z z 0 0 0))
where
  z = PureI 0

```

Unit tests for Aggregated

```

unitAggregatedEqualGT :: Assertion
unitAggregatedEqualGT = do
  assertBool "comparing seconds"
    ((Seconds 3) > (Seconds 2))
  assertBool "comparing microseconds"
    ((Microseconds 3000) > (Microseconds 2000))
  assertBool "comparing nanoseconds"
    ((Nanoseconds 3000000) > (Nanoseconds 2000000))

```

```

assertBool "comparing bytes"
  ((Bytes 2048) > (Bytes 1024))
assertBool "comparing doubles"
  ((PureD 2.34) > (PureD 1.42))
assertBool "comparing integers"
  ((PureI 2) > (PureI 1))
assertBool "comparing severities"
  ((SeverityError) > (SeverityWarning))
unitAggregatedEqualLT :: Assertion
unitAggregatedEqualLT = do
  assertBool "comparing seconds"
    ((Seconds 2) < (Seconds 3))
  assertBool "comparing microseconds"
    ((Microseconds 2000) < (Microseconds 3000))
  assertBool "comparing nanoseconds"
    ((Nanoseconds 2000000) < (Nanoseconds 3000000))
  assertBool "comparing bytes"
    ((Bytes 1024) < (Bytes 2048))
  assertBool "comparing doubles"
    ((PureD 1.34) < (PureD 2.42))
  assertBool "comparing integers"
    ((PureI 1) < (PureI 2))
  assertBool "comparing severities"
    ((SeverityInfo) < (SeverityNotice))
unitAggregatedDiffGT :: Assertion
unitAggregatedDiffGT = do
  assertBool "comparing time (s vs. s)"
    ((Microseconds 3000000) > (Seconds 2))
  assertBool "comparing time (s vs. ns)"
    ((Microseconds 30) > (Nanoseconds 29999))
  assertBool "comparing nanoseconds"
    ((Nanoseconds 3000000) > (Microseconds 2900))
  assertBool "comparing bytes"
    ((Bytes 2048) > (PureI 1024))
  assertBool "comparing doubles"
    ((PureD 2.34) > (PureI 1))
  assertBool "comparing integers"
    ((PureI 2) > (PureD 1.42))
unitAggregatedDiffLT :: Assertion
unitAggregatedDiffLT = do
  assertBool "comparing time (s vs. s)"
    ((Microseconds 2999999) < (Seconds 3))
  assertBool "comparing time (s vs. ns)"
    ((Microseconds 30) < (Nanoseconds 30001))
  assertBool "comparing nanoseconds"
    ((Nanoseconds 3000000) < (Microseconds 3001))
  assertBool "comparing bytes"
    ((PureI 1024) < (Bytes 2048))
  assertBool "comparing doubles"
    ((PureD 2.34) < (PureI 3))
  assertBool "comparing integers"

```

```
((PureI 2) < (PureD 3.42))
```

2.4.2 Cardano.BM.Test.STM

```
module Cardano.BM.Test.STM (
  tests
) where
import Test.Tasty
import Test.Tasty.QuickCheck
tests :: TestTree
tests = testGroup "Observing STM actions" [
  testProperty "minimal" prop_STM_observer
]
prop_STM_observer :: Bool
prop_STM_observer = True
```

2.4.3 Cardano.BM.Test.Trace

```
tests :: TestTree
tests = testGroup "Testing Trace" [
  unit_tests
  , testCase "forked traces stress testing" stressTraceInFork
# ifdef ENABLE_OBSERVABLES
  , testCase "stress testing: ObservableTrace vs. NoTrace" timingObservableVsUntimed
# endif
  , testCaseInfo "demonstrating logging" simpleDemo
  , testCaseInfo "demonstrating nested named context logging" exampleWithNamedContexts
]
unit_tests :: TestTree
unit_tests = testGroup "Unit tests" [
  testCase "opening messages should not be traced" unitNoOpeningTrace
-- , testCase "hierarchy of traces" unitHierarchy
  , testCase "forked traces" unitTraceInFork
  , testCase "hierarchy of traces with NoTrace" $
    unitHierarchy' [Neutral, NoTrace, (ObservableTrace observablesSet)]
    onlyLevelOneMessage
  , testCase "hierarchy of traces with DropOpening" $
    unitHierarchy' [Neutral, DropOpening, (ObservableTrace observablesSet)]
    notObserveOpen
  , testCase "hierarchy of traces with UntimedTrace" $
    unitHierarchy' [Neutral, UntimedTrace, UntimedTrace]
    observeNoMeasures
  , testCase "changing the minimum severity of a trace at runtime"
    unitTraceMinSeverity
  , testCase "changing the minimum severity of a named context at runtime"
    unitNamedMinSeverity
  , testCase "appending names" unitAppendName
  , testCase "create subtrace which duplicates messages" unitTraceDuplicate
```

```

,testCase "testing name filtering" unitNameFiltering
,testCase "testing throwing of exceptions" unitExceptionThrowing
,testCase "NoTrace: check lazy evaluation" unitTestLazyEvaluation
,testCase "private messages should not be logged into private files" unitLoggingPrivate
]
where
  observablesSet = [MonotonicClock, MemoryStats]
  notObserveOpen :: [LogObject a] → Bool
  notObserveOpen = all (\case {LogObject _ _ (ObserveOpen _) → False; _ → True})
  notObserveClose :: [LogObject a] → Bool
  notObserveClose = all (\case {LogObject _ _ (ObserveClose _) → False; _ → True})
  notObserveDiff :: [LogObject a] → Bool
  notObserveDiff = all (\case {LogObject _ _ (ObserveDiff _) → False; _ → True})
  onlyLevelOneMessage :: [LogObject Text] → Bool
  onlyLevelOneMessage = \case
    [LogObject _ _ (LogMessage "Message from level 1.")] → True
    _ → False
  observeNoMeasures :: [LogObject a] → Bool
  observeNoMeasures obs = notObserveOpen obs ∧ notObserveClose obs ∧ notObserveDiff obs

```

Helper routines

```

data TraceConfiguration = TraceConfiguration
  { tcConfig      :: Configuration
  , tcOutputKind :: OutputKind Text
  , tcName        :: LoggerName
  , tcSubTrace    :: SubTrace
  }

setupTrace :: TraceConfiguration → IO (Trace IO Text)
setupTrace (TraceConfiguration cfg outk name subTr) = do
  let logTrace = case outk of
    TVarList tvar → TracerT.natTrace liftIO $ traceInTVarIOConditionally tvar cfg
  setSubTrace cfg name (Just subTr)
  appendName name logTrace

```

Simple demo of logging.

```

simpleDemo :: IO String
simpleDemo = do
  cfg ← defaultConfigTesting
  logTrace :: Trace IO String ← Setup.setupTrace (Right cfg) "test"
  putStrLn "\n"
  logDebug logTrace "This is how a Debug message looks like."
  logInfo logTrace "This is how an Info message looks like."
  logNotice logTrace "This is how a Notice message looks like."
  logWarning logTrace "This is how a Warning message looks like."
  logError logTrace "This is how an Error message looks like."
  logCritical logTrace "This is how a Critical message looks like."

```

```

logAlert    logTrace "This is how an Alert message looks like."
logEmergency logTrace "This is how an Emergency message looks like."
return ""

```

Example of using named contexts with **Trace**

```

exampleWithNamedContexts :: IO String
exampleWithNamedContexts = do
  cfg ← defaultConfigTesting
  Setup.withTrace cfg "test" $ λ(logTrace :: Trace IO Text) → do
    putStrLn "\n"
    logInfo logTrace "entering"
    logTrace0 ← appendName "simple-work-0" logTrace
    work0 ← complexWork0 cfg logTrace0 "0"
    logTrace1 ← appendName "complex-work-1" logTrace
    work1 ← complexWork1 cfg logTrace1 "42"

    Async.wait work0
    Async.wait work1
    -- the named context will include "complex" in the logged message
    logInfo logTrace "done."
    threadDelay 100000
    -- force garbage collection to allow exceptions to be thrown
    performMajorGC
    threadDelay 100000

  return ""
where
  complexWork0 _ tr msg = Async.async $ logInfo tr ("let's see (0): " 'append' msg)
  complexWork1 cfg tr msg = Async.async $ do
    logInfo tr ("let's see (1): " 'append' msg)
    trInner ← appendName "inner-work-1" tr
    let observablesSet = [MonotonicClock]
    setSubTrace cfg "test.complex-work-1.inner-work-1.STM-action" $
      Just $ ObservableTrace observablesSet
  # ifdef ENABLE_OBSERVABLES
    _ ← STMObserver.bracketObserveIO cfg trInner Debug "STM-action" setVar _
  # endif
    logInfo trInner "let's see: done."

```

Show effect of turning off observables

```

# ifdef ENABLE_OBSERVABLES
runTimedAction :: Configuration → Trace IO Text → Int → IO Measurable
runTimedAction cfg logTrace reps = do
  runid ← newUnique
  t0 ← getMonoClock
  forM_ [(1 :: Int)..reps] $ const $ observeAction logTrace
  t1 ← getMonoClock
  return $ diffTimeObserved (CounterState runid t0) (CounterState runid t1)

```



```

where
  observeAction trace = do
    _ ← MonadObserver.bracketObserveIO cfg trace Debug " " action
    return ()
    action = return $ forM [1 :: Int..100] $ \x → [x] ++ (init $ reverse [1 :: Int..10000])
timingObservableVsUntimed :: Assertion
timingObservableVsUntimed = do
  cfg ← defaultConfigTesting
  msgs1 ← STM.newTVarIO []
  traceObservable ← setupTrace $ TraceConfiguration cfg
    (TVarList msgs1)
    "observables"
    (ObservableTrace observablesSet)
  msgs2 ← STM.newTVarIO []
  traceUntimed ← setupTrace $ TraceConfiguration cfg
    (TVarList msgs2)
    "no timing"
    UntimedTrace
  msgs3 ← STM.newTVarIO []
  traceNoTrace ← setupTrace $ TraceConfiguration cfg
    (TVarList msgs3)
    "no trace"
    NoTrace
  t_observable ← runTimedAction cfg traceObservable 100
  t_untimed ← runTimedAction cfg traceUntimed 100
  t_notrace ← runTimedAction cfg traceNoTrace 100
  assertBool
    ("Untimed consumed more time than ObservableTrace " ++ (show [t_untimed,t_observable]))
    True
  assertBool
    ("NoTrace consumed more time than ObservableTrace" ++ (show [t_notrace,t_observable]))
    True
  assertBool
    ("NoTrace consumed more time than Untimed" ++ (show [t_notrace,t_untimed]))
    True
where
  observablesSet = [MonotonicClock,GhcRtsStats,MemoryStats,IOStats,ProcessStats]
#endif

```

Control tracing in a hierarchy of **Traces**

We can lay out traces in a hierarchical manner, that the children forward traced items to the parent **Trace**. A **NoTrace** introduced in this hierarchy will cut off a branch from messaging to the root.

```

unitHierarchy :: Assertion
unitHierarchy = do
  cfg ← defaultConfigTesting
  msgs ← STM.newTVarIO []
  trace0 ← setupTrace $ TraceConfiguration cfg (TVarList msgs) "test" Neutral

```

```

logInfo trace0 "This should have been displayed!"
-- subtrace of trace which traces nothing
setSubTrace cfg "test.inner" (Just NoTrace)
trace1 ← appendName "inner" trace0
logInfo trace1 "This should NOT have been displayed!"
setSubTrace cfg "test.inner.innermost" (Just Neutral)
trace2 ← appendName "innermost" trace1
logInfo trace2 "This should NOT have been displayed also due to the trace one level above"
-- acquire the traced objects
res ← STM.readTVarIO msgs
-- only the first message should have been traced
assertBool
  ("Found more or less messages than expected: " ++ show res)
  (length res == 1)

```

Change a trace's minimum severity

A trace is configured with a minimum severity and filters out messages that are labelled with a lower severity. This minimum severity of the current trace can be changed.

```

unitTraceMinSeverity :: Assertion
unitTraceMinSeverity = do
  cfg ← defaultConfigTesting
  msgs ← STM.newTVarIO []
  trace ← setupTrace $ TraceConfiguration cfg (TVarList msgs) "test min severity" Neutral
  logInfo trace "Message #1"
  -- raise the minimum severity to Warning
  setMinSeverity cfg Warning
  msev ← Cardano.BM.Configuration.minSeverity cfg
  assertBool ("min severity should be Warning, but is " ++ (show msev))
    (msev == Warning)
  -- this message will not be traced
  logInfo trace "Message #2"
  -- lower the minimum severity to Info
  setMinSeverity cfg Info
  -- this message is traced
  logInfo trace "Message #3"
  -- acquire the traced objects
  res ← STM.readTVarIO msgs
  -- only the first and last messages should have been traced
  assertBool
    ("Found more or less messages than expected: " ++ show res)
    (length res == 2)
  assertBool
    ("Found Info message when Warning was minimum severity: " ++ show res)
    (all
      (λcase
        LogObject _ (LOMeta _ _ Info _) (LogMessage "Message #2") → False

```

```

    _ → True)
  res)

```

Define a subtrace's behaviour to duplicate all messages

The **SubTrace** will duplicate all messages that pass through it. Each message will be in its own named context.

```

unitTraceDuplicate :: Assertion
unitTraceDuplicate = do
  cfg ← defaultConfigTesting
  msgs ← STM.newTVarIO []
  trace0 ← setupTrace $ TraceConfiguration cfg (TVarList msgs) "test-duplicate" Neutral
  logInfo trace0 "Message #1"
  -- create a subtrace which duplicates all messages
  setSubTrace cfg "test-duplicate.orig" $ Just (TeeTrace "test-duplicate.dup")
  trace ← appendName "orig" trace0
  -- this message will be duplicated
  logInfo trace "You will see me twice!"
  -- acquire the traced objects
  res ← STM.readTVarIO msgs
  -- only the first and last messages should have been traced
  assertBool
    ("Found more or less messages than expected: " ++ show res)
    (length res == 3)

```

Change the minimum severity of a named context

A trace of a named context can be configured with a minimum severity, such that the trace will filter out messages that are labelled with a lower severity.

```

unitNamedMinSeverity :: Assertion
unitNamedMinSeverity = do
  cfg ← defaultConfigTesting
  msgs ← STM.newTVarIO []
  trace0 ← setupTrace $ TraceConfiguration cfg (TVarList msgs) "test-named-severity" Neutral
  trace ← appendName "sev-change" trace0
  logInfo trace "Message #1"
  -- raise the minimum severity to Warning
  setSeverity cfg "test-named-severity.sev-change" (Just Warning)
  msev ← Cardano.BM.Configuration.inspectSeverity cfg "test-named-severity.sev-change"
  assertBool ("min severity should be Warning, but is " ++ (show msev))
    (msev == Just Warning)
  -- this message will not be traced
  logInfo trace "Message #2"
  -- lower the minimum severity to Info
  setSeverity cfg "test-named-severity.sev-change" (Just Info)
  -- this message is traced
  logInfo trace "Message #3"
  -- acquire the traced objects

```

```

res ← STM.readTVarIO msgs
-- only the first and last messages should have been traced
assertBool
  ("Found more or less messages than expected: " ++ show res)
  (length res == 2)
assertBool
  ("Found Info message when Warning was minimum severity: " ++ show res)
  (all
    (λcase
      LogObject _ (LOMeta _ _ Info _) (LogMessage "Message #2") → False
      _ → True)
    res)

unitHierarchy' :: [SubTrace] → ([LogObject Text] → Bool) → Assertion
unitHierarchy' subtraces f = do
  cfg ← liftIO Cardano.BM.Configuration o Model.empty
  let (t1 : t2 : t3 : _) = cycle subtraces
  msgs ← STM.newTVarIO []
  -- create trace of type 1
  trace1 ← setupTrace $ TraceConfiguration cfg (TVarList msgs) "test" t1
  logInfo trace1 "Message from level 1."
  -- subtrace of type 2
  setSubTrace cfg "test.inner" (Just t2)
  trace2 ← appendName "inner" trace1
  logInfo trace2 "Message from level 2."
  -- subsubtrace of type 3
  setSubTrace cfg "test.inner.innermost" (Just t3)
  # ifdef ENABLE_OBSERVABLES
  _ ← STMObserver.bracketObserveIO cfg trace2 Debug "test.inner.innermost" setVar_
  # endif
  logInfo trace2 "Message from level 3."
  -- acquire the traced objects
  res ← STM.readTVarIO msgs
  -- only the first message should have been traced
  assertBool
    ("Found more or less messages than expected: " ++ show res)
    (f res)

```

Logging in parallel

```

unitTraceInFork :: Assertion
unitTraceInFork = do
  cfg ← defaultConfigTesting
  msgs ← STM.newTVarIO []
  trace ← setupTrace $ TraceConfiguration cfg (TVarList msgs) "test" Neutral
  trace0 ← appendName "work0" trace
  trace1 ← appendName "work1" trace
  work0 ← work trace0
  threadDelay 5000

```

```

work1 ← work trace1
Async.wait $ work0
Async.wait $ work1
res ← STM.readTVarIO msgs
let names@(_:namesTail) = map loName res
-- each trace should have its own name and log right after the other
assertBool
  ("Consecutive loggernames are not different: " ++ show names)
  (and $ zipWith (≠) names namesTail)
where
work :: Trace IO Text → IO (Async.Async ())
work trace = Async.async $ do
  logInfoDelay trace "1"
  logInfoDelay trace "2"
  logInfoDelay trace "3"
logInfoDelay :: Trace IO Text → Text → IO ()
logInfoDelay trace msg =
  logInfo trace msg >>
  threadDelay 10000

```

Stress testing parallel logging

```

stressTraceInFork :: Assertion
stressTraceInFork = do
  cfg ← defaultConfigTesting
  msgs ← STM.newTVarIO []
  trace ← setupTrace $ TraceConfiguration cfg (TVarList msgs) "test" Neutral
  let names = map (λa → ("work-" <> pack (show a))) [1..(10::Int)]
  ts ← forM names $ λname → do
    trace' ← appendName name trace
    work trace'
  forM_ ts Async.wait
  res ← STM.readTVarIO msgs
  let resNames = map loName res
  let frequencyMap = fromListWith (+) [(x,1) | x ← resNames]
  -- each trace should have traced totalMessages' messages
  assertBool
    ("Frequencies of logged messages according to loggername: " ++ show frequencyMap)
    (all (λname → (lookup ("test." <> name) frequencyMap) ≡ Just totalMessages) names)
  where
work :: Trace IO Text → IO (Async.Async ())
work trace = Async.async $ forM_ [1..totalMessages] $ (logInfo trace) ∘ pack ∘ show
totalMessages :: Int
totalMessages = 10

```

Dropping **ObserveOpen** messages in a subtrace

```

unitNoOpeningTrace :: Assertion
unitNoOpeningTrace = do

```

```

    cfg ← defaultConfigTesting
    msgs ← STM.newTVarIO []
# ifdef ENABLE_OBSERVABLES
    logTrace ← setupTrace $ TraceConfiguration cfg (TVarList msgs) "test" DropOpening
    _ ← STMObserver.bracketObserveIO cfg logTrace Debug "setTVar" setVar_
# endif
    res ← STM.readTVarIO msgs
    assertBool
      ("Found non-expected ObserveOpen message: " ++ show res)
      (all (\case {LogObject _ _ (ObserveOpen _) → False; _ → True}) res)

```

Assert maximum length of log context name

The name of the log context cannot grow beyond a maximum number of characters, currently the limit is set to 80.

```

unitAppendName :: Assertion
unitAppendName = do
  cfg ← defaultConfigTesting
  msgs ← STM.newTVarIO []
  trace0 ← setupTrace $ TraceConfiguration cfg (TVarList msgs) "test" Neutral
  trace1 ← appendName bigName trace0
  trace2 ← appendName bigName trace1
  forM_ [trace0, trace1, trace2] $ (flip logInfo msg)
  res ← reverse < $ > STM.readTVarIO msgs
  let loggernames = map loName res
  assertBool
    ("AppendName did not work properly. The loggernames for the messages are: " ++
     show loggernames)
    (loggernames ≡ [ "test"
                    , "test." <> bigName
                    , "test." <> bigName <> "." <> bigName
                    ])
  where
    bigName = T.replicate 30 "abcdefghijklmnopqrstuvwxy"
    msg = "Hello!"

# ifdef ENABLE_OBSERVABLES
setVar_ :: STM.STM Integer
setVar_ = do
  t ← STM.newTVar 0
  STM.writeTVar t 42
  res ← STM.readTVar t
  return res
# endif

```

Testing log context name filters

```

unitNameFiltering :: Assertion
unitNameFiltering = do

```

```

let contextName = "test.sub.1"
let loname = "sum"-- would be part of a "LogValue loname 42"
let filter1 = [(Drop (Exact "test.sub.1"), Unhide [])]
assertBool ("Dropping a specific name should filter it out and thus return False")
  (False == evalFilters filter1 contextName)
let filter2 = [(Drop (EndsWith ".1"), Unhide [])]
assertBool ("Dropping a name ending with a specific text should filter out the context")
  (False == evalFilters filter2 contextName)
let filter3 = [(Drop (StartsWith "test."), Unhide [])]
assertBool ("Dropping a name starting with a specific text should filter out the context")
  (False == evalFilters filter3 contextName)
let filter4 = [(Drop (Contains ".sub."), Unhide [])]
assertBool ("Dropping a name starting containing a specific text should filter out the context")
  (False == evalFilters filter4 contextName)
let filter5 = [(Drop (StartsWith "test."),
  Unhide [(Exact "test.sub.1")])]
assertBool ("Dropping all and unhiding a specific name should the context name allow pass")
  (True == evalFilters filter5 contextName)
let filter6 = [(Drop (StartsWith "test."),
  Unhide [(EndsWith ".sum"),
  (EndsWith ".other")])]
assertBool ("Dropping all and unhiding some names, the LogObject should pass the filter")
  (True == evalFilters filter6 (contextName <> "." <> loname))
let filter7 = [(Drop (StartsWith "test."),
  Unhide [(EndsWith ".product")])]
assertBool ("Dropping all and unhiding an inexistant named value, the LogObject should not pass")
  (False == evalFilters filter7 (contextName <> "." <> loname))
let filter8 = [(Drop (StartsWith "test."),
  Unhide [(Exact "test.sub.1")]),
  (Drop (StartsWith "something.else."),
  Unhide [(EndsWith ".this")])]
assertBool ("Disjunction of filters that should pass")
  (True == evalFilters filter8 contextName)
let filter9 = [(Drop (StartsWith "test."),
  Unhide [(Exact ".that")]),
  (Drop (StartsWith "something.else."),
  Unhide [(EndsWith ".this")])]
assertBool ("Disjunction of filters that should not pass")
  (False == evalFilters filter9 contextName)

```

Exception throwing

Exceptions encountered should be thrown.

```

unitExceptionThrowing :: Assertion
unitExceptionThrowing = do
  action ← work msg
  res ← Async.waitCatch action
  assertBool
    ("Exception should have been rethrown")

```

```

    (isLeft res)
  where
    msg :: Text
    msg = error "faulty message"
    work :: Text → IO (Async.Async ())
    work message = Async.async $ do
      cfg ← defaultConfigTesting
      trace ← Setup.setupTrace (Right cfg) "test"
      logInfo trace message
      threadDelay 10000

```

Check lazy evaluation of trace

Exception should not be thrown when type of `Trace` is `NoTrace`.

```

unitTestLazyEvaluation :: Assertion
unitTestLazyEvaluation = do
  action ← work msg
  res ← Async.waitCatch action
  assertBool
    ("Exception should not have been rethrown when type of Trace is NoTrace")
    (isRight res)
  where
    msg :: Text
    msg = error "faulty message"
    work :: Text → IO (Async.Async ())
    work message = Async.async $ do
      cfg ← defaultConfigTesting
      trace0 ← Setup.setupTrace (Right cfg) "test"
      setSubTrace cfg "test.work" (Just NoTrace)
      trace ← appendName "work" trace0
      logInfo trace message

```

Check that private messages do not end up in public log files.

```

unitLoggingPrivate :: Assertion
unitLoggingPrivate = do
  tmpDir ← getTemporaryDirectory
  let privateFile = tmpDir </> "private.log"
      publicFile = tmpDir </> "public.log"
  conf ← empty
  setDefaultBackends conf [KatipBK]
  setSetupBackends conf [KatipBK]
  setDefaultScribes conf [ "FileTextSK::" <> pack privateFile
    , "FileTextSK::" <> pack publicFile
  ]
  setSetupScribes conf [ScribeDefinition
    { scKind = FileTextSK

```



```

    ,scName      = pack privateFile
    ,scPrivacy   = ScPrivate
    ,scRotation  = Nothing
  }
  ,ScribeDefinition
  { scKind      = FileTextSK
  , scName      = pack publicFile
  , scPrivacy   = ScPublic
  , scRotation  = Nothing
  }
]
Setup.withTrace conf "test" $ \trace → do
  -- should log in both files
  logInfo trace message
  -- should only log in private file
  logInfoS trace message

  countPublic ← length ∘ lines < $ > readFile publicFile
  countPrivate ← length ∘ lines < $ > readFile privateFile
  -- delete files
  forM_ [privateFile, publicFile] removeFile
  assertBool
    ("Confidential file should contain 2 lines and it contains " ++ show countPrivate ++ " "
     "Public file should contain 1 line and it contains " ++ show countPublic ++ ".\n"
    )
    (countPublic ≡ 1 ∧ countPrivate ≡ 2)
where
  message :: Text
  message = "Just a message"

```

2.4.4 Testing configuration

Test declarations

```

tests :: TestTree
tests = testGroup "config tests" [
  propertyTests
  , unitTests
]
propertyTests :: TestTree
propertyTests = testGroup "Properties" [
  testProperty "minimal" prop_Configuration_minimal
]
unitTests :: TestTree
unitTests = testGroup "Unit tests" [
  testCase "static representation" unitConfigurationStaticRepresentation
  , testCase "parsed representation" unitConfigurationParsedRepresentation
  , testCase "parsed configuration" unitConfigurationParsed
  , testCase "include EKG if defined" unitConfigurationCheckEKGpositive
  , testCase "not include EKG if not def" unitConfigurationCheckEKGnegative
]

```

```
,testCase "check scribe caching" unitConfigurationCheckScribeCache
,testCase "test ops on Configuration" unitConfigurationOps
]
```

Property tests

```
prop_Configuration_minimal :: Bool
prop_Configuration_minimal = True
```

Unit tests

The configuration file only indicates that EKG is listening on port nnnnn. Infer that *EKGViewBK* needs to be started as a backend.

```
unitConfigurationCheckEKGpositive :: Assertion
unitConfigurationCheckEKGpositive = do
# ifdef ENABLE_EKG
    return ()
# else
    tmp ← getTemporaryDirectory
    let c = [ "rotation:"
      , "  rpLogLimitBytes: 5000000"
      , "  rpKeepFilesNum: 10"
      , "  rpMaxAgeHours: 24"
      , "minSeverity: Info"
      , "defaultBackends:"
      , "  - KatipBK"
      , "setupBackends:"
      , "  - KatipBK"
      , "defaultScribes:"
      , "  - StdoutSK"
      , "  - stdout"
      , "setupScribes:"
      , "  - scName: stdout"
      , "  scRotation: null"
      , "  scKind: StdoutSK"
      , "hasEKG: 18321"
      , "options:"
      , "  test:"
      , "    value: nothing"
    ]
    fp = tmp </> "test_ekgv_config.yaml"
    writeFile fp $ unlines c
    repr ← parseRepresentation fp
    assertBool "expecting EKGViewBK to be setup" $
      EKGViewBK ∈ (setupBackends repr)
# endif
```

If there is no port defined for EKG, then do not start it even if present in the config.

```

unitConfigurationCheckEKGnegative :: Assertion
unitConfigurationCheckEKGnegative = do
# ifdef ENABLE_EKG
    return ()
# else
    tmp ← getTemporaryDirectory
    let c = [ "rotation:"
      , "  rpLogLimitBytes: 5000000"
      , "  rpKeepFilesNum: 10"
      , "  rpMaxAgeHours: 24"
      , "minSeverity: Info"
      , "defaultBackends:"
      , "  - KatipBK"
      , "  - EKGViewBK"
      , "setupBackends:"
      , "  - KatipBK"
      , "  - EKGViewBK"
      , "defaultScribes:"
      , "- - StdoutSK"
      , "  - stdout"
      , "setupScribes:"
      , "- scName: stdout"
      , "  scRotation: null"
      , "  scKind: StdoutSK"
      , "###hasEKG: 18321"
      , "options:"
      , "  test:"
      , "    value: nothing"
      ]
    fp = tmp < / > "test_ekgv_config.yaml"
    writeFile fp $ unlines c
    repr ← parseRepresentation fp
    assertBool "EKGViewBK shall not be setup" $
      ¬ $ EKGViewBK ∈ (setupBackends repr)
    assertBool "EKGViewBK shall not receive messages" $
      ¬ $ EKGViewBK ∈ (defaultBackends repr)
# endif

```

```

unitConfigurationStaticRepresentation :: Assertion
unitConfigurationStaticRepresentation =
    let r = Representation
      { minSeverity = Info
      , rotation = Just $ RotationParameters
          { rpLogLimitBytes = 5000000
          , rpMaxAgeHours = 24
          , rpKeepFilesNum = 10
          }
      , setupScribes =
          [ ScribeDefinition { scName = "stdout"
            , scKind = StdoutSK
            , scPrivacy = ScPublic

```

```

        ,scRotation = Nothing}
    ]
    ,defaultScribes = [(StdoutSK, "stdout")]
    ,setupBackends = [EKGVViewBK, KatipBK]
    ,defaultBackends = [KatipBK]
    ,hasGUI = Just 12789
    ,hasEKG = Just 18321
    ,options =
        HM.fromList [("test1", (HM.singleton "value" "object1"))
            ,("test2", (HM.singleton "value" "object2"))]
    }
in
encode r @? =
    (intercalate "\n"
        [ "rotation:"
        , "  rpLogLimitBytes: 5000000"
        , "  rpKeepFilesNum: 10"
        , "  rpMaxAgeHours: 24"
        , "defaultBackends:"
        , "- KatipBK"
        , "setupBackends:"
        , "- EKGVViewBK"
        , "- KatipBK"
        , "hasGUI: 12789"
        , "defaultScribes:"
        , "- - StdoutSK"
        , "  - stdout"
        , "options:"
        , "  test2:"
        , "    value: object2"
        , "  test1:"
        , "    value: object1"
        , "setupScribes:"
        , "- scName: stdout"
        , "  scRotation: null"
        , "  scKind: StdoutSK"
        , "  scPrivacy: ScPublic"
        , "hasEKG: 18321"
        , "minSeverity: Info"
        , "-- to force a line feed at the end of the file"
        ]
    )
unitConfigurationParsedRepresentation :: Assertion
unitConfigurationParsedRepresentation = do
    repr ← parseRepresentation "test/config.yaml"
    encode repr @? =
        (intercalate "\n"
            [ "rotation:"
            , "  rpLogLimitBytes: 5000000"
            , "  rpKeepFilesNum: 10"
            , "  rpMaxAgeHours: 24"

```

```

,"defaultBackends:"
,"- KatipBK"
,"setupBackends:"
,"- AggregationBK"
,"- EKGViewBK"
,"- KatipBK"
,"hasGUI: null"
,"defaultScribes:"
,"- - StdoutSK"
,"  - stdout"
,"options:"
,"  mapSubtrace:"
,"    iohk.benchmarking:"
,"      tag: ObservableTrace"
,"      contents:"
,"        - GhcRtsStats"
,"        - MonotonicClock"
,"    iohk.deadend: NoTrace"
,"  mapSeverity:"
,"    iohk.startup: Debug"
,"    iohk.background.process: Error"
,"    iohk.testing.uncritical: Warning"
,"  mapAggregatedkinds:"
,"    iohk.interesting.value: EwmaAK {alpha = 0.75}"
,"    iohk.background.process: StatsAK"
,"  cfokey:"
,"    value: Release-1.0.0"
,"  mapMonitors:"
,"    chain.creation.block:"
,"      - monitor: ((time > (23 s)) Or (time < (17 s)))"
,"      - actions:"
,"        - AlterMinSeverity \"chain.creation\" Debug"
,"        ! '#aggregation.critproc.observable':"
,"      - monitor: (mean >= (42))"
,"      - actions:"
,"        - CreateMessage \"exceeded\" \"the observable has been too long too high"
,"        - AlterGlobalMinSeverity Info"
,"  mapScribes:"
,"    iohk.interesting.value:"
,"      - StdoutSK::stdout"
,"      - FileTextSK::testlog"
,"    iohk.background.process: FileTextSK::testlog"
,"  mapBackends:"
,"    iohk.interesting.value:"
,"      - EKGViewBK"
,"      - AggregationBK"
,"setupScribes:"
,"- scName: testlog"
,"  scRotation:"
,"    rpLogLimitBytes: 25000000"
,"    rpKeepFilesNum: 3"

```

```

    , "    rpMaxAgeHours: 24"
    , "    scKind: FileTextSK"
    , "    scPrivacy: ScPrivate"
    , "- scName: stdout"
    , "    scRotation: null"
    , "    scKind: StdoutSK"
    , "    scPrivacy: ScPublic"
    , "hasEKG: 12789"
    , "minSeverity: Info"
    , "-- to force a line feed at the end of the file
  ]
)
unitConfigurationParsed :: Assertion
unitConfigurationParsed = do
  cfg ← setup "test/config.yaml"
  cfgInternal ← readMVar $ getCG cfg
  cfgInternal @? = ConfigurationInternal
    { cgMinSeverity      = Info
    , cgMapSeverity      = HM.fromList [ ("iohk.startup", Debug)
    , ("iohk.background.process", Error)
    , ("iohk.testing.uncritical", Warning)
    ]
    , cgMapSubtrace      = HM.fromList [ ("iohk.benchmarking",
    , ObservableTrace [GhcRtsStats, MonotonicClock])
    , ("iohk.deadend", NoTrace)
    ]
    , cgOptions          = HM.fromList
    [ ("mapSubtrace",
    , HM.fromList [ ("iohk.benchmarking",
    , Object (HM.fromList [ ("tag", String "ObservableTrace")
    , ("contents", Array $ V.fromList
    [ String "GhcRtsStats"
    , String "MonotonicClock" ]))]
    , ("iohk.deadend", String "NoTrace"))
    , ("mapMonitors", HM.fromList [ ("chain.creation.block", Array $ V.fromList
    [ Object (HM.fromList [ ("monitor", String "((time > (23 s)) Or (time < (1
    , Object (HM.fromList [ ("actions", Array $ V.fromList
    [ String "AlterMinSeverity \"chain.creation\" Debug" ]))]
    , ("#aggregation.critproc.observable", Array $ V.fromList
    [ Object (HM.fromList [ ("monitor", String "(mean >= (42))"))
    , Object (HM.fromList [ ("actions", Array $ V.fromList
    [ String "CreateMessage \"exceeded\" \"the observable has been t
    , String "AlterGlobalMinSeverity Info" ]))]
    , ("mapSeverity", HM.fromList [ ("iohk.startup", String "Debug")
    , ("iohk.background.process", String "Error")
    , ("iohk.testing.uncritical", String "Warning"))
    , ("mapAggregatedkinds", HM.fromList [ ("iohk.interesting.value",
    , String "EwmaAK {alpha = 0.75}")
    , ("iohk.background.process",
    , String "StatsAK"))
    , ("cfokey", HM.fromList [ ("value", String "Release-1.0.0"))

```

```
(("mapScribes", HM.fromList [("iohk.interesting.value",  
    Array $ V.fromList [String "StdoutSK::stdout"  
        , String "FileTextSK::testlog"])  
    , ("iohk.background.process", String "FileTextSK::testlog")]))  
, ("mapBackends", HM.fromList [("iohk.interesting.value",  
    Array $ V.fromList [String "EKGViewBK"  
        , String "AggregationBK"  
            ]))]  
  
cgMapBackend = HM.fromList [("iohk.interesting.value"  
    , [EKGViewBK  
        , AggregationBK  
            ]  
    )  
]  
  
cgDefBackendKs = [KatipBK]  
cgSetupBackends = [  
    AggregationBK  
    ,  
    EKGViewBK  
    , KatipBK]  
  
cgMapScribe = HM.fromList [("iohk.interesting.value",  
    ["StdoutSK::stdout", "FileTextSK::testlog"])  
    , ("iohk.background.process", ["FileTextSK::testlog"])]  
]  
  
cgMapScribeCache = HM.fromList [("iohk.interesting.value",  
    ["StdoutSK::stdout", "FileTextSK::testlog"])  
    , ("iohk.background.process", ["FileTextSK::testlog"])]  
]  
  
cgDefScribes = ["StdoutSK::stdout"]  
cgSetupScribes = [ScribeDefinition  
    {scKind     = FileTextSK  
    , scName      = "testlog"  
    , scPrivacy   = ScPrivate  
    , scRotation  = Just $ RotationParameters  
        {rpLogLimitBytes = 25000000  
        , rpMaxAgeHours  = 24  
        , rpKeepFilesNum = 3  
            }  
    }  
    , ScribeDefinition  
    {scKind = StdoutSK  
    , scName = "stdout"  
    , scPrivacy = ScPublic  
    , scRotation = Nothing  
    }  
    ]  
  
cgMapAggregatedKind = HM.fromList [("iohk.interesting.value", EwmaAK {alpha = 0.75})  
    , ("iohk.background.process", StatsAK)]  
]  
  
cgDefAggregatedKind = StatsAK
```

```

, cgMonitors      = HM.fromList [ ("chain.creation.block", ((OR (Compare "time" ((>), (Agg.Sec
, ["AlterMinSeverity \"chain.creation\" Debug"]
)
)
), ("#aggregation.critproc.observable", (Compare "mean" ((≥), (Agg.PureI 4
, ["CreateMessage \"exceeded\" \"the observable has been too long
, "AlterGlobalMinSeverity Info"]
)
)
)
]
, cgPortEKG       = 12789
, cgPortGUI       = 0
}

```

Test caching and inheritance of Scribes.

```

unitConfigurationCheckScribeCache :: Assertion
unitConfigurationCheckScribeCache = do
  configuration ← empty
  let defScribes = ["FileTextSK::node.log"]
  setDefaultScribes configuration defScribes
  let scribes12 = ["StdoutSK::stdout", "FileTextSK::out.txt"]
  setScribes configuration "name1.name2" $ Just scribes12
  scribes1234 ← getScribes configuration "name1.name2.name3.name4"
  scribes1 ← getScribes configuration "name1"
  scribes1234cached ← getCacheScribes configuration "name1.name2.name3.name4"
  scribesXcached ← getCacheScribes configuration "nameX"
  assertBool "Scribes for name1.name2.name3.name4 must be the same as name1.name2" $
    scribes1234 == scribes12
  assertBool "Scribes for name1 must be the default ones" $
    scribes1 == defScribes
  assertBool "Scribes for name1.name2.name3.name4 must have been cached" $
    scribes1234cached == Just scribes1234
  assertBool "Scribes for nameX must not have been cached since getScribes was not called" $
    scribesXcached == Nothing

```

Test operations on Configuration.

```

unitConfigurationOps :: Assertion
unitConfigurationOps = do
  configuration ← defaultConfigStdout
  defBackends ← getDefaultBackends configuration
  setDefaultAggregatedKind configuration $ EwmaAK 0.01
  -- since loggename does not exist the default must be inherited
  defAggregatedKind ← getAggregatedKind configuration "non-existent loggename"
  setAggregatedKind configuration "name1" $ Just StatsAK
  name1AggregatedKind ← getAggregatedKind configuration "name1"
  setEKGport configuration 11223
  ekgPort ← getEKGport configuration
  setGUIport configuration 1080

```



```

guiPort ← getGUIport configuration
assertBool "Default backends" $
  defBackends ≡ [KatipBK]
assertBool "Default aggregated kind" $
  defAggregatedKind ≡ EwmaAK 0.01
assertBool "Specific name aggregated kind" $
  name1AggregatedKind ≡ StatsAK
assertBool "Set EKG port" $
  ekgPort ≡ 11223
assertBool "Set GUI port" $
  guiPort ≡ 1080

```

2.4.5 Rotator

```

tests :: TestTree
tests = testGroup "testing Trace" [
  property_tests
]
property_tests :: TestTree
property_tests = testGroup "Property tests" [
  testProperty "rotator: file naming" propNaming
# ifdef POSIX
  , testProperty "rotator: cleanup" $ propCleanup $ rot n
# endif
]
# ifdef POSIX
where
  n = 5
  rot num = RotationParameters
    { rpLogLimitBytes = 10000000 -- 10 MB
    , rpMaxAgeHours   = 24
    , rpKeepFilesNum  = num
    }
# endif

```

Check that the generated file name has only 15 digits added to the base name.

```

propNaming :: FilePath → Property
propNaming name = ioProperty $ do
  filename ← nameLogFile name
  return $ length filename == length name + 15

```

Test cleanup of rotator.

This test creates a random number of files with the same name but with different dates and afterwards it calls the `cleanupRotator` function which removes old log files keeping only `rpKeepFilesNum` files and deleting the others.

```

# ifdef POSIX
data LocalFilePath = Dir FilePath
  deriving (Show)
instance Arbitrary LocalFilePath where
  arbitrary = do
    start ← QC.sized $ λn → replicateM (n + 1) (QC.elements $ [ 'a' .. 'z' ])
    x ← QC.sized $ λn → replicateM n (QC.elements $ [ 'a' .. 'd' ] ++ "/" )
    pure $ Dir $ start ++ removeAdjacentAndLastSlashes x
  shrink (Dir path) = map (Dir ∘ removeAdjacentAndLastSlashes ∘ (intercalate "/" )) $
    product' $ map (filter (≠ "/")) $ map QC.shrink (splitOn "/" path)
  where
    product' :: [[a]] → [[a]]
    product' = mapM (λx → x >> return)
removeAdjacentAndLastSlashes :: FilePath → FilePath
removeAdjacentAndLastSlashes = concat ∘ filter (≠ "/") ∘ groupBy (\_b → b ≠ '/')
data SmallAndLargeInt = SL Int
  deriving (Show)
instance Arbitrary SmallAndLargeInt where
  arbitrary = do
    QC.oneof [smallGen
      ,largeGen
    ]
  where
    smallGen :: QC.Gen SmallAndLargeInt
    smallGen = do
      QC.Small x ← (QC.arbitrary :: QC.Gen (QC.Small Int))
      pure $ SL $ abs x
    largeGen :: QC.Gen SmallAndLargeInt
    largeGen = do
      let maxBoundary = 0010000000000000 -- 10 years for the format which is used
          minBoundary = 000000000010000 -- 1 hour for the format which is used
      x ← QC.choose (minBoundary, maxBoundary)
      pure $ SL x
  shrink _ = []
data NumFiles = NF Int deriving (Show)
instance Arbitrary NumFiles where
  arbitrary = QC.oneof [return (NF 0), return (NF 1), return (NF 5), return (NF 7)]
propCleanup :: RotationParameters → LocalFilePath → NumFiles → SmallAndLargeInt → Property
propCleanup rotationParams (Dir filename) (NF nFiles) (SL maxDev) = QC.withMaxSuccess 20 $ ioProperty
  tmpDir0 ← getTemporaryDirectory
  let tmpDir = tmpDir0 </> "rotatorTest.base"
  let path = tmpDir </> filename
  -- generate nFiles different dates
  now ← getCurrentTime
  let tsnow = formatTime defaultTimeLocale tsformat now
  deviations ← replicateM nFiles $ QC.generate $ QC.choose (1, maxDev + 1)
  -- TODO if generated within the same sec we have a problem
  let dates = map show $ scanl (+) (read tsnow) deviations
      files = map (λa → path ++ ('-' : a)) dates

```

```

    sortedFiles = reverse $ sort files
    keepFilesNum = fromIntegral $ rpKeepFilesNum rotationParams
    toBeKept = reverse $ take keepFilesNum sortedFiles
    createDirectoryIfMissing True $ takeDirectory path
    forM_ (files) $ \f → openFile f WriteMode
    cleanupRotator rotationParams path
    filesRemained ← listLogFiles path
    let kept = case filesRemained of
        Nothing → [ ]
        Just l → NE.toList l
    removeDirectoryRecursive tmpDir
    return $ kept === toBeKept
# endif

```

Index

- Aggregated, 42
 - instance of Semigroup, 42
 - instance of Show, 43
- AggregatedExpanded, 84
- AggregatedKind, 43
 - EwmaAK, 43
 - StatsAK, 43
- AggregatedMessage, 48
- Aggregation, 84
 - instance of IsBackend, 85
 - instance of IsEffectuator, 85
- appendName, 24
- Backend, 44
- BackendKind, 44
 - AggregationBK, 44
 - EKGViewBK, 44
 - KatipBK, 44
 - MonitoringBK, 44
 - SwitchboardBK, 44
- condTracing, 19
- condTracingM, 19
- Counter, 45
- Counters
 - Dummy
 - readCounters, 30
 - Linux
 - readCounters, 30
- CounterState, 46
- CounterType, 45
- debugTracer, 18
- diffCounters, 46
- diffTimeObserved, 29
- EKGView, 80
 - instance of IsBackend, 82
 - instance of IsEffectuator, 82
- evalFilters, 25
- EWMA, 42
- ewma, 89
- getMonoClock, 29
- getOptionOrDefault, 54
- IsBackend, 43
- IsEffectuator, 43
- LOContent, 47
- Log, 73
 - instance of IsBackend, 74
 - instance of IsEffectuator, 73
- logAlert, 26
- logAlertS, 26
- logCritical, 26
- logCriticalS, 26
- logDebug, 26
- logDebugS, 26
- logEmergency, 26
- logEmergencyS, 26
- logError, 26
- logErrorS, 26
- LoggerName, 47
- logInfo, 26
- logInfoS, 26
- LogMessage, 48
- LogNamed, 52
- logNotice, 26
- logNoticeS, 26
- LogObject, 47
- LogValue, 48
- logWarning, 26
- logWarningS, 26
- LOMeta, 47
- mainTraceConditionally, 69
- Measurable, 38
 - instance of Num, 39
 - instance of Show, 40
- modifyName, 24
- Monitor, 89
 - instance of IsBackend, 90
 - instance of IsEffectuator, 90
- nameCounter, 46
- natTrace, 19
- nominalTimeToMicroseconds, 29
- nullTracer, 18
- ObservableInstance, 48

- GhcRtsStats, 48
- IOStats, 48
- MemoryStats, 48
- MonotonicClock, 48
- NetStats, 48
- ProcessStats, 48
- ObserveClose, 48
- ObserveDiff, 48
- ObserveOpen, 48
- OutputKind, 48
 - TVarList, 48
- parseRepresentation, 44
- Port, 44
- PrivacyAnnotation, 48
 - Confidential, 48
 - Public, 48
- readRTSStats, 29
- Representation, 44
- RotationParameters, 50
 - rpKeepFilesNum, 50
 - rpLogLimitBytes, 50
 - rpMaxAgeHours, 50
- ScribeDefinition, 49
 - scKind, 49
 - scName, 49
 - scPrivacy, 49
 - scRotation, 49
- ScribeId, 49
- ScribeKind
 - FileJsonSK, 49
 - FileTextSK, 49
 - StderrSK, 49
 - StdoutSK, 49
- ScribePrivacy, 49
 - ScPrivate, 49
 - ScPublic, 49
- setupTrace, 27
- Severity, 50
 - Alert, 50
 - Critical, 50
 - Debug, 50
 - Emergency, 50
 - Error, 50
 - Info, 50
 - instance of FromJSON, 50
 - Notice, 50
 - Warning, 50
- shutdown, 28
- singletonStats, 42
- Stats, 40
 - instance of Semigroup, 41
- stats2Text, 41
- stdoutTrace, 25
- stdoutTracer, 18
- SubTrace, 50
 - DropOpening, 50
 - FilterTrace, 50
 - NameOperator, 50
 - NameSelector, 50
 - Neutral, 50
 - NoTrace, 50
 - ObservableTrace, 50
 - TeeTrace, 50
 - UntimedTrace, 50
- Switchboard, 68
 - instance of IsBackend, 70
 - instance of IsEffectuator, 69
 - setupBackends, 72
- Trace, 51
- traceInTVar, 25
- traceInTVarIO, 25
- traceNamedItem, 26
- traceNamedObject, 24
- Tracer, 18
 - instance of Contravariant, 18
- tracingWith, 18
- updateAggregation, 88
- withTrace, 28